
A Greedy PDE Router for Blending Neural Operators and Classical Methods

Anonymous Author(s)

Affiliation

Address

email

1 **Revision note.** All text, tables and figures added in this revision are typeset in red (new figures
2 carry a red frame); everything in black is unchanged from the submitted version. The iteration-based
3 experiments of the submitted version on 31^2 and 63^2 grids with the LSTM router, and their appendix
4 analyses, have been removed: they are superseded by the wall-clock study of Sections E and 7,
5 which covers the same solvers and equations (plus anisotropic and variable-coefficient diffusion) on
6 128^2 – 512^2 grids and also reports the iteration-based metrics. Only the Dirichlet experiment of the
7 submitted version is kept (Section D).

8 **Status of the confirmatory re-run (draft for review).** In response to three rounds of review, every
9 wall-clock result is being regenerated under a revised, frozen protocol (Section E.1): compiled
10 classical kernels, a corrected forcing sampler, per-operation costs measured in the form in which
11 they are charged, the deployed rule equal to Algorithm 1 on macro-actions, HINTS with its two-
12 dimensional period $\tau = 15$, residual-decay controls, classical baselines timed in the same session
13 as the policies, a drift guard, two-sided tests with bootstrap intervals, a fresh test seed and three
14 warmed-up timed replays. At the time of this draft the following are *final*: the theory corrections
15 (Sections 4 and 5 and their proofs) and the wording of every claim, which is macro-driven and will
16 be recomputed from the complete tables. Everything marked **[pending]** in the text or shown as a
17 **[results pending]** table is still running, in this order: the 128^2 pairwise cells with their baselines;
18 the 128^2 ensembles; all 256^2 and 512^2 cells; usage traces, seed trials, overheads, the assumption,
19 screening and theorem checks and the discretisation study; variable-coefficient diffusion at 128^2 and
20 256^2 (Section E.5); and the decision-granularity ablation (Section E.6). Summary ranges quoted
21 in the abstract and in Section 7 are computed from the cells available at build time. Where a first
22 confirmatory run (before the protocol revision) had produced a number that is quoted in the text, it is
23 labelled as such.

Abstract

24 When solving PDEs, classical numerical solvers are often computationally ex-
25 pensive, while machine learning methods can suffer from spectral bias, failing to
26 capture high-frequency components. Designing an optimal hybrid iterative solver–
27 where, at each iteration, a solver is selected from an ensemble of solvers to leverage
28 their complementary strengths–poses a challenging combinatorial problem. While
29 greedy selection is desirable for its constant-factor approximation guarantee to the
30 optimal solution under Lipschitz and **shared-eigenbasis** assumptions, it requires
31 knowledge of the true error at each step, which is unavailable in practice. We
32 address this by proposing an approximate greedy router that efficiently mimics
33 a greedy approach to solver selection. Empirical results on the Poisson, convec-
34 tion–diffusion and anisotropic diffusion equations show that our method reduces
35 the area-under-the-curve (AUC) of the error trajectory relative to single-solver
36 baselines and existing hybrid approaches such as HINTS in every setting, reaching
37 comparable error levels in substantially fewer iterations with more stable error
38 decay. **With a cost-aware instantiation of the greedy rule, a lightweight router, and**
39 **a band-limited scale-equivariant corrector, we also measure end-to-end wall-clock**

time on 128^2 to 512^2 grids with compiled classical kernels, four equations (Poisson, convection–diffusion, anisotropic and variable-coefficient diffusion) and classical solvers from Jacobi to geometric multigrid, against HINTS, the best fixed schedule and the best residual-decay rule chosen per setting on the test set, and multigrid, preconditioned Krylov and direct solves timed in the same session. At truncation-level accuracy the learned router reproduces, without any search, the best static schedule found by a test-set search (one corrector call, then smoothing), to within $0.95\times$ – $0.98\times$ of it, and is $1.32\times$ – $3.58\times$ faster than HINTS with its published period; below truncation level it tracks the cost-aware oracle and beats HINTS and the one-shot schedule by large factors, but it is faster than the best fixed schedule in only 1 of 5 pairings at 128^2 , and a compiled geometric multigrid remains the strongest classical method there; routing over ensembles of classical solvers is free at truncation level and, at the coarse decision granularity that cheap compiled sweeps impose, gains nothing over the best single solver at 128^2 [256², 512² and finer-granularity cells pending].

1 Introduction

Natural phenomena and engineered systems are often governed by ordinary and partial differential equations (PDEs). Solving these equations enables a variety of tasks - predicting the evolution of the system through simulation (e.g., forecasting weather using the Navier-Stokes equations) [Kalnay, 2003, Bauer et al., 2015, Staniforth and Côté, 1991], addressing control problems (e.g., optimizing heat shield design for spacecrafts using heat transfer equations) [Anderson, 1989, Tröltzsch, 2010], and tackling inverse problems based on external measurements (e.g., reconstructing brain activity from EEG data using electrophysiological models) [Baillet et al., 2001, Grech et al., 2008].

Traditionally, PDEs are solved using finite difference methods [Smith, 1985, LeVeque, 2007], which discretize the spatial and/or temporal domain and approximate the solution by solving a system of equations at the discrete grid points. Similarly, finite element methods [Hughes, 2003, Bathe, 2006, Brenner, 2008] construct a system of equations based on local basis functions, and then rely on numerical linear algebra techniques—such as the Jacobi and Gauss-Seidel method [Saad, 2003]—to compute the solution. These iterative methods, however, **must be re-run from scratch for every new** initial condition, boundary condition, or forcing function, as even minor changes to any of these parameters require re-solving the entire system of equations from scratch.

These challenges have motivated the development of neural operators [Kovachki et al., 2023]—a class of machine learning models that aim to learn the solution operator directly, enabling fast inference across a range of parameters. Neural operators lift the classical linear layer in neural networks to an operator—typically a kernel integral operator—between function spaces. Despite strong approximation guarantees [Chen and Chen, 1995], neural operators are prone to spectral bias [Liu and Cai, 2021, Liu et al., 2024, You et al., 2024, Khodakarami et al., 2025, Xu et al., 2025], similar to traditional neural networks [Rahaman et al., 2019, Xu et al., 2019, Luo et al., 2019]. As a result, they favor learning low-frequency components of the solution function and often struggle to capture high-frequency features that iterative solvers excel at.

Recognizing this complementarity, Zhang et al. [2024] proposed HINTS, a hybrid solver that interleaves a neural operator pass every τ^{th} iteration of a classical solver, with τ fixed a priori. While HINTS reports significant empirical gains in convergence and accuracy over the standalone neural operator and Jacobi solver, this fixed schedule can be detrimental: a poorly timed neural correction may increase the error and undo recent progress.

Although state-dependent solver schedules are desirable, this naturally suggest a reinforcement learning formulation for learning routing policies. However, in our setting, the structure of iterative PDE solvers allows for a simpler approach: each solver induces a known update, and its effect can be evaluated through error-based signals. Leveraging this, we adopt a greedy rule that selects, at each iteration, the solver with the largest immediate error reduction. Such a rule can be near-optimal when the final error satisfies (weak) supermodularity—so that applying a beneficial solver earlier is at least as valuable as applying it later. To support this approach, we make the following contributions:

- We present a general hybrid PDE solver in which a routing rule chooses, at each iteration, from a set of classical solvers and neural operators. With oracle access to the true error at

every step, we show that a greedy routing rule achieves a constant-factor approximation to the optimal strategy for linear PDEs, provided the updates are error-reducing (Lipschitz with constant < 1) and zero-preserving—conditions under which weak supermodularity holds—and share a unitary eigenbasis, which gives the postfix monotonicity the guarantee also needs; for other ensembles the guarantee carries a measured slack.

- Given that the true error is not observed at test time, a convex surrogate loss is introduced, which, when minimized, enables the learned model to imitate the greedy router without access to true error information. This alignment is guaranteed through Bayes consistency.
- We empirically demonstrate that our approximate greedy solver achieves faster convergence than HINTS and individual solvers on the Poisson, convection-diffusion, anisotropic and variable-coefficient diffusion equations, in iterations and in end-to-end wall-clock time on 128^2 to 512^2 grids with compiled classical kernels and with the classical baselines timed in the same session. At truncation-level accuracy the router reproduces the best static schedule (one corrector call, then smoothing) without a search and is $1.32\times$ – $3.58\times$ faster than HINTS with its published period; below truncation level it beats HINTS and the one-shot schedule by large factors, whereas against the best fixed schedule and a residual-decay rule, both selected on the test set, the outcome depends on the pairing (1 of 5 pairings at 128^2), and geometric multigrid alone remains faster there. The construction extends to ensembles of classical solvers, where we find no gain at the decision granularity the compiled kernels impose.

2 Background

Consider the following linear PDE:

$$\mathcal{L}_x^a(u) = f, x \in D; \quad \mathcal{B}_x(u) = b, x \in \partial D \quad (1)$$

where $\mathcal{L}_x^a$ is a differential operator with respect to the spatial variable $x \in D$ parameterized by the coefficient function a , f is a forcing function, u is the PDE solution, $\mathcal{B}_x$ is the boundary operator, and b is the boundary term. Such PDEs can be solved numerically using various discretization strategies, such as finite differences [Smith, 1985, LeVeque, 2007], finite element, and spectral methods [Boyd, 2001, Canuto et al., 2006]. Here, we focus on finite differences, which replace derivatives with difference quotients (e.g., $\partial_x u(x) \approx (u(x+h) - u(x))/h$) on a uniform grid.

Formally, let h be the grid size and $G_h(D)$ be a uniform grid with spacing h over D . For a function $g : D \rightarrow \mathbb{R}$, we denote its restriction to $G_h(D)$ by $g_h \in \mathbb{R}^N$, where $N = |G_h(D)|$. If $\mathcal{L}_h^a \in \mathbb{R}^{N \times N}$ is the discretized operator, the PDE reduces to a linear system of equations

$$\mathcal{L}_h^a u_h = f_h, \quad (2)$$

with the boundary conditions incorporated in $\mathcal{L}_h^a$. Unlike much of the neural operator literature, which retains function space formulations for discretization invariance, we follow the conventions of classical numerical analysis and represent functions and operators as vectors and matrices. Discretization invariance is not essential here, since we focus on fixed-grid problems.

2.1 Iterative PDE Solvers

Direct methods like Gaussian elimination and the Thomas algorithm can be computationally expensive in high-dimensional domains when solving systems like Equation (2). In contrast, iterative methods like Jacobi and Gauss-Seidel offer computational speedups by iteratively updating the solution, gradually converging to the true solution for the PDE. The general iterative update is

$$u^{(t+1)} = u^{(t)} + C \left(f_h - \mathcal{L}_h^a u^{(t)} \right), \quad (3)$$

where $u^{(t)}$ is the t^{th} iterate of the solution and C is a preconditioning matrix. Jacobi uses $C = D^{-1}$, where D is the diagonal of $\mathcal{L}_h^a$, and Gauss-Seidel uses $C = (D + L)^{-1}$, where L denotes the strict lower triangular part. However, these methods tend to damp low frequency parts of the error slowly.

2.2 Neural Operators

Suppose we observe samples $\{(a_i, f_i, u_i)\}_{i=1}^N$ such that $(a_i, f_i) \sim \mathcal{P}$ are i.i.d. samples, and $u_i = \mathcal{G}^*(a_i, f_i)$ be generated by a deterministic solution operator $\mathcal{G}^*$. A neural operator (NO) $\mathcal{G}_\theta$ seeks to approximate $\mathcal{G}^*$ by minimizing the expected squared $L^2(D)$ error:

$$\mathcal{R}_{\text{NO}}(\mathcal{G}_\theta) = \mathbb{E}_{(a,f) \sim \mathcal{P}} \left[\int_D \|\mathcal{G}^*(a, f)(x) - \mathcal{G}_\theta(a, f)(x)\|_2^2 dx \right]$$

Neural operators use discretization-invariant layers. Prominent instantiations include Fourier Neural Operators (FNO), which use spectral transforms [Li et al., 2020a], Graph Neural Operators, which perform graph-based aggregation over sampled points [Li et al., 2020b], and DeepONets, which employ a trunk-branch decomposition to map input functions to output functions [Lu et al., 2021].

2.3 Greedy Optimization

Consider maximizing $g(S)$ over $S \subseteq \Omega$ with $|S| \leq T$ where $g : 2^\Omega \rightarrow \mathbb{R}$ is a set function defined on subsets of a set Ω . An exhaustive search is infeasible, so an efficient alternative is the greedy algorithm. It builds the solution subset iteratively by adding ω^* that yields the largest marginal gain, i.e. $\omega^* = \arg \max_{\omega \in \Omega \setminus S} g(S \cup \omega)$. When g is non-negative, monotone (adding elements never decreases the value), and submodular (diminishing returns), the greedy algorithm achieves a $(1 - e^{-1})$ approximation to the optimal solution [Nemhauser et al., 1978]. Formally, a function g is submodular if $g(A \cup \{\omega\}) - g(A) \geq g(B \cup \{\omega\}) - g(B)$ for all $A \subseteq B \subseteq \Omega$ and $\omega \in \Omega \setminus B$; intuitively, delaying an addition cannot increase its benefit. For set function minimization problems, supermodularity, where the inequality flips, plays an analogous role, and the greedy rule achieves constant-factor approximation guarantees [Liberty and Sviridenko, 2017].

The suboptimality of the greedy algorithm has been extensively studied for both set-function minimization [Bounia and Koriche, 2023] and maximization [Das and Kempe, 2018, Feige et al., 2011, Bian et al., 2017, Harshaw et al., 2019] when standard assumptions—non-negativity, monotonicity, and sub/supermodularity—are weakened. In contrast, results on greedy sequence maximization [Streeter and Golovin, 2008, Alaei et al., 2021, Zhang et al., 2015, Bernardini et al., 2020, Van Over et al., 2024, Tschitschek et al., 2017]—where the ordering of the elements affects the function—have led to sequential analogues of submodularity and monotonicity. We leverage these tools to analyze the suboptimality of the greedy solution to minimizing the final error.

3 General framework for Hybrid Solvers

To solve Equation (2), consider the following hybrid iterative update:

$$u_h^{(t+1)} = u_h^{(t)} + C_{S_t} \left(f_h - \mathcal{L}_h^a u_h^{(t)} \right) \quad (4)$$

where $\mathcal{C} = \{C_j\}_{j=1}^K$ is a set of preconditioning functions and $S_t \in [K] = \{1, \dots, K\}$ indexes the function chosen at step t . Here, we use “preconditioning function” broadly to include classical linear updates ($C_j(x) = C_j x$ where C_j is a matrix), and learned models, such as neural operators. This update generalizes the classical iterative update in Equation (3), by allowing C_j to be non-linear and enabling the preconditioning function to be adaptively chosen at every step. $\mathcal{C}$ can also accommodate parameterized solver families (e.g., different Jacobi relaxation weights), enabling adaptive selection of solver parameters. As the number of solvers K grows, the chance of selecting a more effective update increases. Furthermore, HINTS [Zhang et al., 2024] is a special case of this hybrid solver with $K = 2$, where C_1 is a neural operator and C_2 is a classical preconditioner. Its routing rule is given by $S_t = \mathbf{1}_{(t+1) \bmod \tau > 0} + 1$, which selects C_1 every τ steps (first at iteration τ , after $\tau - 1$ classical steps, as in Algorithm S3 of Zhang et al. [2024]) and C_2 otherwise.

Let $e_h^{(t)} = u_h - u_h^{(t)}$ denote the error at step t . Then,

$$e_h^{(t+1)} = u_h - u_h^{(t)} - C_{S_t} \left(\mathcal{L}_h^a u_h - \mathcal{L}_h^a u_h^{(t)} \right) = (I - C_{S_t} \circ \mathcal{L}_h^a) (e_h^{(t)})$$

where I is the identity map and $I - C_j \circ \mathcal{L}_h^a$ is the error propagation function for solver j . Here, “ $\circ$ ” denotes function composition (matrix multiplication for linear updates). The objective is to select a

sequence $S = (S_1, \dots, S_T)$ that minimizes the error norm after T steps or $\|e_h^{(T)}\|_2^2$:

$$\min_{S_t \in [K], |S| \leq T} h(S) := \left\| (I - C_{S_{|S|}} \circ \mathcal{L}_h^a) \circ \dots \circ (I - C_{S_1} \circ \mathcal{L}_h^a) \left(e_h^{(0)} \right) \right\|_2^2 \quad (5)$$

with compositions applied from right to left, so that the S_1 update acts on the initial error.

4 Greedy Algorithm

Equation (5) defines a combinatorial optimization problem with a search space exponential in T , making exact optimization intractable for large T or K . As a starting point, we consider an “omniscient” greedy algorithm that assumes access to the true initial error $e_h^{(0)}$. This is unrealistic since knowledge of $e_h^{(0)}$ could directly recover the solution via $u_h = u_h^{(0)} + e_h^{(0)}$, but it provides a useful benchmark. In Section 5, we relax this assumption using a practical learning strategy that is Bayes consistent with this omniscient rule: **in the infinite-sample limit it recovers the greedy decision at every state it is trained on, which is a per-decision statement; the sequential guarantee below is stated for the rule with access to the true error.**

As discussed in Section 2.3, when supermodularity and monotonicity hold, the greedy rule—such as the one described in Algorithm 1—enjoys constant-factor approximation guarantees. However, classical results focus on *set* functions, with only recent extensions made to sequences. Building on this line, we introduce a sequence-based notion of weak supermodularity.

Algorithm 1 Greedy Algorithm for a Hybrid PDE solver

Require: $\{C_j\}_{j=1}^K, T, \mathcal{L}_h^a, e_h^{(0)}$
 $S^0 \leftarrow \emptyset$
for $t < T$ **do**
 $S^{t+1} \leftarrow S^t \oplus \arg \min_{j \in [K]} \|(I - C_j \circ \mathcal{L}_h^a)(e_h^{(t)})\|_2^2$
 $e_h^{(t+1)} \leftarrow (I - C_{S_{t+1}} \circ \mathcal{L}_h^a)(e_h^{(t)})$
end for
return S^T

Let Ω^* denote the space of sequences with elements in Ω . For $S, S' \in \Omega^*$, we denote their concatenation as $S \oplus S'$. S is a prefix of S' or $S \preceq S'$ if $S' = S \oplus L$ for some $L \in \Omega^*$. A function $g : \Omega^* \rightarrow \mathbb{R}$ is prefix monotonically non-increasing if $g(S \oplus S') \leq g(S)$ for all $S, S' \in \Omega^*$, and postfix monotonically non-increasing if $g(S' \oplus S) \leq g(S)$ for all $S, S' \in \Omega^*$. A prefix non-increasing function g is sequence supermodular if, for all $S', S \in \Omega^* : S \preceq S'$, it holds that

$$g(S) - g(S \oplus \omega) \geq g(S') - g(S' \oplus \omega), \quad \forall \omega \in \Omega \quad (6)$$

However, $h(S)$ (defined in Equation (5)) may not satisfy this property in general. Therefore, we introduce weak sequence supermodularity. A prefix non-increasing function g is weakly supermodular with respect to $S' \in \Omega^*$ if, for any $S \in \Omega^*$, there exists $\alpha(S') \geq 1$ such that

$$g(S) - g(S \oplus S') \leq \alpha(S') \sum_{i \in [|S'|]} g(S) - g(S \oplus S'_i) \quad (7)$$

The parameter $\alpha(S')$ or the supermodularity ratio quantifies deviation from exact sequence supermodularity. Expanding the $g(S) - g(S \oplus S')$ as a telescoping sum $\sum_{i=1}^{|S'|} g(S \oplus (S'_1, \dots, S'_{i-1})) - g(S \oplus (S'_1, \dots, S'_i))$ shows that the marginal decrease from appending S'_i after its predecessors is controlled by the effect of appending S'_i directly to S . Thus, postponing the inclusion of S'_i cannot yield a significantly larger benefit than adding it earlier. The supermodularity ratio $\alpha(S)$ quantifies the extent to which future gains from delays may exceed immediate gains.

Postfix monotonicity enters the analysis only through the greedy prefixes, so we also allow a slack: g is μ -postfix monotone with respect to O along a sequence S^T if $g(S^t \oplus O) \leq \mu g(O)$ for every prefix $S^t, t < T$; a postfix non-increasing function satisfies this with $\mu = 1$, and μ can be computed exactly for short horizons (Section E.10). Having introduced these notions, we now characterize the suboptimality of greedy solutions of weakly supermodular and postfix monotonic sequence functions in Theorem 4.1.

Theorem 4.1. Let $g : \Omega^* \rightarrow \mathbb{R}$ be a non-negative, prefix non-increasing function that is weakly supermodular with respect to the optimal solution $O = \arg \min_{S \in \Omega^T} g(S)$ with a supermodularity ratio of $\alpha(O)$, and let the greedy solution of length T be S^T . Suppose g is μ -postfix monotone with respect to O along S^T ($\mu = 1$ if g is postfix non-increasing). If $\phi_T(\alpha) = (1 - \frac{1}{\alpha T})^T$, then

$$g(S^T) \leq \mu (1 - \phi_T(\alpha(O))) g(O) + \phi_T(\alpha(O)) g(\emptyset), \quad \text{i.e.,} \quad g(\emptyset) - g(S^T) \geq (1 - \phi_T(\alpha(O))) (g(\emptyset) - \mu g(O)).$$

The proof of Theorem 4.1 appears in Section A.2. As $T \rightarrow \infty$, the factor $1 - \phi_T(\alpha(O))$ decreases to $1 - e^{-1/\alpha(O)}$, so the worst-case performance of the greedy rule saturates rather than degrades with horizon, provided $\alpha(O)$ and μ stay bounded as T grows (both are defined through the optimal sequence of length T and can depend on it). The bound concerns the reduction of the error over T steps relative to the best sequence of T steps: with $\alpha = \mu = 1$ it guarantees a $(1 - e^{-1})$ fraction of the optimal reduction, uniformly over ensembles, but it is not a convergence rate—a single convergent solver applied T times already reduces the error by ρ^{2T} , which for large T is far below the $e^{-1}g(\emptyset)$ term. In the settings of Section E the greedy sequence is in fact optimal or within a few percent of the exact optimum over short horizons on most instances (Section E.10), much better than the worst case. Additionally, larger $\alpha(O)$, which indicates higher reward for delayed inclusions, loosens the suboptimality bound. Theorem 4.1 requires that the sequence objective h be weakly supermodular with respect to the optimal solution and postfix monotone—properties established in Proposition 4.2. We defer the proof to Section A.3.

Proposition 4.2. Suppose that for all $j \in [K]$, the error propagation function $I - C_j \circ \mathcal{L}_h^a$ is ρ_j -Lipschitz continuous with $\rho_j \leq 1$, that $(I - C_j \circ \mathcal{L}_h^a)(0_N) = 0_N$, and that $\sum_{i=1}^T (1 - \rho_{O_i}^2) > 0$, i.e., at least one action of the optimal sequence O is a strict contraction. Then, the function h is weakly supermodular with respect to the optimal solution O , with

$$\alpha(O) = \max \left\{ \frac{1}{T - \sum_{i=1}^T \rho_{O_i}^2}, 1 \right\}$$

Furthermore, if the error propagation maps are linear and simultaneously diagonalizable by a unitary matrix with eigenvalues of modulus at most one (the setting of Theorem 4.3), h is also postfix non-increasing; invertibility alone does not suffice (Section A.3).

The conditions of Proposition 4.2 are natural. The corrector of Section 7 leaves the modes outside its band untouched and has $\rho = 1$ exactly; the proposition still applies because the smoother in O is a strict contraction. For classical solvers, the Lipschitz constant is $\|I_N - C_j \mathcal{L}_h^a\|$, which is typically less than 1 for well-posed linear elliptic PDEs (i.e., the update is damping errors). For neural networks, Lipschitz continuity can be enforced via weight regularization [Gouk et al., 2021] and a sufficiently trained model should approximate $(\mathcal{L}_h^a)^{-1}$ well enough to make the Lipschitz constant small. The requirement $(I - C_j \circ \mathcal{L}_h^a)(0_N) = 0_N$ is both natural and desirable: it precludes spurious updates when the residual $f_h - \mathcal{L}_h^a u_h^{(t)}$ is 0. This holds by design for classical schemes, and it can be enforced for any learned model by excluding bias terms.

The form of $\alpha(O)$ in Proposition 4.2 highlights that the suboptimality factor in Theorem 4.1 is governed by the collective contraction factors of the solvers chosen by the optimal solution: as $\sum_{i=1}^T \rho_{O_i}^2 \rightarrow T$, $\alpha(O)$ grows, resulting in a weaker bound. Postfix monotonicity ($\mu = 1$) holds when the error propagation maps share a unitary eigenbasis—constant-coefficient periodic operators with Jacobi-type smoothers and band-limited spectral correctors—but it can fail for non-commuting contractions even when all of them are invertible, and a Gauss–Seidel sweep or a learned corrector need not be invertible at all. For the non-commuting solvers of our experiments no a priori bound on μ is available; Section E.10 computes μ , $\alpha(O)$ and the bound of Theorem 4.1 by exhaustive enumeration for short horizons ($T \leq 4$ macro-actions) and finds $\mu \leq [\text{pending}]$ and $\alpha(O) = 1$ on every ensemble of the study, including those with Gauss–Seidel, SOR and multigrid; at the horizons of tens to hundreds of macro-actions at which the oracle actually runs, only the worst-case constant of Proposition 4.2 is available, and it is loose.

Theorem 4.1 thus indicates that the approximation guarantee is strongest when the sequence is nearly supermodular ($\alpha(O) \approx 1$). Generally, if all error propagation maps share an eigenbasis, supermodularity is established. This occurs, for example, for linear, constant-coefficient PDEs with periodic boundary conditions when the solver ensemble consists of (damped) Jacobi iterations and

band-limited spectral correctors such as a single-layer linear Fourier neural operator, all of which are diagonal in the Fourier basis; lexicographic Gauss–Seidel and SOR are not, because their sweep order breaks translation invariance. The proof of Theorem 4.3 is deferred to Section A.5.

Proposition 4.3. *Let $\|I_N - C_j \mathcal{L}_h^a\| \leq 1$ for all $j \in [K]$ and $(I - C_j \mathcal{L}_h^a) = P \Lambda_j P^{-1}$ with a common unitary matrix P (orthogonal in the real symmetric case; the Jacobi matrix of the convection–diffusion operator is a real circulant with complex eigenvalues, hence unitarily but not orthogonally diagonalisable). Then, h is supermodular.*

Cost-aware greedy routing. Algorithm 1 treats every update as equally expensive, whereas in practice a neural-operator evaluation and a classical sweep have very different costs, and the quantity of interest is the error reached per unit of wall-clock time. We make the greedy rule cost-aware by changing its action set, not its comparison: let c_j denote the (measured) cost of one application of C_j , take the cost u of one neural-operator call as the unit, and replace each preconditioner by the macro-action $C_j^{(m_j)}$ that applies C_j for $m_j = \max\{1, \text{round}(u/c_j)\}$ consecutive iterations, so that every action not dearer than the unit costs u up to rounding ($m_j c_j / u$ within 0.5/ m_j of 1). For a linear update, m consecutive sweeps are again a stationary iteration, $(I - C_j \circ \mathcal{L}_h^a)^m = I - C_j^{(m)} \mathcal{L}_h^a$ with $C_j^{(m)} = \sum_{i=0}^{m-1} (I - C_j \mathcal{L}_h^a)^i C_j$, and for a nonlinear C_j the macro-action is the m -fold composition of its error propagation map. Hence $\{C_j^{(m_j)}\}_{j=1}^K$ is again a set of preconditioning functions in the sense of Section 4: if $I - C_j \circ \mathcal{L}_h^a$ is ρ_j -Lipschitz and zero-preserving, then so is its m_j -fold composition (with constant $\rho_j^{m_j}$), a shared eigenbasis is inherited by powers, and Theorem 4.1, Propositions 4.2 and 4.3, and the consistency result of Section 5 apply verbatim to an ensemble of macro-actions, all of which are equalised to the unit cost. The deployed rule is exactly Algorithm 1 on these macro-actions—the plain comparison of the errors after one macro-action—and a horizon of T macro-actions corresponds to a wall-clock budget of $\approx Tu$; it selects, at every decision epoch, the action with the largest error reduction per unit cost up to the rounding of the m_j . The one exception is an action dearer than the unit ($c_j > u$, hence $m_j = 1$), whose error reduction is compared per unit of cost, $(\|e_j\|/\|e^{(t)}\|)^{u/c_j}$: this is a geometric-rate heuristic rather than an equal-cost step, and the analysis does not cover it. With the compiled kernels of Section E every point smoother is cheaper than the unit on every grid and a multigrid cycle is cheaper than the unit at 128^2 and 256^2 , so the heuristic is only used for multigrid at 512^2 (and, in the granularity ablation of Section E.6, for the corrector itself); every other cell of the study runs Algorithm 1 on macro-actions, which is also the rule that the exhaustive check of Section E.10 verifies. (An earlier version applied the exponent $u/(m_j c_j)$ to every action, including the rounding error of the cheap ones; that rule is not Algorithm 1, and it was changed so that the theory covers what is deployed.) Section E uses this cost-aware instantiation throughout.

5 Approximate Greedy Router

The results in Section 4 indicate that the error reduction from the greedy solution closely matches that of the optimal sequence, but this is predicated on having access to the initial error, $e_h^{(0)}$. Inaccurate error estimates can result in poor solver decisions, causing errors to amplify. To remedy this, we learn a router r that selects solvers myopically, as in Algorithm 1, without access to true errors.

We adopt the following learning setup. Let $\mathcal{A}, \mathcal{F}, \mathcal{U} \subseteq \mathbb{R}^N$ denote the spaces of coefficient, forcing, and solution functions on the grid $G_h(D)$. We assume an application-specific data distribution $\mathcal{P}_{\mathcal{A} \times \mathcal{F}}$ over $\mathcal{A} \times \mathcal{F}$ that reflects test time conditions. During training, (a_h, f_h) is drawn from $\mathcal{P}_{\mathcal{A} \times \mathcal{F}}$, and a high-accuracy reference solution u_h is computed, providing the true per-step error. The router r is learned offline on this data; at test time, it operates without access to true errors.

At iteration t , router r selects a solver using the coefficient $a_h \in \mathcal{A}$, forcing $f_h \in \mathcal{F}$, and the current iterate $u_h^{(t)} \in \mathcal{U}$. If $r(a_h, f_h, u_h^{(t)}) = j$, the next iterate is computed with solver j , resulting in an error of $\|(I - C_j \circ \mathcal{L}_h^a)(e_h^{(t)})\|_2^2$. Learning such a router requires minimizing the following loss:

$$l_{\text{route}}(r, a_h, f_h, u_h^{(t)}, u_h) = \sum_{j=1}^K \left\| (I - C_j \circ \mathcal{L}_h^a) (u_h - u_h^{(t)}) \right\|_2^2 \mathbf{1}_{r(a_h, f_h, u_h^{(t)})=j} \quad (8)$$

with risk $\mathcal{R}_{\text{route}}(r) = \mathbb{E}_{a_h, f_h \sim \mathcal{P}_{\mathcal{A} \times \mathcal{F}}} [l_{\text{route}}(r, a_h, f_h, u_h^{(t)}, u_h)]$. For analysis, we fix the iterate generation via teacher-forcing [Williams and Zipser, 1989, Lamb et al., 2016]: during training, the iterates fed to the router are produced by an oracle greedy rollout, not by the router’s own past choices. Formally, with $u_h^{(0)} = 0_N$, for $t \geq 1$,

$$j_t^* \in \operatorname{argmin}_{j \in [K]} \left\| (I - C_j \circ \mathcal{L}_h^a) \left(e_h^{(t-1)} \right) \right\|_2^2, \quad u_h^{(t)} = u_h^{(t-1)} + C_{j_t^*} \left(f_h - \mathcal{L}_h^a u_h^{(t-1)} \right)$$

Thus, $\{u_h^{(t)}\}$ is a deterministic function of (a_h, f_h) . At each t , l_{route} is evaluated on $r(a_h, f_h, u_h^{(t)})$ which takes in the teacher-forced iterate. Finally, the greedy choice j_{t+1}^* is used to advance $u_h^{(t+1)}$. Thus the input distribution seen by r depends on (a_h, f_h) , not on the router’s predictions. However, full teacher forcing induces a distributional mismatch at test time (exposure bias). In experiments, we mitigate this with scheduled sampling [Bengio et al., 2015]; see Section C for details.

Minimizing $\mathcal{R}_{\text{route}}(r)$ yields the following Bayes-Optimal Router:

$$r^*(a_h, f_h, u_h^{(t)}) \in \operatorname{argmin}_{j \in [K]} \left\| (I - C_j \circ \mathcal{L}_h^a) \left(e_h^{(t)} \right) \right\|_2^2 \quad (9)$$

which aligns with the update in Algorithm 1. Hence, l_{route} is consistent with learning the greedy rule.

5.1 Surrogate Loss

Since Equation (8) is discontinuous and non-convex, direct minimization is intractable in practice. We instead introduce a surrogate loss that (1) is convex, enabling efficient optimization; (2) upper bounds the original loss; and (3) is Bayes consistent, ensuring the Bayes optimal decision of the original loss is preserved upon minimization. Formally, a surrogate ϕ is considered to be Bayes consistent with respect to the loss l if

$$\lim_{n \rightarrow \infty} \mathcal{R}_\phi(f_n) - \mathcal{R}_\phi^* = 0 \implies \lim_{n \rightarrow \infty} \mathcal{R}_l(f_n) - \mathcal{R}_l^* = 0$$

where $\mathcal{R}_l(f) = \mathbb{E} [l(f(X), Y)]$ and $\mathcal{R}_l^* = \inf_f \mathcal{R}_l(f)$. In other words, in the limit of infinite data, if the risk of a sequence of learned hypotheses $\{f_n\}$ converges to the optimal risk under ϕ , it also converges to the optimal risk with respect to the original loss l .

To define a surrogate loss for the routing problem, consider a set of scoring functions $\mathbf{g} = \{g_j\}_{j=1}^K$ with $g_j : \mathcal{A} \times \mathcal{F} \times \mathcal{U} \rightarrow \mathbb{R}$ and we define the router as $r(a, f, u^{(t)}) = \operatorname{argmax}_{j \in [K]} g_j(a, f, u^{(t)})$. For example, $\mathbf{g}$ can be a neural network with K outputs. Then, minimizing the following surrogate loss yields the same decision as minimizing Equation (8):

$$\Psi(\mathbf{g}, a_h, f_h, u_h^{(t)}, u_h) = - \sum_{j=1}^K \sum_{k=1}^K \tilde{c}_k(a_h, u_h^{(t)}, u_h) \mathbf{1}_{k \neq j} \log \left(\frac{\exp(g_j(a_h, f_h, u_h^{(t)}))}{\sum_{m=1}^K \exp(g_m(a_h, f_h, u_h^{(t)}))} \right) \quad (10)$$

where $\tilde{c}_j(a_h, u_h^{(t)}, u_h) = \|(I - C_j \circ \mathcal{L}_h^a)(u_h - u_h^{(t)})\|_2^2$. The Ψ -risk is denoted by $\mathcal{R}_\Psi(\mathbf{g}) = \mathbb{E}_{a_h, f_h \sim \mathcal{P}_{\mathcal{A} \times \mathcal{F}}} [\Psi(\mathbf{g}, a_h, f_h, u_h^{(t)}, u_h)]$. The convexity of Ψ with respect to the score vector $\mathbf{g}$ follows from the convexity of the negative log-softmax function in its inputs (the surrogate is convex in the scores, not in the parameters of a network that produces them). Moreover, Ψ upper bounds l_{route} up to a constant factor, and we refer the reader to Section B.2 for the proof. Finally, Theorem 5.1 shows that Ψ achieves Bayes consistency with respect to l_{route} ; the proof can be found in Section B.3.

Theorem 5.1. *Let the per-step costs be normalised state by state, $\hat{c}_j = \tilde{c}_j / \sum_{k=1}^K \tilde{c}_k$, at every state at which $\sum_k \tilde{c}_k > 0$, i.e., at which at least one solver does not annihilate the error (states at which every solver annihilates it carry no routing loss and are discarded), and let $\hat{l}_{\text{route}}$ and $\hat{\Psi}$ denote the routing loss and the surrogate with $\tilde{c}_j$ replaced by $\hat{c}_j$; the normalisation does not change the Bayes-optimal decision equation 9 at any state. Then, for any collection of solvers $\{C_j\}_{j=1}^K$ and linear discrete operator $\mathcal{L}_h^a$, $\hat{\Psi}$ is a Bayes consistent surrogate for $\hat{l}_{\text{route}}$.*

Theorem 5.1 is a cost-sensitive analogue of the classical Bayes-consistency of multiclass cross-entropy for 0-1 loss: in the infinite-sample limit, minimizing cross-entropy recovers the true conditional

class probabilities, so the induced decision is Bayes optimal. Our result extends this to cross-entropy with instance-dependent weights $\sum_{k \neq j}^K \hat{c}_k(a_h, u_h^{(t)}, u_h)$. The normalisation is what the router is trained with (Section E.1) and what makes the calibration constants meaningful: an earlier version assumed uniform absolute bounds $E_{\min} \leq \max_j \tilde{c}_j \leq \bar{E}$, which cannot hold along a convergent solve, where every $\tilde{c}_j \rightarrow 0$; with the normalised weights the only hypothesis is that at least one solver leaves a positive fraction of the error at every visited state, which Section E.10 checks. Under this condition, minimizing $\hat{\Psi}$ recovers the Bayes-optimal router of Equation (9), i.e., the greedy solution of Equation (5). Following the work of Mao et al. [2024], the proof uses standard conditional-risk calibration: we relate the excess risks of l_{route} and Ψ and take the infinite-sample limit. This is a per-decision statement on the state distribution the router is trained on: it says that minimising Ψ recovers the greedy decision at each visited state, not that the induced sequence, its time to tolerance, or its distance to the optimal sequence are controlled; the latter would require a sequential argument about the router’s own state distribution, which we do not provide and mitigate empirically with DAGger (Section E). The deployed router, moreover, does not see $(a_h, f_h, u_h^{(t)})$ but a few scalar features of the residual history (Section E.1); with respect to those features the Bayes-optimal router minimises the conditional expectation of the normalised costs given the features, which can differ from the full-information greedy decision when states with different error spectra share the same features, and the normalisation preserves the decision state by state but not necessarily after that averaging. Theorem 5.1 therefore concerns the full-information router, and the price of the feature compression is measured rather than proved, by the agreement between the deployed router and the oracle (Section E).

6 Related Works

Hybrid PDE Solvers: Early data-driven solvers sought convergence guarantees by predicting parameters—e.g., preconditioning matrices, multi-grid smoothers or restriction matrices—within iterative schemes [Taghibakhshi et al., 2021, Caldana et al., 2024, Kopaničáková and Karniadakis, 2025, Huang et al., 2022, Katrutsa et al., 2020]. These works, however, do not leverage the neural surrogates’ ability to generalize across varying coefficients and forcings highlighted in Section 2.2 and thus, offer only modest speedups over classical solvers. HINTS [Zhang et al., 2024] introduced hybrid solvers that interleave a classical method with a pre-trained DeepONet on a fixed schedule (e.g., 24 Jacobi steps, then one DeepONet correction), achieving faster convergence than the standalone numerical solver. Subsequent works extend this idea to new geometries [Kahana et al., 2023] and analyze error mode damping, replacing DeepONet with alternatives such as MIONet [Jin et al., 2022] [Hu and Jin, 2025] or FNO [Li et al., 2020a, Cui et al., 2022]. However, these hybrid solvers are limited in two ways: they, firstly, only combine the trained surrogate with a *single* numerical solver, and secondly rely on a fixed, heuristic schedule. Our proposed method addresses these shortcomings, resulting in significant empirical improvements (Section 7).

Model Routing: Routing [Shnitzer et al., 2023, Hu et al., 2024, Ding et al., 2024, Huang et al., 2025] selects, for each input, a model from a fixed set to optimize a task metric under cost or latency constraints. Simple heuristics—e.g., thresholding a cheap model’s uncertainty estimates [Chuang et al., 2024, 2025]—often poorly balance the cost-accuracy tradeoff, motivating many systems to instead learn a router that maps inputs to the best model [Hari and Thomson, 2023, Mohammadshahi et al., 2024, Šakota et al., 2024]. We similarly learn a router, but over numerical and neural solvers for PDEs at the iteration level. Our work is the first to apply routing to the problem of learning hybrid PDE solvers, in contrast to prior approaches that use fixed schedules. Additionally, by exploiting the algebraic structure of PDE solvers, we derive theoretical guarantees for our routing strategy (Section 4), unlike typical model-routing settings.

Mixture of Experts: Classical mixture-of-experts (MoE) [Jacobs et al., 1991, Jordan and Jacobs, 1994] uses a gating network to combine the outputs of multiple experts, trained jointly with the gate. In modern LLMs, MoE instead performs sparse, token-level routing to a subset of in-layer experts, enabling large capacity without proportional compute, typically with load-balancing mechanisms [Shazeer et al., 2017, Lepikhin et al., 2020, Fedus et al., 2022, Jiang et al., 2024]. Our method can be regarded as gating of a different kind: “experts” (solvers) are not jointly trained with the router, and a single router is reused across iterations, unlike MoE which commonly employs layer-specific gates.

7 Experiments

We evaluate the approximate greedy router on four PDEs posed on the unit domain $D = [0, 1]^2$ with periodic boundary conditions:

Poisson: $-\Delta u = f$, **ConvDiff:** $-\Delta u + 20 \partial_{x_1} u + 20 \partial_{x_2} u = f$, **AnisoDiff:** $-\epsilon \partial_{x_1}^2 u - \partial_{x_2}^2 u = f$, $\epsilon = 10^{-2}$, **VarCoef**

where f is centred to satisfy the compatibility condition $\int_D f(\mathbf{x}) d\mathbf{x} = 0$ and is drawn from a zero-mean hierarchical Gaussian random field with covariance operator $\alpha(-\Delta + \beta I)^{-\gamma}$ and sample-dependent (α, β, γ) (Section C). Anisotropic diffusion is the standard problem on which point relaxation fails to smooth: an error mode oscillatory in x_1 and smooth in x_2 is damped by only $(1 - \epsilon)/(1 + \epsilon) \approx 0.98$ per Jacobi sweep. Variable-coefficient diffusion uses a fixed smooth coefficient field $a(\mathbf{x}) = \exp(\kappa z(\mathbf{x}))$ with $\max a / \min a = 10$: its discrete operator is not diagonalised by the FFT, so no direct spectral solve exists, the corrector’s band-limited inverse has no closed form and is fitted, and the fitted corrector is only approximately accurate (Section E.5). The domain is discretised on 128^2 , 256^2 and 512^2 grids; Dirichlet boundary conditions are not covered by this study (Section D). Because iteration counts favour methods whose iterations are expensive, our primary metric is the end-to-end wall-clock time to a target relative error $\varepsilon = \|e_h^{(t)}\| / \|u_h\|$ on 64, 32 and 16 test instances per grid, with $\varepsilon = h^2$ —a grid-dependent tolerance that lies below the measured discretisation error of every test instance, so that further algebraic refinement is very unlikely to improve the computed solution (Section E.1)—highlighted; every comparison is paired over instances and tested with a two-sided Wilcoxon signed-rank test with a Holm correction. We also report the iteration-based metric of the original submission, the area under the error curve, $\text{AUC} = \sum_{t=1}^T \|e_h^{(t)}\| / \|u_h\|$ over $T = 300$ iterations. Section E gives the full protocol and collects every result per equation.

Setup. The classical solvers are Jacobi, damped Jacobi ($\omega = 0.67$), Gauss–Seidel (GS), symmetric GS (SymGS), SOR ($\omega = 1.5$), line Gauss–Seidel (anisotropic diffusion only: every grid line along the strongly coupled axis solved exactly, the classical remedy for anisotropy) and geometric multigrid (V(2,2)-cycles with GS smoothing; line-GS smoothing for anisotropic diffusion), all implemented as compiled single-threaded C stencil kernels (a GS sweep costs $61 \mu\text{s}$ and a V-cycle $326 \mu\text{s}$ at 128^2 , measured in the form in which the benchmark charges them), so the classical baselines are not slowed by the implementation. The neural operator is a DeepONet corrector that acts on the residual through a 64×64 sensor grid with a Fourier trunk basis and is applied scale-equivariantly: its deployed form is a linear, Fourier-structured map (a least-squares fit of the band-limited inverse of the discrete operator, which it reproduces to $\approx 10^{-6}$ on the constant-coefficient equations), so one call solves the smooth band of the spectrum at the cost of about 32 Jacobi or 11 GS sweeps at 128^2 (and [pending] Jacobi sweeps at 512^2 , since its cost is nearly grid-independent) and never injects error outside its band. The router is made *cost-aware* through the cost-equalised macro-actions of Section 4 and selects from $6+K$ scalar features (log residual, its recent decay rate, the previous action and call counts) with a two-layer MLP whose decision costs $16\text{--}19 \mu\text{s}$ in situ, independently of N ; it is trained with the surrogate of Section 5.1 to imitate the cost-aware greedy oracle on oracle rollouts followed by DAGger rounds. We compare against the classical solver alone; HINTS with the period $\tau=15$ that Zhang et al. [2024] use in all of their two-dimensional experiments ($\tau=25$, their one-dimensional period, is reported in Section E); the best fixed schedule chosen *a posteriori* on the test set for each cell from HINTS with $\tau \in \{2, 5, 10, 15, 25, 50\}$, phase-shifted HINTS (first call at iteration 0) and the static one-shot schedule (corrector once at iteration 0, then the solver); the best of three residual-decay rules—deterministic controls at the router’s granularity that call the corrector again whenever the residual’s contraction over the last macro-action falls below a threshold, also chosen on the test set; the cost-aware greedy oracle (Algorithm 1 with the true error; not deployable); and the classical methods a practitioner would use on these problems without a corrector—multigrid alone (line-smoothed for anisotropic diffusion), multigrid-preconditioned CG or BiCGSTAB, and the FFT or sparse direct solve—timed inside the same replay loop as the policies. The forcings are sampled from the hierarchical GRF with the Hermitian symmetry of a real field and without Nyquist modes (Section E.1). All configurations were fixed on a development test set, and every number reported is a single confirmatory evaluation on a fresh test set with three timed replays in random order and a drift guard per instance (Section E.9); paired comparisons use two-sided Wilcoxon tests with a Holm correction and bootstrap intervals.

Paired solvers. Table 1 reports, for every pairing $\{\text{NO}, \text{solver}\}$, equation and grid, the median wall-clock time to $\varepsilon = h^2$ of HINTS, of the best fixed schedule, and of the learned router, together

with the classical methods without corrector. Two regimes appear. *At truncation level* the whole solve is one corrector call followed by a few sweeps: the router, the cost-aware oracle and the best static schedule (the corrector at iteration 0, then smoothing) take the same decisions, so the router is within a few percent of the best fixed schedule found by a test-set search—paired ratios $0.95\times\text{--}0.98\times$ at 128^2 , [pending]–[pending] at 256^2 and [pending]–[pending] at 512^2 , the difference being its own decision overhead—without any search, and it is $1.32\times\text{--}3.58\times$ faster than HINTS with its published period $\tau=15$ ([pending]–[pending] at 256^2 , [pending]–[pending] at 512^2 ; $1.57\times\text{--}4.14\times$ against $\tau=25$), whose first call comes $\tau - 1$ sweeps too late; the best residual-decay rule, which also calls the corrector first, is within $0.95\times\text{--}0.96\times$ of the router there; the classical solver alone is $195\times\text{--}664\times$ slower on the isotropic equations and never reaches h^2 on the anisotropic one. *Below truncation level* the schedule matters and the picture is mixed: down to $\varepsilon = 10^{-8}$ the router is $0.93\times\text{--}3.07\times$ faster than HINTS ($\tau=25$) and one to two orders of magnitude faster than the one-shot schedule, because it calls the corrector again when the smooth error has re-emerged (Figure 1), but it beats the best fixed schedule in only 1 of the 5 pairings at 128^2 (paired ratios $0.83\times\text{--}1.24\times$) and the best residual-decay rule in 5 ($1.28\times\text{--}1.92\times$): the best static schedule there is a phase-shifted HINTS with a short period, which calls the corrector every 5–10 sweeps, more often than a cost-equalised macro-action of 11 Gauss–Seidel or 32 Jacobi sweeps allows the router to (Section E.6 examines this granularity limit). The router tracks the cost-aware oracle within $1.06\times$ in median time; the stability of its decisions over five independent training seeds is reported in Section E [pending]. Against the classical methods without corrector, the compiled multigrid is a strong competitor: at 128^2 the stationary hybrids match it at h^2 (router $0.97\times\text{--}1.09\times$ of a stand-alone V-cycle solve) and are slower than it below h^2 ($0.04\times\text{--}0.92\times$), whereas routing over {NO, multigrid} is $0.99\times\text{--}0.99\times$ and $1.14\times\text{--}1.14\times$ of multigrid alone at h^2 and 10^{-8} ; the router is $0.95\times\text{--}1.07\times$ of multigrid-preconditioned CG or BiCGSTAB at h^2 , and on the constant-coefficient equations the FFT direct solve remains $2.98\times\text{--}2.98\times$ faster than any iterative method including ours, whereas on variable-coefficient diffusion, which has no such solve, the router is compared with the sparse direct solve (Section E.5). The iteration-based metrics tell the same story: at 128^2 the router’s AUC over $T = 300$ iterations is orders of magnitude below HINTS’ and the solver-only baseline’s for every pairing and equation (paired t -tests on log AUC, $p < 10^{-3}$; Tables 8, 14, 20 and 26).

Solver ensembles. The same construction routes over ensembles $\text{NO} \cup \mathcal{W}$ of several classical solvers. Since the greedy rule is a per-decision comparison, a larger ensemble can only help where the error passes through regimes that favour *different* members; an exhaustive oracle-level screening of every subset of up to three members (Section E.7) shows that this never happens at $\varepsilon = h^2$ —one corrector call and one macro-action of the cheapest smoother already reach truncation level—and, in the development runs, happened at 10^{-8} through a nonstationary Jacobi schedule (undamped Jacobi on the band-edge modes, damped Jacobi on the near-Nyquist modes). We evaluate the learned router over the nested sets {Jacobi}, {Jacobi (0.67)}, {Jacobi, Jacobi (0.67)} and {Jacobi, Jacobi (0.67), GS} and the four larger sets of Table 30 on every equation, with the single-member routers evaluated in the same session on the same instances. At 128^2 the confirmatory run shows no separation: at $\varepsilon = 10^{-8}$ the ensemble router is significantly faster than the best single-member router in [pending] of [pending] ensemble cells, and the ensemble *oracle* is itself [pending]–[pending] of the best single-member oracle, because with the compiled kernels a macro-action is about forty Jacobi sweeps and a single damped-Jacobi macro-action already removes the whole band the corrector leaves at 128^2 ; the nonstationary schedule of the development runs, where a macro-action was about ten sweeps, has nothing left to exploit. Whether the gain returns at a finer decision granularity (Section E.6) and on the 256^2 grids, where the band is twice as wide, is [pending]. Enlarging the set is free at h^2 (every ensemble router within [pending]–[pending] of the best single-member router in work units), whereas at 10^{-8} the ensemble routers are [pending]–[pending] of it: the greedy oracle is not monotone in its action set, and its imitation degrades among near-tied members (Section E.7). Finally, we verify numerically that the assumptions of Sections 4 and 5 hold in all these settings, in the energy norm for the non-normal sweeps (Section E.10).

Every result per equation and grid, the ensemble study, the accounting of overheads and training cost, and the verification of the theoretical assumptions are collected in Section E; Section D reports the Dirichlet experiment of the original submission.

Table 1: Median wall-clock time to truncation-level relative error $\varepsilon = h^2$ for every solver pairing $\{\text{NO, solver}\}$, equation and grid ($64, 32$ and 16 test instances at $128^2, 256^2$ and 512^2): HINTS with the published period $\tau=25$, HINTS with the best period chosen per cell on the test set itself (subscript), and the learned router (paired median speedup over HINTS-25 / over the best period in parentheses; bold: faster than both). The classical solver alone is 10^2 – $10^3 \times$ slower than the router on the isotropic equations and never reaches h^2 on the anisotropic one (Section E). The last rows of each block give geometric multigrid alone and the best multigrid-preconditioned Krylov method, without corrector, on the same instances. Anisotropic diffusion is not run at 512^2 : its corrector keeps full resolution along x_1 , and the resulting $65,536 \times 65,536$ branch matrix does not fit in memory (Section E); multigrid is not run on it either (Section E.1). $\dagger n$: n instances did not reach h^2 within the iteration cap.

Equation	Pairing	128×128			256×256			512×512		
		HINTS-15	best schedule	router (ours)	HINTS-15	best schedule	router (ours)	HINTS-15	best schedule	router (ours)
Poisson	Jacobi	1.02 ms	718 $\mu\text{s}_{\tau=15, t_0=0}$	758 μs (1.37 $\times$ / 0.96 $\times$)	–	–	–	–	–	–
	Jacobi (0.67)	981 μs	711 $\mu\text{s}_{\tau=5, t_0=0}$	743 μs (1.32 $\times$ / 0.98 $\times$)	–	–	–	–	–	–
	GS	1.58 ms	711 $\mu\text{s}_{\tau=15, t_0=0}$	752 μs (2.06 $\times$ / 0.95 $\times$)	–	–	–	–	–	–
	SymGS	1.99 ms	726 $\mu\text{s}_{\text{one-shot}}$	772 μs (2.53 $\times$ / 0.95 $\times$)	–	–	–	–	–	–
	SOR (1.5)	3.03 ms	711 $\mu\text{s}_{\tau=25, t_0=0}$	764 μs (3.58 $\times$ / 0.96 $\times$)	–	–	–	–	–	–
	Multigrid	933 μs	821 $\mu\text{s}_{\tau=25, t_0=0}$	874 μs (0.99 $\times$ / 0.97 $\times$)	–	–	–	–	–	–
	Multigrid alone (no corrector)			1.01 ms	–	–	–	–	–	–
	PCG (MG) (no corrector)			978 μs	–	–	–	–	–	–
	FFT direct solve (exact)			246 μs	–	–	–	–	–	–

8 Discussion

We have introduced an adaptive method for selecting solvers that efficiently minimizes the final error of a PDE in iterative methods, opening many directions for future work. Our current DeepONet is trained without accounting for its downstream role in correction term prediction; jointly training the ML solver and router could yield larger gains. Another avenue is to use reinforcement learning to learn a cost-aware routing strategy that optimizes error under compute budgets; [Sections 4 and E](#) show that a myopic cost-aware variant of the greedy rule—[Algorithm 1 on cost-equalised macro-actions](#)—already reproduces the best static schedule at truncation level and gives mixed results below it, and RL could extend it to non-myopic budgeted routing. RL would additionally enable extending the ensemble routing to *continuous* parameterization of solvers. Finally, framing routing as an online learning problem could enable adaptation to unknown deployment conditions.

References

- Eugenia Kalnay. *Atmospheric modeling, data assimilation and predictability*. Cambridge university press, 2003.
- Peter Bauer, Alan Thorpe, and Gilbert Brunet. The quiet revolution of numerical weather prediction. *Nature*, 525(7567):47–55, 2015.
- Andrew Staniforth and Jean Côté. Semi-lagrangian integration schemes for atmospheric models—a review. *Monthly weather review*, 119(9):2206–2223, 1991.
- John David Anderson. *Hypersonic and high temperature gas dynamics*. Aiaa, 1989.
- Fredi Tröltzsch. *Optimal control of partial differential equations: theory, methods, and applications*, volume 112. American Mathematical Soc., 2010.
- Sylvain Baillet, John C Mosher, and Richard M Leahy. Electromagnetic brain mapping. *IEEE Signal processing magazine*, 18(6):14–30, 2001.
- Roberta Grech, Tracey Cassar, Joseph Muscat, Kenneth P Camilleri, Simon G Fabri, Michalis Zervakis, Petros Xanthopoulos, Vangelis Sakkalis, and Bart Vanrumste. Review on solving the inverse problem in eeg source analysis. *Journal of neuroengineering and rehabilitation*, 5:1–33, 2008.
- Gordon D Smith. *Numerical solution of partial differential equations: finite difference methods*. Oxford university press, 1985.
- Randall J LeVeque. *Finite difference methods for ordinary and partial differential equations: steady-state and time-dependent problems*. SIAM, 2007.




Figure 1: Corrector usage on the 128×128 grid: fraction of test runs that execute a corrector call at each iteration, for the learned router, the cost-aware oracle and HINTS ($\tau=25$), for every solver pairing (columns) and all three equations (rows); the same figure for the 256^2 and 512^2 grids is in Section E.

- 541 Thomas JR Hughes. *The finite element method: linear static and dynamic finite element analysis*.
542 Courier Corporation, 2003.
- 543 Klaus-Jürgen Bathe. *Finite element procedures*. Klaus-Jurgen Bathe, 2006.
- 544 Susanne C Brenner. *The mathematical theory of finite element methods*. Springer, 2008.
- 545 Yousef Saad. *Iterative methods for sparse linear systems*. SIAM, 2003.
- 546 Nikola Kovachki, Zongyi Li, Burigede Liu, Kamyar Azizzadenesheli, Kaushik Bhattacharya, Andrew
547 Stuart, and Anima Anandkumar. Neural operator: Learning maps between function spaces with
548 applications to pdes. *Journal of Machine Learning Research*, 24(89):1–97, 2023.
- 549 Tianping Chen and Hong Chen. Universal approximation to nonlinear operators by neural networks
550 with arbitrary activation functions and its application to dynamical systems. *IEEE transactions on*
551 *neural networks*, 6(4):911–917, 1995.
- 552 Lizuo Liu and Wei Cai. Multiscale deepnet for nonlinear operators in oscillatory function spaces for
553 building seismic wave responses. *arXiv preprint arXiv:2111.04860*, 2021.
- 554 Xinliang Liu, Bo Xu, Shuhao Cao, and Lei Zhang. Mitigating spectral bias for the multiscale operator
555 learning. *Journal of Computational Physics*, 506:112944, 2024.
- 556 Zhilin You, Zhenli Xu, and Wei Cai. Mscalefno: Multi-scale fourier neural operator learning for
557 oscillatory function spaces. *arXiv preprint arXiv:2412.20183*, 2024.
- 558 Siavash Khodakarami, Vivek Oommen, Aniruddha Bora, and George Em Karniadakis. Mitigating
559 spectral bias in neural operators via high-frequency scaling for physical systems. *arXiv preprint*
560 *arXiv:2503.13695*, 2025.

561 Zhi-Qin John Xu, Lulu Zhang, and Wei Cai. On understanding and overcoming spectral biases of
562 deep neural network learning methods for solving pdes. *arXiv preprint arXiv:2501.09987*, 2025.

563 Nasim Rahaman, Aristide Baratin, Devansh Arpit, Felix Draxler, Min Lin, Fred Hamprecht, Yoshua
564 Bengio, and Aaron Courville. On the spectral bias of neural networks. In *International conference*
565 *on machine learning*, pages 5301–5310. PMLR, 2019.

566 Zhi-Qin John Xu, Yaoyu Zhang, Tao Luo, Yanyang Xiao, and Zheng Ma. Frequency principle:
567 Fourier analysis sheds light on deep neural networks. *arXiv preprint arXiv:1901.06523*, 2019.

568 Tao Luo, Zheng Ma, Zhi-Qin John Xu, and Yaoyu Zhang. Theory of the frequency principle for
569 general deep neural networks. *arXiv preprint arXiv:1906.09235*, 2019.

570 Enrui Zhang, Adar Kahana, Alena Kopaničáková, Eli Turkel, Rishikesh Ranade, Jay Pathak, and
571 George Em Karniadakis. Blending neural operators and relaxation methods in pde numerical
572 solvers. *Nature Machine Intelligence*, pages 1–11, 2024.

573 John P Boyd. *Chebyshev and Fourier spectral methods*. Courier Corporation, 2001.

574 Claudio Canuto, M. Yousuff Hussaini, Alfio Quarteroni, and Thomas A. Zang. *Spectral Methods:*
575 *Fundamentals in Single Domains*. Springer, 2006. ISBN 9783540307264.

576 Zongyi Li, Nikola Kovachki, Kamyar Azizzadenesheli, Burigede Liu, Kaushik Bhattacharya, Andrew
577 Stuart, and Anima Anandkumar. Fourier neural operator for parametric partial differential equations.
578 *arXiv preprint arXiv:2010.08895*, 2020a.

579 Zongyi Li, Nikola Kovachki, Kamyar Azizzadenesheli, Burigede Liu, Andrew Stuart, Kaushik
580 Bhattacharya, and Anima Anandkumar. Multipole graph neural operator for parametric partial
581 differential equations. *Advances in Neural Information Processing Systems*, 33:6755–6766, 2020b.

582 Lu Lu, Pengzhan Jin, Guofei Pang, Zhongqiang Zhang, and George Em Karniadakis. Learning
583 nonlinear operators via deepnet based on the universal approximation theorem of operators.
584 *Nature machine intelligence*, 3(3):218–229, 2021.

585 George L Nemhauser, Laurence A Wolsey, and Marshall L Fisher. An analysis of approximations for
586 maximizing submodular set functions—i. *Mathematical programming*, 14(1):265–294, 1978.

587 Edo Liberty and Maxim Sviridenko. Greedy minimization of weakly supermodular set functions.
588 In *Approximation, Randomization, and Combinatorial Optimization. Algorithms and Techniques*
589 *(APPROX/RANDOM 2017)*, pages 19–1. Schloss Dagstuhl–Leibniz-Zentrum für Informatik, 2017.

590 Louenas Bounia and Frederic Koriche. Approximating probabilistic explanations via supermodular
591 minimization. In *Uncertainty in Artificial Intelligence*, pages 216–225. PMLR, 2023.

592 Abhimanyu Das and David Kempe. Approximate submodularity and its applications: Subset selection,
593 sparse approximation and dictionary selection. *Journal of Machine Learning Research*, 19(3):
594 1–34, 2018.

595 Uriel Feige, Vahab S Mirrokni, and Jan Vondrák. Maximizing non-monotone submodular functions.
596 *SIAM Journal on Computing*, 40(4):1133–1153, 2011.

597 Andrew An Bian, Joachim M Buhmann, Andreas Krause, and Sebastian Tschieschek. Guarantees for
598 greedy maximization of non-submodular functions with applications. In *International conference*
599 *on machine learning*, pages 498–507. PMLR, 2017.

600 Chris Harshaw, Moran Feldman, Justin Ward, and Amin Karbasi. Submodular maximization beyond
601 non-negativity: Guarantees, fast algorithms, and applications. In *International Conference on*
602 *Machine Learning*, pages 2634–2643. PMLR, 2019.

603 Matthew Streeter and Daniel Golovin. An online algorithm for maximizing submodular functions.
604 *Advances in Neural Information Processing Systems*, 21, 2008.

605 Saeed Alaei, Ali Makhdoumi, and Azarakhsh Malekian. Maximizing sequence-submodular functions
606 and its application to online advertising. *Management Science*, 67(10):6030–6054, 2021.

607 Zhenliang Zhang, Edwin KP Chong, Ali Pezeshki, and William Moran. String submodular functions
608 with curvature constraints. *IEEE Transactions on Automatic Control*, 61(3):601–616, 2015.

609 Sara Bernardini, Fabio Fagnani, and Chiara Piacentini. Through the lens of sequence submodularity.
610 In *Proceedings of the International Conference on Automated Planning and Scheduling*, volume 30,
611 pages 38–47, 2020.

612 Brandon Van Over, Bowen Li, Edwin KP Chong, and Ali Pezeshki. A performance bound for
613 the greedy algorithm in a generalized class of string optimization problems. *arXiv preprint*
614 *arXiv:2409.05020*, 2024.

615 Sebastian Tschischek, Adish Singla, and Andreas Krause. Selecting sequences of items via
616 submodular maximization. In *Proceedings of the AAAI Conference on Artificial Intelligence*,
617 volume 31, 2017.

618 Henry Gouk, Eibe Frank, Bernhard Pfahringer, and Michael J Cree. Regularisation of neural networks
619 by enforcing lipschitz continuity. *Machine Learning*, 110(2):393–416, 2021.

620 Ronald J Williams and David Zipser. A learning algorithm for continually running fully recurrent
621 neural networks. *Neural computation*, 1(2):270–280, 1989.

622 Alex M Lamb, Anirudh Goyal ALIAS PARTH GOYAL, Ying Zhang, Saizheng Zhang, Aaron C
623 Courville, and Yoshua Bengio. Professor forcing: A new algorithm for training recurrent networks.
624 *Advances in neural information processing systems*, 29, 2016.

625 Samy Bengio, Oriol Vinyals, Navdeep Jaitly, and Noam Shazeer. Scheduled sampling for sequence
626 prediction with recurrent neural networks. *Advances in neural information processing systems*, 28,
627 2015.

628 Anqi Mao, Mehryar Mohri, and Yutao Zhong. Regression with multi-expert deferral. In *International*
629 *Conference on Machine Learning*, pages 34738–34759. PMLR, 2024.

630 Ali Taghibakhshi, Scott MacLachlan, Luke Olson, and Matthew West. Optimization-based algebraic
631 multigrid coarsening using reinforcement learning. *Advances in neural information processing*
632 *systems*, 34:12129–12140, 2021.

633 Matteo Caldana, Paola F Antonietti, et al. A deep learning algorithm to accelerate algebraic multigrid
634 methods in finite element solvers of 3d elliptic pdes. *Computers & Mathematics with Applications*,
635 167:217–231, 2024.

636 Alena Kopaničáková and George Em Karniadakis. Deepnet based preconditioning strategies for
637 solving parametric linear systems of equations. *SIAM Journal on Scientific Computing*, 47(1):
638 C151–C181, 2025.

639 Ru Huang, Ruipeng Li, and Yuanzhe Xi. Learning optimal multigrid smoothers via neural networks.
640 *SIAM Journal on Scientific Computing*, 45(3):S199–S225, 2022.

641 Alexandr Katrutsa, Talgat Daulbaev, and Ivan Oseledets. Black-box learning of multigrid parameters.
642 *Journal of Computational and Applied Mathematics*, 368:112524, 2020.

643 Adar Kahana, Enrui Zhang, Somdatta Goswami, George Karniadakis, Rishikesh Ranade, and Jay
644 Pathak. On the geometry transferability of the hybrid iterative numerical solver for differential
645 equations. *Computational Mechanics*, 72(3):471–484, 2023.

646 Pengzhan Jin, Shuai Meng, and Lu Lu. Mionet: Learning multiple-input operators via tensor product.
647 *SIAM Journal on Scientific Computing*, 44(6):A3490–A3514, 2022.

648 Jun Hu and Pengzhan Jin. A hybrid iterative method based on mionet for pdes: Theory and numerical
649 examples. *Mathematics of Computation*, 2025.

650 Chen Cui, Kai Jiang, Yun Liu, and Shi Shu. Fourier neural solver for large sparse linear algebraic
651 systems. *Mathematics*, 10(21):4014, 2022.

652 Tal Shnitzer, Anthony Ou, Mirian Silva, Kate Soule, Yuekai Sun, Justin Solomon, Neil Thompson,
653 and Mikhail Yurochkin. Large language model routing with benchmark datasets. *arXiv preprint*
654 *arXiv:2309.15789*, 2023.

655 Qitian Jason Hu, Jacob Bieker, Xiuyu Li, Nan Jiang, Benjamin Keigwin, Gaurav Ranganath, Kurt
656 Keutzer, and Shriyash Kaustubh Upadhyay. Routerbench: A benchmark for multi-llm routing
657 system. *arXiv preprint arXiv:2403.12031*, 2024.

658 Dujian Ding, Ankur Mallick, Chi Wang, Robert Sim, Subhabrata Mukherjee, Victor Ruhle, Laks VS
659 Lakshmanan, and Ahmed Hassan Awadallah. Hybrid llm: Cost-efficient and quality-aware query
660 routing. *arXiv preprint arXiv:2404.14618*, 2024.

661 Zhongzhan Huang, Guoming Ling, Yupei Lin, Yandong Chen, Shanshan Zhong, Hefeng Wu, and
662 Liang Lin. Routereval: A comprehensive benchmark for routing llms to explore model-level
663 scaling up in llms. *arXiv preprint arXiv:2503.10657*, 2025.

664 Yu-Neng Chuang, Prathusha Kameswara Sarma, Parikshit Gopalan, John Boccio, Sara Bolouki,
665 Xia Hu, and Helen Zhou. Learning to route llms with confidence tokens. *arXiv preprint*
666 *arXiv:2410.13284*, 2024.

667 Yu-Neng Chuang, Leisheng Yu, Guanchu Wang, Lizhe Zhang, Zirui Liu, Xuanning Cai, Yang Sui,
668 Vladimir Braverman, and Xia Hu. Confident or seek stronger: Exploring uncertainty-based
669 on-device llm routing from benchmarking to generalization. *arXiv preprint arXiv:2502.04428*,
670 2025.

671 Surya Narayanan Hari and Matt Thomson. Tryage: Real-time, intelligent routing of user prompts to
672 large language models. *arXiv preprint arXiv:2308.11601*, 2023.

673 Alireza Mohammadshahi, Arshad Rafiq Shaikh, and Majid Yazdani. Routoo: Learning to route to
674 large language models effectively. *arXiv preprint arXiv:2401.13979*, 2024.

675 Marija Šakota, Maxime Peyrard, and Robert West. Fly-swat or cannon? cost-effective language
676 model choice via meta-modeling. In *Proceedings of the 17th ACM International Conference on*
677 *Web Search and Data Mining*, pages 606–615, 2024.

678 Robert A Jacobs, Michael I Jordan, Steven J Nowlan, and Geoffrey E Hinton. Adaptive mixtures of
679 local experts. *Neural computation*, 3(1):79–87, 1991.

680 Michael I Jordan and Robert A Jacobs. Hierarchical mixtures of experts and the em algorithm. *Neural*
681 *computation*, 6(2):181–214, 1994.

682 Noam Shazeer, Azalia Mirhoseini, Krzysztof Maziarsz, Andy Davis, Quoc Le, Geoffrey Hinton, and
683 Jeff Dean. Outrageously large neural networks: The sparsely-gated mixture-of-experts layer. *arXiv*
684 *preprint arXiv:1701.06538*, 2017.

685 Dmitry Lepikhin, HyoukJoong Lee, Yuanzhong Xu, Dehao Chen, Orhan Firat, Yanping Huang,
686 Maxim Krikun, Noam Shazeer, and Zhifeng Chen. Gshard: Scaling giant models with conditional
687 computation and automatic sharding. *arXiv preprint arXiv:2006.16668*, 2020.

688 William Fedus, Barret Zoph, and Noam Shazeer. Switch transformers: Scaling to trillion parameter
689 models with simple and efficient sparsity. *Journal of Machine Learning Research*, 23(120):1–39,
690 2022.

691 Albert Q Jiang, Alexandre Sablayrolles, Antoine Roux, Arthur Mensch, Blanche Savary, Chris
692 Bamford, Devendra Singh Chaplot, Diego de las Casas, Emma Bou Hanna, Florian Bressand, et al.
693 Mixtral of experts. *arXiv preprint arXiv:2401.04088*, 2024.

694 Anqi Mao, Mehryar Mohri, and Yutao Zhong. Cross-entropy loss functions: Theoretical analysis and
695 applications. In *International conference on Machine learning*, pages 23803–23828. pmlr, 2023.

696 Michael C Mozer. A focused backpropagation algorithm for temporal pattern recognition. In
697 *Backpropagation*, pages 137–169. Psychology Press, 2013.

- 698 Anthony J Robinson and Frank Fallside. *The utility driven dynamic error propagation network*,
699 volume 11. University of Cambridge Department of Engineering Cambridge, 1987.
- 700 Paul J Werbos. Generalization of backpropagation with application to a recurrent gas market model.
701 *Neural networks*, 1(4):339–356, 1988.
- 702 Lu Lu, Xuhui Meng, Shengze Cai, Zhiping Mao, Somdatta Goswami, Zhongqiang Zhang, and
703 George Em Karniadakis. A comprehensive and fair comparison of two neural operators (with prac-
704 tical extensions) based on FAIR data. *Computer Methods in Applied Mechanics and Engineering*,
705 393:114778, 2022.
- 706 Stéphane Ross, Geoffrey Gordon, and Drew Bagnell. A reduction of imitation learning and structured
707 prediction to no-regret online learning. In *Proceedings of the Fourteenth International Confer-*
708 *ence on Artificial Intelligence and Statistics*, pages 627–635. JMLR Workshop and Conference
709 Proceedings, 2011.

710 A Proofs for Section 4

711 A.1 Proof of Proposition A.1

712 **Proposition A.1.** *Any prefix monotonically non-increasing sequence supermodular function g is*
713 *weakly supermodular with respect to all sequences $S \in \Omega^*$ with $\alpha(S) = 1$*

714 *Proof.* This proof is adapted from Liberty and Sviridenko [2017].

715 If g is sequence supermodular then,

$$\begin{aligned}
& g(S) - g(S \oplus S') \\
&= \sum_{i=1}^{|S'|} g(S \oplus (S'_1, \dots, S'_{i-1})) - g(S \oplus (S'_1, \dots, S'_i)) \\
&\stackrel{(a)}{\leq} \sum_{i=1}^{|S'|} g(S) - g(S \oplus S'_i) \\
&\leq |S'| \max_{i \in [|S'|]} g(S) - g(S \oplus S'_i)
\end{aligned}$$

716 (a) by supermodularity

717 □

718 A.2 Proof of Theorem 4.1

719 **Theorem 4.1.** *Let $g : \Omega^* \rightarrow \mathbb{R}$ be a weakly supermodular function with respect to the optimal*
720 *solution $O = \arg \min_{S \in \Omega^T} g(S)$ with a supermodularity ratio of $\alpha(O)$ and μ -postfix monotonicity*
721 *with respect to O along the greedy solution S^T of length T . If $\phi_T(\alpha) = (1 - \frac{1}{\alpha T})^T$, then*

$$g(S^T) \leq \mu(1 - \phi_T(\alpha(O)))g(O) + \phi_T(\alpha(O))g(\emptyset)$$

722 *Proof.* This proof strategy is inspired by Streeter and Golovin [2008] and Liberty and Sviridenko
723 [2017]

$$\begin{aligned}
g(S^t) - \mu g(O) &\stackrel{(a)}{\leq} g(S^t) - g(S^t \oplus O) \\
&= \sum_{i=1}^{|O|} g(S^t \oplus (o_1, \dots, o_{i-1})) - g(S^t \oplus (o_1, \dots, o_i)) \\
&\stackrel{(b)}{\leq} \alpha(O) \sum_{i=1}^{|O|} g(S^t) - g(S^t \oplus o_i) \\
&\leq \alpha(O)|O| \max_{i \in [|O|]} g(S^t) - g(S^t \oplus o_i) \\
&\leq \alpha(O)Tg(S^t) - \alpha(O)T \min_{\omega \in \Omega} g(S^t \oplus \omega) \\
&= \alpha(O)Tg(S^t) - \alpha(O)Tg(S^{t+1})
\end{aligned}$$

724 (a) by μ -postfix monotonicity along S^T , i.e., $g(S^t \oplus O) \leq \mu g(O)$; (b) by weak supermodularity with
725 respect to O (applied to the prefix S^t).

726 After rearranging the inequality, we get:

$$\begin{aligned}
g(S^{t+1}) &\leq \frac{1}{\alpha(O)T} (\mu g(O) + (\alpha(O)T - 1)g(S^t)) \\
&= \frac{\mu}{\alpha(O)T} g(O) + \left(1 - \frac{1}{\alpha(O)T}\right) g(S^t)
\end{aligned}$$

727 When recursively applying this inequality, we get:

$$\begin{aligned}
g(S^T) &\leq \frac{\mu}{\alpha(O)T} g(O) \sum_{i=0}^{T-1} \left(1 - \frac{1}{\alpha(O)T}\right)^i + \left(1 - \frac{1}{\alpha(O)T}\right)^T g(\emptyset) \\
&= \mu \left(1 - \left(1 - \frac{1}{\alpha(O)T}\right)^T\right) g(O) + \left(1 - \frac{1}{\alpha(O)T}\right)^T g(\emptyset)
\end{aligned}$$

728

□

729 A.3 Proof of Proposition 4.2

730 **Proposition 4.2.** Suppose that for all $j \in [K]$, the error propagation function $I - C_j \circ \mathcal{L}_h^a$ is
731 ρ_j -Lipschitz continuous with $\rho_j \leq 1$, that $(I - C_j \circ \mathcal{L}_h^a)(0_N) = 0_N$, and that $\sum_{i=1}^T (1 - \rho_{O_i}^2) > 0$,
732 i.e., at least one action of the optimal sequence O is a strict contraction. Then, the function h is
733 weakly supermodular with respect to the optimal solution O , with

$$\alpha(O) = \max \left\{ \frac{1}{T - \sum_{i=1}^T \rho_{O_i}^2}, 1 \right\}$$

734 Furthermore, if the error propagation maps are linear and simultaneously diagonalizable by a
735 unitary matrix with eigenvalues of modulus at most one (the setting of Theorem 4.3), h is also postfix
736 non-increasing; invertibility alone does not suffice (Section A.3).

737 *Proof.* For brevity, we use the notation $(g_1 \circ \dots \circ g_T)(x) = \circ_{t=1}^T g_t(x)$, where composition is
738 applied from right to left so that g_T acts first. In this proof, we use a few properties of Lipschitz
739 continuous functions:

740 • **Property 1:** If g is ρ -Lipschitz continuous and $g(0) = 0$, then $\|g(x)\|_2 = \|g(x) - 0\|_2 =$
741 $\|g(x) - g(0)\|_2 \leq \rho \|x - 0\|_2 = \rho \|x\|_2$

- 742 • **Property 2:** If g_1 and g_2 are Lipschitz continuous functions with Lipschitz constants of ρ_1
 743 and ρ_2 respectively, the Lipschitz constant of $g_1 + g_2$ and $g_1 - g_2$ is $\rho_1 + \rho_2$.
- 744 • **Property 3:** If g_1 and g_2 are Lipschitz continuous functions with Lipschitz constants of ρ_1
 745 and ρ_2 respectively, the Lipschitz constant of $g_1 \circ g_2$ is $\rho_1 \rho_2$.

746 In order to prove weakly- α -supermodularity, we must first prove prefix monotonicity.

747 **Prefix monotonicity:** Let $S \preceq S'$ where $S' = S \oplus N$.

$$\begin{aligned} h(S') &= \left\| \circ_{t=|S'|}^1 (I - C_{S'_t} \circ \mathcal{L}_h^a) \left(e_h^{(0)} \right) \right\|_2^2 \\ &= \left\| \circ_{t=|S'|}^{|S|+1} (I - C_{S'_t} \circ \mathcal{L}_h^a) \circ \circ_{t=|S|}^1 (I - C_{S'_t} \circ \mathcal{L}_h^a) \left(e_h^{(0)} \right) \right\|_2^2 \\ &\stackrel{(a)}{\leq} \left(\prod_{t=|S|+1}^{|S'|} \rho_{S'_t}^2 \right) \left\| \circ_{t=|S|}^1 (I - C_{S_t} \circ \mathcal{L}_h^a) \left(e_h^{(0)} \right) \right\|_2^2 \\ &\leq h(S) \end{aligned}$$

748 (a) by Property 1

749 **Weak supermodularity:** We will upper bound $\alpha(O)$ by providing an upper bound for $h(S) - h(S \oplus$
 750 $O)$ and a lower bound for $\sum_{i=1}^T h(S) - h(S \oplus O_i)$. Since $h \geq 0$, $h(S) - h(S \oplus O) \leq h(S) =$
 751 $\left\| \circ_{t=|S|}^1 (I - C_{S_t} \circ \mathcal{L}_h^a) \left(e_h^{(0)} \right) \right\|_2^2$. (An earlier version of this step used $\|x\|^2 - \|Gx\|^2 \leq \|x - Gx\|^2$,
 752 which is false, e.g. for $G = \frac{1}{2}I$; the direct bound is sharper by a factor 4.)

753 To lower bound $\sum_{i=1}^{|O|} h(S) - h(S \oplus O_i)$

$$\begin{aligned} \sum_{i=1}^{|O|} h(S) - h(S \oplus O_i) &= \sum_{i=1}^{|O|} \left\| \circ_{t=|S|}^1 (I - C_{S_t} \circ \mathcal{L}_h^a) \left(e_h^{(0)} \right) \right\|_2^2 - \left\| (I - C_{O_i} \circ \mathcal{L}_h^a) \circ \circ_{t=|S|}^1 (I - C_{S_t} \circ \mathcal{L}_h^a) \left(e_h^{(0)} \right) \right\|_2^2 \\ &\geq \sum_{i=1}^{|O|} \left\| \circ_{t=|S|}^1 (I - C_{S_t} \circ \mathcal{L}_h^a) \left(e_h^{(0)} \right) \right\|_2^2 - \rho_{O_i}^2 \left\| \circ_{t=|S|}^1 (I - C_{S_t} \circ \mathcal{L}_h^a) \left(e_h^{(0)} \right) \right\|_2^2 \\ &= \left\| \circ_{t=|S|}^1 (I_N - C_{S_t} \circ \mathcal{L}_h^a) \left(e_h^{(0)} \right) \right\|_2^2 \left(T - \sum_{i=1}^T \rho_{O_i}^2 \right) \end{aligned}$$

754 If $h(S) = 0$ then, by prefix monotonicity, every term above vanishes and the weak-supermodularity
 755 inequality holds for any $\alpha \geq 1$; assume $h(S) > 0$. The lower bound just obtained is positive because
 756 $T - \sum_i \rho_{O_i}^2 > 0$ by hypothesis (this is the only place where a strict contraction is needed; $\rho_j \leq 1$
 757 suffices everywhere else). Finally,

$$\begin{aligned} \frac{h(S) - h(S \oplus O)}{\sum_{i=1}^T (h(S) - h(S \oplus O_i))} &\leq \max \left\{ \frac{\left\| \circ_{t=|S|}^1 (I - C_{S_t} \circ \mathcal{L}_h^a) \left(e_h^{(0)} \right) \right\|_2^2}{\left\| \circ_{t=|S|}^1 (I_N - C_{S_t} \circ \mathcal{L}_h^a) \left(e_h^{(0)} \right) \right\|_2^2 \left(T - \sum_{i=1}^T \rho_{O_i}^2 \right)}, 1 \right\} \\ &= \max \left\{ \frac{1}{T - \sum_{i=1}^T \rho_{O_i}^2}, 1 \right\} \end{aligned}$$

758 **Postfix monotonicity.** Let $S' = S \oplus N$ and suppose that $I - C_j \mathcal{L}_h^a = P \Lambda_j P^*$ for all j with a
 759 common unitary P and $|\lambda_{ji}| \leq 1$ (the products below are of $|\lambda_{ji}|^2$ and $|z_i|^2$; we write them without
 760 moduli for the real case). By Theorem A.2,

$$h(S') = \sum_{i=1}^N z_i^2 \prod_{j=1}^K \lambda_{ji}^{2m_j(N)} \prod_{j=1}^K \lambda_{ji}^{2m_j(S)} \leq \sum_{i=1}^N z_i^2 \prod_{j=1}^K \lambda_{ji}^{2m_j(N)} = h(N),$$

761 since every factor $\lambda_{ji}^{2m_j(S)}$ lies in $[0, 1]$. The shared eigenbasis cannot be replaced by invertibility of
 762 the individual maps: for

$$A = \begin{bmatrix} 0.0674240 & -0.1864331 \\ -0.1473150 & -0.8707252 \end{bmatrix}, \quad B = \begin{bmatrix} 0.7575738 & 0.4816284 \\ -0.1369844 & 0.3257283 \end{bmatrix}, \quad x = \begin{bmatrix} 0.2812107 \\ -0.5538228 \end{bmatrix},$$

763 both matrices are invertible with spectral norm 0.9, yet $\|BAx\|_2 = 0.330 > 0.225 = \|Bx\|_2$, i.e.,
 764 $h((A) \oplus (B)) > h((B))$. (An earlier version of this proof bounded $\|G^{-1}\|$ by the reciprocal of the
 765 Lipschitz constant of G , which is not valid.)

766

□

767 A.4 Proof of Theorem A.2

768 **Lemma A.2.** Let $(I - C_j \mathcal{L}_h^a) = P \Lambda_j P^{-1}$ where P is a unitary matrix and $\Lambda_j =$
 769 $\text{diag}(\lambda_{j1}, \dots, \lambda_{jN})$. If $P^{-1} e_h^{(0)} = z$, then the following equality holds (written for a real or-
 770 thogonal P and real eigenvalues; in general z_i^2 and $\lambda_{ji}^{2m_j}$ are to be read as $|z_i|^2$ and $|\lambda_{ji}|^{2m_j}$ here and
 771 in the proofs of Propositions 4.2 and 4.3, which use only $\|Px\| = \|x\|$ and $|\lambda_{ji}| \leq 1$):

$$h(S) = \sum_{i=1}^N z_i^2 \prod_{j=1}^K \lambda_{ji}^{2m_j(S)} \quad (11)$$

772 where $m_j(S) = \sum_{t=1}^{|S|} \mathbf{1}_{S_t=j}$

Proof.

$$\begin{aligned} h(S) &= \left\| \prod_{t=|S|}^1 (I_N - C_{S_t} \mathcal{L}_h^a) e_h^{(0)} \right\|^2 \\ &= \left\| \prod_{t=|S|}^1 (P \Lambda_{S_t} P^{-1}) e_h^{(0)} \right\|^2 \\ &= \left\| P \prod_{t=|S|}^1 \Lambda_{S_t} P^{-1} e_h^{(0)} \right\|^2 \\ &= \left\| \prod_{t=|S|}^1 \Lambda_{S_t} P^{-1} e_h^{(0)} \right\|^2 \\ &= \sum_{i=1}^N z_i^2 \prod_{t=|S|}^1 \lambda_{S_t i}^2 \\ &= \sum_{i=1}^N z_i^2 \prod_{j=1}^K \lambda_{ji}^{2m_j(S)} \end{aligned}$$

773

□

774 A.5 Proof of Theorem 4.3

775 **Proposition 4.3.** Let $\|I_N - C_j \mathcal{L}_h^a\| \leq 1$ for all $j \in [K]$ and $(I - C_j \mathcal{L}_h^a) = P \Lambda_j P^{-1}$ with a common
 776 unitary matrix P . Then, h is supermodular.

777 *Proof.* Let $S \preceq S'$ where $S' = S \oplus B$. By Theorem A.2,

$$h(S) = \sum_{i=1}^N z_i^2 \prod_{j=1}^K \lambda_{ji}^{2m_j(S)}$$

778 where $m_j(S) = \sum_{t=1}^{|S|} \mathbf{1}_{S_t=j}$ be the number of times a sequence S calls the solver j . Recall that h is
 779 considered sequence supermodular if $\forall S', S \in \Omega^*$ such that $S \preceq S'$, it holds that

$$h(S) - h(S \oplus \omega) \geq h(S') - h(S' \oplus \omega)$$

780

$$\begin{aligned} h(S) - h(S \oplus \omega) &= \sum_{i=1}^N z_i^2 \prod_{j=1}^K \lambda_{ji}^{2m_j(S)} - \sum_{i=1}^N z_i^2 \lambda_{\omega i}^2 \prod_{k=1}^K \lambda_{ki}^{2m_j(S)} \\ &= \sum_{i=1}^N (1 - \lambda_{\omega i}^2) z_i^2 \prod_{j=1}^K \lambda_{ji}^{2m_j(S)} \end{aligned}$$

781 Similarly,

$$\begin{aligned} h(S') - h(S' \oplus \omega) &= \sum_{i=1}^N (1 - \lambda_{\omega i}^2) z_i^2 \prod_{j=1}^K \lambda_{ji}^{2m_j(S')} \\ &\stackrel{(a)}{=} \sum_{i=1}^N (1 - \lambda_{\omega i}^2) z_i^2 \prod_{j=1}^K \lambda_{ji}^{2(m_j(S) + m_j(B))} \\ &\stackrel{(b)}{\leq} \sum_{i=1}^N (1 - \lambda_{\omega i}^2) z_i^2 \prod_{j=1}^K \lambda_{ji}^{2m_j(S)} \\ &= h(S) - h(S \oplus \omega) \end{aligned}$$

782 (a) Since $m_j(S') = \sum_{t=1}^{|S'|} \mathbf{1}_{S'_t=j} = \sum_{t=1}^{|S|} \mathbf{1}_{S'_t=j} + \sum_{t=|S|+1}^{|S'|} \mathbf{1}_{S'_t=j} = \sum_{t=1}^{|S|} \mathbf{1}_{S_t=k} + \sum_{t=1}^{|B|} \mathbf{1}_{B_t=j} =$
 783 $m_j(S) + m_j(B)$, (b) since $|\lambda_{ji}| \leq 1$ (as $\|I_N - C_j \mathcal{L}_h^a\| \leq 1$ and P is unitary) and $m_j(B) \geq 0$ for
 784 all $j \in [K]$

785

□

786 B Proofs for Section 5

787 B.1 Proof of Theorem B.1

788 **Lemma B.1.** For any set of preconditioning functions $\mathcal{C}$, any discrete operator $\mathcal{L}_h^a$, any router r , any

789 $a_h, f_h, u_h^{(t)}, u_h \in \mathcal{A} \times \mathcal{F} \times \mathcal{U} \times \mathcal{U}$, the following equality holds true:

$$\begin{aligned} l_{route} \left(r, a_h, f_h, u_h^{(t)}, u_h \right) &= \sum_{j=1}^K \sum_{k=1}^K \left\| (I - C_k \circ \mathcal{L}_h^a) \left(u_h - u_h^{(t)} \right) \right\|_2^2 \mathbf{1}_{k \neq j} \mathbf{1}_{r(a_h, f_h, u_h^{(t)}) \neq j} \\ &\quad - (K - 2) \sum_{j=1}^K \left\| (I - C_j \circ \mathcal{L}_h^a) \left(u_h - u_h^{(t)} \right) \right\|_2^2 \end{aligned}$$

790 *Proof.* Note that $\sum_{j=1}^K \mathbf{1}_{r(a_h, f_h, u_h^{(t)}) \neq j} = K - 1$

$$\begin{aligned}
l_{\text{route}} \left(r, a_h, f_h, u_h^{(t)}, u_h \right) &= \sum_{j=1}^K \left\| (I - C_j \circ \mathcal{L}_h^a) \left(u_h - u_h^{(t)} \right) \right\|_2^2 \mathbf{1}_{r(a_h, f_h, u_h^{(t)}) = j} \\
&= \sum_{j=1}^K \left\| (I - C_j \circ \mathcal{L}_h^a) \left(u_h - u_h^{(t)} \right) \right\|_2^2 - \sum_{j=1}^K \left\| (I - C_j \circ \mathcal{L}_h^a) \left(u_h - u_h^{(t)} \right) \right\|_2^2 \mathbf{1}_{r(a_h, f_h, u_h^{(t)}) \neq j} \\
&= \sum_{j=1}^K \left\| (I - C_j \circ \mathcal{L}_h^a) \left(u_h - u_h^{(t)} \right) \right\|_2^2 - \sum_{j=1}^K \left\| (I - C_j \circ \mathcal{L}_h^a) \left(u_h - u_h^{(t)} \right) \right\|_2^2 \mathbf{1}_{r(a_h, f_h, u_h^{(t)}) \neq j} \\
&\quad + (K - 1) \sum_{k=1}^K \left\| (I - C_k \circ \mathcal{L}_h^a) \left(u_h - u_h^{(t)} \right) \right\|_2^2 - (K - 1) \sum_{k=1}^K \left\| (I - C_k \circ \mathcal{L}_h^a) \left(u_h - u_h^{(t)} \right) \right\|_2^2 \\
&= \sum_{j=1}^K \left\| (I - C_j \circ \mathcal{L}_h^a) \left(u_h - u_h^{(t)} \right) \right\|_2^2 - \sum_{j=1}^K \left\| (I - C_j \circ \mathcal{L}_h^a) \left(u_h - u_h^{(t)} \right) \right\|_2^2 \mathbf{1}_{r(a_h, f_h, u_h^{(t)}) \neq j} \\
&\quad + \sum_{j=1}^K \mathbf{1}_{r(a_h, f_h, u_h^{(t)}) \neq j} \sum_{k=1}^K \left\| (I - C_k \circ \mathcal{L}_h^a) \left(u_h - u_h^{(t)} \right) \right\|_2^2 \\
&\quad - (K - 1) \sum_{k=1}^K \left\| (I - C_k \circ \mathcal{L}_h^a) \left(u_h - u_h^{(t)} \right) \right\|_2^2
\end{aligned}$$

791

$$\begin{aligned}
&= \sum_{j=1}^K \left(\sum_{k=1}^K \left\| (I - C_k \circ \mathcal{L}_h^a) \left(u_h - u_h^{(t)} \right) \right\|_2^2 - \left\| (I - C_j \circ \mathcal{L}_h^a) \left(u_h - u_h^{(t)} \right) \right\|_2^2 \right) \mathbf{1}_{r(a_h, f_h, u_h^{(t)}) \neq j} \\
&\quad - (K - 2) \sum_{k=1}^K \left\| (I - C_k \circ \mathcal{L}_h^a) \left(u_h - u_h^{(t)} \right) \right\|_2^2 \\
&= \sum_{j=1}^K \sum_{k=1}^K \left\| (I - C_k \circ \mathcal{L}_h^a) \left(u_h - u_h^{(t)} \right) \right\|_2^2 \mathbf{1}_{k \neq j} \mathbf{1}_{r(a_h, f_h, u_h^{(t)}) \neq j} \\
&\quad - (K - 2) \sum_{k=1}^K \left\| (I - C_k \circ \mathcal{L}_h^a) \left(u_h - u_h^{(t)} \right) \right\|_2^2
\end{aligned}$$

792

□

793 **B.2 Proof of Theorem B.2**

794 **Proposition B.2.** For any router r defined by $r(a, f, u^{(t)}) = \operatorname{argmax}_{j \in [K]} g_j(a, f, u^{(t)})$, any $a_h \in \mathcal{A}$,
795 $f_h \in \mathcal{F}$, and $u_h^{(t)}, u_h \in \mathcal{U}$, the routing loss l_{route} satisfies:

$$\log(2) l_{\text{route}} \left(r, a_h, f_h, u_h^{(t)}, u_h \right) \leq \Psi(\mathbf{g}, a_h, f_h, u_h^{(t)}, u_h)$$

796 *Proof.* By Theorem B.1, we know that

$$\begin{aligned}
\log(2)l_{\text{route}}\left(r, a_h, f_h, u_h^{(t)}, u_h\right) &= \log(2) \sum_{j=1}^K \sum_{k=1}^K \left\| (I - C_k \circ \mathcal{L}_h^a) \left(u_h - u_h^{(t)} \right) \right\|_2^2 \mathbf{1}_{k \neq j} \mathbf{1}_{r(a_h, f_h, u_h^{(t)}) \neq j} \\
&\quad - \log(2) (K-2) \sum_{j=1}^K \left\| (I - C_j \circ \mathcal{L}_h^a) \left(u_h - u_h^{(t)} \right) \right\|_2^2 \\
&\leq \log(2) \sum_{j=1}^K \sum_{k=1}^K \left\| (I - C_k \circ \mathcal{L}_h^a) \left(u_h - u_h^{(t)} \right) \right\|_2^2 \mathbf{1}_{k \neq j} \mathbf{1}_{r(a_h, f_h, u_h^{(t)}) \neq j} \\
&\stackrel{(a)}{\leq} - \sum_{j=1}^K \sum_{k=1}^K \left\| (I - C_k \circ \mathcal{L}_h^a) \left(u_h - u_h^{(t)} \right) \right\|_2^2 \mathbf{1}_{k \neq j} \log \left(\frac{\exp(g_j(a, f, u^{(t)}))}{\sum_{k=1}^K \exp(g_k(a, f, u^{(t)}))} \right) \\
&= \Psi(\mathbf{g}, a_h, f_h, u_h^{(t)}, u_h)
\end{aligned}$$

797 (a) if $r(a_h, f_h, u_h^{(t)}) \neq j$, $\frac{\exp(g_j(a, f, u^{(t)}))}{\sum_{k=1}^K \exp(g_k(a, f, u^{(t)}))} \leq 0.5$ (with equality only at a tie for the largest
798 score), which implies that $-\log \left(\frac{\exp(g_j(a, f, u^{(t)}))}{\sum_{k=1}^K \exp(g_k(a, f, u^{(t)}))} \right) \geq \log(2) \mathbf{1}_{r(a_h, f_h, u_h^{(t)}) \neq j}$

799

□

800 B.3 Proof of Theorem 5.1

801 **Theorem 5.1.** *Let the per-step costs be normalised state by state, $\hat{c}_j = \tilde{c}_j / \sum_{k=1}^K \tilde{c}_k$, at every state at
802 which $\sum_k \tilde{c}_k > 0$, i.e., at which at least one solver does not annihilate the error (states at which every
803 solver annihilates it carry no routing loss and are discarded), and let $\hat{l}_{\text{route}}$ and $\hat{\Psi}$ denote the routing
804 loss and the surrogate with $\tilde{c}_j$ replaced by $\hat{c}_j$; the normalisation does not change the Bayes-optimal
805 decision equation 9 at any state. Then, for any collection of solvers $\{C_j\}_{j=1}^K$ and linear discrete
806 operator $\mathcal{L}_h^a$, $\hat{\Psi}$ is a Bayes consistent surrogate for $\hat{l}_{\text{route}}$.*

807 *Proof.* Throughout the proof $\tilde{c}_j$ stands for the normalised weight $\hat{c}_j$ of the statement, so that $\sum_j \tilde{c}_j =$
808 1 at every state considered; consequently $\tilde{c}_j \leq \bar{E} := 1$ for all j and $\sum_k B_k = (K-1) \sum_j \tilde{c}_j =$
809 $(K-1)E_{\min}$ with $E_{\min} := 1$. With these two constants the argument below, written for general
810 bounds $\bar{E}$ and $E_{\min}$, is unconditional. For a given a_h, f_h , let u_h be $\mathcal{G}_h(a_h, f_h)$ where $\mathcal{G}_h$ denotes
811 the solution operator acting on the grid G_h . Furthermore, let's consider routers of the form

$$r(a, f, u^{(t)}) = \operatorname{argmax}_{j \in [K]} g_j(a, f, u^{(t)})$$

812 For a given $a_h, f_h, u_h^{(t)} \in \mathcal{A} \times \mathcal{F} \times \mathcal{U}$, let the optimal loss under l_{route} be
813 $l_{\text{route}}^*(a_h, f_h, u_h^{(t)}) = \inf_{\tilde{r}} l_{\text{route}}(\tilde{r}, a_h, f_h, u_h^{(t)}, \mathcal{G}_h(a_h, f_h))$. Similarly, let the optimal loss un-
814 der Ψ be $\Psi^*(a_h, f_h, u_h^{(t)}) = \inf_{\tilde{\mathbf{g}}} \Psi(\tilde{\mathbf{g}}, a_h, f_h, u_h^{(t)}, \mathcal{G}_h(a_h, f_h))$. Let $B_j(a_h, f_h, u_h^{(t)}) =$
815 $\sum_{k=1}^K \left\| (I - C_k \circ \mathcal{L}_h^a) \left(\mathcal{G}_h(a_h, f_h) - u_h^{(t)} \right) \right\|_2^2 \mathbf{1}_{k \neq j}$.

$$\begin{aligned}
& l_{\text{route}} \left(r, a_h, f_h, u_h^{(t)}, \mathcal{G}_h(a_h, f_h) \right) - l_{\text{route}}^* \left(a_h, f_h, u_h^{(t)} \right) \\
& \stackrel{(a)}{=} \sum_{j=1}^K \sum_{k=1}^K \left\| (I - C_k \circ \mathcal{L}_h^a) \left(u_h - u_h^{(t)} \right) \right\|_2^2 \mathbf{1}_{k \neq j} \mathbf{1}_{r(a_h, f_h, u_h^{(t)}) \neq j} - (K-2) \sum_{j=1}^K \left\| (I - C_j \circ \mathcal{L}_h^a) \left(u_h - u_h^{(t)} \right) \right\|_2^2 \\
& \quad - \inf_{\tilde{r}} \sum_{j=1}^K \sum_{k=1}^K \left\| (I - C_k \circ \mathcal{L}_h^a) \left(u_h - u_h^{(t)} \right) \right\|_2^2 \mathbf{1}_{k \neq j} \mathbf{1}_{\tilde{r}(a_h, f_h, u_h^{(t)}) \neq j} + (K-2) \sum_{j=1}^K \left\| (I - C_j \circ \mathcal{L}_h^a) \left(u_h - u_h^{(t)} \right) \right\|_2^2 \\
& = \sum_{j=1}^K B_j(a_h, f_h, u_h^{(t)}) \mathbf{1}_{r(a_h, f_h, u_h^{(t)}) \neq j} - \inf_{\tilde{r}} \sum_{j=1}^K B_j(a_h, f_h, u_h^{(t)}) \mathbf{1}_{\tilde{r}(a_h, f_h, u_h^{(t)}) \neq j} \\
& = \sum_{k=1}^K B_k(a_h, f_h, u_h^{(t)}) \left(\sum_{j=1}^K \frac{B_j(a_h, f_h, u_h^{(t)})}{\sum_{k=1}^K B_k(a_h, f_h, u_h^{(t)})} \mathbf{1}_{r(a_h, f_h, u_h^{(t)}) \neq j} - \inf_{\tilde{r}} \sum_{j=1}^K \frac{B_j(a_h, f_h, u_h^{(t)})}{\sum_{k=1}^K B_k(a_h, f_h, u_h^{(t)})} \mathbf{1}_{\tilde{r}(a_h, f_h, u_h^{(t)}) \neq j} \right)
\end{aligned}$$

816 (a) by Theorem B.1

817 Let $\mathcal{X} = \mathcal{A} \times \mathcal{F} \times \mathcal{U}$ and $\mathcal{Y} = [K]$. Let $\mathcal{P}_{\mathcal{X}}$ denote the degenerate distribution supported at the
818 point $(a_h, f_h, u_h^{(t)})$. We define the conditional distribution - $P(Y = j \mid X = (a_h, f_h, u_h^{(t)})) =$
819 $\frac{B_j(a_h, f_h, u_h^{(t)})}{\sum_{k=1}^K B_k(a_h, f_h, u_h^{(t)})}$ for $j \in [K]$. The risk and optimal risk of 0 - 1 loss under this distribution can
820 be written as:

$$\begin{aligned}
\mathcal{R}_{0-1}(r) &= \sum_{j=1}^K \frac{B_j(a_h, f_h, u_h^{(t)})}{\sum_{k=1}^K B_k(a_h, f_h, u_h^{(t)})} \mathbf{1}_{r(a_h, f_h, u_h^{(t)}) \neq j} \\
\mathcal{R}_{0-1}^* &= \inf_r \sum_{j=1}^K \frac{B_j(a_h, f_h, u_h^{(t)})}{\sum_{k=1}^K B_k(a_h, f_h, u_h^{(t)})} \mathbf{1}_{r(a_h, f_h, u_h^{(t)}) \neq j}
\end{aligned}$$

821 If $r(a_h, f_h, u_h^{(t)}) = \arg \max_{j \in [K]} g_j(a_h, f_h, u_h^{(t)})$ for all $x \in \mathcal{X}$, then the risk and optimal risk of
822 cross entropy loss $(l_{ce}(\mathbf{g}, x, y) - \log \left(\frac{\exp(g_y(x))}{\sum_{k=1}^K \exp(g_k(x))} \right))$ under this distribution can be written as:

$$\begin{aligned}
\mathcal{R}_{ce}(\mathbf{g}) &= - \sum_{j=1}^K \frac{B_j(a_h, f_h, u_h^{(t)})}{\sum_{k=1}^K B_k(a_h, f_h, u_h^{(t)})} \log \left(\frac{\exp(g_j(a_h, f_h, u_h^{(t)}))}{\sum_{k=1}^K \exp(g_k(a_h, f_h, u_h^{(t)}))} \right) \\
\mathcal{R}_{ce}^* &= \inf_{\mathbf{g}} - \sum_{j=1}^K \frac{B_j(a_h, f_h, u_h^{(t)})}{\sum_{k=1}^K B_k(a_h, f_h, u_h^{(t)})} \log \left(\frac{\exp(g_j(a_h, f_h, u_h^{(t)}))}{\sum_{k=1}^K \exp(g_k(a_h, f_h, u_h^{(t)}))} \right)
\end{aligned}$$

823 From Theorem 3.1 of [Mao et al., 2023], $\mathcal{R}_{0-1}(r) - \mathcal{R}_{0-1}^* \leq \Gamma^{-1}(\mathcal{R}_{ce}(\mathbf{g}) - \mathcal{R}_{ce}^*)$ if
824 $r(a_h, f_h, u_h^{(t)}) = \arg \max_{j \in [K]} g_j(a_h, f_h, u_h^{(t)})$ where $\Gamma(z) = \frac{1+z}{2} \log(1+z) + \frac{1-z}{2} \log(1-z)$.
825 Then,

$$\begin{aligned}
& l_{\text{route}} \left(r, a_h, f_h, u_h^{(t)}, \mathcal{G}_h(a_h, f_h) \right) - l_{\text{route}}^* \left(a_h, f_h, u_h^{(t)} \right) \\
&= \sum_{k=1}^K B_k(a_h, f_h, u_h^{(t)}) \left(\sum_{j=1}^K \frac{B_j(a_h, f_h, u_h^{(t)})}{\sum_{k=1}^K B_k(a_h, f_h, u_h^{(t)})} \mathbf{1}_{r(a_h, f_h, u_h^{(t)}) \neq j} - \inf_{\tilde{r}} \sum_{j=1}^K \frac{B_j(a_h, f_h, u_h^{(t)})}{\sum_{k=1}^K B_k(a_h, f_h, u_h^{(t)})} \mathbf{1}_{\tilde{r}(a_h, f_h, u_h^{(t)}) \neq j} \right) \\
&\leq \sum_{k=1}^K B_k(a_h, f_h, u_h^{(t)}) \Gamma^{-1} \left(- \sum_{j=1}^K \frac{B_j(a_h, f_h, u_h^{(t)})}{\sum_{k=1}^K B_k(a_h, f_h, u_h^{(t)})} \log \left(\frac{\exp(g_j(a_h, f_h, u_h^{(t)}))}{\sum_{k=1}^K \exp(g_k(a_h, f_h, u_h^{(t)}))} \right) \right. \\
&\quad \left. - \inf_{\mathbf{g}} - \sum_{j=1}^K \frac{B_j(a_h, f_h, u_h^{(t)})}{\sum_{k=1}^K B_k(a_h, f_h, u_h^{(t)})} \log \left(\frac{\exp(g_j(a_h, f_h, u_h^{(t)}))}{\sum_{k=1}^K \exp(g_k(a_h, f_h, u_h^{(t)}))} \right) \right) \\
&\stackrel{(a)}{\leq} \bar{E} K (K-1) \Gamma^{-1} \left(- \sum_{j=1}^K \frac{B_j(a_h, f_h, u_h^{(t)})}{\sum_{k=1}^K B_k(a_h, f_h, u_h^{(t)})} \log \left(\frac{\exp(g_j(a_h, f_h, u_h^{(t)}))}{\sum_{k=1}^K \exp(g_k(a_h, f_h, u_h^{(t)}))} \right) \right. \\
&\quad \left. - \inf_{\mathbf{g}} - \sum_{j=1}^K \frac{B_j(a_h, f_h, u_h^{(t)})}{\sum_{k=1}^K B_k(a_h, f_h, u_h^{(t)})} \log \left(\frac{\exp(g_j(a_h, f_h, u_h^{(t)}))}{\sum_{k=1}^K \exp(g_k(a_h, f_h, u_h^{(t)}))} \right) \right) \\
&\stackrel{(b)}{\leq} \bar{E} K (K-1) \Gamma^{-1} \left(- \sum_{j=1}^K \frac{B_j(a_h, f_h, u_h^{(t)})}{(K-1)E_{\min}} \log \left(\frac{\exp(g_j(a_h, f_h, u_h^{(t)}))}{\sum_{k=1}^K \exp(g_k(a_h, f_h, u_h^{(t)}))} \right) \right. \\
&\quad \left. - \inf_{\mathbf{g}} - \sum_{j=1}^K \frac{B_j(a_h, f_h, u_h^{(t)})}{(K-1)E_{\min}} \log \left(\frac{\exp(g_j(a_h, f_h, u_h^{(t)}))}{\sum_{k=1}^K \exp(g_k(a_h, f_h, u_h^{(t)}))} \right) \right) \\
&= \bar{E} K (K-1) \Gamma^{-1} \left(\frac{\Psi(\mathbf{g}, a_h, f_h, u_h^{(t)}, \mathcal{G}_h(a_h, f_h)) - \Psi^*(a_h, f_h, u_h^{(t)})}{(K-1)E_{\min}} \right)
\end{aligned}$$

826 (a) since $\left\| (I - C_j \circ \mathcal{L}_h^a)(e_h^{(t)}) \right\|_2^2 < \bar{E}$ for all $j \in [K]$, (b) since Γ^{-1} is non-decreasing and

827 $\exists j \in [K]$ such that $\left\| (I - C_j \circ \mathcal{L}_h^a)(e_h^{(t)}) \right\|_2^2 > E_{\min}$

828 Finally,

$$\begin{aligned}
& \lim_{n \rightarrow \infty} \mathcal{R}_{\text{route}}(r_n) - \mathcal{R}_{\text{route}}^* \\
&\stackrel{(a)}{=} \lim_{n \rightarrow \infty} \mathbb{E}_{a_h, f_h \sim \mathcal{P}_{\mathcal{A} \times \mathcal{F}}} \left[l_{\text{route}} \left(r_n, a_h, f_h, u_h^{(t)}, \mathcal{G}_h(a_h, f_h) \right) - l_{\text{route}}^* \left(a_h, f_h, u_h^{(t)} \right) \right] \\
&\leq \lim_{n \rightarrow \infty} \mathbb{E}_{a_h, f_h \sim \mathcal{P}_{\mathcal{A} \times \mathcal{F}}} \left[\bar{E} K (K-1) \Gamma^{-1} \left(\frac{\Psi(\tilde{\mathbf{g}}, a_h, f_h, u_h^{(t)}, \mathcal{G}_h(a_h, f_h)) - \Psi^*(a_h, f_h, u_h^{(t)})}{(K-1)E_{\min}} \right) \right] \\
&\stackrel{(b)}{\leq} \lim_{n \rightarrow \infty} \bar{E} K (K-1) \Gamma^{-1} \left(\frac{\mathbb{E}_{a_h, f_h \sim \mathcal{P}_{\mathcal{A} \times \mathcal{F}}} \left[\Psi(\mathbf{g}_n, a_h, f_h, u_h^{(t)}, \mathcal{G}_h(a_h, f_h)) - \Psi^*(a_h, f_h, u_h^{(t)}) \right]}{(K-1)E_{\min}} \right)
\end{aligned}$$

$$\begin{aligned}
&= \lim_{n \rightarrow \infty} \bar{E} K (K - 1) \Gamma^{-1} \left(\frac{\mathcal{R}_\Psi(\mathbf{g}_n) - \mathcal{R}_\Psi^*}{(K - 1)E_{min}} \right) \\
&\stackrel{(c)}{=} \bar{E} K (K - 1) \Gamma^{-1} \left(\frac{\lim_{n \rightarrow \infty} \mathcal{R}_\Psi(\mathbf{g}_n) - \mathcal{R}_\Psi^*}{(K - 1)E_{min}} \right) \\
&= \bar{E} K (K - 1) \Gamma^{-1}(0) \\
&\stackrel{(d)}{=} 0
\end{aligned}$$

830 (a) $\mathcal{R}_{\text{route}}^* = \mathbb{E}_{a_h, f_h \sim \mathcal{P}_{\mathcal{A} \times \mathcal{F}}} \left[l_{\text{route}}^* \left(a_h, f_h, u_h^{(t)} \right) \right]$ since the infimum is taken over all measurable
831 functions, (b) by Jensen's inequality since Γ^{-1} is concave, (c) by continuity of Γ^{-1} at 0, (d) $\Gamma^{-1}(0) =$
832 0

833

□

834 C Training Details

835 The periodic experiments of Section 7 use the corrector and the lightweight router described in
836 Section E.1; the data distribution below is shared by all experiments, and the DeepONet and LSTM
837 router described afterwards are those of the original submission's pipeline, whose Dirichlet experiment
838 has been withdrawn (Section D); they are kept for the record.

839 Data for both DeepONet and the routers is sampled from a zero-mean Hierarchical Gaussian Random
840 Field on the periodic domain with covariance operator $\alpha(-\Delta + \beta I)^{-\gamma}$, where α, β, γ are sampled
841 as:

$$\begin{aligned}
\alpha &\sim \text{Log-Uniform}(0.01, 100) \\
\beta &\sim \text{Log-Uniform}(0.1, 1000) \\
\gamma &\sim \text{Uniform}(\{0.5, 1.0, 1.5, 2.0, 2.5, 3.0, 4.0\})
\end{aligned}$$

842 Given (α, β, γ) , samples are generated in the Fourier space. For each non-zero frequency mode
843 k , we draw an independent complex coefficient from a Gaussian distribution with mean 0 and
844 variance $\alpha(4\pi^2\|k\|_2^2 + \beta)^{-\gamma}$. We enforce a Hermitian symmetry to obtain a real-valued field, set the
845 zero-frequency (DC) mode to 0 to obtain a zero mean field, and apply the inverse Discrete Fourier
846 Transform to obtain the field in physical space.

847 For each sample, we compute reference solutions with a least squares solver and treat them as ground
848 truth.

849 This data is used to trained our DeepONet models and LSTM routers. All the models were imple-
850 mented using PyTorch and all the models were trained on one Nvidia A40 GPU.

851 Table 2 contains all hyperparameter details for the DeepONet. DeepONet took 30 minutes to train.
852 We then use the model with the best validation loss.

Table 2: Hyperparameter settings for DeepONet

Hyperparameter	Value
Learning rate	1e-3
Branch Dimension	128
Hidden dimension for branch net	256
No. of hidden layers in branch net	4
Hidden dimension for trunk net	256
No. of hidden layers in trunk net	4
Gradient Clipping Norm	1.0
Weight Decay	0.005
Batch size	256
Training samples	10000
Validation samples	2000
Epochs	1000

853 The routers are LSTM models. Along with the inputs specified in Section 5, namely $a_h, f_h, u_h^{(t)}$,
854 we also supply the routers with the iteration index t , current residual $f_h - \mathcal{L}_h^a u_h^{(t)}$, and the routing
855 decisions from the previous iteration. While the iteration index and residual are deterministic
856 functions of $(a_h, f_h, u_h^{(t)})$ and align with the theory presented in Section 5, incorporating previous
857 routing decisions is a mild deviation from the idealized setting. Nevertheless, this modification is
858 motivated by the same considerations that justify the use of scheduled sampling: under deployment,
859 routing decisions are predicted autoregressively and influence future inputs, whereas training under
860 teacher forcing assumes independence from past predictions. Including previous routing decisions
861 bridges the gap induced by exposure bias by providing the model with trajectory-level information.

862 All the routers are trained with scheduled sampling. We use a warm-up of e_w epochs with teacher-
863 forcing probability $p_{tf}(e) = ss_{start}$. After the warm-up, the p_{tf} decays geometrically by a factor of
864 $\gamma_{tf} < 1$ per epoch and is floored by ss_{end} :

$$p_{tf}(e) = \begin{cases} ss_{start} & e \leq e_w \\ \max(ss_{start}\gamma_{tf}^{e-e_w}, ss_{end}) & e > e_w \end{cases}$$

865 At each time step, with probability $p_{tf}(e)$, we feed the teacher-forced greedy iterate; otherwise, we
866 feed the router’s own predicted iterate.

867 Since LSTMs on long rollouts can suffer from exploding/vanishing gradients, we use truncated
868 backpropagation through time (TBPTT) [Mozier, 2013, Robinson and Fallside, 1987, Werbos, 1988]:
869 the forward pass unrolls the entire trajectory, but gradients are propagated only through the most
870 recent $w_{bptt}(e)$ steps at epoch e . Hidden states are passed forward between segments, while earlier
871 segments are treated as stop-gradient.

872 We employ a curriculum learning approach analogous to scheduled sampling. Let T_{max} be the
873 horizon (300 iterations). With a warm-up of e_w epochs,

$$w_{bptt}(e) = \begin{cases} w_{start} & e \leq e_w \\ \min\left(T_{max}, w_{start}\gamma_{bptt}^{\lfloor \frac{e-e_w}{f_{bptt}} \rfloor}\right) & e > e_w \end{cases} \quad (12)$$

874 so the window grows geometrically by a factor of $\gamma_{bptt} > 1$ every f_{bptt} epochs and is capped at the
875 full trajectory length.

876 Table 3 contains all hyperparameter details for the LSTM routers. The routers for the Paired Solver
877 experiment took a maximum of 4 hours and 30 minutes to train while the routers for the Solver
878 Ensemble experiment took a maximum of 6 hours to train. We then use the model with the best
879 validation loss for testing. The architecture and dataset size is larger for the Solver ensemble
880 experiments to encourage the model to learn some of the nuanced differences between the classes.

Table 3: Hyperparameter settings for routers in Paired solver and Solver Ensemble Experiments

Hyperparameter	Paired Solver	Solver Ensemble
Learning rate	1e-3	1e-3
Hidden dimension	256	512
No. of hidden layers	4	3
Gradient Clipping Norm	1.0	1.0
Weight Decay	0.005	0.005
Batch size	32	32
Training samples	256	1024
Validation samples	32	128
Epochs	200	200
ss_{start}	1.0	1.0
γ_{tf}	0.95	0.95
ss_{end}	0.0	0.0
w_{start}	50	50
γ_{bptt}	1.25	1.25
e_w	10	10
f_{bptt}	1	1

D Additional Experimental Results

The original submission’s Dirichlet experiment (31×31 grid, LSTM router, MLP-branch DeepONet), which this appendix previously kept, has been withdrawn. A review of its evaluation code found that HINTS’ predictions had been mean-centred before comparison with the (uncentred) Dirichlet reference solution while the other methods’ predictions had not, which penalises HINTS by an amount we cannot reconstruct without the original per-instance predictions; in addition, its symmetric-GS rows used the mis-scaled SSOR reference noted in Section E.1. The evaluation code has been corrected (`results.py`), but the experiment has not been rerun, so the wall-clock study of Section E, which is restricted to periodic boundary conditions, is the paper’s only experimental evidence. Extending the corrector to Dirichlet boundaries requires a sine-transform trunk in place of the Fourier trunk and is left to future work.

E Cost-Aware Routing and Wall-Clock Evaluation

The experiments of Section 7 compare methods per iteration, which favors methods whose iterations are expensive. This appendix evaluates the end-to-end *wall-clock time to a target accuracy* of the learned router against HINTS, the classical solver alone and the classical methods a practitioner would reach for on these problems (geometric multigrid and preconditioned Krylov methods), for all five classical solvers and three PDEs, on 128^2 , 256^2 and 512^2 grids, using the cost-aware instantiation of the greedy rule introduced at the end of Section 4. Besides the Poisson and convection–diffusion equations of Section 7 we add *variable-coefficient diffusion* (Section E.5), on which no Fourier diagonalisation exists and the corrector must be learned from data, and *anisotropic diffusion*, $-\epsilon \partial_{x_1}^2 u - \partial_{x_2}^2 u = f$ with $\epsilon = 10^{-2}$ (diffusion tensor $\text{diag}(\epsilon, 1)$), on the same periodic domain with the same forcing distribution. It is the standard model problem on which point-relaxation methods fail to smooth: an error component oscillatory in x_1 and smooth in x_2 is damped by only $(1 - \epsilon)/(1 + \epsilon) \approx 0.98$ per Jacobi sweep, and geometric multigrid with point smoothing loses its grid-independent convergence (line relaxation or semi-coarsening would be needed). Section E.1 describes the protocol; Sections E.2 to E.5 collect all results per equation (every grid, every tolerance, every baseline, with significance tests and seed trials); Section E.7 studies ensembles of one, two and three classical solvers; Section E.8 accounts for overheads and training cost; and Section E.10 verifies numerically that the assumptions of Theorem 4.1, Propositions 4.2–4.3 and Theorem 5.1 hold in these settings.

E.1 Protocol

Fast, matched implementations. A wall-clock comparison is meaningful only if the classical baseline is not handicapped by implementation. All classical operations are therefore implemented as compiled, single-threaded C kernels on the periodic 5-point stencil (`stencil.c`, built with `cc -O3`): the residual together with its norm (the stopping test, in one pass), fused (damped) Jacobi sweeps, in-place lexicographic Gauss–Seidel, SOR and symmetric SOR sweeps (forward and backward), a lexicographic line Gauss–Seidel sweep in which every grid line along the strongly coupled axis is solved exactly by a cyclic Thomas algorithm (the line relaxation that anisotropic problems call for), and the full-weighting and bilinear transfers of geometric multigrid, whose V(2,2)-cycles use Gauss–Seidel smoothing—or line Gauss–Seidel smoothing for the anisotropic equation—with rediscretised coarse operators and a direct solve on the 32×32 coarsest grid (error contraction ≈ 0.03 per cycle on the isotropic equations and, with line smoothing, on the anisotropic one; point-smoothed multigrid contracts the anisotropic error by only ≈ 0.8 per cycle). At 128^2 a residual evaluation with its norm costs $14 \mu\text{s}$, a Gauss–Seidel sweep $61 \mu\text{s}$, a V-cycle $326 \mu\text{s}$ and one corrector call $672 \mu\text{s}$, all measured in the exact per-iteration form in which the benchmark charges them (Table 7); an earlier version of this study used numpy/scipy sweeps that were $4\text{--}5\times$ slower and correspondingly overstated the router’s margins (Section E.9). The compiled kernels agree with the numpy/scipy implementations they replace to 2×10^{-14} (relative), the line sweep agrees with a sparse block-triangular solve to 10^{-13} , and their operators and one-step iterates match the dense reference implementations of the original submission to the reference’s own precision (relative error below 10^{-7} , the reference assembling its matrices in single precision; 10^{-12} where the reference is exact); the comparison is an assertion in `validate_fast_pde.py`. In the course of this check we found that the original dense symmetric-GS reference mis-implemented the SSOR scaling (the factor $\omega/(2 - \omega)$ was added to D^{-1}

instead of multiplying it); it has been corrected, and the wall-clock study uses the standard SSOR throughout. Reference solutions are computed by FFT diagonalization of the same discrete operator, i.e., they are exact solutions of the discrete system, so true errors are available for evaluation at no charged cost; the FFT solve is also reported as a baseline in every table (it is a direct method that exists because these test problems have constant coefficients and periodic boundaries, and it is faster than every iterative method, including ours; the variable-coefficient problem of Section E.5 has no such solve). All timings are single-threaded CPU measurements (Apple M4 Pro, numpy/scipy/PyTorch, one thread). The benchmark guards against machine drift, which on this machine is caused by macOS background services (media analysis and Spotlight indexing, which run when the machine is otherwise idle and slow the memory-bound corrector call by up to $3\times$ for minutes at a time); before the timed replays of every test instance it re-times a reference operation (one corrector call) and waits, in 5 s steps and for up to 5 min, while that time is more than 10% slower than at the start of the session; an instance whose reference is still too slow after the wait is timed anyway and re-timed at the end of the cell, after waiting up to 15 min more. Over the instance timings of the study the guard waited 720 times, instances were re-timed, remained above the tolerance (largest accepted deviation 221.9%), and the per-instance ratios and load averages are stored with the results (an earlier confirmatory run had been contaminated by a 3.5 h stall of the machine inside one cell, which the harness could not detect). Two accounting checks are reported in Section E.8: the charged time of a router solve is 99–99% of the outer elapsed time of its replay, and its live time to h^2 is 0.99–1.06 times its work-unit time.

Corrector. The neural corrector is a DeepONet applied to the residual $r = f_h - \mathcal{L}_h^a u_h^{(t)}$ and evaluated on a 64×64 sensor grid, i.e., on the coarser grid obtained by band-limited (FFT) restriction to the modes $|k_1|, |k_2| \leq 31$ at every grid size (for the anisotropic equation the sensor grid is $N \times N/4$, i.e., full resolution along x_1 , so that the band contains every mode the point smoothers cannot damp; the sensor grid is the only PDE-specific hyperparameter of the pipeline; the variable-coefficient equation uses the isotropic 64×64 grid); its output is prolonged back to the fine grid by exact trigonometric interpolation, so the correction is band-limited by construction and never injects error into the high-frequency band that the classical smoother owns (Figure 2). The trunk net is the fixed orthonormal Fourier basis of that band ($p = 3969$ basis functions, as in POD-DeepONet [Lu et al., 2022]), and the branch net is a linear layer on the 4096 sensor values; it is fit by least squares on 32,000–64,000 exact (residual, error) pairs of the discrete operator drawn from a mixture of spectral shapes (forcing-like, residual-like, white and randomly tilted spectra), and reaches a validation relative error of $\approx 2 \times 10^{-6}$ on every mode of the band. The corrector is applied scale-equivariantly, $C_{\text{NO}}(r) = \|Rr\| \mathcal{G}_\theta(Rr/\|Rr\|)$, which preserves $C_{\text{NO}}(0) = 0$ (Proposition 4.2). Its deployed form is therefore a linear, Fourier-structured map—a least-squares approximation of the exact band-limited inverse $P(\mathcal{L}_h)_{\text{band}}^{-1} R$ of the discrete operator. Replacing it by that exact inverse changes nothing (development run): on the GS pairing at 128^2 the cost-aware oracle’s median work-unit time to 10^{-8} was 4.472 ms with either corrector and its decision sequence identical on 13 of 16 instances. The gains reported below are thus those of an accurate band-limited spectral correction applied at the right moments; the learned corrector stands in for a spectral inverse that, on these constant-coefficient periodic problems, could also be computed directly. Its cost is dominated by a 4096×4096 matrix–vector product and is therefore almost independent of N ($672 \mu\text{s}$ at 128^2 , [pending] at 512^2 , Table 31), i.e., about 32 Jacobi or 11 GS sweeps at 128^2 and [pending] Jacobi sweeps at 512^2 (Tables 7, 13 and 19). The same corrector is used by HINTS and by all routing policies.

Cost-equalised macro-actions. Per-iteration costs c_j are measured once per pairing and grid on the benchmark machine, in the exact form in which the timed replay charges them—the update given the current residual, then the residual of the new iterate and its norm, i.e. the stopping test—as the median over seven interleaved blocks of the block median of 60 repetitions (the minimum and the spread across blocks are stored with the costs). The unit of cost u is one corrector call; every operation is applied $m_j = \max\{1, \text{round}(u/c_j)\}$ times per decision, and these macro-actions are compared by their plain error, i.e. the deployed rule is Algorithm 1 on macro-actions; only an operation dearer than the unit ($c_j > u$, hence $m_j = 1$) is compared per unit of cost, $(\|e_j\|/\|e^{(t)}\|)^{u/c_j}$, and the surrogate weights $\tilde{c}_j$ use the same exponent. With the compiled kernels every point smoother is cheaper than the unit on every grid and a multigrid cycle is cheaper than the unit at 128^2 and 256^2 , so the per-unit-cost comparison is used only for multigrid at 512^2 and, in the granularity ablation of Section E.6, for the corrector itself; everywhere else the deployed rule is exactly Algorithm 1 on macro-actions, which is also the rule the exhaustive check of Section E.10 verifies. Tables 7, 13 and 19 list c_j , u and m_j for

every cell; the equalisation is exact up to rounding ($m_j c_j / u$ lies within $0.5/m_j$ of 1). Ensembles reuse the cached pairwise costs of their members, so every member has the same macro-action size in the ensemble as in its pairwise run. The cost-aware oracle runs Algorithm 1 on these macro-actions with the true error; the *greedy oracle* rows run the cost-agnostic Algorithm 1 on single operations, as in the main text. (An earlier version of this study applied the exponent $u/(m_j c_j)$ to every action, including the rounding error of cheap ones, which is not Algorithm 1; the deployed rule was changed so that the theory covers it.)

Lightweight router. The learned router selects macro-actions from $6 + K$ scalar features, all $O(1)$ given the residual norm that any stopping test computes: the log relative residual, its per-iteration change over the last macro-action and an exponential moving average thereof, the log epoch index, a one-hot encoding of the previous macro-action, and the log counts of corrector calls so far and of iterations since the last call. It is a two-hidden-layer MLP (64 units) evaluated in numpy, costing $\approx 5 \mu\text{s}$ per decision in a warm loop and $16\text{--}19 \mu\text{s}$ in situ, where every decision after a macro-action finds cold caches (measured inside the timed replays, Section E.8), independently of N . It is trained with the surrogate of Section 5.1, with the per-state weights $\tilde{c}_k$ normalised by $\sum_k \tilde{c}_k$ (which does not change the Bayes decision), on states visited by oracle rollouts with ε -exploration (probability 0.1), followed by two to three DAgger rounds [Ross et al., 2011] in which the current router’s own rollouts are relabelled by the oracle—the batch analogue of the scheduled sampling used in Section C. Training uses 32–128 instances per round (fewer on the larger grids) drawn from the same GRF distribution as the test set but with a different seed; the router of every pairing is trained in seconds to a few minutes. Ensemble routers, which must additionally learn to choose among classical members, use a larger imitation budget (256 rollouts at 128^2 and 128 at 256^2 , six DAgger rounds, hidden width 128, 300 epochs): with the pairwise recipe the router over $\{\text{Jacobi}, \text{Jacobi} (0.67)\}$ imitated its oracle imperfectly at tight tolerances (convection–diffusion, 128^2 , $\varepsilon = 10^{-8}$: 3.34 against the oracle’s 2.67 work-unit ms), which the larger budget removes (2.70). At 512^2 the routers are trained from 32 oracle rollouts of at most 400 iterations (stopping error 10^{-8}); for the convection–diffusion Jacobi pairing, where the first such router missed h^2 within the cap on 4 of 16 instances, the reported router uses 64 rollouts of at most 800 iterations—the only per-cell change to the recipe.

Baselines and metric. HINTS applies the corrector every τ -th iteration, first at iteration τ (Section 3). Zhang et al. [2024] use the proportion $1/15$ (one DeepONet call per 14 relaxation sweeps) in all of their two- and three-dimensional experiments and $1/25$ only in the one-dimensional Poisson experiment; since all our problems are two-dimensional, the reference schedule is $\tau = 15$, and $\tau = 25$ is reported as well (an earlier version of this study used it as the reference; the corresponding margins are $1.57 \times \text{--} 4.14 \times$ at h^2 and $1.13 \times \text{--} 2.72 \times$ at 10^{-8} on the 128^2 pairwise cells). We further report the *best fixed schedule*, selected *per cell and per tolerance* on the test instances themselves (an optimistic baseline), from the family of HINTS with $\tau \in \{2, 5, 10, 15, 25, 50\}$, phase-shifted HINTS with the first call at iteration 0 and $\tau \in \{5, 10, 15, 25, 50\}$, and the *one-shot schedule*—the corrector once at iteration 0, then the classical solver only—which is the static rule closest to what the router learns at truncation level and therefore the natural ablation for state-dependent routing ($\tau = 1$, the corrector at every iteration, never reduces the error outside the corrector’s band and is omitted); and the best of three *residual-decay rules*, deterministic controls that decide at the router’s granularity from the same residual norm that the stopping test computes: the corrector at iteration 0, then the solver’s macro-action, and the corrector again whenever the per-iteration contraction of the residual over the last macro-action was worse than $\theta \in \{0.9, 0.95, 0.98\}$ (the threshold is selected per cell and tolerance on the test set, like the best fixed schedule). The classical methods without corrector run on the same stencil, with the same accounting, *inside the same replay loop as the policies*—same instance, same warm-up, same random order, same three replays: multigrid V(2,2) alone; conjugate gradients, CG preconditioned by symmetric GS and by a symmetric V-cycle (forward and backward GS, both compiled) for Poisson; BiCGSTAB, multigrid-preconditioned BiCGSTAB and restarted GMRES(20) for convection–diffusion (GMRES is accounted at the granularity of its restart cycles and is not run at 512^2); point-smoothed and line-smoothed multigrid, CG and its two preconditioned variants for anisotropic diffusion; and the sparse direct solve with a cached factorisation together with the multigrid and CG variants for variable-coefficient diffusion. The Krylov methods are SciPy implementations driven by the compiled matrix–vector product and preconditioners: their per-iteration true errors come from an untimed pass with a callback, and the timed pass runs exactly the same number of iterations with a timestamp-only callback (a few hundred nanoseconds per iteration) and verifies that its final iterate reaches the error the untimed

1048 pass recorded (largest ratio of the two final errors over the study: 1.00). An earlier version of this
 1049 study timed the classical baselines in separate sessions; the same computation differed by 4–9%
 1050 between sessions on identical instances (paired, $p < 10^{-4}$), so every cross-session p -value was
 1051 uninformative, which is why all baselines are now timed in-session. All configuration choices of
 1052 the study were made on a development test set (GRF seed 72, with the earlier numpy kernels, the
 1053 earlier forcing sampler and $\tau = 25$ as reference; archived as `results_dev/`); the results reported
 1054 here are a single confirmatory evaluation of the frozen configurations on a fresh test set (seed 73)
 1055 with three timed replays per instance, and Section E.9 places the two side by side. The test forcings
 1056 are drawn from the hierarchical GRF of Section 7 as real white noise filtered in Fourier space, which
 1057 realises exactly the stated covariance with the Hermitian symmetry of a real field (the development
 1058 runs used a sampler that drew independent complex coefficients on the half-spectrum, which halved
 1059 the variance of the modes on the two coordinate axes); the Nyquist modes are excluded, so every
 1060 forcing is a trigonometric polynomial with $|k| < N/2$ in both directions, and the instances are drawn
 1061 sequentially so that the first n instances of any script are the benchmark’s. For each test instance (64
 1062 at 128^2 , 32 at 256^2 , 16 at 512^2) we record, after every iteration, the true relative error $\|e_h^{(t)}\|/\|u_h\|$
 1063 (untimed, benchmark-only) and the cumulative charged wall-clock time of a replay that executes
 1064 only the chosen operations (residual and its norm, update, and—for the learned router—every feature
 1065 and decision computation; oracle decisions are not charged). The reported time to a tolerance is
 1066 the charged time at the first iteration whose true error is below it, identically for every method: it
 1067 is a benchmark of the work needed to reach an error level, not of a deployed stopping rule (which
 1068 would stop on the residual norm that every method computes anyway). We report the median first
 1069 time at which the error drops below a tolerance ε ; paired speedups—the median over instances of
 1070 the per-instance time ratio, which is the estimand behind every speedup, bold mark and test in this
 1071 appendix (medians of raw times are shown for scale and can order differently from the paired median
 1072 on small samples)—with bootstrap 95% intervals (2000 resamples over instances) and two-sided
 1073 Wilcoxon signed-rank tests on the paired log time ratios (the direction is read from the median ratio;
 1074 “slower” marks a median below one); the study performs several hundred such tests, so the speedup
 1075 tables mark as significant only comparisons that survive a Holm correction at level 0.05 over all
 1076 tests of the table. A run that does not reach ε within the iteration cap (60,000 at 128^2 , 40,000 at
 1077 256^2 , 8,000 at 512^2 ; 5,000 units for the classical baselines, 3,000 at 512^2) is marked $\dagger$ with the
 1078 number of censored instances, reported as a lower bound when it is the majority, and enters every
 1079 paired ratio and test at the time it spent up to the cap, on whichever side it occurs (a router failure
 1080 therefore counts against the router; a ratio against a majority-censored baseline is marked $\geq$ as a
 1081 lower bound)—a labelled lower-bound analysis rather than a survival analysis. Timed replays are
 1082 repeated three times per instance in a random order over policies and baselines, each replay preceded
 1083 by an untimed warm-up (one corrector call and one sweep, or the baseline’s own operation) so that
 1084 every method starts from the same cache state—the corrector’s matrix exceeds the CPU caches,
 1085 and without the warm-up a policy’s first call was up to 0.1 ms slower or faster depending on which
 1086 policy had run before it; the median over replays is used. The measured per-operation costs set the
 1087 macro-action sizes and define *work units* (measured per-iteration costs times executed operations,
 1088 plus the measured decision cost per decision), a timer-free measure of the decision sequence used
 1089 wherever runs from different sessions must be compared; all reported live times are measured directly.
 1090 We highlight $\varepsilon = h^2$ as a nominal grid-dependent tolerance and verified on the confirmatory instances
 1091 themselves ([**pending**] instances over all equations and grids) that it lies below their discretisation
 1092 error: comparing the exact discrete solution with the exact discrete solution on a four times finer grid
 1093 for the same forcing (`discretization_error.py`; for variable coefficients the coefficient is one
 1094 grid-independent function), the relative discretisation error has median [**pending**] h^2 over
 1095 the equations and grids and is at least [**pending**] h^2 on every instance, so driving the algebraic error
 1096 below h^2 is very unlikely to improve the computed solution materially (the two error components can
 1097 partly cancel, so “never” would be too strong); we also report 10^{-3} and 10^{-8} , and the iteration-based
 1098 metrics of the main text ($\text{AUC} = \sum_{t=1}^T \|e_h^{(t)}\|/\|u_h\|$ over $T = 300$ iterations, compared by paired
 1099 t -tests on log AUC) on the 128^2 grid. Every router was additionally retrained from scratch with five
 1100 independent seeds (different training instances, exploration and initialisation; the corrector and the
 1101 test set are fixed) and re-run on the test set; because those trials ran in separate sessions, whose live
 1102 timings are not comparable with the main run’s, they are compared in work units.

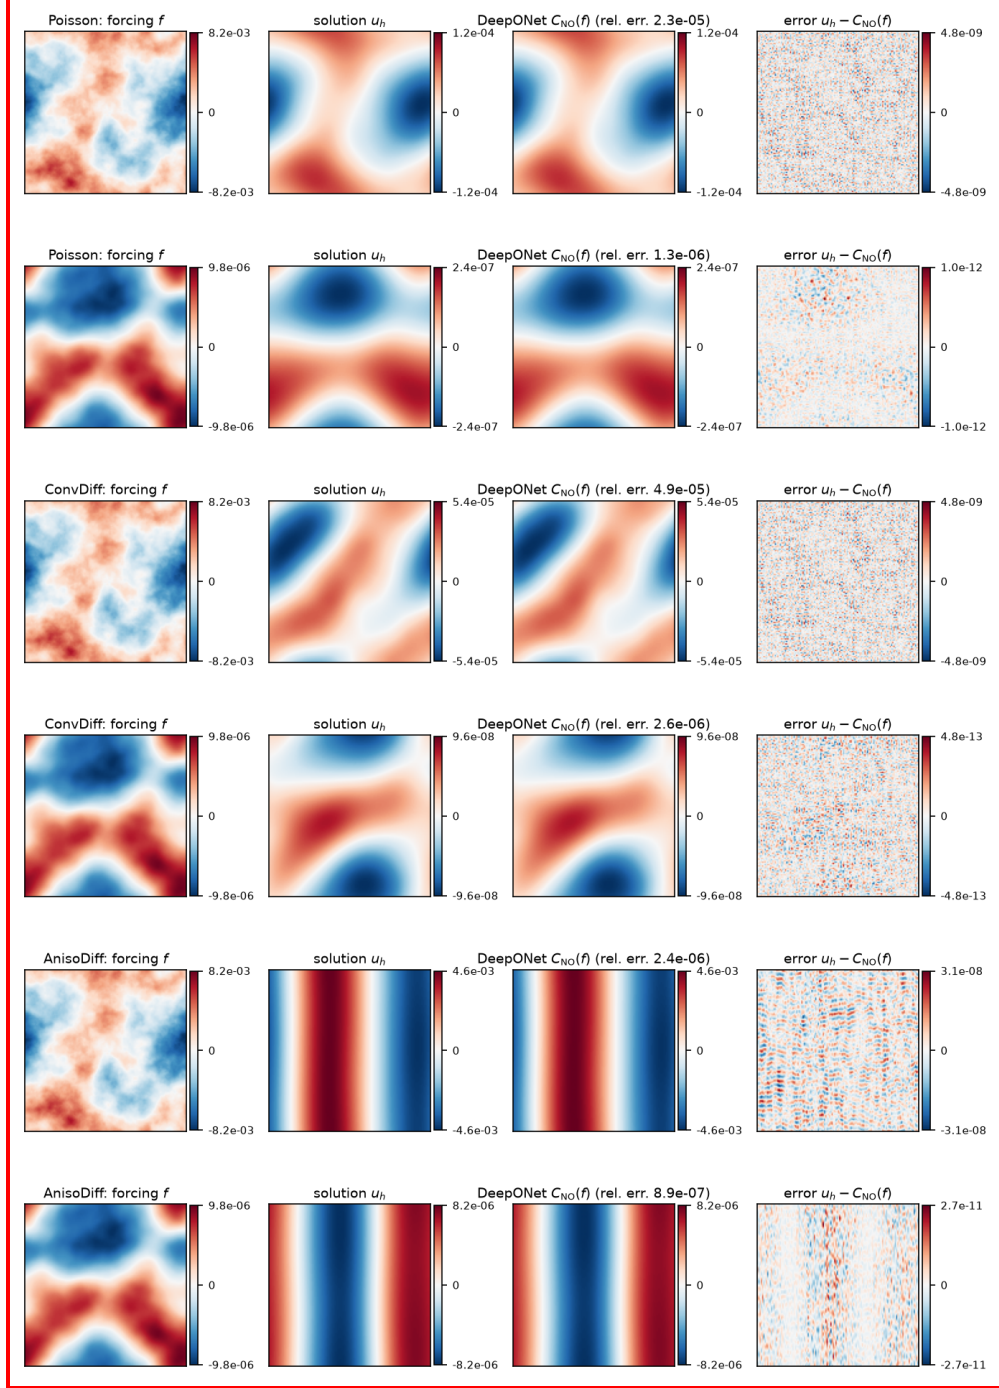


Figure 2: One-shot corrector predictions on the 128×128 grid (two test instances per equation): forcing, exact discrete solution, the DeepONet correction applied to the initial residual ($u_h^{(0)} = 0$), and the remaining error, which is confined to the modes outside the corrector's band (the relative error of the prediction is 10^{-6} – 10^{-5} on the isotropic equations and 10^{-6} on anisotropic diffusion, whose solution is dominated by modes that are smooth in x_2).




Figure 3: Relative error versus charged wall-clock time on the 128×128 grid for one representative test instance of every pairing (columns) and equation (rows); the dashed line is $\varepsilon = h^2$. The learned router and the cost-aware oracle are indistinguishable.

1103 E.2 Poisson

1104 Table 4 collects the median time to 10^{-3} , h^2 and 10^{-8} for every pairing, every policy and the classical
 1105 baselines on all three grids; Table 5 the paired speedups of the learned router with their p -values;
 1106 Figure 4 and table 6 where the iterations go; Table 7 the measured costs and macro-action sizes;
 1107 Table 8 the iteration-based metrics; and Table 9 the seed trials.

1108 **Two regimes.** At $\varepsilon = h^2$ the whole solve consists of one corrector call and a few sweeps (Table 6):
 1109 the router, the cost-aware oracle, phase-shifted HINTS and the one-shot schedule take the same
 1110 decisions and their times coincide up to timer resolution, so the router’s paired ratio against the best
 1111 fixed schedule is close to one (Table 5; bold marks the comparisons that remain significant after the
 1112 Holm correction), while HINTS with its published period, whose first call comes after $\tau - 1$ wasted
 1113 sweeps, is several times slower, and the solver alone hundreds to thousands of times slower. Below
 1114 h^2 the remaining error is the band the corrector cannot see, and the schedules separate: the router
 1115 calls the corrector again only when the smooth error has re-emerged (two or three calls down to
 1116 10^{-8} ; Figure 4), the one-shot schedule never calls it again and pays for hundreds of extra sweeps,
 1117 and every periodic schedule keeps paying for calls that no longer help, so at 10^{-8} the router beats
 1118 HINTS and the one-shot schedule by large factors, whereas against the best fixed schedule and the
 1119 best residual-decay rule—both selected on the test set—the outcome depends on the pairing (Table 5;
 1120 the counts over all pairings are in Section 7, and Section E.6 traces the losses to the macro-action
 1121 granularity). The cost-agnostic greedy oracle (Algorithm 1 on single operations) makes the same
 1122 decisions as its cost-aware counterpart on the path to h^2 (they differ later, in when the corrector is
 1123 called again); the learned router reproduces the cost-aware oracle’s decisions almost exactly (same
 1124 median iteration and call counts, median times within a few percent), and the five retrained routers
 1125 are compared with the main router in Table 9 [pending].

1126 **Classical baselines and larger grids.** Against the classical methods without corrector (bottom
 1127 block of Table 4; timed in the same session as the router of every pairing), the compiled multigrid is
 1128 the strongest competitor: at 128^2 and $\varepsilon = h^2$ the stationary hybrids are $0.97 \times - 1.09 \times$ of a stand-alone
 1129 V-cycle solve in paired ratio, and below h^2 they are slower than it ($0.04 \times - 0.92 \times$), because one
 1130 corrector call removes the modes $|k| \leq 31$ at the price of one matrix–vector product whereas a V(2,2)-
 1131 cycle contracts the whole spectrum by ≈ 0.03 and keeps doing so; routing over {NO, multigrid}
 1132 calls the corrector once and lets the cycles remove the rest ($0.99 \times - 0.99 \times$ of multigrid alone at h^2 and
 1133 $1.14 \times - 1.14 \times$ at 10^{-8}), and against multigrid-preconditioned CG the stationary hybrids are $0.95 \times -$

1134 $1.07 \times$ at h^2 and $0.04 \times -0.79 \times$ at 10^{-8} . The FFT direct solve, exact for this constant-coefficient
1135 periodic problem, is faster than every iterative method at every grid. At 256^2 and 512^2 **[pending]**
1136 the same two regimes are expected to hold with shrinking margins for the Jacobi and Gauss–Seidel
1137 pairings, because the corrector’s fixed band $|k| \leq 31$ leaves the smoother all modes $k \geq 32$, whose
1138 slowest member contracts by only 0.98 per Jacobi sweep at $N = 512$ (against 0.71 at $N = 128$):
1139 every schedule then needs the same ≈ 80 sweeps after the first call, the wall-clock time is dominated
1140 by sweeps, and the corrector—whose cost is nearly grid-independent—is worth fewer and fewer of
1141 them, so the best classical method at 512^2 is multigrid, alone or with one corrector call. This is the
1142 expected limit of a corrector whose band does not grow with the grid.

Table 4: Poisson: median wall-clock time to relative error $\varepsilon \in \{10^{-3}, h^2, 10^{-8}\}$ on the 128^2 , 256^2 and 512^2 grids (64, 32 and 16 test instances) for every pairing {NO, solver} and policy, and for the classical baselines without corrector. “Best fixed schedule” is the best of HINTS ($\tau \in \{2, 5, 10, 15, 25, 50\}$), phase-shifted HINTS ($\tau \in \{5, 10, 15, 25, 50\}$) and the one-shot schedule, and “Best residual-decay rule” the best threshold $\theta \in \{0.9, 0.95, 0.98\}$, both chosen per cell and tolerance on the test set itself (the chosen schedule is the subscript). The greedy oracle is the cost-agnostic Algorithm 1, the cost-aware oracle is Algorithm 1 on cost-equalised macro-actions; both use the true error and are not deployable. Bold marks the router where its median is the smallest among the deployable methods of its pairing (solver alone, every fixed schedule, the one-shot and residual-decay rules); the oracles are in italics; $\dagger n$: n instances did not reach the tolerance within the iteration cap ($>$: lower bound).

Pairing	Method	128 × 128		
		$\varepsilon=10^{-3}$	$\varepsilon=h^2$	$\varepsilon=10^{-8}$
{NO, Jacobi}	Solver only	233.0 ms	329.3 ms	630.1 ms
	HINTS ($\tau=15$)	1.01 ms	1.02 ms	143.9 ms
	Best fixed schedule	703 $\mu s_{\tau=5, t_0=0}$	718 $\mu s_{\tau=15, t_0=0}$	66.5 ms _{one-shot}
	Best residual-decay rule	698 $\mu s_{\theta=0.9}$	715 $\mu s_{\theta=0.9}$	86.8 ms _{$\theta=0.98$}
	One-shot schedule	708 μs	729 μs	66.5 ms
	Greedy oracle (Alg. 1)	698 μs	717 μs	45.4 ms
	Cost-aware oracle	703 μs	717 μs	46.2 ms
	Learned router (ours)	733 μs	758 μs	46.0 ms
{NO, Jacobi (0.67)}	Solver only	350.5 ms	493.4 ms	938.5 ms
	HINTS ($\tau=15$)	981 μs	981 μs	1.94 ms
	Best fixed schedule	690 $\mu s_{\tau=10, t_0=0}$	711 $\mu s_{\tau=5, t_0=0}$	1.74 ms _{$\tau=15, t_0=0$}
	Best residual-decay rule	689 $\mu s_{\theta=0.95}$	711 $\mu s_{\theta=0.98}$	2.69 ms _{$\theta=0.9$}
	One-shot schedule	700 μs	717 μs	76.7 ms
	Greedy oracle (Alg. 1)	693 μs	709 μs	1.72 ms
	Cost-aware oracle	686 μs	709 μs	1.99 ms
	Learned router (ours)	725 μs	743 μs	2.09 ms
{NO, GS}	Solver only	308.0 ms	433.3 ms	823.3 ms
	HINTS ($\tau=15$)	1.52 ms	1.58 ms	4.52 ms
	Best fixed schedule	676 $\mu s_{\tau=10, t_0=0}$	711 $\mu s_{\tau=15, t_0=0}$	2.02 ms _{$\tau=10, t_0=0$}
	Best residual-decay rule	684 $\mu s_{\theta=0.98}$	717 $\mu s_{\theta=0.9}$	2.76 ms _{$\theta=0.9$}
	One-shot schedule	683 μs	718 μs	107.7 ms
	Greedy oracle (Alg. 1)	686 μs	721 μs	2.45 ms
	Cost-aware oracle	687 μs	707 μs	2.06 ms
	Learned router (ours)	721 μs	752 μs	2.14 ms
{NO, SymGS}	Solver only	234.5 ms	329.7 ms	627.3 ms
	HINTS ($\tau=15$)	1.96 ms	1.99 ms	4.01 ms
	Best fixed schedule	676 $\mu s_{\tau=25, t_0=0}$	726 $\mu s_{\text{one-shot}}$	1.85 ms _{$\tau=5, t_0=0$}
	Best residual-decay rule	685 $\mu s_{\theta=0.9}$	720 $\mu s_{\theta=0.95}$	3.29 ms _{$\theta=0.9$}
	One-shot schedule	685 μs	726 μs	78.1 ms
	Greedy oracle (Alg. 1)	695 μs	737 μs	1.86 ms
	Cost-aware oracle	693 μs	735 μs	2.03 ms
	Learned router (ours)	720 μs	772 μs	2.11 ms
{NO, SOR (1.5)}	Solver only	105.4 ms	148.6 ms	282.7 ms
	HINTS ($\tau=15$)	1.53 ms	3.03 ms	4.56 ms
	Best fixed schedule	682 $\mu s_{\tau=25, t_0=0}$	711 $\mu s_{\tau=25, t_0=0}$	2.70 ms _{$\tau=15, t_0=0$}
	Best residual-decay rule	685 $\mu s_{\theta=0.95}$	721 $\mu s_{\theta=0.95}$	4.02 ms _{$\theta=0.9$}
	One-shot schedule	688 μs	723 μs	55.0 ms
	Greedy oracle (Alg. 1)	687 μs	732 μs	3.03 ms
	Cost-aware oracle	696 μs	733 μs	2.69 ms
	Learned router (ours)	716 μs	764 μs	2.85 ms
{NO, Multigrid}	Solver only	680 μs	901 μs	1.90 ms
	HINTS ($\tau=15$)	678 μs	933 μs	1.89 ms
	Best fixed schedule	651 $\mu s_{\tau=15, t_0=0}$	821 $\mu s_{\tau=25, t_0=0}$	1.65 ms _{$\tau=10, t_0=0$}
	Best residual-decay rule	654 $\mu s_{\theta=0.98}$	824 $\mu s_{\theta=0.98}$	1.65 ms _{$\theta=0.95$}
	One-shot schedule	666 μs	822 μs	1.65 ms
	Greedy oracle (Alg. 1)	657 μs	816 μs	1.65 ms
	Cost-aware oracle	672 μs	801 μs	1.67 ms
	Learned router (ours)	684 μs	874 μs	1.70 ms
<i>Classical baselines without corrector (same stencil, same accounting)</i>				
	Multigrid V(2,2) alone (point GS)	692 μs	1.01 ms	2.01 ms
	CG	4.48 ms	5.41 ms	7.59 ms
	PCG (SymGS)	6.33 ms	9.56 ms	16.5 ms
	PCG (multigrid)	706 μs	978 μs	1.70 ms
	FFT direct solve (exact; constant coefficients only)	246 μs	246 μs	246 μs

Table 5: Poisson: paired median speedup [bootstrap 95% interval] of the learned router at $\varepsilon = h^2$ and 10^{-8} over HINTS ($\tau=15$), over the best fixed period for that cell, over multigrid alone and over multigrid-preconditioned CG, with the one-sided Wilcoxon p -value that the router is faster (or, where its median is larger, that it is slower). Bold: $\geq 1.1 \times$ with $p < 0.01$.

Pairing	ε	128 × 128						
		vs. HINTS-15	vs. best schedule	vs. best decay rule	vs. one-shot	vs. multigrid	vs. MG-Krylov	vs. FFT solve
{NO, Jacobi}	h^2	1.37 × [1.34, 1.41] ($< 10^{-10}$)	0.96 × [0.95, 0.99] (slower, 0.095)	0.95 × [0.94, 0.97] (slower, 0.030)	0.95 × [0.95, 0.97] (slower, 2×10^{-10})	0.98 × [0.95, 1.03] (slower, 0.952)	0.99 × [0.96, 1.03] (slower, 0.738)	0.32 × [0.31, 0.33] (slower, $< 10^{-10}$)
	10^{-8}	3.07 × [2.86, 3.34] ($< 10^{-10}$)	1.24 × [1.04, 1.58] (10^{-1})	1.92 × [1.86, 1.96] ($< 10^{-10}$)	1.24 × [1.04, 1.58] (10^{-1})	0.94 × [0.93, 0.95] (slower, $< 10^{-10}$)	0.94 × [0.92, 0.96] (slower, $< 10^{-10}$)	0.01 × [0.00, 0.01] (slower, $< 10^{-10}$)
{NO, Jacobi (0.67)}	h^2	1.32 × [1.24, 1.37] ($< 10^{-10}$)	0.98 × [0.96, 1.03] (slower, 0.158)	0.95 × [0.94, 0.96] (slower, 9×10^{-10})	0.95 × [0.94, 0.96] (slower, 4×10^{-10})	1.09 × [1.03, 1.15] (10^{-1})	1.07 × [1.03, 1.13] (2×10^{-10})	0.34 × [0.31, 0.34] (slower, $< 10^{-10}$)
	10^{-8}	0.93 × [0.92, 0.94] (slower, $< 10^{-10}$)	0.83 × [0.81, 0.86] (slower, 10^{-10})	1.29 × [1.28, 1.30] ($< 10^{-10}$)	36.9 × [33.1, 41.6] ($< 10^{-10}$)	0.92 × [0.84, 0.96] (slower, 10^{-10})	0.79 × [0.76, 0.81] (slower, $< 10^{-10}$)	0.12 × [0.12, 0.12] (slower, $< 10^{-10}$)
{NO, GS}	h^2	2.06 × [1.83, 2.34] ($< 10^{-10}$)	0.95 × [0.94, 0.96] (slower, 9×10^{-10})	0.95 × [0.94, 0.96] (slower, 5×10^{-10})	0.94 × [0.93, 0.95] (slower, 10^{-10})	0.98 × [0.96, 1.03] (slower, 0.805)	0.99 × [0.97, 1.03] (slower, 0.820)	0.32 × [0.27, 0.33] (slower, $< 10^{-10}$)
	10^{-8}	2.08 × [1.64, 2.15] ($< 10^{-10}$)	0.91 × [0.90, 0.92] (slower, 5×10^{-10})	1.28 × [1.26, 1.31] (0×10^{-10})	48.8 × [41.6, 66.8] ($< 10^{-10}$)	0.79 × [0.71, 0.82] (slower, $< 10^{-10}$)	0.66 × [0.59, 0.74] (slower, $< 10^{-10}$)	0.11 × [0.09, 0.11] (slower, $< 10^{-10}$)
{NO, SymGS}	h^2	2.53 × [2.38, 2.70] ($< 10^{-10}$)	0.95 × [0.94, 0.96] (slower, $< 10^{-10}$)	0.95 × [0.94, 0.96] (slower, $< 10^{-10}$)	0.95 × [0.94, 0.96] (slower, $< 10^{-10}$)	1.01 × [0.98, 1.06] (0.113)	1.01 × [0.98, 1.04] (0.183)	0.31 × [0.28, 0.33] (slower, $< 10^{-10}$)
	10^{-8}	1.80 × [1.83, 1.92] ($< 10^{-10}$)	0.86 × [0.84, 0.88] (slower, 2×10^{-10})	1.25 × [1.21, 1.29] ($< 10^{-10}$)	37.8 × [31.2, 54.9] ($< 10^{-10}$)	0.82 × [0.80, 0.84] (slower, $< 10^{-10}$)	0.70 × [0.65, 0.76] (slower, $< 10^{-10}$)	0.11 × [0.11, 0.12] (slower, $< 10^{-10}$)
{NO, SOR (1.5)}	h^2	3.58 × [2.93, 4.32] ($< 10^{-10}$)	0.96 × [0.95, 0.97] (slower, 10^{-10})	0.96 × [0.95, 0.96] (slower, 3×10^{-10})	0.95 × [0.95, 0.97] (slower, 10^{-10})	0.97 × [0.94, 1.00] (slower, 0.025)	0.95 × [0.91, 1.00] (slower, 0.003)	0.32 × [0.25, 0.34] (slower, $< 10^{-10}$)
	10^{-8}	1.63 × [1.25, 1.83] ($< 10^{-10}$)	0.98 × [0.97, 0.99] (slower, 0.016)	1.42 × [1.38, 1.46] ($< 10^{-10}$)	14.4 × [14.6, 22.6] ($< 10^{-10}$)	0.66 × [0.57, 0.70] (slower, $< 10^{-10}$)	0.58 × [0.48, 0.64] (slower, $< 10^{-10}$)	0.09 × [0.07, 0.10] (slower, $< 10^{-10}$)
{NO, Multigrid}	h^2	0.99 × [0.97, 0.99] (slower, 0.003)	0.97 × [0.96, 0.97] (slower, $< 10^{-10}$)	0.96 × [0.96, 0.97] (slower, $< 10^{-10}$)	0.96 × [0.96, 0.97] (slower, 5×10^{-10})	0.99 × [0.98, 1.00] (slower, 0.011)	0.99 × [0.98, 1.00] (slower, 0.022)	0.28 × [0.24, 0.33] (slower, $< 10^{-10}$)
	10^{-8}	1.14 × [1.10, 1.18] (0×10^{-10})	0.97 × [0.97, 0.98] (slower, 5×10^{-10})	0.97 × [0.97, 0.97] (slower, 2×10^{-10})	0.97 × [0.97, 0.97] (slower, 10^{-10})	1.14 × [1.10, 1.18] (1×10^{-10})	0.97 × [0.96, 0.98] (slower, 10^{-10})	0.14 × [0.13, 0.15] (slower, $< 10^{-10}$)




Figure 4: Poisson: fraction of test runs that execute a corrector call at each iteration, for the learned router, the cost-aware oracle and HINTS ($\tau=15$), for every pairing (columns) and grid (rows; all test instances of each grid).

Table 6: Poisson: median iterations (and corrector calls) to $\varepsilon = h^2$ for HINTS ($\tau=15$), the cost-aware oracle and the learned router on every grid.

Pairing	128 × 128		
	HINTS-15	oracle	router
{NO, Jacobi}	15 (1)	2 (1)	2 (1)
{NO, Jacobi (0.67)}	15 (1)	2 (1)	2 (1)
{NO, GS}	16 (1)	2 (1)	2 (1)
{NO, SymGS}	15 (1)	2 (1)	2 (1)
{NO, SOR (1.5)}	30 (2)	2 (1)	2 (1)
{NO, Multigrid}	3 (0)	2 (1)	2 (1)

Table 7: Poisson: measured cost of one classical iteration c_{solver} and of one corrector call c_{NO} (residual evaluation included), and the resulting macro-action size m_{solver} (sweeps per decision; the corrector is always applied once).

Pairing	128×128		
	c_{solver}	c_{NO}	m_{solver}
{NO, Jacobi}	21 μs	672 μs	32
{NO, Jacobi (0.67)}	21 μs	672 μs	32
{NO, GS}	61 μs	672 μs	11
{NO, SymGS}	92 μs	672 μs	7
{NO, SOR (1.5)}	61 μs	672 μs	11
{NO, Multigrid}	326 μs	672 μs	2

Table 8: Poisson, 128×128 : iteration-based metrics of the main text—relative error after $T = 300$ iterations and AUC over the T iterations (mean and standard error over 64 instances; p : one-sided paired t -test that the method’s AUC exceeds the router’s). Bold: router better than both solver-only and HINTS.

Method	AUC	p
{NO, Jacobi}		
Solver only	2.69e+02 (7.2e-01)	$< 10^{-10}$
HINTS ($\tau=15$)	1.39e+01 (2.7e-03)	$< 10^{-10}$
Learned router (ours)	6.08e-03 (1.5e-03)	-
Cost-aware oracle	6.08e-03 (1.5e-03)	-
{NO, Jacobi (0.67)}		
Solver only	2.78e+02 (5.7e-01)	$< 10^{-10}$
HINTS ($\tau=15$)	1.39e+01 (2.4e-03)	$< 10^{-10}$
Learned router (ours)	2.13e-03 (5.0e-04)	-
Cost-aware oracle	2.13e-03 (5.0e-04)	-
{NO, GS}		
Solver only	2.44e+02 (1.0e+00)	$< 10^{-10}$
HINTS ($\tau=15$)	1.38e+01 (6.1e-03)	$< 10^{-10}$
Learned router (ours)	1.71e-03 (4.1e-04)	-
Cost-aware oracle	1.71e-03 (4.1e-04)	-
{NO, SymGS}		
Solver only	2.05e+02 (1.2e+00)	$< 10^{-10}$
HINTS ($\tau=15$)	1.37e+01 (1.1e-02)	$< 10^{-10}$
Learned router (ours)	1.26e-03 (3.0e-04)	-
Cost-aware oracle	1.26e-03 (3.0e-04)	-
{NO, SOR (1.5)}		
Solver only	1.75e+02 (1.2e+00)	$< 10^{-10}$
HINTS ($\tau=15$)	1.35e+01 (1.4e-02)	$< 10^{-10}$
Learned router (ours)	2.79e-03 (6.7e-04)	-
Cost-aware oracle	2.80e-03 (6.7e-04)	-

Table 9: Poisson, 128×128 : five independent router-training seeds per pairing—mean and standard deviation over seeds of the median work-unit time to h^2 , the share of test instances on which the seed router’s decision sequence to h^2 coincides with the main router’s, the range over seeds of the paired speedup over HINTS ($\tau=15$), and the number of seeds significantly faster ($p < 0.01$) than both HINTS ($\tau=15$) and the best fixed period.

[results pending: confirmatory run in progress]




Figure 5: Poisson, 128×128 : routing decisions (dark: corrector call) of the learned router (top row) and the cost-aware oracle (bottom row) over the first 60 iterations of every test instance, for every pairing.

1143 E.3 Convection–diffusion

1144 The convection–diffusion results (Tables 10 to 15 and fig. 6) mirror the Poisson ones: at h^2 the router
 1145 is within a few percent of the best static schedule and $1.5\text{--}4\times$ faster than HINTS with the published
 1146 two-dimensional period, at 10^{-8} it beats HINTS and the one-shot schedule by large factors while
 1147 the comparison with the best fixed schedule and the best residual-decay rule is pairing-dependent
 1148 (Table 11), and it reproduces the cost-aware oracle; the nonsymmetric operator changes only the
 1149 details (the undamped-Jacobi pairing stalls at the near-Nyquist band below 10^{-6} for every policy,
 1150 and GMRES(20) is two orders of magnitude slower than the router).

Table 10: Convection–diffusion: median wall-clock time to relative error $\varepsilon \in \{10^{-3}, h^2, 10^{-8}\}$ on the 128^2 , 256^2 and 512^2 grids for every pairing and policy, and for the classical baselines without corrector (conventions as in Table 4).

[results pending: confirmatory run in progress]

Table 11: Convection–diffusion: paired median speedup [bootstrap 95% interval] of the learned router at $\varepsilon = h^2$ and 10^{-8} over HINTS ($\tau=15$), the best fixed period, multigrid alone and multigrid-preconditioned BiCGSTAB, with two-sided Wilcoxon p -values (conventions as in Table 5).

[results pending: confirmatory run in progress]




Figure 6: Convection–diffusion: corrector usage per iteration of the learned router, the cost-aware oracle and HINTS ($\tau=15$) for every pairing (columns) and grid (rows).

Table 12: Convection–diffusion: median iterations (and corrector calls) to $\varepsilon = h^2$ on every grid.

[results pending: confirmatory run in progress]

Table 13: Convection–diffusion: measured per-iteration costs and macro-action sizes.

[results pending: confirmatory run in progress]

Table 14: Convection–diffusion, 128×128 : iteration-based metrics (conventions as in Table 8).

[results pending: confirmatory run in progress]

Table 15: Convection–diffusion, 128×128 : five-seed router-training trials (conventions as in Table 9).

[results pending: confirmatory run in progress]




Figure 7: Convection–diffusion, 128×128 : routing decisions (dark: corrector call) of the learned router (top row) and the cost-aware oracle (bottom row) over the first 60 iterations of every test instance.

1151 E.4 Anisotropic diffusion

1152 The anisotropic equation is where the two components of the hybrid are most complementary and
 1153 where the point smoothers are weakest: they cannot damp the x_1 -oscillatory modes at all (Table 16:
 1154 none of the five point smoothers alone reaches h^2 within the iteration cap on either grid), whereas the
 1155 corrector, whose sensor grid keeps full resolution along x_1 , removes exactly those modes together
 1156 with the smooth ones in a single call and leaves the smoother only the x_2 -oscillatory band, which
 1157 it damps at the isotropic rate. The corrector’s $N \times N/4$ band is thus the spectral counterpart of

1158 semi-coarsening, and the honest classical comparison is the method a practitioner would use on
 1159 this problem: *line* Gauss–Seidel along x_2 (which contracts the x_1 -oscillatory modes by ≈ 0.35
 1160 per sweep where point Jacobi contracts them by 0.98) and line-smoothed multigrid, which regains
 1161 grid-independent convergence (≈ 0.03 per cycle against ≈ 0.8 with point smoothing). Both are
 1162 included, line GS as the sixth pairing {NO, line GS} and, together with point-smoothed multigrid and
 1163 the CG variants, as classical baselines timed in the same session [pending]. For the point-smoother
 1164 pairings the router reaches h^2 in the time of one corrector call plus a few sweeps, as does the one-shot
 1165 schedule, while HINTS with the published period is $1.7\text{--}3.9\times$ slower (Table 17; first confirmatory
 1166 run, revision 1); the best fixed schedule, which is exactly one call at $t = 0$ followed by sweeps, is
 1167 $3\text{--}7\%$ faster than the router at h^2 , the cost of the router’s per-step decision. At 10^{-8} the router was
 1168 significantly faster than the best fixed schedule in two of the five point-smoother pairings (Jacobi and
 1169 damped Jacobi) and at parity in the other three, within $4\text{--}7\%$ of the cost-aware oracle; these numbers
 1170 are being regenerated under the revised protocol [pending], and the five-seed trials for this equation
 1171 are [pending] (Table 21). The 512^2 grid is not run: the anisotropic corrector keeps full resolution
 1172 along x_1 ($N \times N/4$ sensors), so at 512^2 its least-squares branch would be a dense $65,536 \times 65,536$
 1173 matrix (17 GB in single precision), which the current design cannot hold; a coarser x_1 resolution
 1174 would leave the point smoothers modes they cannot damp (Section E.7 discusses exactly this regime).

Table 16: Anisotropic diffusion: median wall-clock time to relative error $\varepsilon \in \{10^{-3}, h^2, 10^{-8}\}$ on the 128^2 (64 instances, six pairings including line GS) and 256^2 (16 instances) grids for every policy, and for the classical baselines without corrector (conventions as in Table 4; the point-smoother-only runs are capped at 60,000 and 40,000 iterations and never reach h^2).

[results pending: confirmatory run in progress]

Table 17: Anisotropic diffusion: paired median speedup [bootstrap 95% interval] of the learned router at $\varepsilon = h^2$ and 10^{-8} over HINTS ($\tau=15$) and the best fixed period, with two-sided Wilcoxon p -values and bootstrap intervals; the classical baselines (point- and line-smoothed multigrid, CG variants) are timed in the same session.

[results pending: confirmatory run in progress]




Figure 8: Anisotropic diffusion: corrector usage per iteration of the learned router, the cost-aware oracle and HINTS ($\tau=15$) for every pairing (columns) and grid (rows).

Table 18: Anisotropic diffusion: median iterations (and corrector calls) to $\varepsilon = h^2$.

[results pending: confirmatory run in progress]

Table 19: Anisotropic diffusion: measured per-iteration costs and macro-action sizes.

[results pending: confirmatory run in progress]

Table 20: Anisotropic diffusion, 128×128 : iteration-based metrics (conventions as in Table 8).

[results pending: confirmatory run in progress]

Table 21: Anisotropic diffusion, 128×128 : five-seed router-training trials (conventions as in Table 9).

[results pending: confirmatory run in progress]




Figure 9: Anisotropic diffusion, 128×128 : routing decisions (dark: corrector call) of the learned router (top row) and the cost-aware oracle (bottom row) over the first 60 iterations of every test instance.

1175 E.5 Variable-coefficient diffusion

1176 The three equations above have constant coefficients on a periodic domain, so their discrete operators
 1177 are diagonalised by the FFT: the exact band-limited inverse that the corrector approximates has a
 1178 closed form, and a direct FFT solve is faster than any iterative method (Tables 4 and 10). To evaluate
 1179 the hybrid where neither exists we add $-\nabla \cdot (a(\mathbf{x})\nabla u) = f$ on the same periodic domain with a
 1180 fixed, smooth, positive coefficient field $a(\mathbf{x}) = \exp(\kappa z(\mathbf{x}))$, where z is a band-limited random field
 1181 ($|k| \leq 4$, one fixed realisation, scaled by its maximum on a fixed 1024^2 reference grid, so that a is
 1182 one grid-independent function sampled on every grid) and $\kappa = \frac{1}{2} \ln 10$, so that $\max a / \min a = 10$;
 1183 the forcing distribution is unchanged. The discretisation is the conservative 5-point scheme with face-
 1184 averaged coefficients (symmetric positive semidefinite, constants in its null space); reference solutions
 1185 come from a sparse LU factorisation with one pinned unknown, which is also reported as a direct-
 1186 solve baseline (its solve time with the factorisation cached, the factorisation itself costing ≈ 0.5 s at
 1187 128^2). Geometric multigrid rediscrètes the coarse operators with block-averaged coefficients and
 1188 contracts the error by ≈ 0.1 per V(2,2)-cycle. The corrector is the same least-squares DeepONet
 1189 as before—a linear map on the 64×64 band—but fitted to (residual, error) pairs of *this* operator,
 1190 whose band-limited inverse $R(\mathcal{L}_h)^{-1}P$ is not diagonal in the Fourier basis and has no closed form.
 1191 It is nevertheless a fixed 4096×3969 matrix that could be assembled directly from 4096 sparse
 1192 solves, one per sensor basis function; the least-squares fit on random pairs is a cheaper approximation
 1193 of the same matrix, with validation error on the band of order 10^{-2} rather than 10^{-6} . What this
 1194 experiment tests is therefore routing with an inexact corrector on an operator without Fourier structure
 1195 (no shared eigenbasis with the smoothers), not that the correction must be learned: because the
 1196 corrector is inexact it is called repeatedly along a solve, and the router’s decisions matter beyond
 1197 the first call. The coefficient field is fixed rather than sampled, so this experiment tests a corrector
 1198 learned for one operator without an analytic inverse, not generalisation across coefficient fields; the
 1199 latter would require the coefficient-conditioned branch of the original DeepONet and is left to future
 1200 work. Tables 22 to 27 and fig. 10 report the same quantities as for the other equations on the 128^2
 1201 and 256^2 grids.

Table 22: Variable-coefficient diffusion: median wall-clock time to relative error $\varepsilon \in \{10^{-3}, h^2, 10^{-8}\}$ on the 128^2 and 256^2 grids for every pairing and policy, and for the classical baselines without corrector, including the sparse direct solve (conventions as in Table 4).

[results pending: confirmatory run in progress]

Table 23: Variable-coefficient diffusion: paired median speedup [bootstrap 95% interval] of the learned router at $\varepsilon = h^2$ and 10^{-8} over HINTS ($\tau=15$), the best fixed schedule, the one-shot schedule, multigrid alone and multigrid-preconditioned CG, with Holm-corrected one-sided Wilcoxon significance (conventions as in Table 5).

[results pending: confirmatory run in progress]

(figure pending)

Figure 10: Variable-coefficient diffusion: corrector usage per iteration of the learned router, the cost-aware oracle and HINTS ($\tau=15$) for every pairing (columns) and grid (rows).

Table 24: Variable-coefficient diffusion: median iterations (and corrector calls) to $\varepsilon = h^2$.

[results pending: confirmatory run in progress]

Table 25: Variable-coefficient diffusion: measured per-iteration costs and macro-action sizes.

[results pending: confirmatory run in progress]

Table 26: Variable-coefficient diffusion, 128×128 : iteration-based metrics (conventions as in Table 8).

[results pending: confirmatory run in progress]

Table 27: Variable-coefficient diffusion: five-seed router-training trials (conventions as in Table 9).

[results pending: confirmatory run in progress]

1202 E.6 Decision granularity

1203 The cost-aware router decides once per macro-action, and the macro-action of a cheap solver is as
 1204 expensive as one corrector call—32 Jacobi or 11 Gauss–Seidel sweeps at 128^2 with the compiled
 1205 kernels. Below truncation level the best fixed schedule is a phase-shifted HINTS that calls the
 1206 corrector every 5–10 sweeps, a schedule the cost-aware oracle cannot express because its decisions
 1207 are too coarse, which is why the router loses to it in several pairings at 10^{-8} (Tables 5, 11 and 17).
 1208 The granularity is a free parameter of the construction of Section 4: the unit of cost can be any
 1209 fraction of a corrector call, actions dearer than the unit (here the corrector itself) being compared per
 1210 unit of cost. Table 28 repeats the 128^2 study with the unit set to one quarter of a corrector call (about
 1211 a quarter of those sweeps per decision; the corrector is then the one action dearer than the unit and is
 1212 scored per unit of cost), routers retrained with the same recipe, on the same confirmatory instances,
 1213 for every pairing and for the nested two-member Jacobi ensemble.

Table 28: Decision granularity at 128×128 : median time to h^2 and 10^{-8} of the cost-aware oracle and the learned router with the unit of cost equal to one corrector call (main results) and to one quarter of it, and the router’s paired speedup over the best fixed schedule at 10^{-8} in both cases.

[results pending: confirmatory run in progress]

1214 E.7 Solver ensembles

1215 **When can a second classical solver help?** The cost-aware oracle is a per-decision comparison
 1216 of macro-actions, so a router over $\text{NO} \cup \mathcal{W}$ can only beat the best pairwise router of its members
 1217 if, along the way to the tolerance, the error passes through regimes in which *different* members
 1218 give the largest reduction per unit cost. We screened this ceiling exhaustively before training any
 1219 router: for every subset $\mathcal{W}$ of up to three members of {Jacobi, Jacobi (0.67), GS, SymGS, SOR
 1220 (1.5), multigrid} (multigrid on the isotropic equations only), every equation and grid, we ran the
 1221 cost-aware oracle over $\text{NO} \cup \mathcal{W}$ on the test instances and compared its work-unit time to that of
 1222 the oracle over each member alone (Table 29). Two facts emerge. At $\varepsilon = h^2$ no ensemble beats
 1223 the best single solver anywhere: one corrector call followed by one or two macro-actions of the
 1224 cheapest smoother already reaches truncation-level accuracy, so there is no second regime to exploit.
 1225 At $\varepsilon = 10^{-8}$, after the corrector has removed the smooth band, the smoother has to damp the whole
 1226 band $|k| > 31$, whose sub-bands favour different solvers: undamped Jacobi is the cheapest per unit
 1227 cost on the modes near the band edge (it contracts them by $\cos(2\pi \cdot 32/N)$ -like factors that damped
 1228 Jacobi cannot match), but it does not damp the near-Nyquist modes at all, which damped Jacobi and
 1229 GS remove in a few sweeps. In the development screening ([**provisional**]: Table 29 is computed
 1230 with the development-run macro-action sizes; the confirmatory screening with the compiled kernels’
 1231 sizes is pending) the oracle over {Jacobi, Jacobi (0.67)}—a *nonstationary* Jacobi iteration whose
 1232 damping parameter is chosen per macro-action—was therefore faster than every single solver, by
 1233 13–28% on four of the six (equation, grid) settings at 128^2 and 256^2 , and adding a third member (GS)
 1234 changed nothing, i.e., the gain saturates. Whether this ceiling survives the coarser macro-actions of
 1235 the compiled kernels is exactly what the nested-ensemble runs below measure. At 512^2 multigrid
 1236 dominates every subset (its cycle removes all sub-bands at once), so the oracle over any ensemble
 1237 containing it equals the oracle over {NO, multigrid}: adding members neither helps nor hurts. The
 1238 “difficulty” that matters for ensembles is thus not the equation but the width of the spectrum left to the
 1239 classical members after the corrector, which grows with the grid (fixed band) and with the tolerance.

1240 **Nested ensembles.** Table 30 evaluates the learned router over the nested sets {Jacobi},
 1241 {Jacobi (0.67)}, {Jacobi, Jacobi (0.67)} and {Jacobi, Jacobi (0.67), GS} on every equation at 128^2
 1242 and 256^2 , and at 128^2 also over the four ensembles of Section 7 ({Jacobi, GS} up to all five solvers);
 1243 the single-member routers are the pairwise routers of Sections E.2 to E.4, evaluated in the same
 1244 session on the same instances, and all members keep the macro-action sizes of their pairwise runs.
 1245 The table reports live and work-unit times to h^2 , 10^{-6} and 10^{-8} for the router and its oracle, and
 1246 the paired speedup at 10^{-8} over the best single-member router (and of the ensemble oracle over
 1247 the best single-member oracle) with its p -value; work units are the primary measure, since they are
 1248 determined by the decision sequence alone. The 256^2 cells are [**pending**]; at 128^2 three things are
 1249 visible. *No separation at 128^2* : at $\varepsilon = 10^{-8}$ the ensemble router is significantly faster than every
 1250 single-member router in [**pending**] of the [**pending**] ensemble cells, and the ensemble *oracle* is
 1251 [**pending**]-[**pending**] of the best single-member oracle in work units, i.e., the greedy rule itself finds
 1252 no second regime to exploit. This differs from the development runs, in which the two-member Jacobi
 1253 set was 8–15% faster than damped Jacobi alone on convection–diffusion and anisotropic diffusion
 1254 (Table 33), and the difference is the decision granularity. With the numpy sweeps a Jacobi macro-
 1255 action was about ten sweeps, and the oracle alternated undamped Jacobi—the cheapest member on
 1256 the band-edge modes $|k| \approx 32$, which it contracts by about 0.5 per sweep at 128^2 but which leaves
 1257 the Nyquist modes untouched—with damped Jacobi, which removes them. With the compiled sweeps
 1258 a macro-action is about forty sweeps: a single damped-Jacobi macro-action contracts the band-edge
 1259 modes by $0.665^{39} \approx 10^{-7}$ and the Nyquist modes by 0.34^{39} , so the whole band the corrector leaves
 1260 is gone after one or two macro-actions and there is nothing left for a nonstationary schedule to gain.
 1261 Section E.6 repeats these cells with a four times finer unit and tests directly whether the gain returns
 1262 [**pending**]; the 256^2 grids, where the band left to the smoothers is twice as wide and needs several
 1263 macro-actions, are the other place it could. *Saturation*: adding GS to the Jacobi pair changes the
 1264 oracle’s schedule by at most a few percent (the three-member oracles match the two-member ones
 1265 within 1–2% in work units), because GS never wins a cost-aware comparison against the Jacobi pair.
 1266 *Monotonicity*: at h^2 every ensemble router is within [**pending**]-[**pending**] of the best single-member
 1267 router in work units, so enlarging the set is free at truncation level; at 10^{-8} the ensemble routers are
 1268 [**pending**]-[**pending**] of it, and [**pending**] of the [**pending**] cells are more than 5% slower, all of
 1269 them sets containing both GS and SymGS on convection–diffusion or anisotropic diffusion, where
 1270 the router’s imitation of its oracle degrades among near-tied members (the corresponding oracles are

1271 within 6% of the best single oracle). A greedy rule is not monotone in its action set in general, and its
 1272 learned imitation is less so; in our settings the losses vanish at truncation level and reach **[pending]** at
 1273 10^{-8} .

Table 29: Oracle-level screening of ensembles (**[provisional: development-run costs and instances]**): for every equation and grid, the fastest single solver and the fastest ensemble of up to three members under the cost-aware oracle (median work-unit time over 16 test instances), at $\varepsilon = h^2$ and 10^{-8} ; the ratio is the speedup of the ensemble over the best single solver, and the last column is the ratio for {Jacobi, Jacobi (0.67)}.

[results pending: confirmatory run in progress]

Table 30: Nested ensembles: median time to $\varepsilon \in \{h^2, 10^{-6}, 10^{-8}\}$ of the learned router (live and work units) and of the cost-aware oracle (work units) over $\text{NO} \cup \mathcal{W}$, for the single members and the nested two- and three-member sets, all evaluated in one session on the same instances (64 at 128^2 , 32 at 256^2 , 16 for anisotropic diffusion at 256^2). The last columns give the paired speedup at 10^{-8} over the best single-member router (resp. oracle) with the one-sided Wilcoxon p -value; bold: $\geq 1.1 \times$ with $p < 0.01$.

[results pending: confirmatory run in progress]

1274 E.8 Overheads and amortisation

1275 Table 31 lists the measured cost of every operation as a function of N , including one decision of the
 1276 LSTM router of Section 7 (full-state input of size $4N^2$, hidden size 256, 4 layers): at 128^2 one LSTM
 1277 decision costs **[pending]** ms, about **[pending]** Jacobi sweeps, which is why the wall-clock study uses
 1278 the scalar-feature router ($\approx 5 \mu\text{s}$, independent of N). In situ, inside the timed replays, the router’s
 1279 decisions cost 16–19 μs each (the first decision of a solve and every decision after a macro-action
 1280 find cold caches; the $5 \mu\text{s}$ figure is a warm-loop median), i.e. 2.6–3.3% of the router’s time to h^2
 1281 (2–2 decisions) and 0.9–2.1% of its time to 10^{-8} ; the paired difference between the router and the
 1282 cost-aware oracle at h^2 , whose decisions are identical and uncharged, is 22–42 μs per solve and is
 1283 exactly this decision cost (an earlier version of this appendix put the decisions at “well under 1%” on
 1284 the basis of the warm-loop figure). The corrector calls account for 50–70% of the time to h^2 (one or
 1285 two calls). Two further accounting checks: the charged time of a router solve is 99–99% of the outer
 1286 elapsed time of its replay, the remainder being the Python driver’s loop and timer overhead, which
 1287 every method pays per iteration; and the router’s live time to h^2 is 0.99–1.06 times its work-unit
 1288 time, the per-operation costs being measured in the same form as the replay charges them. Table 32
 1289 accounts for training: the corrector (shared with HINTS) is fit in minutes, the router in seconds to a
 1290 few minutes per pairing, and the break-even number of solves after which the router’s own training
 1291 cost is recovered relative to HINTS ($\tau=15$) with the same corrector is in the thousands at 128^2 and
 1292 in the hundreds on the larger grids.

Table 31: Measured single-thread cost of one operation as a function of the grid size (Poisson; median of repeated calls). Multigrid uses a direct solve below 64^2 and is not listed there.

[results pending: confirmatory run in progress]

Table 32: Training-cost amortisation for the GS pairing: corrector training (data generation and least-squares fit, CPU), router training (oracle rollouts and DAGger rounds), median wall-clock saved per solve at $\varepsilon = h^2$ relative to HINTS ($\tau=15$) and to the solver alone, and the number of solves after which the router’s training cost is recovered relative to HINTS.

[results pending: confirmatory run in progress]

E.9 Development versus confirmatory runs

Several choices of this study were made after looking at test results: the corrector band and the anisotropic sensor grid, the router of the convection–diffusion Jacobi pairing at 512^2 (retrained with more rollouts after the first router missed h^2 on 4 of 16 development instances, Section E.1), and the imitation budget of the ensemble routers (raised after the pairwise recipe imitated its oracle imperfectly on one development cell, Section E.7); the optimistic best-schedule baseline is tuned on the test set by design and labelled as such. All of these were made on the development set (seed 72), with the earlier numpy/scipy sweeps, the earlier forcing sampler and $\tau = 25$ as the HINTS reference (Section E.1), which is why the table compares against $\tau = 25$ in both runs. Every result in this appendix comes from a single evaluation of the frozen configurations on the confirmatory set (seed 73) with the compiled kernels and three timed replays. Table 33 lists, for every cell present in both, the development and the confirmatory speedups: the two differ both through the instances and through the kernels—the compiled sweeps are $4\text{--}5\times$ faster, so the corrector is worth $2\text{--}3\times$ more sweeps than in the development runs and the margins over fixed schedules at h^2 shrink accordingly—and the table shows exactly how much of the development-phase margin survives.

Table 33: Development (seed 72, numpy/scipy sweeps, one timed replay) versus confirmatory (seed 73, compiled kernels, three timed replays) paired median speedups of the learned router, for every pairwise cell at $\varepsilon = h^2$ against HINTS ($\tau=25$, the reference of the development runs) and against the best fixed schedule available in that run, and for every nested ensemble at 10^{-8} in work units against the best single-member router and against its oracle; two-sided Wilcoxon p -values in parentheses.

Equation	N	Pairing / ensemble	development (seed 72, numpy kernels)		confirmatory (seed 73, compiled kernels)	
			vs. HINTS-25 or best single	vs. best schedule or oracle	vs. HINTS-25 or best single	vs. best schedule or oracle
Poisson	128^2	{NO, Jacobi}	$2.68\times (< 10^{-10})$	$1.21\times (4\times 10^{-8})$	$1.62\times (< 10^{-10})$	$0.96\times$ (slower, 0.095)
	128^2	{NO, Jacobi (0.67)}	$2.90\times (< 10^{-10})$	$1.28\times (4\times 10^{-9})$	$1.57\times (< 10^{-10})$	$0.98\times$ (slower, 0.158)
	128^2	{NO, GS}	$6.75\times (< 10^{-10})$	$1.94\times (< 10^{-10})$	$2.89\times (< 10^{-10})$	$0.95\times$ (slower, 9×10^{-10})
	128^2	{NO, SymGS}	$9.21\times (< 10^{-10})$	$2.34\times (< 10^{-10})$	$3.73\times (< 10^{-10})$	$0.95\times$ (slower, $< 10^{-10}$)
	128^2	{NO, SOR (1.5)}	$10.3\times (< 10^{-10})$	$3.22\times (< 10^{-10})$	$4.14\times (< 10^{-10})$	$0.96\times$ (slower, 10^{-6})
	128^2	{NO, Multigrid}	$1.96\times (< 10^{-10})$	$1.88\times (< 10^{-10})$	$0.99\times$ (slower, 0.042)	$0.97\times$ (slower, $< 10^{-10}$)

E.10 Verification of the theoretical assumptions

Theorem 4.1 and Propositions 4.2–4.3 assume that every error propagation map $G_j = I - C_j \circ \mathcal{L}_h^a$ is Lipschitz with constant $\rho_j \leq 1$ (with at least one strict contraction in the optimal sequence, respectively a shared unitary eigenbasis), zero-preserving, and—for postfix monotonicity—simultaneously diagonalisable by a unitary matrix; Theorem 5.1 assumes that at every state at least one solver leaves a positive fraction of the error. Table 34 checks these conditions numerically for every operation of the study—finite diagnostics (a power-iteration estimate of the Lipschitz constant, commutators on random vectors, exhaustive search over short horizons), none of which is a certificate over the whole input space—on the mean-free subspace (the constant mode is the null space of the periodic operator and is excluded from the data and from every iteration): the Euclidean Lipschitz constant $\|G_j\|_2$ (power iteration on $G_j^\top G_j$; exact from the Fourier symbol for Jacobi), the Lipschitz constant $\|G_j\|_A = \|A_s^{1/2} G_j A_s^{-1/2}\|_2$ in the energy norm of the symmetric part A_s of the discrete operator, the same for the macro-action $G_j^{m_j}$, the spectral radius, the smallest observed contraction $\sigma_{\min} = \min \|G_j x\|/\|x\|$, $\|C_j(0)\|$, and the commutator $\|G_j G_{\text{NO}} - G_{\text{NO}} G_j\|$ on random vectors. Table 36 evaluates the quantities of Theorem 4.1 and Theorem 5.1 along the cost-aware oracle’s own path to $\varepsilon = 10^{-8}$ (a horizon of tens of macro-actions; to h^2 the path has one or two steps).

What holds, and in which norm. (i) Every classical solver is zero-preserving exactly and contracts in the sense that matters for convergence: the spectral radius of G_j is below 1 in every setting (at most **[pending]**), and so is the energy-norm Lipschitz constant of the solver and of its macro-action (at most **[pending]**). The *Euclidean* Lipschitz constant, however, exceeds 1 for Gauss–Seidel, SOR and symmetric GS (up to **[pending]** for SOR), whose iteration matrices are non-normal; the proof of Proposition 4.2 uses only the triangle inequality and the composition rule for Lipschitz constants, so it holds verbatim with $\|\cdot\|_2$ replaced by the energy norm, in which the condition is satisfied by all classical solvers in all settings—but that statement concerns the greedy rule and objective defined with the energy norm, whereas the oracle and the router of this study minimise the Euclidean error. For the evaluated Euclidean rule, the strict-contraction premise of Proposition 4.2 therefore

1334 holds for the Jacobi-type solvers and fails for the non-normal sweeps, whose Euclidean expansion
 1335 is confined to a single sweep (their spectral radii and macro-action contraction show the iteration
 1336 is convergent). For (damped) Jacobi both constants agree (the maps are Fourier-diagonal); the
 1337 values are close to 1 because the slowest modes (the smoothest ones, and the near-Nyquist mode
 1338 for undamped Jacobi) are barely damped by a point smoother—which is exactly why the corrector
 1339 is needed. (ii) The corrector’s map is zero-preserving, contracts every band mode to **[pending]** or
 1340 better (on average to 10^{-4} – 10^{-6}) and leaves the modes outside the band untouched, so its Lipschitz
 1341 constant is exactly 1 (not < 1): Theorem 4.2 in its strict form does not apply to the corrector, but
 1342 Theorem 4.3 does, and its second condition—a shared unitary eigenbasis—holds exactly for the
 1343 exact band-limited inverse together with (damped) Jacobi, all of which are diagonal in the Fourier
 1344 basis, and holds for the learned corrector up to its fit error (measured commutators $\approx 10^{-6}$, at the
 1345 corrector’s single-precision accuracy; replacing the learned corrector by the exact inverse does not
 1346 change the oracle’s decisions, Section E.1), so for the ensembles $\{\text{NO}, \text{Jacobi}\}$, $\{\text{NO}, \text{Jacobi (0.67)}\}$
 1347 and $\{\text{NO}, \text{Jacobi}, \text{Jacobi (0.67)}\}$ the objective is supermodular and postfix monotone up to that error
 1348 and Theorem 4.1 applies with $\alpha = 1$; for GS, SOR, symmetric GS and multigrid the commutators
 1349 are small but nonzero (10^{-3} – 10^{-1}) and no guarantee is claimed. (iii) Invertibility, which an earlier
 1350 version of the theory used for postfix monotonicity, is neither needed for the corrected result nor
 1351 sufficient in general (Section A.3); for the record, a Gauss–Seidel or symmetric-GS sweep annihilates
 1352 an error concentrated at the first grid point ($\sigma_{\min} = 0$), SOR leaves $(1 - \omega)$ of it ($\sigma_{\min} = 0.72$ – 0.83
 1353 in Table 34), undamped Jacobi annihilates the mode $(N/4, N/4)$, and the corrector annihilates the
 1354 band. Postfix monotonicity ($\mu = 1$) is therefore guaranteed only for the Jacobi-type ensembles,
 1355 and is measured for the others. (iv) Table 35 checks Theorem 4.1 by exhaustive enumeration (in
 1356 floating point, with errors clipped at a relative floor of 10^{-11}) for short horizons on the evaluated
 1357 ensembles, in the Euclidean objective the oracle optimises and for the deployed cost-aware rule:
 1358 for every pairing, for $\{\text{Jacobi}, \text{Jacobi (0.67)}\}$, for the two non-commuting sweeps $\{\text{GS}, \text{SOR}\}$, for
 1359 $\{\text{GS}, \text{multigrid}\}$ and for all five point smoothers, the optimal sequence O of T macro-actions ($T = 4$,
 1360 or 3 for $K = 6$) is found by exhaustive search on 8 test instances (4 for $K = 6$), and μ (over
 1361 the greedy prefixes), $\alpha(O)$ (over the greedy prefixes, and over every sequence of length below T)
 1362 and the bound itself are computed. In every ensemble $\mu \leq \text{[pending]}$ and $\alpha(O) = 1$ along the
 1363 greedy prefixes; the bound holds on all **[pending]** instance–ensemble pairs (on **[pending]** of them the
 1364 exhaustive optimum is at the round-off floor, where the bound holds by construction); the deployed
 1365 rule differs from Algorithm 1 on macro-actions on **[pending]** of them; no prefix violates the weak-
 1366 supermodularity inequality with a nonpositive denominator (**[pending]** such prefixes); and the greedy
 1367 sequence coincides with the exact optimum or is within a few percent of it on most instances (the
 1368 exceptions are myopic misses of up to **[pending]**× on single instances, most often for $\{\text{GS}, \text{SOR}\}$ on
 1369 convection–diffusion); over all short prefixes the ratio exceeds 1 only for the SOR pairing, and only on
 1370 prefixes the greedy rule never visits. The check supports the optimality of the greedy rule over a few
 1371 macro-actions; the horizons at which the oracle runs to 10^{-8} are 10–400 macro-actions (Table 36),
 1372 where the only available constant is the worst-case bound of Proposition 4.2. The worst-case $\alpha(O)$
 1373 of Theorem 4.2, which only uses the Lipschitz constants, is loose in comparison: along the oracle’s
 1374 path to 10^{-8} (Table 36) $\sum_i \rho_{O_i}^2$ is close to T because each ρ is dominated by the slowest mode,
 1375 so the bound ranges from **[pending]** (the multigrid pairing, whose cycle contracts everything) to
 1376 10^2 – 10^4 for the point smoothers, and is infinite on the few instances where $\sum_i \rho_{O_i}^2 \geq T$. (The last
 1377 two columns of Table 36 evaluate the ratio of Equation (7) with the oracle’s own suffix in place of
 1378 the unknown long-horizon O ; this is only a diagnostic, and its maximum is attained trivially at the
 1379 last prefix.) (v) Theorem 5.1 is stated for the per-state normalised weights $\tilde{c}_j / \sum_k \tilde{c}_k$ that the router
 1380 is trained with; its one hypothesis is that at every visited state at least one solver leaves a positive
 1381 fraction of the current error, i.e. $\max_j \tilde{c}_j / \|e^{(t)}\|^2 > 0$ (last column of Table 36; the column before it
 1382 shows that the per-step errors never exceed the current error), so the surrogate’s Bayes consistency
 1383 applies at every visited state.

Table 34: Numerical check of the assumptions of Propositions 4.2–4.3 for every operation of the study (mean-free subspace). **[Provisional: development run; the confirmatory run regenerates this table with the compiled kernels’ macro-action sizes.]** $\|\cdot\|_2$: Euclidean Lipschitz constant of the error propagation map (power-iteration estimate; exact from the symbol for Jacobi); $\|\cdot\|_A$: Lipschitz constant in the energy norm of the symmetric part of $\mathcal{L}_h$, for one iteration and for the macro-action of m_j iterations; ρ : spectral radius; $\sigma_{\min}$: smallest observed contraction (a unit error at the first grid point for the triangular sweeps, the minimum over Fourier modes otherwise; for the corrector, the largest remaining fraction of a band mode); $\|C_j(0)\|$: zero preservation; last column: commutator with the corrector’s map on random unit vectors (Proposition 4.3). For (damped) Jacobi the constants are exact from the Fourier symbol (the value with the Nyquist modes included; without them undamped Jacobi has $\|G\|_2 = \rho = 0.9994$ at 128^2).

[results pending: confirmatory run in progress]

Table 35: Exhaustive short-horizon check of Theorem 4.1 (128×128 , n test instances per row, T macro-actions, K actions including the corrector, deployed cost-aware rule): the optimal sequence O is found by exhaustive search; $\mu = \max_{t < T} g(S^t \oplus O)/g(O)$ over the greedy prefixes; $\alpha(O)$ over the greedy prefixes and over all sequences of length below T ; whether $g(S^T) \leq \mu(1 - \phi_T)g(O) + \phi_T g(\emptyset)$ holds; the ratio $g(S^T)/g(O)$ of the greedy to the optimal error (median and maximum over instances); the number of instances on which the greedy sequence is the optimum; and the numbers of instances on which the exhaustive optimum is at the round-off floor (bound trivially true), on which the deployed rule coincides with Algorithm 1, and of prefixes at which the weak-supermodularity inequality fails with a nonpositive denominator. Errors below 10^{-11} relative are clipped to that floor.

[results pending: confirmatory run in progress]

Table 36: Quantities of Theorem 4.1, Proposition 4.2 and Theorem 5.1 along the cost-aware oracle’s path to $\varepsilon = 10^{-8}$ (8 test instances; medians unless stated): number of macro-actions T , $\sum_i \rho_{O_i}^2$ with the energy-norm constants of the chosen macro-actions, the per-instance margin $T - \sum_i \rho_{O_i}^2$ (median) and the resulting per-instance $\alpha(O)$ of Proposition 4.2 (median and range over the instances on which the margin is positive; ∞ otherwise), the empirical supermodularity ratio $\hat{\alpha}$ of Equation (7) evaluated with S the oracle’s prefix and S' its suffix (maximum and median over prefixes and instances), and the normalised bounds of Theorem 5.1: the largest ratio $\max_j \tilde{c}_j / \|e^{(t)}\|^2$ over visited states ($\bar{E}$) and the smallest value of $\max_j \tilde{c}_j / \|e^{(t)}\|^2$ ($E_{\min}$).

[results pending: confirmatory run in progress]

1384 F Limitations

1385 A primary limitation of our approach is the higher on-time computational cost requires to learn the
 1386 greedy routing strategy. In contrast, the simpler baselines such as HINTS incur no training overhead.
 1387 While this cost is amortized during inference time and leads to improved solution quality, it may be
 1388 prohibitive in settings with limited computational resources or when rapid deployment is required.
 1389 On the deployment side, Section E shows that the gains survive cost accounting: with a corrector
 1390 that covers the smooth half of the spectrum and a classical smoother that covers the rest, one or
 1391 two corrector calls suffice at truncation level, where the router reproduces the best static schedule
 1392 without a search; below truncation level its advantage over the best fixed schedule is not uniform
 1393 across pairings, and a compiled multigrid alone is faster there. The corrector we use is specific to
 1394 the constant-coefficient periodic setting (a fixed Fourier trunk basis on a coarse sensor grid), its cost
 1395 grows with the fourth power of the sensor-grid size, and undamped Jacobi remains limited far below
 1396 truncation level by the near-Nyquist modes that neither component damps. Our wall-clock study is
 1397 restricted to single-thread CPU execution; with a corrector band that does not grow with the grid,

1398 the gains over fixed schedules shrink at 512^2 for the point smoothers, whereas pairing the corrector
1399 with geometric multigrid is expected to restore them (Section E; [512² **pending**]). Ensembles of
1400 classical solvers pay off only where the error passes through regimes that favour different members;
1401 at truncation level a single well-chosen smoother suffices. For the constant-coefficient periodic test
1402 problems the discrete operator is diagonalised by the FFT, and the resulting direct solve, reported
1403 in every table, is faster than every iterative method we measured, including ours ($2.98\times$ – $2.98\times$
1404 against the best router); the hybrid is meant for problems on which no such direct diagonalisation
1405 exists—where the band-limited inverse has to be assembled or fitted rather than read off a symbol,
1406 and the corrector is then inexact—which is why Section E also reports a variable-coefficient diffusion
1407 problem.

1408 Additionally, the effectiveness of the learned router depends on the neural operator exhibiting
1409 heterogeneous performance across samples. When the neural operators produces predictions of
1410 relatively uniform quality, the benefit of adaptive routing diminishes and it is more practical to set a
1411 fixed schedule based on the quality of the neural operator. In practice, achieving such heterogeneity
1412 can be sensitive to training choices, and neural operators themselves can be finicky to train.

1413 G LLM Usage

1414 LLMs were used substantially in preparing this paper, and we disclose this precisely. The text of
1415 the original submission was drafted by the authors, with ChatGPT and Gemini used for feedback
1416 on grammar and clarity. The revision that responds to the reviews was produced with Claude Code
1417 (Anthropic), operated and directed by the authors: the wall-clock benchmark pipeline (the compiled
1418 stencil kernels, the cost-aware macro-action environment, the scalar-feature router and its training
1419 loop, the corrector fitting, the baseline, table and figure generators, and the assumption and theorem
1420 checkers) was written largely by the assistant under the authors’ direction and was reviewed, tested
1421 and run by the authors; the corrections to the theory in this revision (the μ -slack in Theorem 4.1, the
1422 corrected constant and hypotheses of Proposition 4.2, the counterexample to postfix monotonicity,
1423 the unitary-eigenbasis condition, and the restatement of Theorem 5.1 for normalised weights) were
1424 drafted with the assistant’s help in response to reviewer findings and were verified by the authors;
1425 and much of the new text of Section 7 and Section E was drafted by the assistant from the generated
1426 tables and edited by the authors. The theorems and proof strategies of the original submission predate
1427 the use of the assistant. The research questions, the framing of the method, and the decisions about
1428 what to run and what to claim are the authors’. Independent reviews of the manuscript and code were
1429 also obtained from LLM-based reviewers; their findings are addressed in this revision.

1430 NeurIPS Paper Checklist

1431 1. Claims

1432 Question: Do the main claims made in the abstract and introduction accurately reflect the
1433 paper’s contributions and scope?

1434 Answer: [Yes]

1435 Justification: The abstract and introduction state the two-regime empirical outcome (parity
1436 with the best static schedule at truncation level; mixed results against the best fixed schedule
1437 below it; multigrid alone faster below truncation level) with counts and ranges computed
1438 from the tables, and the theoretical claims with their hypotheses (unitary eigenbasis for
1439 postfix monotonicity, measured slack otherwise; per-decision consistency for the full-
1440 information router). At the time of this draft several confirmatory cells are still being
1441 regenerated and are marked as pending in the text.

1442 Guidelines:

- 1443 • The answer [N/A] means that the abstract and introduction do not include the claims
1444 made in the paper.
- 1445 • The abstract and/or introduction should clearly state the claims made, including the
1446 contributions made in the paper and important assumptions and limitations. A [No] or
1447 [N/A] answer to this question will not be perceived well by the reviewers.
- 1448 • The claims made should match theoretical and experimental results, and reflect how
1449 much the results can be expected to generalize to other settings.
- 1450 • It is fine to include aspirational goals as motivation as long as it is clear that these goals
1451 are not attained by the paper.

1452 2. Limitations

1453 Question: Does the paper discuss the limitations of the work performed by the authors?

1454 Answer: [Yes]

1455 Justification: The paper discusses key limitations in Appendix E.

1456 Guidelines:

- 1457 • The answer [N/A] means that the paper has no limitation while the answer [No] means
1458 that the paper has limitations, but those are not discussed in the paper.
- 1459 • The authors are encouraged to create a separate “Limitations” section in their paper.
- 1460 • The paper should point out any strong assumptions and how robust the results are to
1461 violations of these assumptions (e.g., independence assumptions, noiseless settings,
1462 model well-specification, asymptotic approximations only holding locally). The authors
1463 should reflect on how these assumptions might be violated in practice and what the
1464 implications would be.
- 1465 • The authors should reflect on the scope of the claims made, e.g., if the approach was
1466 only tested on a few datasets or with a few runs. In general, empirical results often
1467 depend on implicit assumptions, which should be articulated.
- 1468 • The authors should reflect on the factors that influence the performance of the approach.
1469 For example, a facial recognition algorithm may perform poorly when image resolution
1470 is low or images are taken in low lighting. Or a speech-to-text system might not be
1471 used reliably to provide closed captions for online lectures because it fails to handle
1472 technical jargon.
- 1473 • The authors should discuss the computational efficiency of the proposed algorithms
1474 and how they scale with dataset size.
- 1475 • If applicable, the authors should discuss possible limitations of their approach to
1476 address problems of privacy and fairness.
- 1477 • While the authors might fear that complete honesty about limitations might be used by
1478 reviewers as grounds for rejection, a worse outcome might be that reviewers discover
1479 limitations that aren’t acknowledged in the paper. The authors should use their best
1480 judgment and recognize that individual actions in favor of transparency play an impor-
1481 tant role in developing norms that preserve the integrity of the community. Reviewers
1482 will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [Yes]

Justification: All theoretical results include clearly stated assumptions in the statement of theorems. Complete proofs are provided in Appendices A and B with appropriate cross-referencing from the main text.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: The paper provides sufficient detail to reproduce all experiments, including training procedures and implementation details described in the main text and Appendix C. Additionally, we release code with a comprehensive README to support the reproducibility of our results.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).

1537 (d) We recognize that reproducibility may be tricky in some cases, in which case
1538 authors are welcome to describe the particular way they provide for reproducibility.
1539 In the case of closed-source models, it may be that access to the model is limited in
1540 some way (e.g., to registered users), but it should be possible for other researchers
1541 to have some path to reproducing or verifying the results.

1542 5. Open access to data and code

1543 Question: Does the paper provide open access to the data and code, with sufficient instruc-
1544 tions to faithfully reproduce the main experimental results, as described in supplemental
1545 material?

1546 Answer: [Yes]

1547 Justification: We provide open access to code with detailed instructions in a comprehensive
1548 README, including environment setup and commands to train routers and reproduce
1549 the main experimental results. These files will automatically generate the data needed for
1550 training. Data generation has also been documented in Appendix C.

1551 Guidelines:

- 1552 • The answer [N/A] means that paper does not include experiments requiring code.
- 1553 • Please see the NeurIPS code and data submission guidelines ([https://neurips.cc/
1554 public/guides/CodeSubmissionPolicy](https://neurips.cc/public/guides/CodeSubmissionPolicy)) for more details.
- 1555 • While we encourage the release of code and data, we understand that this might not
1556 be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not
1557 including code, unless this is central to the contribution (e.g., for a new open-source
1558 benchmark).
- 1559 • The instructions should contain the exact command and environment needed to run to
1560 reproduce the results. See the NeurIPS code and data submission guidelines ([https:
1561 //neurips.cc/public/guides/CodeSubmissionPolicy](https://neurips.cc/public/guides/CodeSubmissionPolicy)) for more details.
- 1562 • The authors should provide instructions on data access and preparation, including how
1563 to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- 1564 • The authors should provide scripts to reproduce all experimental results for the new
1565 proposed method and baselines. If only a subset of experiments are reproducible, they
1566 should state which ones are omitted from the script and why.
- 1567 • At submission time, to preserve anonymity, the authors should release anonymized
1568 versions (if applicable).
- 1569 • Providing as much information as possible in supplemental material (appended to the
1570 paper) is recommended, but including URLs to data and code is permitted.

1571 6. Experimental setting/details

1572 Question: Does the paper specify all the training and test details (e.g., data splits, hyperpa-
1573 rameters, how they were chosen, type of optimizer) necessary to understand the results?

1574 Answer: [Yes]

1575 Justification: The paper specifies all key training and evaluation details in the Section 6 and
1576 Appendix C.

1577 Guidelines:

- 1578 • The answer [N/A] means that the paper does not include experiments.
- 1579 • The experimental setting should be presented in the core of the paper to a level of detail
1580 that is necessary to appreciate the results and make sense of them.
- 1581 • The full details can be provided either with the code, in appendix, or as supplemental
1582 material.

1583 7. Experiment statistical significance

1584 Question: Does the paper report error bars suitably and correctly defined or other appropriate
1585 information about the statistical significance of the experiments?

1586 Answer: [Yes]

1587 Justification: The paper reports standard deviations in the main results tables, and further
1588 provides statistical significance via t-test p values in Appendix D.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: The paper provides details on compute resources in Appendix C.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: This paper does not involve any human subjects and all data has been simulated. This research focuses on numerical methods for physical systems and does not raise concerns related to safety, security, discrimination, surveillance, deception & harassment, environment, human rights, and bias & fairness.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: The paper discusses broader impacts in the introduction section. No direct negative societal impacts are anticipated.

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: The paper does not release any models or datasets at all. It focuses on methodological contributions for numerical solvers using simulated data. Therefore, safeguards are not applicable.

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [N/A]

Justification: The paper does not use any external datasets or pretrained models. All experiments are conducted using simulated data and models trained in-house.

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, `paperswithcode.com/datasets` has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [Yes]

Justification: We release code for our method which is documented with a detailed README.

Guidelines:

- The answer [N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [N/A]

Justification: The paper does not involve crowdsourcing or any research with human subjects. All experiments are conducted using simulated data.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects

1744 Question: Does the paper describe potential risks incurred by study participants, whether
1745 such risks were disclosed to the subjects, and whether Institutional Review Board (IRB)
1746 approvals (or an equivalent approval/review based on the requirements of your country or
1747 institution) were obtained?

1748 Answer: [N/A]

1749 Justification: This paper does not involve human subjects. All experiments are conducted
1750 using simulated data. Therefore, IRB approval is not applicable.

1751 Guidelines:

- 1752 • The answer [N/A] means that the paper does not involve crowdsourcing nor research
1753 with human subjects.
- 1754 • Depending on the country in which research is conducted, IRB approval (or equivalent)
1755 may be required for any human subjects research. If you obtained IRB approval, you
1756 should clearly state this in the paper.
- 1757 • We recognize that the procedures for this may vary significantly between institutions
1758 and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the
1759 guidelines for their institution.
- 1760 • For initial submissions, do not include any information that would break anonymity (if
1761 applicable), such as the institution conducting the review.

1762 16. Declaration of LLM usage

1763 Question: Does the paper describe the usage of LLMs if it is an important, original, or
1764 non-standard component of the core methods in this research? Note that if the LLM is used
1765 only for writing, editing, or formatting purposes and does *not* impact the core methodology,
1766 scientific rigor, or originality of the research, declaration is not required.

1767 Answer: [N/A]

1768 Justification: LLMs are primarily used for writing and editing. Further details are provided
1769 in Appendix F.

1770 Guidelines:

- 1771 • The answer [N/A] means that the core method development in this research does not
1772 involve LLMs as any important, original, or non-standard components.
- 1773 • Please refer to our LLM policy in the NeurIPS handbook for what should or should not
1774 be described.