
A Greedy PDE Router for Blending Neural Operators and Classical Methods

Anonymous Author(s)

Affiliation

Address

email

1 **Revision note.** All text, tables and figures added in this revision are typeset in red (new figures
2 carry a red frame); everything in black is unchanged from the submitted version. The iteration-based
3 experiments of the submitted version on 31^2 and 63^2 grids with the LSTM router, and their appendix
4 analyses, have been removed: they are superseded by the wall-clock study of Sections E and 7,
5 which covers the same solvers and equations (plus anisotropic diffusion) on 128^2 – 512^2 grids and
6 also reports the iteration-based metrics. Only the Dirichlet experiment of the submitted version is
7 kept (Section D).

Abstract

8 When solving PDEs, classical numerical solvers are often computationally ex-
9 pensive, while machine learning methods can suffer from spectral bias, failing to
10 capture high-frequency components. Designing an optimal hybrid iterative solver–
11 where, at each iteration, a solver is selected from an ensemble of solvers to leverage
12 their complementary strengths–poses a challenging combinatorial problem. While
13 greedy selection is desirable for its constant-factor approximation guarantee to the
14 optimal solution under Lipschitz assumptions, it requires knowledge of the true
15 error at each step, which is unavailable in practice. We address this by proposing
16 an approximate greedy router that efficiently mimics a greedy approach to solver
17 selection. Empirical results on the Poisson, convection–diffusion and anisotropic
18 diffusion equations show that our method consistently reduces final error and area-
19 under-the-curve (AUC) of the error trajectory relative to single-solver baselines
20 and existing hybrid approaches such as HINTS, reaching comparable error levels
21 in substantially fewer iterations with more stable error decay. With a cost-aware
22 instantiation of the greedy rule, a lightweight router, and a band-limited scale-
23 equivariant corrector, these gains carry over to end-to-end wall-clock time on 128^2
24 to 512^2 grids: in every one of the 15 solver–equation pairings at 128^2 (five classical
25 solvers; Poisson, convection–diffusion and anisotropic diffusion) the learned router
26 reaches truncation-level accuracy $238\times$ – $1983\times$ faster than the classical solver
27 alone (which never reaches it on the anisotropic equation) and $2.68\times$ – $11.8\times$ faster
28 than HINTS, it beats HINTS on every grid even when the latter’s period is tuned
29 per setting, and it is faster than geometric multigrid and multigrid-preconditioned
30 Krylov methods on the isotropic equations; routing over ensembles of classical
31 solvers is never worse than the best pair and strictly better where the error passes
32 through regimes that favour different solvers.

1 Introduction

34 Natural phenomena and engineered systems are often governed by ordinary and partial differential
35 equations (PDEs). Solving these equations enables a variety of tasks - predicting the evolution of the
36 system through simulation (e.g., forecasting weather using the Navier-Stokes equations) [Kalnay,
37 2003, Bauer et al., 2015, Staniforth and Côté, 1991], addressing control problems (e.g., optimizing

38 heat shield design for spacecrafts using heat transfer equations) [Anderson, 1989, Tröltzsch, 2010],
 39 and tackling inverse problems based on external measurements (e.g., reconstructing brain activity
 40 from EEG data using electrophysiological models) [Baillet et al., 2001, Grech et al., 2008].

41 Traditionally, PDEs are solved using finite difference methods [Smith, 1985, LeVeque, 2007], which
 42 discretize the spatial and/or temporal domain and approximate the solution by solving a system
 43 of equations at the discrete grid points. Similarly, finite element methods [Hughes, 2003, Bathe,
 44 2006, Brenner, 2008] construct a system of equations based on local basis functions, and then
 45 rely on numerical linear algebra techniques—such as the Jacobi and Gauss-Seidel method [Saad,
 46 2003]—to compute the solution. These iterative methods, however, lack generalization across different
 47 initial conditions, boundary conditions, or forcing functions, as even minor changes to any of these
 48 parameters require re-solving the entire system of equations from scratch.

49 These challenges have motivated the development of neural operators [Kovachki et al., 2023]—a class
 50 of machine learning models that aim to learn the solution operator directly, enabling fast inference
 51 across a range of parameters. Neural operators lift the classical linear layer in neural networks to an
 52 operator—typically a kernel integral operator—between function spaces. Despite strong approximation
 53 guarantees [Chen and Chen, 1995], neural operators are prone to spectral bias [Liu and Cai, 2021, Liu
 54 et al., 2024, You et al., 2024, Khodakarami et al., 2025, Xu et al., 2025], similar to traditional neural
 55 networks [Rahaman et al., 2019, Xu et al., 2019, Luo et al., 2019]. As a result, they favor learning
 56 low-frequency components of the solution function and often struggle to capture high-frequency
 57 features that iterative solvers excel at.

58 Recognizing this complementarity, Zhang et al. [2024] proposed HINTS, a hybrid solver that inter-
 59 leaves a neural operator pass every τ^{th} iteration of a classical solver, with τ fixed a priori. While
 60 HINTS reports significant empirical gains in convergence and accuracy over the standalone neural
 61 operator and Jacobi solver, this fixed schedule can be detrimental: a poorly timed neural correction
 62 may increase the error and undo recent progress.

63 Although state-dependent solver schedules are desirable, this naturally suggest a reinforcement
 64 learning formulation for learning routing policies. However, in our setting, the structure of iterative
 65 PDE solvers allows for a simpler approach: each solver induces a known update, and its effect can be
 66 evaluated through error-based signals. Leveraging this, we adopt a greedy rule that selects, at each
 67 iteration, the solver with the largest immediate error reduction. Such a rule can be near-optimal when
 68 the final error satisfies (weak) supermodularity—so that applying a beneficial solver earlier is at least
 69 as valuable as applying it later. To support this approach, we make the following contributions:

- 70 • We present a general hybrid PDE solver in which a routing rule chooses, at each iteration,
 71 from a set of classical solvers and neural operators. With oracle access to the true error at
 72 every step, we show that a greedy routing rule achieves a constant-factor approximation to
 73 the optimal strategy for linear PDEs, provided the updates are error-reducing (Lipschitz with
 74 constant < 1) and zero-preserving—conditions under which weak supermodularity holds.
- 75 • Given that the true error is not observed at test time, a convex surrogate loss is introduced,
 76 which, when minimized, enables the learned model to imitate the greedy router without
 77 access to true error information. This alignment is guaranteed through Bayes consistency.
- 78 • We empirically demonstrate that our approximate greedy solver achieves faster convergence
 79 than HINTS and individual solvers on the Poisson, convection-diffusion and anisotropic
 80 diffusion equations, **in iterations and in end-to-end wall-clock time on 128^2 to 512^2 grids,**
 81 **where it is also faster than geometric multigrid and multigrid-preconditioned Krylov methods**
 82 **and extends to ensembles of classical solvers.**

83 2 Background

84 Consider the following linear PDE:

$$\mathcal{L}_x^a(u) = f, x \in D; \quad \mathcal{B}_x(u) = b, x \in \partial D \quad (1)$$

85 where $\mathcal{L}_x^a$ is a differential operator with respect to the spatial variable $x \in D$ parameterized by the
 86 coefficient function a , f is a forcing function, u is the PDE solution, $\mathcal{B}_x$ is the boundary operator, and
 87 b is the boundary term. Such PDEs can be solved numerically using various discretization strategies,
 88 such as finite differences [Smith, 1985, LeVeque, 2007], finite element, and spectral methods [Boyd,

2001, Canuto et al., 2006]. Here, we focus on finite differences, which replace derivatives with difference quotients (e.x., $\partial_x u(x) \approx (u(x+h) - u(x))/h$) on a uniform grid.

Formally, let h be the grid size and $G_h(D)$ be a uniform grid with spacing h over D . For a function $g : D \rightarrow \mathbb{R}$, we denote its restriction to $G_h(D)$ by $g_h \in \mathbb{R}^N$, where $N = |G_h(D)|$. If $\mathcal{L}_h^a \in \mathbb{R}^{N \times N}$ is the discretized operator, the PDE reduces to a linear system of equations

$$\mathcal{L}_h^a u_h = f_h, \quad (2)$$

with the boundary conditions incorporated in $\mathcal{L}_h^a$. Unlike much of the neural operator literature, which retains function space formulations for discretization invariance, we follow the conventions of classical numerical analysis and represent functions and operators as vectors and matrices. Discretization invariance is not essential here, since we focus on fixed-grid problems.

2.1 Iterative PDE Solvers

Direct methods like Gauss Elimination and Thomas Algorithm can be computationally expensive in high-dimensional domains when solving systems like Equation (2). In contrast, iterative methods like Jacobi and Gauss-Seidel offer computational speedups by iteratively updating the solution, gradually converging to the true solution for the PDE. The general iterative update is

$$u^{(t+1)} = u^{(t)} + C \left(f_h - \mathcal{L}_h^a u^{(t)} \right), \quad (3)$$

where $u^{(t)}$ is the t^{th} iterate of the solution and C is a preconditioning matrix. Jacobi uses $C = D^{-1}$, where D is the diagonal of $\mathcal{L}_h^a$, and Gauss-Seidel uses $C = (D + L)^{-1}$, where L denotes the strict lower triangular part. However, these methods tend to damp low frequency parts of the error slowly.

2.2 Neural Operators

Suppose we observe samples $\{(a_i, f_i, u_i)\}_{i=1}^N$ such that $(a_i, f_i) \sim \mathcal{P}$ are i.i.d. samples, and $u_i = \mathcal{G}^*(a_i, f_i)$ be generated by a deterministic solution operator $\mathcal{G}^*$. A neural operator (NO) $\mathcal{G}_\theta$ seeks to approximate $\mathcal{G}^*$ by minimizing the expected squared $L^2(D)$ error:

$$\mathcal{R}_{\text{NO}}(\mathcal{G}_\theta) = \mathbb{E}_{(a,f) \sim \mathcal{P}} \left[\int_D \|u_i(x) - \mathcal{G}_\theta(a_i, f_i)(x)\|_2^2 dx \right]$$

Neural operators use discretization-invariant layers. Prominent instantiations include Fourier Neural Operators (FNO), which use spectral transforms [Li et al., 2020a], Graph Neural Operators, which perform graph-based aggregation over sampled points [Li et al., 2020b], and DeepONets, which employ a trunk-branch decomposition to map input functions to output functions [Lu et al., 2021].

2.3 Greedy Optimization

Consider maximizing $g(S)$ over $S \subseteq \Omega$ with $|S| \leq T$ where $g : 2^\Omega \rightarrow \mathbb{R}$ is a set function defined on subsets of a set Ω . An exhaustive search is infeasible, so an efficient alternative is the greedy algorithm. It builds the solution subset iteratively by adding ω^* that yields the largest marginal gain, i.e. $\omega^* = \arg \max_{\omega \in \Omega \setminus S} g(S \cup \omega)$. When g is non-negative, monotone (adding elements never decreases the value), and submodular (diminishing returns), the greedy algorithm achieves a $(1 - e^{-1})$ approximation to the optimal solution [Nemhauser et al., 1978]. Formally, a function g is submodular if $g(A \cup \{\omega\}) - g(A) \geq g(B \cup \{\omega\}) - g(B)$ for all $A \subseteq B \subseteq \Omega$ and $\omega \in \Omega \setminus B$; intuitively, delaying an addition cannot increase its benefit. For set function minimization problems, supermodularity, where the inequality flips, plays an analogous role, and the greedy rule achieves constant-factor approximation guarantees [Liberty and Sviridenko, 2017].

The suboptimality of the greedy algorithm has been extensively studied for both set-function minimization [Bounia and Koriche, 2023] and maximization [Das and Kempe, 2018, Feige et al., 2011, Bian et al., 2017, Harshaw et al., 2019] when standard assumptions—non-negativity, monotonicity, and sub/supermodularity—are weakened. In contrast, results on greedy sequence maximization [Streeter and Golovin, 2008, Alaei et al., 2021, Zhang et al., 2015, Bernardini et al., 2020, Van Over et al., 2024, Tschitschek et al., 2017]—where the ordering of the elements affects the function—have led to sequential analogues of submodularity and monotonicity. We leverage these tools to analyze the suboptimality of the greedy solution to minimizing the final error.

3 General framework for Hybrid Solvers

To solve Equation (2), consider the following hybrid iterative update:

$$u_h^{(t+1)} = u_h^{(t)} + C_{S_t} \left(f_h - \mathcal{L}_h^a u_h^{(t)} \right) \quad (4)$$

where $\mathcal{C} = \{C_j\}_{j=1}^K$ is a set of preconditioning functions and $S_t \in [K] = \{1, \dots, K\}$ indexes the function chosen at step t . Here, we use “preconditioning function” broadly to include classical linear updates ($C_j(x) = C_j x$ where C_j is a matrix), and learned models, such as neural operators. This update generalizes the classical iterative update in Equation (3), by allowing C_j to be non-linear and enabling the preconditioning function to be adaptively chosen at every step. $\mathcal{C}$ can also accommodate parameterized solver families (e.g., different Jacobi relaxation weights), enabling adaptive selection of solver parameters. As the number of solvers K grows, the chance of selecting a more effective update increases. Furthermore, HINTS [Zhang et al., 2024] is a special case of this hybrid solver with $K = 2$, where C_1 is a neural operator and C_2 is a classical preconditioner. Its routing rule is given by $S_t = \mathbf{1}_{t \bmod \tau > 0} + 1$, which selects C_1 every τ steps and C_2 otherwise.

Let $e_h^{(t)} = u_h - u_h^{(t)}$ denote the error at step t . Then,

$$e_h^{(t+1)} = u_h - u_h^{(t)} - C_{S_t} \left(\mathcal{L}_h^a u_h - \mathcal{L}_h^a u_h^{(t)} \right) = (I - C_{S_t} \circ \mathcal{L}_h^a) (e_h^{(t)})$$

where I is the identity map and $I - C_j \circ \mathcal{L}_h^a$ is the error propagation function for solver j . Here, “ $\circ$ ” denotes function composition (matrix multiplication for linear updates). The objective is to select a sequence $S = (S_1, \dots, S_T)$ that minimizes the error norm after T steps or $\|e_h^{(T)}\|_2^2$:

$$\min_{S_t \in [K], |S| \leq T} h(S) := \left\| (I - C_{S_{|S|}} \circ \mathcal{L}_h^a) \circ \dots \circ (I - C_{S_1} \circ \mathcal{L}_h^a) \left(e_h^{(0)} \right) \right\|_2^2 \quad (5)$$

with compositions applied from right to left, so that the S_1 update acts on the initial error.

4 Greedy Algorithm

Equation (5) defines a combinatorial optimization problem with a search space exponential in T , making exact optimization intractable for large T or K . As a starting point, we consider an “omniscient” greedy algorithm that assumes access to the true initial error $e_h^{(0)}$. This is unrealistic since knowledge of $e_h^{(0)}$ could directly recover the solution via $u_h = u_h^{(0)} + e_h^{(0)}$, but it provides a useful benchmark. In Section 5, we relax this assumption using a practical learning strategy that is Bayes consistent with this omniscient rule, thereby recovering the guarantees shown below.

As discussed in Section 2.3, when supermodularity and monotonicity hold, the greedy rule—such as the one described in Algorithm 1—enjoys constant-factor approximation guarantees. However, classical results focus on *set* functions, with only recent extensions made to sequences. Building on this line, we introduce a sequence-based notion of weak supermodularity.

Algorithm 1 Greedy Algorithm for a Hybrid PDE solver

Require: $\{C_j\}_{j=1}^K, T, \mathcal{L}_h^a, e_h^{(0)}$
 $S^0 \leftarrow \emptyset$
for $t < T$ **do**
 $S^{t+1} \leftarrow S^t \oplus \arg \min_{j \in [K]} \|(I - C_j \circ \mathcal{L}_h^a)(e_h^{(t)})\|_2^2$
 $e_h^{(t+1)} \leftarrow (I - C_{S^{t+1}} \circ \mathcal{L}_h^a)(e_h^{(t)})$
end for
return S^T

Let Ω^* denote the space of sequences with elements in Ω . For $S, S' \in \Omega^*$, we denote their concatenation as $S \oplus S'$. S is a prefix of S' or $S \preceq S'$ if $S' = S \oplus L$ for some $L \in \Omega^*$. A function $g : \Omega^* \rightarrow \mathbb{R}$ is prefix monotonically non-increasing if $g(S \oplus S') \leq g(S)$ for all $S, S' \in \Omega^*$, and postfix monotonically non-increasing if $g(S \oplus S') \leq g(S')$ for all $S, S' \in \Omega^*$. A prefix non-increasing function g is sequence supermodular if, for all $S', S \in \Omega^* : S \preceq S'$, it holds that

$$g(S) - g(S \oplus \omega) \geq g(S') - g(S' \oplus \omega), \quad \forall \omega \in \Omega \quad (6)$$

166 However, $h(S)$ (defined in Equation (5)) may not satisfy this property in general. Therefore, we
 167 introduce weak sequence supermodularity. A prefix non-increasing function g is weakly supermodular
 168 with respect to $S' \in \Omega^*$ if, for any $S \in \Omega^{|S'|}$, there exists $\alpha(S') \geq 1$ such that

$$g(S) - g(S \oplus S') \leq \alpha(S') \sum_{i \in [|S'|]} g(S) - g(S \oplus S'_i) \quad (7)$$

169 The parameter $\alpha(S')$ or the supermodularity ratio quantifies deviation from exact sequence super-
 170 modularity. Expanding the $g(S) - g(S \oplus S')$ as a telescoping sum $\sum_{i=1}^{|S'|} g(S \oplus (S'_1, \dots, S'_{i-1})) -$
 171 $g(S \oplus (S'_1, \dots, S'_i))$ shows that the marginal decrease from appending S'_i after its predecessors is
 172 controlled by the effect of appending S'_i directly to S . Thus, postponing the inclusion of S'_i cannot
 173 yield a significantly larger benefit than adding it earlier. The supermodularity ratio $\alpha(S)$ quantifies
 174 the extent to which future gains from delays may exceed immediate gains.

175 Having introduced these notions, we now characterize the suboptimality of greedy solutions of weakly
 176 supermodular and postfix monotonic sequence functions in Theorem 4.1.

177 **Theorem 4.1.** *Let $g : \Omega^* \rightarrow \mathbb{R}$ be a weakly supermodular function with respect to the optimal*
 178 *solution $O = \arg \min_{S \in \Omega^T} h(S)$ with a supermodularity ratio of $\alpha(O)$ and postfix monotonicity. Let*
 179 *the greedy solution of length T be S^T . If $\phi_T(\alpha) = (1 - \frac{1}{\alpha T})^T$, then*

$$g(\emptyset) - g(S^T) \geq (1 - \phi_T(\alpha(O))) (g(\emptyset) - g(O))$$

180 The proof of Theorem 4.1 appears in Section A.2. As $T \rightarrow \infty$, the factor $1 - \phi_T(\alpha(O))$ decreases to
 181 $1 - e^{-1/\alpha(O)}$, so the worst-case performance of the greedy rule saturates rather than degrades with
 182 horizon. Additionally, larger $\alpha(O)$, which indicates higher reward for delayed inclusions, loosens the
 183 suboptimality bound. Theorem 4.1 requires that the sequence objective h be weakly supermodular
 184 with respect to the optimal solution and postfix monotone-properties established in Proposition 4.2.
 185 We defer the proof to Section A.3.

186 **Proposition 4.2.** *Suppose that for all $j \in [K]$, the error propagation function $I - C_j \circ \mathcal{L}_h^a$ is*
 187 *ρ_j -Lipschitz continuous with $\rho_j < 1$, and that $(I - C_j \circ \mathcal{L}_h^a)(0_N) = 0_N$. Then, the function h is*
 188 *weakly supermodular with respect to the optimal solution O , with*

$$\alpha(O) = \max \left\{ \frac{4}{T - \sum_{i=1}^T \rho_{O_i}^2}, 1 \right\}$$

189 Furthermore, if $I - C_j \circ \mathcal{L}_h^a$ is invertible for all $j \in [K]$, h is also postfix non-increasing.

190 The conditions of Proposition 4.2 are natural. For classical solvers, the Lipschitz constant is $\|I_N -$
 191 $C_j \mathcal{L}_h^a\|$, which is typically less than 1 for well-posed linear elliptic PDEs (i.e., the update is damping
 192 errors). For neural networks, Lipschitz continuity can be enforced via weight regularization [Gouk
 193 et al., 2021] and a sufficiently trained model should approximate $(\mathcal{L}_h^a)^{-1}$ well enough to make the
 194 Lipschitz constant small. The requirement $(I - C_j \circ \mathcal{L}_h^a)(0_N) = 0_N$ is both natural and desirable: it
 195 precludes spurious updates when the residual $f_h - \mathcal{L}_h^a u_h^{(t)}$ is 0. This holds by design for classical
 196 schemes, and it can be enforced for any learned model by excluding bias terms.

197 The form of $\alpha(O)$ in Proposition 4.2 highlights that the suboptimality factor in Theorem 4.1 is
 198 governed by the collective contraction factors of the solvers chosen by the optimal solution: as
 199 $\sum_{i=1}^T \rho_{O_i}^2 \rightarrow T$, $\alpha(O)$ grows, resulting in a weaker bound. Invertibility of the error propagation
 200 functions is often satisfied with Jacobi and Gauss-Seidel updates. However, neural networks with di-
 201 mension changes or ReLU activations may yield non-invertible error propagation maps. Nevertheless,
 202 our experiments show that greedy routers remain effective even when invertibility is not met.

203 Theorem 4.1 thus indicates that the approximation guarantee is strongest when the sequence is
 204 nearly supermodular ($\alpha(O) \approx 1$). Generally, if all error propagation maps share an eigenbasis,
 205 supermodularity is established. This occurs, for example, for linear, constant-coefficient PDEs
 206 with periodic boundary conditions where the solver ensemble includes Jacobi, Gauss-Seidel, and a
 207 single-layer linear Fourier Neural Operator. The proof of Theorem 4.3 is deferred to Section A.5.

208 **Proposition 4.3.** *Let $\|I_N - C_j \mathcal{L}_h^a\| \leq 1$ for all $j \in [K]$ and $(I - C_j \mathcal{L}_h^a) = P \Lambda_j P^{-1}$. Then, h is*
 209 *supermodular.*

Cost-aware greedy routing. Algorithm 1 treats every update as equally expensive, whereas in practice a neural-operator evaluation and a classical sweep have very different costs, and the quantity of interest is the error reached per unit of wall-clock time. We make the greedy rule cost-aware without changing the analysis: let c_j denote the (measured) cost of one application of C_j , take the cost u of one neural-operator call as the unit, and replace each preconditioner by the *macro-action* $C_j^{(m_j)}$ that applies C_j for $m_j = \max\{1, \text{round}(u/c_j)\}$ consecutive iterations, so that all actions cheaper than the unit cost approximately u ; an action dearer than the unit (a multigrid cycle, say) is applied once and its error is compared per unit of cost, $\|e\|^{u/(m_j c_j)}$, which reduces to the plain comparison when the costs match. For a linear update, m consecutive sweeps are again a stationary iteration, $(I - C_j \circ \mathcal{L}_h^a)^m = I - C_j^{(m)} \mathcal{L}_h^a$ with $C_j^{(m)} = \sum_{i=0}^{m-1} (I - C_j \mathcal{L}_h^a)^i C_j$, and for a nonlinear C_j the macro-action is the m -fold composition of its error propagation map. Hence $\{C_j^{(m_j)}\}_{j=1}^K$ is again a set of preconditioning functions in the sense of Section 4: if $I - C_j \circ \mathcal{L}_h^a$ is ρ_j -Lipschitz, zero-preserving and invertible, then so is its m_j -fold composition (with constant $\rho_j^{m_j} < 1$), and Theorem 4.1, Propositions 4.2 and 4.3, and the consistency result of Section 5 apply verbatim to the macro-action ensemble. Running Algorithm 1 on macro-actions selects, at every decision epoch, the action with the largest error reduction per unit cost, and a horizon of T macro-actions corresponds to a wall-clock budget of $\approx Tu$; Section E uses this cost-aware instantiation throughout.

5 Approximate Greedy Router

The results in Section 4 indicate that the error reduction from the greedy solution closely matches that of the optimal sequence, but this is predicated on having access to the initial error, $e_h^{(0)}$. Inaccurate error estimates can result in poor solver decisions, causing errors to amplify. To remedy this, we learn a router r that selects solvers myopically, as in Algorithm 1, without access to true errors.

We adopt the following learning setup. Let $\mathcal{A}, \mathcal{F}, \mathcal{U} \subseteq \mathbb{R}^N$ denote the spaces of coefficient, forcing, and solution functions on the grid $G_h(D)$. We assume an application-specific data distribution $\mathcal{P}_{\mathcal{A} \times \mathcal{F}}$ over $\mathcal{A} \times \mathcal{F}$ that reflects test time conditions. During training, (a_h, f_h) is drawn from $\mathcal{P}_{\mathcal{A} \times \mathcal{F}}$, and a high-accuracy reference solution u_h is computed, providing the true per-step error. The router r is learned offline on this data; at test time, it operates without access to true errors.

At iteration t , router r selects a solver using the coefficient $a_h \in \mathcal{A}$, forcing $f_h \in \mathcal{F}$, and the current iterate $u_h^{(t)} \in \mathcal{U}$. If $r(a_h, f_h, u_h^{(t)}) = j$, the next iterate is computed with solver j , resulting in an error of $\|(I - C_j \circ \mathcal{L}_h^a)(e_h^{(t)})\|_2^2$. Learning such a router requires minimizing the following loss:

$$l_{\text{route}}(r, a_h, f_h, u_h^{(t)}, u_h) = \sum_{j=1}^K \left\| (I - C_j \circ \mathcal{L}_h^a) (u_h - u_h^{(t)}) \right\|_2^2 \mathbf{1}_{r(a_h, f_h, u_h^{(t)})=j} \quad (8)$$

with risk $\mathcal{R}_{\text{route}}(r) = \mathbb{E}_{a_h, f_h \sim \mathcal{P}_{\mathcal{A} \times \mathcal{F}}} [l_{\text{route}}(r, a_h, f_h, u_h^{(t)}, u_h)]$. For analysis, we fix the iterate generation via teacher-forcing [Williams and Zipser, 1989, Lamb et al., 2016]: during training, the iterates fed to the router are produced by an oracle greedy rollout, not by the router's own past choices. Formally, with $u_h^{(0)} = 0_N$, for $t \geq 1$,

$$j_t^* \in \operatorname{argmin}_{j \in [K]} \left\| (I - C_j \circ \mathcal{L}_h^a) (e_h^{(t-1)}) \right\|_2^2, \quad u_h^{(t)} = u_h^{(t-1)} + C_{j_t^*} (f_h - \mathcal{L}_h^a u_h^{(t-1)})$$

Thus, $\{u_h^{(t)}\}$ is a deterministic function of (a_h, f_h) . At each t , l_{route} is evaluated on $r(a_h, f_h, u_h^{(t)})$ which takes in the teacher-forced iterate. Finally, the greedy choice j_{t+1}^* is used to advance $u_h^{(t+1)}$. Thus the input distribution seen by r depends on (a_h, f_h) , not on the router's predictions. However, full teacher forcing induces a distributional mismatch at test time (exposure bias). In experiments, we mitigate this with scheduled sampling [Bengio et al., 2015]; see Section C for details.

Minimizing $\mathcal{R}_{\text{route}}(r)$ yields the following Bayes-Optimal Router:

$$r^*(a_h, f_h, u_h^{(t)}) \in \operatorname{argmin}_{j \in [K]} \left\| (I - C_j \circ \mathcal{L}_h^a) (e_h^{(t)}) \right\|_2^2 \quad (9)$$

which aligns with the update in Algorithm 1. Hence, l_{route} is consistent with learning the greedy rule.

5.1 Surrogate Loss

Since Equation (8) is discontinuous and non-convex, direct minimization is intractable in practice. We instead introduce a surrogate loss that (1) is convex, enabling efficient optimization; (2) upper bounds the original loss; and (3) is Bayes consistent, ensuring the Bayes optimal decision of the original loss is preserved upon minimization. Formally, a surrogate ϕ is considered to be Bayes consistent with respect to the loss l if

$$\lim_{n \rightarrow \infty} \mathcal{R}_\phi(f_n) - \mathcal{R}_\phi^* \implies \lim_{n \rightarrow \infty} \mathcal{R}_l(f_n) - \mathcal{R}_l^*$$

where $\mathcal{R}_l(f) = \mathbb{E}[l(f(X), Y)]$ and $\mathcal{R}_l^* = \inf_f \mathcal{R}_l(f)$. In other words, in the limit of infinite data, if the risk of a sequence of learned hypotheses $\{f_n\}$ converges to the optimal risk under ϕ , it also converges to the optimal risk with respect to the original loss l .

To define a surrogate loss for the routing problem, consider a set of scoring functions $\mathbf{g} = \{g_j\}_{j=1}^K$ with $g_j : \mathcal{A} \times \mathcal{F} \times \mathcal{U} \rightarrow \mathbb{R}$ and we define the router as $r(a, f, u^{(t)}) = \arg\max_{j \in [K]} g_j(a, f, u^{(t)})$. For example, $\mathbf{g}$ can be a neural network with K outputs. Then, minimizing the following surrogate loss yields the same decision as minimizing Equation (8):

$$\Psi(\mathbf{g}, a_h, f_h, u_h^{(t)}, u_h) = - \sum_{j=1}^K \sum_{k=1}^K \tilde{c}_k(a_h, u_h^{(t)}, u_h) \mathbf{1}_{k \neq j} \log \left(\frac{\exp(g_j(a_h, f_h, u_h^{(t)}))}{\sum_{m=1}^K \exp(g_m(a_h, f_h, u_h^{(t)}))} \right) \quad (10)$$

where $\tilde{c}_j(a_h, u_h^{(t)}, u_h) = \|(I - C_j \circ \mathcal{L}_h^a)(u_h - u_h^{(t)})\|_2^2$. The Ψ -risk is denoted by $\mathcal{R}_\Psi(\mathbf{g}) = \mathbb{E}_{a_h, f_h \sim \mathcal{P}_{\mathcal{A} \times \mathcal{F}}}[\Psi(\mathbf{g}, a_h, f_h, u_h^{(t)}, u_h)]$. The convexity of Ψ with respect to $\mathbf{g}$ follows from the convexity of log-softmax function in its inputs. Moreover, Ψ upper bounds l_{route} up to a constant factor, and we refer the reader to Section B.2 for the proof. Finally, Theorem 5.1 shows that Ψ achieves Bayes consistency with respect to l_{route} ; the proof can be found in Section B.3.

Theorem 5.1. *Let $\tilde{c}_j(a_h, u_h^{(t)}, u_h) < \bar{E} < \infty$ for all $j \in [K]$. If there exists $j \in [K]$ such that $\tilde{c}_j(a_h, u_h^{(t)}, u_h) > E_{\min} > 0$, then, for any collection of solvers $\{C_j\}_{j=1}^K$ and linear discrete operator $\mathcal{L}_h^a$, Ψ is Bayes consistent surrogate for l_{route} .*

Theorem 5.1 is a cost-sensitive analogue of the classical Bayes-consistency of multiclass cross-entropy for 0-1 loss: in the infinite-sample limit, minimizing cross-entropy recovers the true conditional class probabilities, so the induced decision is Bayes optimal. Our result extends this to cross-entropy with instance-dependent weights $\sum_{k \neq j}^K \tilde{c}_k(a_h, u_h^{(t)}, u_h)$. The uniform upper bound holds when all preconditioning functions are error-damping. It is also reasonable to assume at least one solver cannot annihilate the error in one step, yielding the lower bound. Under these conditions, minimizing Ψ recovers the Bayes-optimal router of Equation (9), i.e., the greedy solution of Equation (5). Following the work of Mao et al. [2024], the proof uses standard conditional-risk calibration: we relate the excess risks of l_{route} and Ψ and take the infinite-sample limit.

6 Related Works

Hybrid PDE Solvers: Early data-driven solvers sought convergence guarantees by predicting parameters—e.g., preconditioning matrices, multi-grid smoothers or restriction matrices—within iterative schemes [Taghibakhshi et al., 2021, Caldana et al., 2024, Kopaničáková and Karniadakis, 2025, Huang et al., 2022, Katrutsa et al., 2020]. These works, however, do not leverage the neural surrogates’ ability to generalize across varying coefficients and forcings highlighted in Section 2.2 and thus, offer only modest speedups over classical solvers. HINTS [Zhang et al., 2024] introduced hybrid solvers that interleave a classical method with a pre-trained DeepONet on a fixed schedule (e.g., 24 Jacobi steps, then one DeepONet correction), achieving faster convergence than the standalone numerical solver. Subsequent works extend this idea to new geometries [Kahana et al., 2023] and analyze error mode damping, replacing DeepONet with alternatives such as MIONet [Jin et al., 2022] [Hu and Jin, 2025] or FNO [Li et al., 2020a, Cui et al., 2022]. However, these hybrid solvers are limited in two ways: they, firstly, only combine the trained surrogate with a *single* numerical solver, and secondly rely on a fixed, heuristic schedule. Our proposed method addresses these shortcomings, resulting in significant empirical improvements (Section 7).

Model Routing: Routing [Shnitzer et al., 2023, Hu et al., 2024, Ding et al., 2024, Huang et al., 2025] selects, for each input, a model from a fixed set to optimize a task metric under cost or latency constraints. Simple heuristics—e.g., thresholding a cheap model’s uncertainty estimates [Chuang et al., 2024, 2025]—often poorly balance the cost-accuracy tradeoff, motivating many systems to instead learn a router that maps inputs to the best model [Hari and Thomson, 2023, Mohammadshahi et al., 2024, Šakota et al., 2024]. We similarly learn a router, but over numerical and neural solvers for PDEs at the iteration level. Our work is the first to apply routing to the problem of learning hybrid PDE solvers, in contrast to prior approaches that use fixed schedules. Additionally, by exploiting the algebraic structure of PDE solvers, we derive theoretical guarantees for our routing strategy (Section 4), unlike typical model-routing settings.

Mixture of Experts: Classical mixture-of-experts (MoE) [Jacobs et al., 1991, Jordan and Jacobs, 1994] uses a gating network to combine the outputs of multiple experts, trained jointly with the gate. In modern LLMs, MoE instead performs sparse, token-level routing to a subset of in-layer experts, enabling large capacity without proportional compute, typically with load-balancing mechanisms [Shazeer et al., 2017, Lepikhin et al., 2020, Fedus et al., 2022, Jiang et al., 2024]. Our method can be regarded as gating of a different kind: “experts” (solvers) are not jointly trained with the router, and a single router is reused across iterations, unlike MoE which commonly employs layer-specific gates.

7 Experiments

We evaluate the approximate greedy router on three PDEs posed on the unit domain $D = [0, 1]^2$ with periodic boundary conditions:

$$\textbf{Poisson: } -\Delta u = f, \quad \textbf{ConvDiff: } -\Delta u + 20 \partial_{x_1} u + 20 \partial_{x_2} u = f, \quad \textbf{AnisoDiff: } -\epsilon \partial_{x_1}^2 u - \partial_{x_2}^2 u = f, \quad \epsilon = 10^{-2},$$

where f is centred to satisfy the compatibility condition $\int_D f(\mathbf{x}) d\mathbf{x} = 0$ and is drawn from a zero-mean hierarchical Gaussian random field with covariance operator $\alpha(-\Delta + \beta I)^{-\gamma}$ and sample-dependent (α, β, γ) (Section C). Anisotropic diffusion is the standard problem on which point relaxation fails to smooth: an error mode oscillatory in x_1 and smooth in x_2 is damped by only $(1 - \epsilon)/(1 + \epsilon) \approx 0.98$ per Jacobi sweep. The domain is discretised on 128^2 , 256^2 and 512^2 grids; results with Dirichlet boundary conditions, obtained with an earlier version of the pipeline, are deferred to Section D.1. Because iteration counts favour methods whose iterations are expensive, our primary metric is the end-to-end *wall-clock time to a target relative error* $\varepsilon = \|e_h^{(t)}\|/\|u_h\|$ on 64, 32 and 16 test instances per grid, with $\varepsilon = h^2$ —the truncation level of the second-order scheme, beyond which further algebraic refinement is not meaningful—highlighted; every comparison is paired over instances and tested with a one-sided Wilcoxon signed-rank test. We also report the iteration-based metrics of the original submission, the final relative error and the area under the error curve, $\text{AUC} = \sum_{t=1}^T \|e_h^{(t)}\|/\|u_h\|$ over $T = 300$ iterations. Section E gives the full protocol and collects every result per equation.

Setup. The classical solvers are Jacobi, damped Jacobi ($\omega = 0.67$), Gauss–Seidel (GS), symmetric GS (SymGS), SOR ($\omega = 1.5$) and geometric multigrid (V(2,2)-cycles with GS smoothing), all implemented as $O(N^2)$ stencil or cached sparse-triangular sweeps that match the dense implementations to machine precision, so the classical baselines are not artificially slowed. The neural operator is a DeepONet corrector that acts on the residual through a 64×64 sensor grid with a Fourier trunk basis and is applied scale-equivariantly: one call solves the smooth band of the spectrum to $\approx 10^{-6}$ at the cost of about 12 Jacobi or 3 GS sweeps at 128^2 (and only 2.5 Jacobi sweeps at 512^2 , since its cost is nearly grid-independent) and never injects high-frequency error. The router is made *cost-aware* through the cost-equalised macro-actions of Section 4 and selects from $6+K$ scalar features (log residual, its recent decay rate, the previous action and call counts) with a two-layer MLP whose decision costs $\approx 5 \mu\text{s}$ independently of N ; it is trained with the surrogate of Section 5.1 to imitate the cost-aware greedy oracle on oracle rollouts followed by Dagger rounds. We compare against the classical solver alone, HINTS with the published period $\tau=25$ and with the best period $\tau \in \{5, 10, 25, 50\}$ chosen *a posteriori* on the test set for each cell, the cost-aware greedy oracle (Algorithm 1 with the true error; not deployable), and the classical methods a practitioner would use on these problems without a corrector, multigrid alone and multigrid-preconditioned CG or BiCGSTAB.

Paired solvers. Table 1 reports, for every pairing $\{\text{NO}, \text{solver}\}$, equation and grid, the median wall-clock time to $\varepsilon = h^2$ of HINTS, of HINTS with the best period, and of the learned router. On the 128^2 grid the router is $238\times$ – $1983\times$ faster than the classical solver alone on the two isotropic equations (on the anisotropic equation no point smoother reaches h^2 at all within 60,000 iterations), $2.68\times$ – $11.8\times$ faster than HINTS ($\tau=25$) and $1.21\times$ – $3.22\times$ faster than the best period in every one of the 15 solver–equation pairings; on 256^2 it is $1.58\times$ – $10.9\times$ and $1.11\times$ – $2.33\times$ faster in all 13 pairings; and on 512^2 still $1.05\times$ – $3.10\times$ and $1.09\times$ – $1.24\times$ faster in all 6 stationary pairings, with the $\{\text{NO}, \text{multigrid}\}$ router the fastest iterative method we measured at that size (all $p < 0.01$; the margins shrink at 512^2 for the Jacobi and GS pairings because the corrector’s fixed band leaves those smoothers a slowly damped band that every schedule has to sweep through). Against the classical methods without corrector, the router with its best stationary pairing is $3.62\times$ – $3.71\times$ faster than geometric multigrid alone and $3.85\times$ – $5.31\times$ faster than the best multigrid-preconditioned Krylov method at 128^2 , and routing over $\{\text{NO}, \text{multigrid}\}$ is itself $1.93\times$ – $2.03\times$ faster than multigrid alone. The router tracks the cost-aware oracle within $1.06\times$ in median time and its decisions are identical across five independent training seeds. Figure 1 shows where the gain comes from: the router calls the corrector once at the very start, where the error is smooth and one call removes four orders of magnitude, spends a solver-dependent number of sweeps on the high-frequency band, and calls it again only if the smooth error re-emerges—between one and three calls in total—while HINTS spends $\tau-1$ sweeps before its first call and keeps calling the corrector when it no longer helps. The iteration-based metrics tell the same story: at 128^2 the router’s AUC over $T = 300$ iterations is orders of magnitude below HINTS’ and the solver-only baseline’s for every pairing and equation (paired t -tests, $p < 10^{-3}$) and its final error tracks the oracle’s (Tables 9, 15 and 21).

Solver ensembles. The same construction routes over ensembles $\text{NO} \cup \mathcal{W}$ of several classical solvers. Since the greedy rule is a per-decision comparison, a larger ensemble can only help where the error passes through regimes that favour *different* members; an exhaustive oracle-level screening of every subset of up to three members (Section E.5) shows that this never happens at $\varepsilon = h^2$ —one corrector call and one macro-action of the cheapest smoother already reach truncation level—but does at tolerances far below h^2 , where the smoother must damp the whole band outside the corrector’s reach: undamped Jacobi is cheapest on the modes near the band edge but cannot damp the near-Nyquist modes, which damped Jacobi removes. We therefore evaluate the learned router over the nested sets $\{\text{Jacobi}\}$, $\{\text{Jacobi} (0.67)\}$, $\{\text{Jacobi}, \text{Jacobi} (0.67)\}$ and $\{\text{Jacobi}, \text{Jacobi} (0.67), \text{GS}\}$ on every equation at 128^2 and 256^2 , with the single-member routers evaluated in the same session on the same instances. Two properties emerge. A router over a larger set is never slower than the best single-member router beyond a few percent, because the oracle ignores a member that never wins a cost-aware comparison and the router imitates it (adding GS changes nothing). And at $\varepsilon = 10^{-8}$ the two-member router is strictly faster than every single-member router in 2 of the 5 ensemble cells ($1.08\times$ – $1.14\times$ over the best single solver, $p < 0.01$): it learns a *nonstationary* damping schedule—undamped Jacobi while the band-edge modes dominate, damped Jacobi once the near-Nyquist modes do—that no single member can reproduce. Finally, we verify numerically that the assumptions of Sections 4 and 5 hold in all these settings, in the energy norm for the non-normal sweeps (Section E.7).

Every result per equation and grid, the ensemble study, the accounting of overheads and training cost, and the verification of the theoretical assumptions are collected in Section E; Section D reports the Dirichlet experiment of the original submission.

8 Discussion

We have introduced an adaptive method for selecting solvers that efficiently minimizes the final error of a PDE in iterative methods, opening many directions for future work. Our current DeepONet is trained without accounting for its downstream role in correction term prediction; jointly training the ML solver and router could yield larger gains. Another avenue is to use reinforcement learning to learn a cost-aware routing strategy that optimizes error under compute budgets; Sections 4 and E show that a myopic cost-aware variant of the greedy rule—Algorithm 1 on cost-equalised macro-actions—already delivers uniform wall-clock gains, and RL could extend it to non-myopic budgeted routing. RL would additionally enable extending the ensemble routing to *continuous* parameterization of solvers. Finally, framing routing as an online learning problem could enable adaptation to unknown deployment conditions.

Table 1: Median wall-clock time to truncation-level relative error $\varepsilon = h^2$ for every solver pairing $\{\text{NO, solver}\}$, equation and grid ($64, 32$ and 16 test instances at $128^2, 256^2$ and 512^2): HINTS with the published period $\tau=25$, HINTS with the best period chosen per cell on the test set itself (subscript), and the learned router (paired median speedup over HINTS-25 / over the best period in parentheses; bold: faster than both). The classical solver alone is 10^2 – $10^3 \times$ slower than the router on the isotropic equations and never reaches h^2 on the anisotropic one (Section E). The last rows of each block give geometric multigrid alone and the best multigrid-preconditioned Krylov method, without corrector, on the same instances. $\dagger n$: n instances did not reach h^2 within the iteration cap.

Equation	Pairing	128 $\times$ 128			256 $\times$ 256			512 $\times$ 512		
		HINTS-25	best τ	router (ours)	HINTS-25	best τ	router (ours)	HINTS-25	best τ	router (ours)
Poisson	Jacobi	2.65 ms	1.10 ms _{$\tau=5$}	987 μs ($2.68 \times / 1.21 \times$)	5.95 ms	2.47 ms _{$\tau=5$}	2.38 ms ($2.52 \times / 1.56 \times$)	87.0 ms	87.0 ms _{$\tau=25$}	78.8 ms ^{†2} ($1.11 \times / 1.11 \times$)
	Jacobi (0.67)	2.71 ms	1.13 ms _{$\tau=5$}	912 μs ($2.90 \times / 1.28 \times$)	6.41 ms	3.43 ms _{$\tau=5$}	2.83 ms ($2.26 \times / 1.73 \times$)	—	—	—
	GS	6.65 ms	1.99 ms _{$\tau=5$}	1.01 ms ($6.75 \times / 1.94 \times$)	22.0 ms	6.07 ms _{$\tau=5$}	3.92 ms ($5.77 \times / 1.70 \times$)	142.8 ms	142.8 ms _{$\tau=25$}	131.5 ms ($1.11 \times / 1.11 \times$)
	SymGS	10.3 ms	2.60 ms _{$\tau=5$}	1.13 ms ($9.21 \times / 2.34 \times$)	37.3 ms	8.71 ms _{$\tau=5$}	4.22 ms ($9.96 \times / 2.27 \times$)	—	—	—
	SOR (1.5)	12.2 ms	3.37 ms _{$\tau=5$}	1.02 ms ($10.3 \times / 3.22 \times$)	40.7 ms	9.06 ms _{$\tau=5$}	3.80 ms ($10.9 \times / 2.33 \times$)	182.7 ms	74.4 ms _{$\tau=5$}	60.7 ms ($3.10 \times / 1.24 \times$)
	Multigrid	3.83 ms	3.81 ms _{$\tau=5$}	2.05 ms ($1.96 \times / 1.88 \times$)	14.7 ms	14.6 ms _{$\tau=5$}	5.97 ms ($2.46 \times / 2.39 \times$)	92.0 ms	90.5 ms _{$\tau=10$}	38.7 ms ($2.08 \times / 2.10 \times$)
	Multigrid alone (no corrector)	—	—	4.39 ms	—	—	14.1 ms	—	—	81.9 ms
ConvDiff	PCG (MG) (no corrector)	—	—	4.35 ms	—	—	15.2 ms	—	—	85.6 ms
	Jacobi	3.44 ms	1.28 ms _{$\tau=5$}	1.14 ms ($2.95 \times / 1.34 \times$)	9.37 ms	4.95 ms _{$\tau=10$}	4.76 ms ($1.92 \times / 1.30 \times$)	158.0 ms	147.4 ms _{$\tau=50$}	144.3 ms ($1.17 \times / 1.14 \times$)
	Jacobi (0.67)	3.40 ms	1.31 ms _{$\tau=5$}	1.10 ms ($3.14 \times / 1.32 \times$)	9.31 ms	5.68 ms _{$\tau=10$}	5.76 ms ($1.80 \times / 1.25 \times$)	—	—	—
	GS	8.27 ms	3.60 ms _{$\tau=5$}	1.35 ms ($6.76 \times / 2.60 \times$)	31.0 ms	10.5 ms _{$\tau=5$}	7.05 ms ($5.52 \times / 1.53 \times$)	214.1 ms	192.7 ms _{$\tau=10$}	179.9 ms ($1.05 \times / 1.09 \times$)
	SymGS	11.1 ms	2.96 ms _{$\tau=5$}	1.28 ms ($8.70 \times / 2.40 \times$)	42.3 ms	10.9 ms _{$\tau=5$}	6.96 ms ($6.26 \times / 1.69 \times$)	—	—	—
	SOR (1.5)	13.6 ms	3.64 ms _{$\tau=5$}	1.59 ms ($8.89 \times / 2.40 \times$)	46.1 ms	10.6 ms _{$\tau=5$}	5.80 ms ($7.99 \times / 1.96 \times$)	206.6 ms	89.3 ms _{$\tau=5$}	83.6 ms ($2.48 \times / 1.18 \times$)
	Multigrid	4.68 ms	4.66 ms _{$\tau=5$}	2.32 ms ($2.07 \times / 2.01 \times$)	16.7 ms	16.5 ms _{$\tau=5$}	7.26 ms ($2.19 \times / 2.20 \times$)	97.6 ms	92.5 ms _{$\tau=10$}	53.7 ms ($1.83 \times / 1.75 \times$)
AnisoDiff	Multigrid alone (no corrector)	—	—	4.52 ms	—	—	18.8 ms	—	—	99.7 ms
	BiCGSTAB (MG) (no corrector)	—	—	6.17 ms	—	—	24.0 ms	—	—	93.5 ms
	Jacobi	2.90 ms	1.15 ms _{$\tau=5$}	911 μs ($3.16 \times / 1.25 \times$)	14.9 ms	10.3 ms _{$\tau=5$}	9.55 ms ($1.58 \times / 1.11 \times$)	—	—	—
	Jacobi (0.67)	2.85 ms	1.18 ms _{$\tau=5$}	850 μs ($3.31 \times / 1.34 \times$)	—	—	—	—	—	—
	GS	6.48 ms	1.74 ms _{$\tau=5$}	852 μs ($7.61 \times / 1.99 \times$)	29.7 ms	12.8 ms _{$\tau=5$}	9.30 ms ($3.18 \times / 1.37 \times$)	—	—	—
	SymGS	10.3 ms	2.40 ms _{$\tau=5$}	871 μs ($11.8 \times / 2.72 \times$)	—	—	—	—	—	—
	SOR (1.5)	6.62 ms	1.80 ms _{$\tau=5$}	869 μs ($7.36 \times / 2.04 \times$)	29.7 ms	12.8 ms _{$\tau=5$}	9.93 ms ($2.98 \times / 1.38 \times$)	—	—	—




Figure 1: Corrector usage on the 128×128 grid: fraction of test runs that execute a corrector call at each iteration, for the learned router, the cost-aware oracle and HINTS ($\tau=25$), for every solver pairing (columns) and all three equations (rows); the same figure for the 256^2 and 512^2 grids is in Section E.

References

- Eugenia Kalnay. *Atmospheric modeling, data assimilation and predictability*. Cambridge university press, 2003.
- Peter Bauer, Alan Thorpe, and Gilbert Brunet. The quiet revolution of numerical weather prediction. *Nature*, 525(7567):47–55, 2015.
- Andrew Staniforth and Jean Côté. Semi-lagrangian integration schemes for atmospheric models—a review. *Monthly weather review*, 119(9):2206–2223, 1991.
- John David Anderson. *Hypersonic and high temperature gas dynamics*. Aiaa, 1989.

409 Fredi Tröltzsch. *Optimal control of partial differential equations: theory, methods, and applications*,
410 volume 112. American Mathematical Soc., 2010.

411 Sylvain Baillet, John C Mosher, and Richard M Leahy. Electromagnetic brain mapping. *IEEE Signal
412 processing magazine*, 18(6):14–30, 2001.

413 Roberta Grech, Tracey Cassar, Joseph Muscat, Kenneth P Camilleri, Simon G Fabri, Michalis
414 Zervakis, Petros Xanthopoulos, Vangelis Sakkalis, and Bart Vanrumste. Review on solving the
415 inverse problem in eeg source analysis. *Journal of neuroengineering and rehabilitation*, 5:1–33,
416 2008.

417 Gordon D Smith. *Numerical solution of partial differential equations: finite difference methods*.
418 Oxford university press, 1985.

419 Randall J LeVeque. *Finite difference methods for ordinary and partial differential equations: steady-
420 state and time-dependent problems*. SIAM, 2007.

421 Thomas JR Hughes. *The finite element method: linear static and dynamic finite element analysis*.
422 Courier Corporation, 2003.

423 Klaus-Jürgen Bathe. *Finite element procedures*. Klaus-Jurgen Bathe, 2006.

424 Susanne C Brenner. *The mathematical theory of finite element methods*. Springer, 2008.

425 Yousef Saad. *Iterative methods for sparse linear systems*. SIAM, 2003.

426 Nikola Kovachki, Zongyi Li, Burigede Liu, Kamyar Azizzadenesheli, Kaushik Bhattacharya, Andrew
427 Stuart, and Anima Anandkumar. Neural operator: Learning maps between function spaces with
428 applications to pdes. *Journal of Machine Learning Research*, 24(89):1–97, 2023.

429 Tianping Chen and Hong Chen. Universal approximation to nonlinear operators by neural networks
430 with arbitrary activation functions and its application to dynamical systems. *IEEE transactions on
431 neural networks*, 6(4):911–917, 1995.

432 Lizuo Liu and Wei Cai. Multiscale deepnet for nonlinear operators in oscillatory function spaces for
433 building seismic wave responses. *arXiv preprint arXiv:2111.04860*, 2021.

434 Xinliang Liu, Bo Xu, Shuhao Cao, and Lei Zhang. Mitigating spectral bias for the multiscale operator
435 learning. *Journal of Computational Physics*, 506:112944, 2024.

436 Zhilin You, Zhenli Xu, and Wei Cai. Mscalefno: Multi-scale fourier neural operator learning for
437 oscillatory function spaces. *arXiv preprint arXiv:2412.20183*, 2024.

438 Siavash Khodakarami, Vivek Oommen, Aniruddha Bora, and George Em Karniadakis. Mitigating
439 spectral bias in neural operators via high-frequency scaling for physical systems. *arXiv preprint
440 arXiv:2503.13695*, 2025.

441 Zhi-Qin John Xu, Lulu Zhang, and Wei Cai. On understanding and overcoming spectral biases of
442 deep neural network learning methods for solving pdes. *arXiv preprint arXiv:2501.09987*, 2025.

443 Nasim Rahaman, Aristide Baratin, Devansh Arpit, Felix Draxler, Min Lin, Fred Hamprecht, Yoshua
444 Bengio, and Aaron Courville. On the spectral bias of neural networks. In *International conference
445 on machine learning*, pages 5301–5310. PMLR, 2019.

446 Zhi-Qin John Xu, Yaoyu Zhang, Tao Luo, Yanyang Xiao, and Zheng Ma. Frequency principle:
447 Fourier analysis sheds light on deep neural networks. *arXiv preprint arXiv:1901.06523*, 2019.

448 Tao Luo, Zheng Ma, Zhi-Qin John Xu, and Yaoyu Zhang. Theory of the frequency principle for
449 general deep neural networks. *arXiv preprint arXiv:1906.09235*, 2019.

450 Enrui Zhang, Adar Kahana, Alena Kopaničáková, Eli Turkel, Rishikesh Ranade, Jay Pathak, and
451 George Em Karniadakis. Blending neural operators and relaxation methods in pde numerical
452 solvers. *Nature Machine Intelligence*, pages 1–11, 2024.

453 John P Boyd. *Chebyshev and Fourier spectral methods*. Courier Corporation, 2001.

454 Claudio Canuto, M. Yousuff Hussaini, Alfio Quarteroni, and Thomas A. Zang. *Spectral Methods:
455 Fundamentals in Single Domains*. Springer, 2006. ISBN 9783540307264.

456 Zongyi Li, Nikola Kovachki, Kamyar Azizzadenesheli, Burigede Liu, Kaushik Bhattacharya, Andrew
457 Stuart, and Anima Anandkumar. Fourier neural operator for parametric partial differential equations.
458 *arXiv preprint arXiv:2010.08895*, 2020a.

459 Zongyi Li, Nikola Kovachki, Kamyar Azizzadenesheli, Burigede Liu, Andrew Stuart, Kaushik
460 Bhattacharya, and Anima Anandkumar. Multipole graph neural operator for parametric partial
461 differential equations. *Advances in Neural Information Processing Systems*, 33:6755–6766, 2020b.

462 Lu Lu, Pengzhan Jin, Guofei Pang, Zhongqiang Zhang, and George Em Karniadakis. Learning
463 nonlinear operators via deepnet based on the universal approximation theorem of operators.
464 *Nature machine intelligence*, 3(3):218–229, 2021.

465 George L Nemhauser, Laurence A Wolsey, and Marshall L Fisher. An analysis of approximations for
466 maximizing submodular set functions—i. *Mathematical programming*, 14(1):265–294, 1978.

467 Edo Liberty and Maxim Sviridenko. Greedy minimization of weakly supermodular set functions.
468 In *Approximation, Randomization, and Combinatorial Optimization. Algorithms and Techniques*
469 (*APPROX/RANDOM 2017*), pages 19–1. Schloss Dagstuhl–Leibniz-Zentrum für Informatik, 2017.

470 Louenas Bounia and Frederic Koriche. Approximating probabilistic explanations via supermodular
471 minimization. In *Uncertainty in Artificial Intelligence*, pages 216–225. PMLR, 2023.

472 Abhimanyu Das and David Kempe. Approximate submodularity and its applications: Subset selection,
473 sparse approximation and dictionary selection. *Journal of Machine Learning Research*, 19(3):
474 1–34, 2018.

475 Uriel Feige, Vahab S Mirrokni, and Jan Vondrák. Maximizing non-monotone submodular functions.
476 *SIAM Journal on Computing*, 40(4):1133–1153, 2011.

477 Andrew An Bian, Joachim M Buhmann, Andreas Krause, and Sebastian Tschieschek. Guarantees for
478 greedy maximization of non-submodular functions with applications. In *International conference*
479 *on machine learning*, pages 498–507. PMLR, 2017.

480 Chris Harshaw, Moran Feldman, Justin Ward, and Amin Karbasi. Submodular maximization beyond
481 non-negativity: Guarantees, fast algorithms, and applications. In *International Conference on*
482 *Machine Learning*, pages 2634–2643. PMLR, 2019.

483 Matthew Streeter and Daniel Golovin. An online algorithm for maximizing submodular functions.
484 *Advances in Neural Information Processing Systems*, 21, 2008.

485 Saeed Alaei, Ali Makhdoumi, and Azarakhsh Malekian. Maximizing sequence-submodular functions
486 and its application to online advertising. *Management Science*, 67(10):6030–6054, 2021.

487 Zhenliang Zhang, Edwin KP Chong, Ali Pezeshki, and William Moran. String submodular functions
488 with curvature constraints. *IEEE Transactions on Automatic Control*, 61(3):601–616, 2015.

489 Sara Bernardini, Fabio Fagnani, and Chiara Piacentini. Through the lens of sequence submodularity.
490 In *Proceedings of the International Conference on Automated Planning and Scheduling*, volume 30,
491 pages 38–47, 2020.

492 Brandon Van Over, Bowen Li, Edwin KP Chong, and Ali Pezeshki. A performance bound for
493 the greedy algorithm in a generalized class of string optimization problems. *arXiv preprint*
494 *arXiv:2409.05020*, 2024.

495 Sebastian Tschieschek, Adish Singla, and Andreas Krause. Selecting sequences of items via
496 submodular maximization. In *Proceedings of the AAAI Conference on Artificial Intelligence*,
497 volume 31, 2017.

498 Henry Gouk, Eibe Frank, Bernhard Pfahringer, and Michael J Cree. Regularisation of neural networks
499 by enforcing lipschitz continuity. *Machine Learning*, 110(2):393–416, 2021.

500 Ronald J Williams and David Zipser. A learning algorithm for continually running fully recurrent
501 neural networks. *Neural computation*, 1(2):270–280, 1989.

502 Alex M Lamb, Anirudh Goyal ALIAS PARTH GOYAL, Ying Zhang, Saizheng Zhang, Aaron C
503 Courville, and Yoshua Bengio. Professor forcing: A new algorithm for training recurrent networks.
504 *Advances in neural information processing systems*, 29, 2016.

505 Samy Bengio, Oriol Vinyals, Navdeep Jaitly, and Noam Shazeer. Scheduled sampling for sequence
506 prediction with recurrent neural networks. *Advances in neural information processing systems*, 28,
507 2015.

508 Anqi Mao, Mehryar Mohri, and Yutao Zhong. Regression with multi-expert deferral. In *International
509 Conference on Machine Learning*, pages 34738–34759. PMLR, 2024.

510 Ali Taghibakhshi, Scott MacLachlan, Luke Olson, and Matthew West. Optimization-based algebraic
511 multigrid coarsening using reinforcement learning. *Advances in neural information processing
512 systems*, 34:12129–12140, 2021.

513 Matteo Caldana, Paola F Antonietti, et al. A deep learning algorithm to accelerate algebraic multigrid
514 methods in finite element solvers of 3d elliptic pdes. *Computers & Mathematics with Applications*,
515 167:217–231, 2024.

516 Alena Kopaničáková and George Em Karniadakis. Deeponet based preconditioning strategies for
517 solving parametric linear systems of equations. *SIAM Journal on Scientific Computing*, 47(1):
518 C151–C181, 2025.

519 Ru Huang, Ruipeng Li, and Yuanzhe Xi. Learning optimal multigrid smoothers via neural networks.
520 *SIAM Journal on Scientific Computing*, 45(3):S199–S225, 2022.

521 Alexandr Katrutsa, Talgat Daulbaev, and Ivan Oseledets. Black-box learning of multigrid parameters.
522 *Journal of Computational and Applied Mathematics*, 368:112524, 2020.

523 Adar Kahana, Enrui Zhang, Somdatta Goswami, George Karniadakis, Rishikesh Ranade, and Jay
524 Pathak. On the geometry transferability of the hybrid iterative numerical solver for differential
525 equations. *Computational Mechanics*, 72(3):471–484, 2023.

526 Pengzhan Jin, Shuai Meng, and Lu Lu. Mionet: Learning multiple-input operators via tensor product.
527 *SIAM Journal on Scientific Computing*, 44(6):A3490–A3514, 2022.

528 Jun Hu and Pengzhan Jin. A hybrid iterative method based on mionet for pdes: Theory and numerical
529 examples. *Mathematics of Computation*, 2025.

530 Chen Cui, Kai Jiang, Yun Liu, and Shi Shu. Fourier neural solver for large sparse linear algebraic
531 systems. *Mathematics*, 10(21):4014, 2022.

532 Tal Shnitzer, Anthony Ou, Mírian Silva, Kate Soule, Yuekai Sun, Justin Solomon, Neil Thompson,
533 and Mikhail Yurochkin. Large language model routing with benchmark datasets. *arXiv preprint
534 arXiv:2309.15789*, 2023.

535 Qitian Jason Hu, Jacob Bieker, Xiuyu Li, Nan Jiang, Benjamin Keigwin, Gaurav Ranganath, Kurt
536 Keutzer, and Shriyash Kaustubh Upadhyay. Routerbench: A benchmark for multi-llm routing
537 system. *arXiv preprint arXiv:2403.12031*, 2024.

538 Dujian Ding, Ankur Mallick, Chi Wang, Robert Sim, Subhabrata Mukherjee, Victor Ruhle, Laks VS
539 Lakshmanan, and Ahmed Hassan Awadallah. Hybrid llm: Cost-efficient and quality-aware query
540 routing. *arXiv preprint arXiv:2404.14618*, 2024.

541 Zhongzhan Huang, Guoming Ling, Yupei Lin, Yandong Chen, Shanshan Zhong, Hefeng Wu, and
542 Liang Lin. Routereval: A comprehensive benchmark for routing llms to explore model-level
543 scaling up in llms. *arXiv preprint arXiv:2503.10657*, 2025.

544 Yu-Neng Chuang, Prathusha Kameswara Sarma, Parikshit Gopalan, John Boccio, Sara Bolouki,
545 Xia Hu, and Helen Zhou. Learning to route llms with confidence tokens. *arXiv preprint
546 arXiv:2410.13284*, 2024.

547 Yu-Neng Chuang, Leisheng Yu, Guanchu Wang, Lizhe Zhang, Zirui Liu, Xuanning Cai, Yang Sui,
548 Vladimir Braverman, and Xia Hu. Confident or seek stronger: Exploring uncertainty-based
549 on-device llm routing from benchmarking to generalization. *arXiv preprint arXiv:2502.04428*,
550 2025.

551 Surya Narayanan Hari and Matt Thomson. Tryage: Real-time, intelligent routing of user prompts to
552 large language models. *arXiv preprint arXiv:2308.11601*, 2023.

553 Alireza Mohammadshahi, Arshad Rafiq Shaikh, and Majid Yazdani. Routoo: Learning to route to
554 large language models effectively. *arXiv preprint arXiv:2401.13979*, 2024.

555 Marija Šakota, Maxime Peyrard, and Robert West. Fly-swat or cannon? cost-effective language
556 model choice via meta-modeling. In *Proceedings of the 17th ACM International Conference on*
557 *Web Search and Data Mining*, pages 606–615, 2024.

558 Robert A Jacobs, Michael I Jordan, Steven J Nowlan, and Geoffrey E Hinton. Adaptive mixtures of
559 local experts. *Neural computation*, 3(1):79–87, 1991.

560 Michael I Jordan and Robert A Jacobs. Hierarchical mixtures of experts and the em algorithm. *Neural*
561 *computation*, 6(2):181–214, 1994.

562 Noam Shazeer, Azalia Mirhoseini, Krzysztof Maziarz, Andy Davis, Quoc Le, Geoffrey Hinton, and
563 Jeff Dean. Outrageously large neural networks: The sparsely-gated mixture-of-experts layer. *arXiv*
564 *preprint arXiv:1701.06538*, 2017.

565 Dmitry Lepikhin, HyoukJoong Lee, Yuanzhong Xu, Dehao Chen, Orhan Firat, Yanping Huang,
566 Maxim Krikun, Noam Shazeer, and Zhifeng Chen. Gshard: Scaling giant models with conditional
567 computation and automatic sharding. *arXiv preprint arXiv:2006.16668*, 2020.

568 William Fedus, Barret Zoph, and Noam Shazeer. Switch transformers: Scaling to trillion parameter
569 models with simple and efficient sparsity. *Journal of Machine Learning Research*, 23(120):1–39,
570 2022.

571 Albert Q Jiang, Alexandre Sablayrolles, Antoine Roux, Arthur Mensch, Blanche Savary, Chris
572 Bamford, Devendra Singh Chaplot, Diego de las Casas, Emma Bou Hanna, Florian Bressand, et al.
573 Mixtral of experts. *arXiv preprint arXiv:2401.04088*, 2024.

574 Anqi Mao, Mehryar Mohri, and Yutao Zhong. Cross-entropy loss functions: Theoretical analysis and
575 applications. In *International conference on Machine learning*, pages 23803–23828. pmlr, 2023.

576 Michael C Mozer. A focused backpropagation algorithm for temporal pattern recognition. In
577 *Backpropagation*, pages 137–169. Psychology Press, 2013.

578 Anthony J Robinson and Frank Fallside. *The utility driven dynamic error propagation network*,
579 volume 11. University of Cambridge Department of Engineering Cambridge, 1987.

580 Paul J Werbos. Generalization of backpropagation with application to a recurrent gas market model.
581 *Neural networks*, 1(4):339–356, 1988.

582 Lu Lu, Xuhui Meng, Shengze Cai, Zhiping Mao, Somdatta Goswami, Zhongqiang Zhang, and
583 George Em Karniadakis. A comprehensive and fair comparison of two neural operators (with prac-
584 tical extensions) based on FAIR data. *Computer Methods in Applied Mechanics and Engineering*,
585 393:114778, 2022.

586 Stéphane Ross, Geoffrey Gordon, and Drew Bagnell. A reduction of imitation learning and structured
587 prediction to no-regret online learning. In *Proceedings of the Fourteenth International Confer-*
588 *ence on Artificial Intelligence and Statistics*, pages 627–635. JMLR Workshop and Conference
589 Proceedings, 2011.

590 A Proofs for Section 4

591 A.1 Proof of Proposition A.1

592 **Proposition A.1.** *Any prefix monotonically non-increasing sequence supermodular function g is*
 593 *weakly supermodular with respect to all sequences $S \in \Omega^*$ with $\alpha(S) = 1$*

594 *Proof.* This proof is adapted from Liberty and Sviridenko [2017].

595 If g is sequence supermodular then,

$$\begin{aligned}
 & g(S) - g(S \oplus S') \\
 &= \sum_{i=1}^{|S'|} g(S \oplus (S'_1, \dots, S'_{i-1})) - g(S \oplus (S'_1, \dots, S'_i)) \\
 &\stackrel{(a)}{\leq} \sum_{i=1}^{|S'|} g(S) - g(S \oplus S'_i) \\
 &\leq |S'| \max_{i \in [|S'|]} g(S) - g(S \oplus S'_i)
 \end{aligned}$$

596 (a) by supermodularity

597 □

598 A.2 Proof of Theorem 4.1

599 **Theorem 4.1.** *Let $g : \Omega^* \rightarrow \mathbb{R}$ be a weakly supermodular function with respect to the optimal*
 600 *solution $O = \arg \min_{S \in \Omega^T} h(S)$ with a supermodularity ratio of $\alpha(O)$ and postfix monotonicity. Let*
 601 *the greedy solution of length T be S^T . If $\phi_T(\alpha) = (1 - \frac{1}{\alpha T})^T$, then*

$$g(S^T) \leq (1 - \phi_T(\alpha(O))) g(O) + \phi_T(\alpha(O)) g(\emptyset)$$

602 *Proof.* This proof strategy is inspired by Streeter and Golovin [2008] and Liberty and Sviridenko
 603 [2017]

$$\begin{aligned}
 g(S^t) - g(O) &\stackrel{(a)}{\leq} g(S^t) - g(S^t \oplus O) \\
 &= \sum_{i=1}^{|O|} g(S^t \oplus (o_1, \dots, o_{i-1})) - g(S^t \oplus (o_1, \dots, o_i)) \\
 &\stackrel{(b)}{\leq} \alpha(O) \sum_{i=1}^{|O|} g(S^t) - g(S^t \oplus o_i) \\
 &\leq \alpha(O) |O| \max_{i \in [|O|]} g(S^t) - g(S^t \oplus o_i) \\
 &\leq \alpha(O) T g(S^t) - \alpha(O) T \min_{\omega \in \Omega} g(S^t \oplus \omega) \\
 &= \alpha(O) T g(S^t) - \alpha(O) T g(S^{t+1})
 \end{aligned}$$

604 (a) by μ - postfix monotonicity, (b) by supermodularity.

605 After rearranging the inequality, we get:

$$\begin{aligned}
 g(S^{t+1}) &\leq \frac{1}{\alpha(O)T} (g(O) - (\alpha(O)T - 1) g(S^t)) \\
 &= \frac{1}{\alpha(O)T} g(O) + \left(1 - \frac{1}{\alpha(O)T}\right) g(S^t)
 \end{aligned}$$

606 When recursively applying this inequality, we get:

$$\begin{aligned} g(S^T) &\leq \frac{1}{\alpha(O)T} g(O) \sum_{i=0}^{T-1} \left(1 - \frac{1}{\alpha(O)T}\right)^i + \left(1 - \frac{1}{\alpha(O)T}\right)^T g(\emptyset) \\ &= \left(1 - \left(1 - \frac{1}{\alpha(O)T}\right)^T\right) g(O) + \left(1 - \frac{1}{\alpha(O)T}\right)^T g(\emptyset) \end{aligned}$$

607

□

608 A.3 Proof of Proposition 4.2

609 **Proposition 4.2.** Suppose that for all $j \in [K]$, the error propagation function $I - C_j \circ \mathcal{L}_h^a$ is
 610 ρ_j -Lipschitz continuous with $\rho_j < 1$, and that $(I - C_j \circ \mathcal{L}_h^a)(0_N) = 0_N$. Then, the function h is
 611 weakly supermodular with respect to the optimal solution O , with

$$\alpha(O) = \max \left\{ \frac{4}{T - \sum_{i=1}^T \rho_{O_i}^2}, 1 \right\}$$

612 Furthermore, if $I - C_j \circ \mathcal{L}_h^a$ is invertible for all $j \in [K]$, h is also postfix monotonically non-increasing.

613 *Proof.* For brevity, we use the notation $(g_1 \circ \dots \circ g_T)(x) = \circ_{t=1}^T g_t(x)$, where composition is
 614 applied from right to left so that g_T acts first. In this proof, we use a few properties of Lipschitz
 615 continuous functions:

- 616 • **Property 1:** If g is ρ -Lipschitz continuous and $g(0) = 0$, then $\|g(x)\|_2 = \|g(x) - 0\|_2 =$
 617 $\|g(x) - g(0)\|_2 \leq \rho \|x - 0\|_2 = \rho \|x\|_2$
- 618 • **Property 2:** If g_1 and g_2 are Lipschitz continuous functions with Lipschitz constants of ρ_1
 619 and ρ_2 respectively, the Lipschitz constant of $g_1 + g_2$ and $g_1 - g_2$ is $\rho_1 + \rho_2$.
- 620 • **Property 3:** If g_1 and g_2 are Lipschitz continuous functions with Lipschitz constants of ρ_1
 621 and ρ_2 respectively, the Lipschitz constant of $g_1 \circ g_2$ is $\rho_1 \rho_2$.

622 In order to prove weakly- α -supermodularity, we must first prove prefix monotonicity.

623 **Prefix monotonicity:** Let $S \preceq S'$ where $S' = S \oplus N$.

$$\begin{aligned} h(S') &= \left\| \circ_{t=|S'|}^1 (I - C_{S'_t} \circ \mathcal{L}_h^a) \left(e_h^{(0)} \right) \right\|_2^2 \\ &= \left\| \circ_{t=|S'|}^{|S|+1} (I - C_{S'_t} \circ \mathcal{L}_h^a) \circ \circ_{t=|S|}^1 (I - C_{S'_t} \circ \mathcal{L}_h^a) \left(e_h^{(0)} \right) \right\|_2^2 \\ &\stackrel{(a)}{\leq} \left(\prod_{t=|S|+1}^{|S'|} \rho_{S'_t}^2 \right) \left\| \circ_{t=|S|}^1 (I - C_{S_t} \circ \mathcal{L}_h^a) \left(e_h^{(0)} \right) \right\|_2^2 \\ &\leq h(S) \end{aligned}$$

624 (a) by Property 1

625 **Weak supermodularity:** We will upper bound $\alpha(O)$ by providing an upper bound for $h(S) -$
626 $h(S \oplus O)$ and a lower bound for $\sum_{i=1}^T h(S) - h(S \oplus O_i)$.

$$\begin{aligned}
h(S) - h(S \oplus O) &= \left\| \circ_{t=|S|}^1 (I - C_{S_t} \circ \mathcal{L}_h^a) \left(e_h^{(0)} \right) \right\|_2^2 - \left\| \circ_{t=|O|}^1 (I - C_{O_t} \circ \mathcal{L}_h^a) \circ \circ_{t=|S|}^1 (I - C_{S_t} \circ \mathcal{L}_h^a) \left(e_h^{(0)} \right) \right\|_2^2 \\
&\stackrel{(a)}{\leq} \left\| \circ_{t=|S|}^1 (I - C_{S_t} \circ \mathcal{L}_h^a) \left(e_h^{(0)} \right) - \circ_{t=|O|}^1 (I - C_{O_t} \circ \mathcal{L}_h^a) \circ \circ_{t=|S|}^1 (I - C_{S_t} \circ \mathcal{L}_h^a) \left(e_h^{(0)} \right) \right\|_2^2 \\
&= \left\| \left(I - \circ_{t=|O|}^1 (I - C_{O_t} \circ \mathcal{L}_h^a) \right) \circ \circ_{t=|S|}^1 (I - C_{S_t} \circ \mathcal{L}_h^a) \left(e_h^{(0)} \right) \right\|_2^2 \\
&\stackrel{(b)}{\leq} \left(1 + \prod_{t=1}^{|O|} \rho_{O_t} \right)^2 \left\| \circ_{t=|S|}^1 (I - C_{S_t} \circ \mathcal{L}_h^a) \left(e_h^{(0)} \right) \right\|_2^2 \\
&\stackrel{(c)}{\leq} 4 \left\| \circ_{t=|S|}^1 (I - C_{S_t} \circ \mathcal{L}_h^a) \left(e_h^{(0)} \right) \right\|_2^2
\end{aligned}$$

627 (a) by reverse triangle property and $h(S) - h(S \oplus S') > 0$ by prefix monotonicity, (b) since the
628 Lipschitz constant of $I - \circ_{t=|S'|}^1 (I - C_{S'_t} \circ \mathcal{L}_h^a)$ is $1 + \prod_{t=1}^{|S'|} \rho_{S'_t}$ by Property 2 and 3, (c) since
629 $\rho_j < 1$

630 To lower bound $\sum_{i=1}^{|O|} h(S) - h(S \oplus O_i)$

$$\begin{aligned}
\sum_{i=1}^{|O|} h(S) - h(S \oplus O_i) &= \sum_{i=1}^{|O|} \left\| \circ_{t=|S|}^1 (I - C_{S_t} \circ \mathcal{L}_h^a) \left(e_h^{(0)} \right) \right\|_2^2 - \left\| (I - C_{O_i} \circ \mathcal{L}_h^a) \circ \circ_{t=|S|}^1 (I - C_{S_t} \circ \mathcal{L}_h^a) \left(e_h^{(0)} \right) \right\|_2^2 \\
&\geq \sum_{i=1}^{|O|} \left\| \circ_{t=|S|}^1 (I - C_{S_t} \circ \mathcal{L}_h^a) \left(e_h^{(0)} \right) \right\|_2^2 - \rho_{O_i}^2 \left\| \circ_{t=|S|}^1 (I - C_{S_t} \circ \mathcal{L}_h^a) \left(e_h^{(0)} \right) \right\|_2^2 \\
&= \left\| \circ_{t=|S|}^1 (I_N - C_{S_t} \circ \mathcal{L}_h^a) \left(e_h^{(0)} \right) \right\|_2^2 \left(T - \sum_{i=1}^T \rho_{O_i}^2 \right)
\end{aligned}$$

631 Finally,

$$\begin{aligned}
\frac{h(S) - h(S \oplus O)}{T \max_i h(S) - h(S \oplus O_i)} &\leq \max \left\{ \frac{4 \left\| \circ_{t=|S|}^1 (I - C_{S_t} \circ \mathcal{L}_h^a) \left(e_h^{(0)} \right) \right\|_2^2}{\left\| \circ_{t=|S|}^1 (I_N - C_{S_t} \circ \mathcal{L}_h^a) \left(e_h^{(0)} \right) \right\|_2^2 \left(T - \sum_{i=1}^T \rho_{O_i}^2 \right)}, 1 \right\} \\
&= \max \left\{ \frac{4}{T - \sum_{i=1}^T \rho_{O_i}^2}, 1 \right\}
\end{aligned}$$

632 **Postfix monotonicity:** Let $S' = S \oplus N$.

$$\begin{aligned}
h(S') &= \left\| \circ_{t=|S'|}^1 (I - C_{S'_t} \circ \mathcal{L}_h) e_h^{(0)} \right\|_2^2 \\
&= \left\| \circ_{t=|N|}^1 (I - C_{N_t} \circ \mathcal{L}_h) \circ \circ_{t=|S|}^1 (I - C_{S_t} \circ \mathcal{L}_h) e_h^{(0)} \right\|_2^2 \\
&\stackrel{(a)}{=} \left\| \circ_{t=|N|}^1 (I - C_{N_t} \circ \mathcal{L}_h) \circ \circ_{t=|S|}^1 (I - C_{S_t} \circ \mathcal{L}_h) \circ \left(\circ_{t=|N|}^1 (I - C_{N_t} \circ \mathcal{L}_h) \right)^{-1} \circ \circ_{t=|N|}^1 (I - C_{N_t} \circ \mathcal{L}_h) e_h^{(0)} \right\|_2^2 \\
&\leq \prod_{t=1}^{|N|} \rho_{N_t}^2 \prod_{t=1}^{|S|} \rho_{S_t}^2 \prod_{t=1}^{|N|} \rho_{N_t}^{-2} \left\| \circ_{t=|N|}^1 (I - C_{N_t} \circ \mathcal{L}_h) e_h^{(0)} \right\|_2^2 \\
&\leq h(N)
\end{aligned}$$

633 (a) due the invertibility of $I_N - C_j \mathcal{L}_h$

634

□

635 **A.4 Proof of Theorem A.2**

636 **Lemma A.2.** Let $(I - C_j \mathcal{L}_h^a) = P \Lambda_j P^{-1}$ where P is an orthogonal matrix and $\Lambda_j =$
 637 $\text{diag}(\lambda_{j1}, \dots, \lambda_{jN})$. If $P^{-1} e_h^{(0)} = z$, then the following equality holds:

$$h(S) = \sum_{i=1}^N z_i^2 \prod_{j=1}^K \lambda_{ji}^{2m_j(S)} \quad (11)$$

638 where $m_j(S) = \sum_{t=1}^{|S|} \mathbf{1}_{S_t=j}$

Proof.

$$\begin{aligned} h(S) &= \left\| \prod_{t=|S|}^1 (I_N - C_{S_t} \mathcal{L}_h^a) e_h^{(0)} \right\|^2 \\ &= \left\| \prod_{t=|S|}^1 (P \Lambda_{S_t} P^{-1}) e_h^{(0)} \right\|^2 \\ &= \left\| P \prod_{t=|S|}^1 \Lambda_{S_t} P^{-1} e_h^{(0)} \right\|^2 \\ &= \left\| \prod_{t=|S|}^1 \Lambda_{S_t} P^{-1} e_h^{(0)} \right\|^2 \\ &= \sum_{i=1}^N z_i^2 \prod_{t=|S|}^1 \lambda_{S_t i}^2 \\ &= \sum_{i=1}^N z_i^2 \prod_{j=1}^K \lambda_{ji}^{2m_j(S)} \end{aligned}$$

639 □

640 **A.5 Proof of Theorem 4.3**

641 **Proposition 4.3.** Let $\|I_N - C_j \mathcal{L}_h^a\| \leq 1$ for all $j \in [K]$ and $(I - C_j \mathcal{L}_h^a) = P \Lambda_j P^{-1}$. Then, h is
 642 supermodular.

643 *Proof.* Let $S \preceq S'$ where $S' = S \oplus B$. By Theorem A.2,

$$h(S) = \sum_{i=1}^N z_i^2 \prod_{j=1}^K \lambda_{ji}^{2m_j(S)}$$

644 where $m_j(S) = \sum_{t=1}^{|S|} \mathbf{1}_{S_t=j}$ be the number of times a sequence S calls the solver j . Recall that h is
 645 considered sequence supermodular if $\forall S', S \in \Omega^*$ such that $S \preceq S'$, it holds that

$$h(S) - h(S \oplus \omega) \geq h(S') - h(S' \oplus \omega)$$

646

$$\begin{aligned} h(S) - h(S \oplus \omega) &= \sum_{i=1}^N z_i^2 \prod_{j=1}^K \lambda_{ji}^{2m_j(S)} - \sum_{i=1}^N z_i^2 \lambda_{\omega i}^2 \prod_{k=1}^K \lambda_{ki}^{2m_k(S)} \\ &= \sum_{i=1}^N (1 - \lambda_{\omega i}^2) z_i^2 \prod_{j=1}^K \lambda_{ji}^{2m_j(S)} \end{aligned}$$

647 Similarly,

$$\begin{aligned}
h(S') - h(S \oplus \omega) &= \sum_{i=1}^N (1 - \lambda_{\omega i}^2) z_i^2 \prod_{j=1}^K \lambda_{ji}^{2m_j(S')} \\
&\stackrel{(a)}{=} \sum_{i=1}^N (1 - \lambda_{\omega i}^2) z_i^2 \prod_{j=1}^K \lambda_{ji}^{2(m_j(S) + m_j(B))} \\
&\stackrel{(b)}{\leq} \sum_{i=1}^N (1 - \lambda_{\omega i}^2) z_i^2 \prod_{j=1}^K \lambda_{ji}^{2m_j(S)} \\
&= h(S) - h(S \oplus \omega)
\end{aligned}$$

648 (a) Since $m_j(S') = \sum_{t=1}^{|S'|} \mathbf{1}_{S'_t=j} = \sum_{t=1}^{|S|} \mathbf{1}_{S'_t=j} + \sum_{t=|S|+1}^{|S'|} \mathbf{1}_{S'_t=j} = \sum_{t=1}^{|S|} \mathbf{1}_{S_t=k} + \sum_{t=1}^{|B|} \mathbf{1}_{B_t=j} =$
649 $m_j(S) + m_j(B)$, (b) since $\rho(I_N - C_j \mathcal{L}_h^a) < 1$ and $m_j(B) \geq 0$ for all $j \in [K]$

650 □

651 B Proofs for Section 5

652 B.1 Proof of Theorem B.1

653 **Lemma B.1.** *For any set of preconditioning functions $\mathcal{C}$, any discrete operator $\mathcal{L}_h^a$, any router r , any*
654 *$a_h, f_h, u_h^{(t)}, u_h \in \mathcal{A} \times \mathcal{F} \times \mathcal{U} \times \mathcal{U}$, the following equality holds true:*

$$\begin{aligned}
l_{\text{route}} \left(r, a_h, f_h, u_h^{(t)}, u_h \right) &= \sum_{j=1}^K \sum_{k=1}^K \left\| (I - C_k \circ \mathcal{L}_h^a) \left(u_h - u_h^{(t)} \right) \right\|_2^2 \mathbf{1}_{k \neq j} \mathbf{1}_{r(a_h, f_h, u_h^{(t)}) \neq j} \\
&\quad - (K-2) \sum_{j=1}^K \left\| (I - C_j \circ \mathcal{L}_h^a) \left(u_h - u_h^{(t)} \right) \right\|_2^2
\end{aligned}$$

655 *Proof.* Note that $\sum_{j=1}^K \mathbf{1}_{r(a_h, f_h, u_h^{(t)}) \neq j} = K - 1$

$$\begin{aligned}
l_{\text{route}} \left(r, a_h, f_h, u_h^{(t)}, u_h \right) &= \sum_{j=1}^K \left\| (I - C_j \circ \mathcal{L}_h^a) \left(u_h - u_h^{(t)} \right) \right\|_2^2 \mathbf{1}_{r(a_h, f_h, u_h^{(t)}) = j} \\
&= \sum_{j=1}^K \left\| (I - C_j \circ \mathcal{L}_h^a) \left(u_h - u_h^{(t)} \right) \right\|_2^2 - \sum_{j=1}^K \left\| (I - C_j \circ \mathcal{L}_h^a) \left(u_h - u_h^{(t)} \right) \right\|_2^2 \mathbf{1}_{r(a_h, f_h, u_h^{(t)}) \neq j} \\
&= \sum_{j=1}^K \left\| (I - C_j \circ \mathcal{L}_h^a) \left(u_h - u_h^{(t)} \right) \right\|_2^2 - \sum_{j=1}^K \left\| (I - C_j \circ \mathcal{L}_h^a) \left(u_h - u_h^{(t)} \right) \right\|_2^2 \mathbf{1}_{r(a_h, f_h, u_h^{(t)}) \neq j} \\
&\quad + (K-1) \sum_{k=1}^K \left\| (I - C_k \circ \mathcal{L}_h^a) \left(u_h - u_h^{(t)} \right) \right\|_2^2 - (K-1) \sum_{k=1}^K \left\| (I - C_k \circ \mathcal{L}_h^a) \left(u_h - u_h^{(t)} \right) \right\|_2^2 \\
&= \sum_{j=1}^K \left\| (I - C_j \circ \mathcal{L}_h^a) \left(u_h - u_h^{(t)} \right) \right\|_2^2 - \sum_{j=1}^K \left\| (I - C_j \circ \mathcal{L}_h^a) \left(u_h - u_h^{(t)} \right) \right\|_2^2 \mathbf{1}_{r(a_h, f_h, u_h^{(t)}) \neq j} \\
&\quad + \sum_{j=1}^K \mathbf{1}_{r(a_h, f_h, u_h^{(t)}) \neq j} \sum_{k=1}^K \left\| (I - C_k \circ \mathcal{L}_h^a) \left(u_h - u_h^{(t)} \right) \right\|_2^2 \\
&\quad - (K-1) \sum_{k=1}^K \left\| (I - C_k \circ \mathcal{L}_h^a) \left(u_h - u_h^{(t)} \right) \right\|_2^2
\end{aligned}$$

656

$$\begin{aligned}
&= \sum_{j=1}^K \left(\sum_{k=1}^K \left\| (I - C_k \circ \mathcal{L}_h^a) (u_h - u_h^{(t)}) \right\|_2^2 - \left\| (I - C_j \circ \mathcal{L}_h^a) (u_h - u_h^{(t)}) \right\|_2^2 \right) \mathbf{1}_{r(a_h, f_h, u_h^{(t)}) \neq j} \\
&\quad - (K-2) \sum_{k=1}^K \left\| (I - C_k \circ \mathcal{L}_h^a) (u_h - u_h^{(t)}) \right\|_2^2 \\
&= \sum_{j=1}^K \sum_{k=1}^K \left\| (I - C_k \circ \mathcal{L}_h^a) (u_h - u_h^{(t)}) \right\|_2^2 \mathbf{1}_{k \neq j} \mathbf{1}_{r(a_h, f_h, u_h^{(t)}) \neq j} \\
&\quad - (K-2) \sum_{k=1}^K \left\| (I - C_k \circ \mathcal{L}_h^a) (u_h - u_h^{(t)}) \right\|_2^2
\end{aligned}$$

657

□

658 B.2 Proof of Theorem B.2

659 **Proposition B.2.** For any router r defined by $r(a, f, u^{(t)}) = \operatorname{argmax}_{j \in [K]} g_j(a, f, u^{(t)})$, any $a_h \in \mathcal{A}$,
660 $f_h \in \mathcal{F}$, and $u_h^{(t)}, u_h \in \mathcal{U}$, the routing loss l_{route} satisfies:

$$\log(2) l_{\text{route}}(r, a_h, f_h, u_h^{(t)}, u_h) \leq \Psi(\mathbf{g}, a_h, f_h, u_h^{(t)}, u_h)$$

661 *Proof.* By Theorem B.1, we know that

$$\begin{aligned}
\log(2) l_{\text{route}}(r, a_h, f_h, u_h^{(t)}, u_h) &= \log(2) \sum_{j=1}^K \sum_{k=1}^K \left\| (I - C_k \circ \mathcal{L}_h^a) (u_h - u_h^{(t)}) \right\|_2^2 \mathbf{1}_{k \neq j} \mathbf{1}_{r(a_h, f_h, u_h^{(t)}) \neq j} \\
&\quad - \log(2) (K-2) \sum_{j=1}^K \left\| (I - C_j \circ \mathcal{L}_h^a) (u_h - u_h^{(t)}) \right\|_2^2 \\
&\leq \log(2) \sum_{j=1}^K \sum_{k=1}^K \left\| (I - C_k \circ \mathcal{L}_h^a) (u_h - u_h^{(t)}) \right\|_2^2 \mathbf{1}_{k \neq j} \mathbf{1}_{r(a_h, f_h, u_h^{(t)}) \neq j} \\
&\stackrel{(a)}{\leq} - \sum_{j=1}^K \sum_{k=1}^K \left\| (I - C_k \circ \mathcal{L}_h^a) (u_h - u_h^{(t)}) \right\|_2^2 \mathbf{1}_{k \neq j} \log \left(\frac{\exp(g_j(a, f, u^{(t)}))}{\sum_{k=1}^K \exp(g_k(a, f, u^{(t)}))} \right) \\
&= \Psi(\mathbf{g}, a_h, f_h, u_h^{(t)}, u_h)
\end{aligned}$$

662 (a) if $r(a_h, f_h, u_h^{(t)}) \neq j$, $\frac{\exp(g_j(a, f, u^{(t)}))}{\sum_{k=1}^K \exp(g_k(a, f, u^{(t)}))} < 0.5$ which implies that

$$663 -\log \left(\frac{\exp(g_j(a, f, u^{(t)}))}{\sum_{k=1}^K \exp(g_k(a, f, u^{(t)}))} \right) \geq \log(2) \mathbf{1}_{r(a_h, f_h, u_h^{(t)}) \neq j}$$

664

□

665 B.3 Proof of Theorem 5.1

666 **Theorem 5.1.** Let $\tilde{c}_j(a_h, u_h^{(t)}, u_h) < \bar{E} < \infty$ for all $j \in [K]$. If there exists $j \in [K]$ such that
667 $\tilde{c}_j(a_h, u_h^{(t)}, u_h) > E_{\min} > 0$, then, for any collection of solvers $\{C_j\}_{j=1}^K$ and linear discrete
668 operator $\mathcal{L}_h^a$, Ψ is Bayes consistent surrogate for l_{route} .

669 *Proof.* For a given a_h, f_h , let u_h be $\mathcal{G}_h(a_h, f_h)$ where $\mathcal{G}_h$ denotes the solution operator acting on the
670 grid G_h . Furthermore, let's consider routers of the form

$$r(a, f, u^{(t)}) = \operatorname{argmax}_{j \in [K]} g_j(a, f, u^{(t)})$$

671 For a given $a_h, f_h, u_h^{(t)} \in \mathcal{A} \times \mathcal{F} \times \mathcal{U}$, let the optimal loss under l_{route} be
672 $l_{\text{route}}^* (a_h, f_h, u_h^{(t)}) = \inf_{\tilde{r}} l_{\text{route}} (\tilde{r}, a_h, f_h, u_h^{(t)}, \mathcal{G}_h(a_h, f_h))$. Similarly, let the optimal loss un-
673 der Ψ be $\Psi^* (a_h, f_h, u_h^{(t)}) = \inf_{\tilde{\mathbf{g}}} \Psi (\tilde{\mathbf{g}}, a_h, f_h, u_h^{(t)}, \mathcal{G}_h(a_h, f_h))$. Let $B_j(a_h, f_h, u_h^{(t)}) =$
674 $\sum_{k=1}^K \left\| (I - C_k \circ \mathcal{L}_h^a) (\mathcal{G}_h(a_h, f_h) - u_h^{(t)}) \right\|_2^2 \mathbf{1}_{k \neq j}$.

$$\begin{aligned}
& l_{\text{route}} (r, a_h, f_h, u_h^{(t)}, \mathcal{G}_h(a_h, f_h)) - l_{\text{route}}^* (a_h, f_h, u_h^{(t)}) \\
& \stackrel{(a)}{=} \sum_{j=1}^K \sum_{k=1}^K \left\| (I - C_k \circ \mathcal{L}_h^a) (u_h - u_h^{(t)}) \right\|_2^2 \mathbf{1}_{k \neq j} \mathbf{1}_{r(a_h, f_h, u_h^{(t)}) \neq j} - (K-2) \sum_{j=1}^K \left\| (I - C_j \circ \mathcal{L}_h^a) (u_h - u_h^{(t)}) \right\|_2^2 \\
& \quad - \inf_{\tilde{r}} \sum_{j=1}^K \sum_{k=1}^K \left\| (I - C_k \circ \mathcal{L}_h^a) (u_h - u_h^{(t)}) \right\|_2^2 \mathbf{1}_{k \neq j} \mathbf{1}_{\tilde{r}(a_h, f_h, u_h^{(t)}) \neq j} + (K-2) \sum_{j=1}^K \left\| (I - C_j \circ \mathcal{L}_h^a) (u_h - u_h^{(t)}) \right\|_2^2 \\
& = \sum_{j=1}^K B_j(a_h, f_h, u_h^{(t)}) \mathbf{1}_{r(a_h, f_h, u_h^{(t)}) \neq j} - \inf_{\tilde{r}} \sum_{j=1}^K B_j(a_h, f_h, u_h^{(t)}) \mathbf{1}_{\tilde{r}(a_h, f_h, u_h^{(t)}) \neq j} \\
& = \sum_{k=1}^K B_k(a_h, f_h, u_h^{(t)}) \left(\sum_{j=1}^K \frac{B_j(a_h, f_h, u_h^{(t)})}{\sum_{k=1}^K B_k(a_h, f_h, u_h^{(t)})} \mathbf{1}_{r(a_h, f_h, u_h^{(t)}) \neq j} - \inf_{\tilde{r}} \sum_{j=1}^K \frac{B_j(a_h, f_h, u_h^{(t)})}{\sum_{k=1}^K B_k(a_h, f_h, u_h^{(t)})} \mathbf{1}_{\tilde{r}(a_h, f_h, u_h^{(t)}) \neq j} \right)
\end{aligned}$$

675 (a) by Theorem B.1

676 Let $\mathcal{X} = \mathcal{A} \times \mathcal{F} \times \mathcal{U}$ and $\mathcal{Y} = [K]$. Let $\mathcal{P}_{\mathcal{X}}$ denote the degenerate distribution supported at the
677 point $(a_h, f_h, u_h^{(t)})$. We define the conditional distribution - $P(Y = j \mid X = (a_h, f_h, u_h^{(t)})) =$
678 $\frac{B_j(a_h, f_h, u_h^{(t)})}{\sum_{k=1}^K B_k(a_h, f_h, u_h^{(t)})}$ for $j \in [K]$. The risk and optimal risk of 0 - 1 loss under this distribution can
679 be written as:

$$\begin{aligned}
\mathcal{R}_{0-1}(r) &= \sum_{j=1}^K \frac{B_j(a_h, f_h, u_h^{(t)})}{\sum_{k=1}^K B_k(a_h, f_h, u_h^{(t)})} \mathbf{1}_{r(a_h, f_h, u_h^{(t)}) \neq j} \\
\mathcal{R}_{0-1}^* &= \inf_r \sum_{j=1}^K \frac{B_j(a_h, f_h, u_h^{(t)})}{\sum_{k=1}^K B_k(a_h, f_h, u_h^{(t)})} \mathbf{1}_{r(a_h, f_h, u_h^{(t)}) \neq j}
\end{aligned}$$

680 If $r(a_h, f_h, u_h^{(t)}) = \arg \max_{j \in [K]} g_j(a_h, f_h, u_h^{(t)})$ for all $x \in \mathcal{X}$, then the he risk and optimal risk of
681 cross entropy loss $(l_{ce}(\mathbf{g}, x, y) - \log \left(\frac{\exp(g_y(x))}{\sum_{k=1}^K \exp(g_k(x))} \right))$ under this distribution can be written as:

$$\begin{aligned}
\mathcal{R}_{ce}(\mathbf{g}) &= - \sum_{j=1}^K \frac{B_j(a_h, f_h, u_h^{(t)})}{\sum_{k=1}^K B_k(a_h, f_h, u_h^{(t)})} \log \left(\frac{\exp(g_j(a_h, f_h, u_h^{(t)}))}{\sum_{k=1}^K \exp(g_k(a_h, f_h, u_h^{(t)}))} \right) \\
\mathcal{R}_{ce}^* &= \inf_{\mathbf{g}} - \sum_{j=1}^K \frac{B_j(a_h, f_h, u_h^{(t)})}{\sum_{k=1}^K B_k(a_h, f_h, u_h^{(t)})} \log \left(\frac{\exp(g_j(a_h, f_h, u_h^{(t)}))}{\sum_{k=1}^K \exp(g_k(a_h, f_h, u_h^{(t)}))} \right)
\end{aligned}$$

682 From Theorem 3.1 of [Mao et al., 2023], $\mathcal{R}_{0-1}(r) - \mathcal{R}_{0-1}^* \leq \Gamma^{-1}(\mathcal{R}_{ce}(\mathbf{g}) - \mathcal{R}_{ce}^*)$ if
683 $r(a_h, f_h, u_h^{(t)}) = \arg \max_{j \in [K]} g_j(a_h, f_h, u_h^{(t)})$ where $\Gamma(z) = \frac{1+z}{2} \log(1+z) + \frac{1-z}{2} \log(1-z)$.
684 Then,

$$\begin{aligned}
& l_{\text{route}} \left(r, a_h, f_h, u_h^{(t)}, \mathcal{G}_h(a_h, f_h) \right) - l_{\text{route}}^* \left(a_h, f_h, u_h^{(t)} \right) \\
&= \sum_{k=1}^K B_k(a_h, f_h, u_h^{(t)}) \left(\sum_{j=1}^K \frac{B_j(a_h, f_h, u_h^{(t)})}{\sum_{k=1}^K B_k(a_h, f_h, u_h^{(t)})} \mathbf{1}_{r(a_h, f_h, u_h^{(t)}) \neq j} - \inf_{\tilde{r}} \sum_{j=1}^K \frac{B_j(a_h, f_h, u_h^{(t)})}{\sum_{k=1}^K B_k(a_h, f_h, u_h^{(t)})} \mathbf{1}_{\tilde{r}(a_h, f_h, u_h^{(t)}) \neq j} \right) \\
&\leq \sum_{k=1}^K B_k(a_h, f_h, u_h^{(t)}) \Gamma^{-1} \left(- \sum_{j=1}^K \frac{B_j(a_h, f_h, u_h^{(t)})}{\sum_{k=1}^K B_k(a_h, f_h, u_h^{(t)})} \log \left(\frac{\exp(g_j(a_h, f_h, u_h^{(t)}))}{\sum_{k=1}^K \exp(g_k(a_h, f_h, u_h^{(t)}))} \right) \right. \\
&\quad \left. - \inf_{\mathbf{g}} - \sum_{j=1}^K \frac{B_j(a_h, f_h, u_h^{(t)})}{\sum_{k=1}^K B_k(a_h, f_h, u_h^{(t)})} \log \left(\frac{\exp(g_j(a_h, f_h, u_h^{(t)}))}{\sum_{k=1}^K \exp(g_k(a_h, f_h, u_h^{(t)}))} \right) \right) \\
&\stackrel{(a)}{\leq} \bar{E} K (K-1) \Gamma^{-1} \left(- \sum_{j=1}^K \frac{B_j(a_h, f_h, u_h^{(t)})}{\sum_{k=1}^K B_k(a_h, f_h, u_h^{(t)})} \log \left(\frac{\exp(g_j(a_h, f_h, u_h^{(t)}))}{\sum_{k=1}^K \exp(g_k(a_h, f_h, u_h^{(t)}))} \right) \right. \\
&\quad \left. - \inf_{\mathbf{g}} - \sum_{j=1}^K \frac{B_j(a_h, f_h, u_h^{(t)})}{\sum_{k=1}^K B_k(a_h, f_h, u_h^{(t)})} \log \left(\frac{\exp(g_j(a_h, f_h, u_h^{(t)}))}{\sum_{k=1}^K \exp(g_k(a_h, f_h, u_h^{(t)}))} \right) \right) \\
&\stackrel{(b)}{\leq} \bar{E} K (K-1) \Gamma^{-1} \left(- \sum_{j=1}^K \frac{B_j(a_h, f_h, u_h^{(t)})}{(K-1)E_{\min}} \log \left(\frac{\exp(g_j(a_h, f_h, u_h^{(t)}))}{\sum_{k=1}^K \exp(g_k(a_h, f_h, u_h^{(t)}))} \right) \right. \\
&\quad \left. - \inf_{\mathbf{g}} - \sum_{j=1}^K \frac{B_j(a_h, f_h, u_h^{(t)})}{(K-1)E_{\min}} \log \left(\frac{\exp(g_j(a_h, f_h, u_h^{(t)}))}{\sum_{k=1}^K \exp(g_k(a_h, f_h, u_h^{(t)}))} \right) \right) \\
&= \bar{E} K (K-1) \Gamma^{-1} \left(\frac{\Psi(\mathbf{g}, a_h, f_h, u_h^{(t)}, \mathcal{G}_h(a_h, f_h)) - \Psi^*(a_h, f_h, u_h^{(t)})}{(K-1)E_{\min}} \right)
\end{aligned}$$

685 (a) since $\left\| (I - C_j \circ \mathcal{L}_h^a)(e_h^{(t)}) \right\|_2^2 < \bar{E}$ for all $j \in [K]$, (b) since Γ^{-1} is non-decreasing and

686 $\exists j \in [K]$ such that $\left\| (I - C_j \circ \mathcal{L}_h^a)(e_h^{(t)}) \right\|_2^2 > E_{\min}$

687 Finally,

$$\begin{aligned}
& \lim_{n \rightarrow \infty} \mathcal{R}_{\text{route}}(r_n) - \mathcal{R}_{\text{route}}^* \\
&\stackrel{(a)}{=} \lim_{n \rightarrow \infty} \mathbb{E}_{a_h, f_h \sim \mathcal{P}_{\mathcal{A} \times \mathcal{F}}} \left[l_{\text{route}} \left(r_n, a_h, f_h, u_h^{(t)}, \mathcal{G}_h(a_h, f_h) \right) - l_{\text{route}}^* \left(a_h, f_h, u_h^{(t)} \right) \right] \\
&\leq \lim_{n \rightarrow \infty} \mathbb{E}_{a_h, f_h \sim \mathcal{P}_{\mathcal{A} \times \mathcal{F}}} \left[\bar{E} K (K-1) \Gamma^{-1} \left(\frac{\Psi(\tilde{\mathbf{g}}, a_h, f_h, u_h^{(t)}, \mathcal{G}_h(a_h, f_h)) - \Psi^*(a_h, f_h, u_h^{(t)})}{(K-1)E_{\min}} \right) \right] \\
&\stackrel{(b)}{\leq} \lim_{n \rightarrow \infty} \bar{E} K (K-1) \Gamma^{-1} \left(\frac{\mathbb{E}_{a_h, f_h \sim \mathcal{P}_{\mathcal{A} \times \mathcal{F}}} \left[\Psi(\mathbf{g}_n, a_h, f_h, u_h^{(t)}, \mathcal{G}_h(a_h, f_h)) - \Psi^*(a_h, f_h, u_h^{(t)}) \right]}{(K-1)E_{\min}} \right)
\end{aligned}$$

$$\begin{aligned}
&= \lim_{n \rightarrow \infty} \bar{E} K (K-1) \Gamma^{-1} \left(\frac{\mathcal{R}_\Psi(\mathbf{g}_n) - \mathcal{R}_\Psi^*}{(K-1)E_{min}} \right) \\
&\stackrel{(c)}{=} \bar{E} K (K-1) \Gamma^{-1} \left(\frac{\lim_{n \rightarrow \infty} \mathcal{R}_\Psi(\mathbf{g}_n) - \mathcal{R}_\Psi^*}{(K-1)E_{min}} \right) \\
&= \bar{E} K (K-1) \Gamma^{-1}(0) \\
&\stackrel{(d)}{=} 0
\end{aligned}$$

689 (a) $\mathcal{R}_{\text{route}}^* = \mathbb{E}_{a_h, f_h \sim \mathcal{P}_{\mathcal{A} \times \mathcal{F}}} \left[l_{\text{route}}^* \left(a_h, f_h, u_h^{(t)} \right) \right]$ since the infimum is taken over all measurable
690 functions, (b) by Jensen's inequality since Γ^{-1} is concave, (c) by continuity of Γ^{-1} at 0, (d) $\Gamma^{-1}(0) =$
691 0

692

□

693 C Training Details

694 The periodic experiments of Section 7 use the corrector and the lightweight router described in
695 Section E.1; the data distribution below is shared by all experiments, and the DeepONet and LSTM
696 router described afterwards are those of the Dirichlet experiment in Section D.1.

697 Data for both DeepONet and the routers is sampled from a zero-mean Hierarchical Gaussian Random
698 Field on the periodic domain with covariance operator $\alpha(-\Delta + \beta I)^{-\gamma}$, where α, β, γ are sampled
699 as:

$$\begin{aligned}
\alpha &\sim \text{Log-Uniform}(0.01, 100) \\
\beta &\sim \text{Log-Uniform}(0.1, 1000) \\
\gamma &\sim \text{Uniform}(\{0.5, 1.0, 1.5, 2.0, 2.5, 3.0, 4.0\})
\end{aligned}$$

700 Given (α, β, γ) , samples are generated in the Fourier space. For each non-zero frequency mode
701 k , we draw an independent complex coefficient from a Gaussian distribution with mean 0 and
702 variance $\alpha(4\pi^2\|k\|_2^2 + \beta)^{-\gamma}$. We enforce a Hermitian symmetry to obtain a real-valued field, set the
703 zero-frequency (DC) mode to 0 to obtain a zero mean field, and apply the inverse Discrete Fourier
704 Transform to obtain the field in physical space.

705 For each sample, we compute reference solutions with a least squares solver and treat them as ground
706 truth.

707 This data is used to trained our DeepONet models and LSTM routers. All the models were imple-
708 mented using PyTorch and all the models were trained on one Nvidia A40 GPU.

709 Table 2 contains all hyperparameter details for the DeepONet. DeepONet took 30 minutes to train.
710 We then use the model with the best validation loss.

Table 2: Hyperparameter settings for DeepONet

Hyperparameter	Value
Learning rate	1e-3
Branch Dimension	128
Hidden dimension for branch net	256
No. of hidden layers in branch net	4
Hidden dimension for trunk net	256
No. of hidden layers in trunk net	4
Gradient Clipping Norm	1.0
Weight Decay	0.005
Batch size	256
Training samples	10000
Validation samples	2000
Epochs	1000

711 The routers are LSTM models. Along with the inputs specified in Section 5, namely $a_h, f_h, u_h^{(t)}$,
712 we also supply the routers with the iteration index t , current residual $f_h - \mathcal{L}_h^a u_h^{(t)}$, and the routing
713 decisions from the previous iteration. While the iteration index and residual are deterministic
714 functions of $(a_h, f_h, u_h^{(t)})$ and align with the theory presented in Section 5, incorporating previous
715 routing decisions is a mild deviation from the idealized setting. Nevertheless, this modification is
716 motivated by the same considerations that justify the use of scheduled sampling: under deployment,
717 routing decisions are predicted autoregressively and influence future inputs, whereas training under
718 teacher forcing assumes independence from past predictions. Including previous routing decisions
719 bridges the gap induced by exposure bias by providing the model with trajectory-level information.

720 All the routers are trained with scheduled sampling. We use a warm-up of e_w epochs with teacher-
721 forcing probability $p_{tf}(e) = ss_{start}$. After the warm-up, the p_{tf} decays geometrically by a factor of
722 $\gamma_{tf} < 1$ per epoch and is floored by ss_{end} :

$$p_{tf}(e) = \begin{cases} ss_{start} & e \leq e_w \\ \max(ss_{start}\gamma_{tf}^{e-e_w}, ss_{end}) & e > e_w \end{cases}$$

723 At each time step, with probability $p_{tf}(e)$, we feed the teacher-forced greedy iterate; otherwise, we
724 feed the router’s own predicted iterate.

725 Since LSTMs on long rollouts can suffer from exploding/vanishing gradients, we use truncated
726 backpropagation through time (TBPTT) [Mozier, 2013, Robinson and Fallside, 1987, Werbos, 1988]:
727 the forward pass unrolls the entire trajectory, but gradients are propagated only through the most
728 recent $w_{bptt}(e)$ steps at epoch e . Hidden states are passed forward between segments, while earlier
729 segments are treated as stop-gradient.

730 We employ a curriculum learning approach analogous to scheduled sampling. Let T_{max} be the
731 horizon (300 iterations). With a warm-up of e_w epochs,

$$w_{bptt}(e) = \begin{cases} w_{start} & e \leq e_w \\ \min\left(T_{max}, w_{start}\gamma_{bptt}^{\lfloor \frac{e-e_w}{f_{bptt}} \rfloor}\right) & e > e_w \end{cases} \quad (12)$$

732 so the window grows geometrically by a factor of $\gamma_{bptt} > 1$ every f_{bptt} epochs and is capped at the
733 full trajectory length.

734 Table 3 contains all hyperparameter details for the LSTM routers. The routers for the Paired Solver
735 experiment took a maximum of 4 hours and 30 minutes to train while the routers for the Solver
736 Ensemble experiment took a maximum of 6 hours to train. We then use the model with the best
737 validation loss for testing. The architecture and dataset size is larger for the Solver ensemble
738 experiments to encourage the model to learn some of the nuanced differences between the classes.

Table 3: Hyperparameter settings for routers in Paired solver and Solver Ensemble Experiments

Hyperparameter	Paired Solver	Solver Ensemble
Learning rate	1e-3	1e-3
Hidden dimension	256	512
No. of hidden layers	4	3
Gradient Clipping Norm	1.0	1.0
Weight Decay	0.005	0.005
Batch size	32	32
Training samples	256	1024
Validation samples	32	128
Epochs	200	200
ss_{start}	1.0	1.0
γ_{tf}	0.95	0.95
ss_{end}	0.0	0.0
w_{start}	50	50
γ_{bptt}	1.25	1.25
e_w	10	10
f_{bptt}	1	1

D Additional Experimental Results

This appendix keeps the one experiment of the original submission that the wall-clock study of Section E does not supersede: the paired-solver comparison under Dirichlet boundary conditions, which was run with the original pipeline (a 31×31 grid, a DeepONet with an MLP branch and an LSTM router trained with scheduled sampling; Section C). The other results of the original submission—iteration-based comparisons on 31^2 and 63^2 periodic grids with the LSTM router, residual and Fourier-mode analyses, and ensemble decision traces—are superseded by the per-equation results, the usage analyses and the ensemble study of Section E and have been removed.

D.1 Results with Dirichlet Boundary Conditions

We replicated the paired-solver experiment of the original submission for solutions with zero Dirichlet boundary conditions, with the original pipeline; the metrics are the final error and AUC over 300 iterations, with paired t -tests against the learned router.

In Table 4, we notice that, similar to PDEs with periodic boundary conditions, the learned greedy router closely matches the performance of the true greedy oracle and outperforms its single solver and HINTS baselines with great statistical significance. Despite the learned router matching the performance of the true greedy oracle, we observe that the learned router doesn’t seem to provide much improvement over single-solver baseline in the case of SOR (1.5)-related solvers. This may be a result of the deepnet correction always being weaker relative to the SOR (1.5) correction in terms of error reduction.

Table 4: Final error and AUC of squared L^2 error (lower is better). Values are mean ($\pm$ standard error (s.e.)) over 128 test instances with dirichlet boundaries; both mean and s.e. of error are reported in $\times 10^{-3}$. If a standard error is not shown, it is $< 10^{-3}$ in the reported units (raw $< 10^{-6}$). Statistical significance is assessed via paired t -tests comparing each baseline (single-solver only and HINTS) against the learned greedy router, using a one-sided alternative that the learned router achieves lower error/AUC. Reported p -values correspond to these tests.

Equation	Poisson				ConvDiff			
Methods	$\ e_h^{(T)}\ \times 10^3$	$\ e_h^{(T)}\ $ p -value	AUC	AUC p -value	$\ e_h^{(T)}\ \times 10^3$	$\ e_h^{(T)}\ $ p -value	AUC	AUC p -value
Jacobi-related solvers								
Jacobi Only	0.704 (2.151)	0.001	0.703 (1.929)	$< 10^{-3}$	0.003 (0.006)	0.057	0.169 (0.430)	$< 10^{-3}$
HINTS-Jacobi	2.674 (7.315)	$< 10^{-3}$	0.869 (2.396)	$< 10^{-3}$	0.474 (0.001)	$< 10^{-3}$	0.199 (0.254)	$< 10^{-3}$
Learned Greedy-Jacobi	0.243 (0.851)	-	0.236 (0.719)	-	0.003 (0.006)	-	0.076 (0.162)	-
True-Greedy-Jacobi	0.042 (0.047)	-	0.099 (0.252)	-	0.003 (0.007)	-	0.042 (0.089)	-
GS-related solvers								
GS only	0.138 (0.431)	0.001	0.408 (1.139)	$< 10^{-3}$	0.003 (0.006)	0.028	0.057 (0.146)	$< 10^{-3}$
HINTS-GS	2.668 (7.316)	$< 10^{-3}$	0.841 (2.360)	$< 10^{-3}$	0.470 (0.001)	$< 10^{-3}$	0.115 (0.137)	$< 10^{-3}$
Learned Greedy-GS	0.032 (0.158)	-	0.115 (0.433)	-	0.003 (0.006)	-	0.024 (0.048)	-
True-Greedy-GS	0.010 (0.015)	-	0.063 (0.166)	-	0.003 (0.006)	-	0.017 (0.036)	-
SymGS-related solvers								
SymGS only	2.924 (8.865)	0.01	0.926 (2.740)	0.006	1.331 (3.687)	0.004	0.413 (1.133)	0.004
HINTS-SymGS	2.748 (7.375)	0.008	0.856 (2.427)	0.008	0.917 (1.015)	0.006	0.332 (0.697)	0.003
Learned Greedy-SymGS	1.333 (5.715)	-	0.416 (1.722)	-	0.628 (1.468)	-	0.196 (0.449)	-
True-Greedy-SymGS	0.342 (0.906)	-	0.147 (0.375)	-	0.184 (0.323)	-	0.067 (0.118)	-
Jacobi (0.67)-related solvers								
Jacobi (0.67) only	1.265 (3.768)	$< 10^{-3}$	0.910 (2.468)	$< 10^{-3}$	0.008 (0.021)	0.003	0.253 (0.642)	$< 10^{-3}$
HINTS-Jacobi (0.67)	2.686 (7.311)	$< 10^{-3}$	0.886 (2.412)	$< 10^{-3}$	0.487 (0.002)	$< 10^{-3}$	0.220 (0.279)	$< 10^{-3}$
Learned Greedy-Jacobi (0.67)	0.134 (0.275)	-	0.215 (0.660)	-	0.004 (0.010)	-	0.127 (0.393)	-
True-Greedy-Jacobi (0.67)	0.096 (0.220)	-	0.145 (0.408)	-	0.004 (0.013)	-	0.061 (0.131)	-
SOR (1.5)-related solvers								
SOR (1.5) only	0.004 (0.011)	0.012	0.145 (0.412)	$< 10^{-3}$	0.003 (0.006)	1.0	0.006 (0.014)	1.0
HINTS-SOR (1.5)	2.650 (7.321)	$< 10^{-3}$	0.791 (2.295)	$< 10^{-3}$	0.468 (0.001)	$< 10^{-3}$	0.014 (0.014)	$< 10^{-3}$
Learned Greedy-SOR (1.5)	0.004 (0.011)	-	0.039 (0.091)	-	0.003 (0.006)	-	0.006 (0.014)	-
True-Greedy-SOR (1.5)	0.004 (0.011)	-	0.028 (0.070)	-	0.003 (0.006)	-	0.006 (0.014)	-

E Cost-Aware Routing and Wall-Clock Evaluation

The experiments of Section 7 compare methods per iteration, which favors methods whose iterations are expensive. This appendix evaluates the end-to-end *wall-clock time to a target accuracy* of the learned router against HINTS, the classical solver alone and the classical methods a practitioner would reach for on these problems (geometric multigrid and preconditioned Krylov methods), for all five classical solvers and three PDEs, on 128^2 , 256^2 and 512^2 grids, using the cost-aware instantiation of the greedy rule introduced at the end of Section 4. Besides the Poisson and convection–diffusion equations of Section 7 we add *anisotropic diffusion*, $-\epsilon \partial_{x_1}^2 u - \partial_{x_2}^2 u = f$ with $\epsilon = 10^{-2}$ (diffusion

tensor $\text{diag}(\epsilon, 1)$), on the same periodic domain with the same forcing distribution. It is the standard model problem on which point-relaxation methods fail to smooth: an error component oscillatory in x_1 and smooth in x_2 is damped by only $(1 - \epsilon)/(1 + \epsilon) \approx 0.98$ per Jacobi sweep, and geometric multigrid with point smoothing loses its grid-independent convergence (line relaxation or semi-coarsening would be needed). Section E.1 describes the protocol; Sections E.2 to E.4 collect all results per equation (every grid, every tolerance, every baseline, with significance tests and seed trials); Section E.5 studies ensembles of one, two and three classical solvers; Section E.6 accounts for overheads and training cost; and Section E.7 verifies numerically that the assumptions of Theorem 4.1, Propositions 4.2–4.3 and Theorem 5.1 hold in these settings.

E.1 Protocol

Fast, matched implementations. A wall-clock comparison is meaningful only if the classical baseline is not handicapped by implementation. All solvers are therefore re-implemented for the periodic constant-coefficient setting as $O(N^2)$ operations: Jacobi and damped Jacobi as 5-point stencil sweeps, Gauss–Seidel, SOR and symmetric GS via cached sparse triangular factorizations, geometric multigrid as V(2,2)-cycles with lexicographic Gauss–Seidel smoothing, full-weighting restriction, bilinear prolongation, rediscrctised coarse operators and a direct solve on the 32×32 coarsest grid (error contraction ≈ 0.03 per cycle on the isotropic equations). Their one-step iterates match the dense implementations used in the rest of the paper to 10^{-12} in double precision. Reference solutions are computed by FFT diagonalization of the same discrete operator, i.e., they are exact solutions of the discrete system, so true errors are available for evaluation at no charged cost; the FFT solve is used only as this reference (it is a direct method that exists because the test problems have constant coefficients and periodic boundaries, and is not a competitor of the iterative methods compared here). All timings are single-threaded CPU measurements (Apple M4 Pro, numpy/scipy/PyTorch, one thread) on an otherwise idle machine.

Corrector. The neural corrector is a DeepONet applied to the residual $r = f_h - \mathcal{L}_h^a u_h^{(t)}$ and evaluated on a 64×64 sensor grid, i.e., on the coarser grid obtained by band-limited (FFT) restriction to the modes $|k_1|, |k_2| \leq 31$ at every grid size (for the anisotropic equation the sensor grid is $N \times N/4$, i.e., full resolution along x_1 , so that the band contains every mode the point smoothers cannot damp; the sensor grid is the only PDE-specific hyperparameter of the pipeline); its output is prolonged back to the fine grid by exact trigonometric interpolation, so the correction is band-limited by construction and never injects error into the high-frequency band that the classical smoother owns (Figure 2). The trunk net is the fixed orthonormal Fourier basis of that band ($p = 3969$ basis functions, as in POD-DeepONet [Lu et al., 2022]), and the branch net is a linear layer on the 4096 sensor values; it is fit by least squares on 32,000–64,000 exact (residual, error) pairs of the discrete operator drawn from a mixture of spectral shapes (forcing-like, residual-like, white and randomly tilted spectra), and reaches a validation relative error of $\approx 2 \times 10^{-6}$ on every mode of the band. The corrector is applied scale-equivariantly, $C_{\text{NO}}(r) = \|Rr\| \mathcal{G}_\theta(Rr/\|Rr\|)$, which preserves $C_{\text{NO}}(0) = 0$ (Proposition 4.2). Its cost is dominated by a 4096×4096 matrix–vector product and is therefore almost independent of N (0.65 ms at 128^2 , 3.6 ms at 512^2 , Table 29), i.e., about 12 Jacobi or 3 GS sweeps at 128^2 and about 2.5 Jacobi sweeps at 512^2 (Tables 8, 14 and 20). The same corrector is used by HINTS and by all routing policies.

Cost-equalised macro-actions. Per-iteration costs c_j (residual evaluation plus the update) are measured once per pairing and grid on the benchmark machine (minimum over seven repeated blocks of the block median, cached), the unit of cost u is one corrector call, and every operation is applied $m_j = \max\{1, \text{round}(u/c_j)\}$ times per decision; for the five stationary solvers this equalises costs to within 8%. An operation dearer than the unit (a multigrid cycle) is applied once and compared per unit of cost, $\|e_j\|^{u/(m_j c_j)}$; the surrogate weights $\tilde{c}_j$ use the same exponent. Ensembles reuse the cached pairwise costs of their members, so every member has the same macro-action size in the ensemble as in its pairwise run. The cost-aware oracle runs Algorithm 1 on these macro-actions with the true error; the *greedy oracle* rows run the cost-agnostic Algorithm 1 on single operations, as in the main text.

Lightweight router. The learned router selects macro-actions from $6 + K$ scalar features, all $O(1)$ given the residual norm that any stopping test computes: the log relative residual, its per-iteration change over the last macro-action and an exponential moving average thereof, the log epoch index, a

one-hot encoding of the previous macro-action, and the log counts of corrector calls so far and of iterations since the last call. It is a two-hidden-layer MLP (64 units) evaluated in numpy, costing $\approx 5 \mu\text{s}$ per decision independently of N . It is trained with the surrogate of Section 5.1, with the per-state weights $\tilde{c}_k$ normalised by $\sum_k \tilde{c}_k$ (which does not change the Bayes decision), on states visited by oracle rollouts with ε -exploration (probability 0.1), followed by two to three DAgger rounds [Ross et al., 2011] in which the current router’s own rollouts are relabelled by the oracle—the batch analogue of the scheduled sampling used in Section C. Training uses 32–128 instances per round (fewer on the larger grids) drawn from the same GRF distribution as the test set but with a different seed; the router of every pairing is trained in seconds to a few minutes. Ensemble routers, which must additionally learn to choose among classical members, use a larger imitation budget (256 rollouts at 128^2 and 128 at 256^2 , six DAgger rounds, hidden width 128, 300 epochs): with the pairwise recipe the router over $\{\text{Jacobi}, \text{Jacobi} (0.67)\}$ imitated its oracle imperfectly at tight tolerances (convection–diffusion, 128^2 , $\varepsilon = 10^{-8}$: 3.34 against the oracle’s 2.67 work-unit ms), which the larger budget removes (2.70). At 512^2 the routers are trained from 32 oracle rollouts of at most 400 iterations (stopping error 10^{-8}); for the convection–diffusion Jacobi pairing, where the first such router missed h^2 within the cap on 4 of 16 instances, the reported router uses 64 rollouts of at most 800 iterations—the only per-cell change to the recipe.

Baselines and metric. HINTS applies the corrector every τ -th iteration; we report $\tau = 25$ (the schedule used in the main text and in Zhang et al. [2024]) and the best of $\tau \in \{5, 10, 25, 50\}$ selected *per cell and per tolerance* on the test instances themselves (an optimistic baseline). The classical baselines without corrector run on the same stencil with the same accounting: multigrid V(2,2) alone; conjugate gradients, CG preconditioned by symmetric GS and by a symmetric V-cycle for Poisson; BiCGSTAB, multigrid-preconditioned BiCGSTAB and restarted GMRES(20) for convection–diffusion (GMRES is accounted at the granularity of its restart cycles and was not run at 512^2 , where it was two orders of magnitude slower than the router on the smaller grids). They are not run for anisotropic diffusion, where point-smoothed multigrid loses its grid-independent contraction (≈ 0.85 per cycle) and would not be informative. For each test instance (64 at 128^2 , 32 at 256^2 , 16 at 512^2) we record, after every iteration, the true relative error $\|e_h^{(t)}\|/\|u_h\|$ (untimed, benchmark-only) and the cumulative charged wall-clock time of a replay that executes only the chosen operations (residual, update, and—for the learned router—every feature and decision computation; oracle decisions are not charged). We report the median first time at which the error drops below a tolerance ε , paired per-instance median speedups, and one-sided Wilcoxon signed-rank tests on the paired log time ratios (alternative: the router is faster); a run that does not reach ε within the iteration cap (60,000 at 128^2 , 40,000 at 256^2 , 8,000 at 512^2) is marked †, reported as a lower bound when it is the majority, counted against the router when it is the router’s, and entered at the time it spent up to the cap on the baseline side of the tests (a lower bound, hence conservative). We highlight $\varepsilon = h^2$: the discretization error of the second-order scheme scales as $O(h^2)$, so driving the algebraic error far below h^2 does not improve the computed solution; we also report 10^{-3} and 10^{-8} , and the iteration-based metrics of the main text (final relative error and $\text{AUC} = \sum_{t=1}^T \|e_h^{(t)}\|/\|u_h\|$ over $T = 300$ iterations) on the 128^2 grid. Every router was additionally retrained from scratch with five independent seeds (different training instances, exploration and initialisation; the corrector and the test set are fixed) and re-run on the test set; because those trials ran in separate sessions, whose live timings are not comparable with the main run’s, they are compared in *work units* (measured per-iteration costs times executed operations, plus $5 \mu\text{s}$ per decision), which are determined by the decision sequence alone.




Figure 2: One-shot corrector predictions on the 128×128 grid (two test instances per equation): forcing, exact discrete solution, the DeepONet correction applied to the initial residual ($u_h^{(0)} = 0$), and the remaining error, which is confined to the modes outside the corrector's band (the relative error of the prediction is 10^{-6} – 10^{-5} on the isotropic equations and 10^{-6} on anisotropic diffusion, whose solution is dominated by modes that are smooth in x_2).




Figure 3: Relative error versus charged wall-clock time on the 128×128 grid for one representative test instance of every pairing (columns) and equation (rows); the dashed line is $\varepsilon = h^2$. The learned router and the cost-aware oracle are indistinguishable.

865 E.2 Poisson

866 Table 5 collects the median time to 10^{-3} , h^2 and 10^{-8} for every pairing, every policy and the classical
867 baselines on all three grids; Table 6 the paired speedups of the learned router with their p -values;
868 Figure 4 and table 7 where the iterations go; Table 8 the measured costs and macro-action sizes;
869 Table 9 the iteration-based metrics; and Table 10 the seed trials.

870 **Uniform gains and their mechanism.** On the 128^2 grid the learned router is the fastest deployable
871 method at $\varepsilon = h^2$ in every pairing: hundreds to thousands of times faster than the solver alone,
872 several times faster than HINTS ($\tau=25$) and still faster than the best fixed period chosen per cell and
873 per tolerance on the test set, with all paired $p < 0.01$ (Table 6); at 10^{-8} the margins over the best
874 fixed period shrink but remain significant in every pairing, and the best period itself varies across
875 pairings and tolerances ($\tau=5$ at truncation level, $\tau \in \{5, 10, 50\}$ at 10^{-8}), so no single fixed schedule
876 is competitive everywhere. Table 7 explains the numbers: with a corrector that removes the smooth
877 band to 10^{-6} in one call, the cost-aware oracle and the router call it first and reach h^2 after one or
878 two further sweeps, i.e., in about two iterations, whereas HINTS reaches the same accuracy only
879 at its first scheduled call (iteration τ), having spent $\tau-1$ sweeps on an error that is almost entirely
880 smooth. Below h^2 the remaining error is the high-frequency band that the corrector cannot see; the
881 router therefore stops calling it (one to three calls in total down to 10^{-8} ; Figure 4), whereas a fixed
882 schedule keeps paying for a call every τ iterations. The cost-agnostic greedy oracle (Algorithm 1 on
883 single operations) makes the same decisions as its cost-aware counterpart at 128^2 , because after the
884 first call the corrector is never better than a single sweep on the remaining high-frequency error; the
885 learned router reproduces the cost-aware oracle’s decisions almost exactly (same median iteration
886 and call counts, median times within a few percent, the difference being its own decision cost and
887 timer noise), and the five retrained routers reproduce the main router’s decision sequence to h^2 on
888 essentially every test instance (Table 10).

889 **Classical baselines and larger grids.** Against the classical methods without corrector (bottom
890 block of Table 5), the router with its best stationary pairing reaches truncation-level accuracy several
891 times faster than multigrid alone (three V-cycles) and than multigrid-preconditioned CG at 128^2
892 and 256^2 , and an order of magnitude faster than unpreconditioned or SymGS-preconditioned CG.
893 The reason is the corrector’s accuracy on the smooth band: one call removes the modes $|k| \leq 31$ to
894 10^{-6} at the price of a single 4096×4096 matrix–vector product, whereas a V(2,2)-cycle contracts
895 the whole spectrum by ≈ 0.03 at the price of four Gauss–Seidel sweeps on the fine grid plus the

coarse-grid work. Routing over $\{\text{NO}, \text{multigrid}\}$ is itself about twice as fast as multigrid alone on every grid: the oracle and the router call the corrector once and let the V-cycle remove the rest, which is the case in which the per-unit-cost comparison of Section 4 matters (the cycle is dearer than the unit). On 256^2 the router remains the fastest iterative method in every pairing ($1.58\times$ – $10.9\times$ over HINTS, $1.11\times$ – $2.33\times$ over the best period across the three equations). At 512^2 the margins shrink but the ordering holds for the stationary pairings, and the $\{\text{NO}, \text{multigrid}\}$ router is the fastest iterative method we measured at this size. With Jacobi and Gauss–Seidel the gains at 512^2 are small because the corrector’s fixed band $|k| \leq 31$ leaves the smoother all modes $k \geq 32$, whose slowest member contracts by only 0.98 per Jacobi sweep at $N = 512$ (against 0.71 at $N = 128$): every policy then needs the same ≈ 80 sweeps after a single corrector call, the wall-clock time is dominated by sweeps, and the schedule barely matters (with Jacobi the router misses h^2 within the 8,000-iteration cap on 2 of 16 instances, marked in the table). This is the expected limit of a corrector whose band does not grow with the grid: the hybrid’s advantage comes from removing the modes the smoother cannot damp, and a corrector with a band that scales with N (at the price of a dense matrix–vector product that grows like the fourth power of the band) or a multigrid member restores the gains, as the multigrid pairing shows.

Table 5: Poisson: median wall-clock time to relative error $\varepsilon \in \{10^{-3}, h^2, 10^{-8}\}$ on the 128^2 , 256^2 and 512^2 grids (64, 32 and 16 test instances) for every pairing $\{\text{NO}, \text{solver}\}$ and policy, and for the classical baselines without corrector. “HINTS (best τ)” is the best of $\tau \in \{5, 10, 25, 50\}$ chosen per cell on the test set itself (the chosen τ is the subscript). The greedy oracle is the cost-agnostic Algorithm 1, the cost-aware oracle is Algorithm 1 on cost-equalised macro-actions; both use the true error and are not deployable. Bold marks the router where it is the fastest deployable method of its pairing; † n : n instances did not reach the tolerance within the iteration cap ($>$: lower bound).

Pairing	Method	128 × 128			256 × 256			512 × 512		
		$\varepsilon=10^{-3}$	$\varepsilon=h^2$	$\varepsilon=10^{-8}$	$\varepsilon=10^{-3}$	$\varepsilon=h^2$	$\varepsilon=10^{-8}$	$\varepsilon=10^{-3}$	$\varepsilon=h^2$	$\varepsilon=10^{-8}$
{NO, Jacobi}	Solver only	806.1 ms	1.15 s	2.14 s	$>7.61 \text{ s}^{132}$	$>7.61 \text{ s}^{132}$	$>7.61 \text{ s}^{132}$	$>9.85 \text{ s}^{116}$	$>9.85 \text{ s}^{116}$	$>9.85 \text{ s}^{116}$
	HINTS ($\tau=25$)	2.61 ms	2.65 ms	226.1 ms	5.90 ms	5.95 ms	85.5 ms	33.1 ms	87.0 ms	309.2 ms ¹⁷
	HINTS (best τ)	1.06 ms _{$\tau=5$}	1.10 ms _{$\tau=5$}	182.8 ms _{$\tau=50$}	2.07 ms _{$\tau=5$}	2.47 ms _{$\tau=5$}	72.9 ms _{$\tau=50$}	9.27 ms _{$\tau=5$}	87.0 ms _{$\tau=25$}	293.9 ms ¹⁷ _{$\tau=50$}
	Greedy oracle (Alg. 1)	803 μs	854 μs	155.1 ms	1.28 ms	2.30 ms	67.2 ms	4.74 ms	78.2 ms	275.5 ms ¹⁷
	Cost-aware oracle	880 μs	942 μs	154.1 ms	1.42 ms	2.32 ms	67.0 ms	4.74 ms	74.2 ms	261.5 ms ¹⁷
	Learned router (ours)	918 μs	987 μs	151.3 ms	1.47 ms	2.38 ms	73.1 ms	4.86 ms	78.8 ms¹²	282.9 ms¹⁷
{NO, Jacobi (0.67)}	Solver only	1.16 s	1.62 s	3.16 s	$>7.74 \text{ s}^{132}$	$>7.74 \text{ s}^{132}$	$>7.74 \text{ s}^{132}$	–	–	–
	HINTS ($\tau=25$)	2.71 ms	2.71 ms	5.16 ms	6.02 ms	6.41 ms	16.8 ms	–	–	–
	HINTS (best τ)	1.11 ms _{$\tau=5$}	1.13 ms _{$\tau=5$}	3.03 ms _{$\tau=10$}	2.10 ms _{$\tau=5$}	3.43 ms _{$\tau=5$}	16.8 ms _{$\tau=25$}	–	–	–
	Greedy oracle (Alg. 1)	818 μs	915 μs	2.83 ms	1.30 ms	2.93 ms	16.2 ms	–	–	–
	Cost-aware oracle	820 μs	922 μs	2.93 ms	1.30 ms	2.77 ms	16.1 ms	–	–	–
	Learned router (ours)	839 μs	912 μs	2.93 ms	1.32 ms	2.83 ms	16.8 ms	–	–	–
{NO, GS}	Solver only	1.29 s	1.83 s	3.45 s	17.80 s	28.83 s	$>31.70 \text{ s}^{132}$	$>29.96 \text{ s}^{116}$	$>29.96 \text{ s}^{116}$	$>29.96 \text{ s}^{116}$
	HINTS ($\tau=25$)	6.43 ms	6.65 ms	13.9 ms	20.4 ms	22.0 ms	43.5 ms	92.5 ms	142.8 ms	439.6 ms
	HINTS (best τ)	1.72 ms _{$\tau=5$}	1.99 ms _{$\tau=5$}	5.55 ms _{$\tau=5$}	4.49 ms _{$\tau=5$}	6.07 ms _{$\tau=5$}	30.4 ms _{$\tau=10$}	20.0 ms _{$\tau=5$}	142.8 ms _{$\tau=25$}	439.6 ms _{$\tau=25$}
	Greedy oracle (Alg. 1)	835 μs	960 μs	4.34 ms	1.36 ms	3.85 ms	25.4 ms	4.65 ms	141.4 ms	447.1 ms
	Cost-aware oracle	827 μs	953 μs	4.37 ms	1.33 ms	3.78 ms	25.6 ms	4.87 ms	130.0 ms	433.0 ms
	Learned router (ours)	850 μs	1.01 ms	4.45 ms	1.36 ms	3.92 ms	26.0 ms	4.70 ms	131.5 ms	417.8 ms
{NO, SymGS}	Solver only	1.08 s	1.52 s	2.91 s	16.41 s	26.53 s	44.30 s	–	–	–
	HINTS ($\tau=25$)	10.2 ms	10.3 ms	20.2 ms	36.1 ms	37.3 ms	71.9 ms	–	–	–
	HINTS (best τ)	2.35 ms _{$\tau=5$}	2.60 ms _{$\tau=5$}	6.94 ms _{$\tau=5$}	7.18 ms _{$\tau=5$}	8.71 ms _{$\tau=5$}	29.8 ms _{$\tau=5$}	–	–	–
	Greedy oracle (Alg. 1)	829 μs	1.09 ms	3.87 ms	1.36 ms	4.18 ms	24.3 ms	–	–	–
	Cost-aware oracle	829 μs	1.13 ms	3.86 ms	1.35 ms	4.19 ms	23.8 ms	–	–	–
	Learned router (ours)	853 μs	1.13 ms	3.88 ms	1.34 ms	4.22 ms	24.4 ms	–	–	–
{NO, SOR (1.5)}	Solver only	424.4 ms	600.9 ms	1.15 s	5.94 s	9.60 s	15.97 s	$>29.74 \text{ s}^{116}$	$>29.74 \text{ s}^{116}$	$>29.74 \text{ s}^{116}$
	HINTS ($\tau=25$)	6.37 ms	12.2 ms	18.5 ms	20.5 ms	40.7 ms	60.9 ms	93.2 ms	182.7 ms	269.4 ms
	HINTS (best τ)	1.72 ms _{$\tau=5$}	3.37 ms _{$\tau=5$}	8.56 ms _{$\tau=10$}	4.57 ms _{$\tau=5$}	9.06 ms _{$\tau=5$}	27.2 ms _{$\tau=10$}	18.7 ms _{$\tau=5$}	74.4 ms _{$\tau=5$}	214.9 ms _{$\tau=10$}
	Greedy oracle (Alg. 1)	814 μs	1.03 ms	7.48 ms	1.31 ms	3.69 ms	22.7 ms	4.56 ms	61.1 ms	186.7 ms
	Cost-aware oracle	800 μs	1.02 ms	6.81 ms	1.30 ms	3.51 ms	22.4 ms	4.62 ms	60.3 ms	185.9 ms
	Learned router (ours)	823 μs	1.02 ms	7.05 ms	1.36 ms	3.80 ms	24.2 ms	4.69 ms	60.7 ms	188.8 ms
{NO, Multigrid}	Solver only	2.62 ms	3.84 ms	7.64 ms	9.61 ms	15.1 ms	28.6 ms	47.3 ms	90.8 ms	140.9 ms
	HINTS ($\tau=25$)	2.62 ms	3.83 ms	7.61 ms	9.73 ms	14.7 ms	28.7 ms	47.2 ms	92.0 ms	140.2 ms
	HINTS (best τ)	2.59 ms _{$\tau=5$}	3.81 ms _{$\tau=5$}	5.89 ms _{$\tau=5$}	9.56 ms _{$\tau=50$}	14.6 ms _{$\tau=5$}	20.5 ms _{$\tau=5$}	47.2 ms _{$\tau=25$}	90.5 ms _{$\tau=10$}	102.2 ms _{$\tau=5$}
	Greedy oracle (Alg. 1)	828 μs	2.01 ms	4.63 ms	1.39 ms	6.14 ms	15.8 ms	4.68 ms	39.5 ms	95.5 ms
	Cost-aware oracle	773 μs	2.02 ms	4.66 ms	1.31 ms	6.09 ms	16.2 ms	4.86 ms	38.6 ms	97.0 ms
	Learned router (ours)	783 μs	2.05 ms	4.72 ms	1.32 ms	5.97 ms	16.6 ms	4.76 ms	38.7 ms	99.5 ms
Classical baselines without corrector (same stencil, same accounting)										
	Multigrid V(2,2) alone	2.75 ms	3.99 ms	7.97 ms	9.27 ms	14.1 ms	27.0 ms	42.8 ms	81.9 ms	126.5 ms
	CG	9.37 ms	11.5 ms	16.3 ms	58.0 ms	73.1 ms	102.5 ms	689.5 ms	920.9 ms	1.18 s
	PCG (SymGS)	20.3 ms	30.8 ms	52.5 ms	146.5 ms	254.7 ms	365.0 ms	1.33 s	2.52 s	3.14 s
	PCG (multigrid)	2.97 ms	4.35 ms	7.36 ms	10.5 ms	15.2 ms	24.5 ms	44.2 ms	85.6 ms	113.2 ms

Table 6: Poisson: paired median speedup of the learned router at $\varepsilon = h^2$ and 10^{-8} over HINTS ($\tau=25$), over the best fixed period for that cell, over multigrid alone and over multigrid-preconditioned CG, with the one-sided Wilcoxon p -value that the router is faster (or, where its median is larger, that it is slower). Bold: $\geq 1.1 \times$ with $p < 0.01$.

Pairing	ε	128 × 128				256 × 256				512 × 512			
		vs. HINTS-25	vs. best τ	vs. multigrid	vs. MG-Krylow	vs. HINTS-25	vs. best τ	vs. multigrid	vs. MG-Krylow	vs. HINTS-25	vs. best τ	vs. multigrid	vs. MG-Krylow
{NO, Jacobi}	h^2	2.68 × ($< 10^{-10}$)	1.21 × (10^{-6})	3.43 × (10^{-7})	3.57 × (10^{-7})	2.52 × (10^{-10})	1.56 × (10^{-7})	6.70 × (10^{-7})	6.30 × (10^{-7})	1.11 × (0.001)	1.11 × (0.001)	1.14 × (0.121)	1.26 × (0.086)
{NO, Jacobi (0.67)}	10^{-8}	1.55 × ($< 10^{-10}$)	1.22 × (10^{-7})	0.06 × (slower, 10^{-7})	0.05 × (slower, 10^{-7})	1.23 × (10^{-10})	1.11 × (10^{-7})	0.37 × (slower, 0.009)	0.35 × (slower, 0.007)	0.90 × (slower, 0.590)	0.89 × (slower, 0.545)	0.45 × (slower, 0.082)	0.40 × (slower, 0.082)
	h^2	2.90 × ($< 10^{-10}$)	1.20 × (10^{-7})	3.62 × ($< 10^{-10}$)	3.85 × ($< 10^{-10}$)	2.36 × (10^{-7})	1.73 × (10^{-7})	5.62 × (10^{-7})	5.29 × (10^{-7})	—	—	—	—
{NO, GS}	10^{-8}	1.78 × ($< 10^{-10}$)	1.09 × (10^{-7})	2.65 × ($< 10^{-10}$)	2.42 × ($< 10^{-10}$)	1.10 × (10^{-7})	1.10 × (10^{-7})	1.65 × (10^{-7})	1.59 × (10^{-7})	—	—	—	—
	h^2	6.75 × ($< 10^{-10}$)	1.94 × ($< 10^{-10}$)	3.39 × ($< 10^{-10}$)	3.52 × ($< 10^{-10}$)	5.77 × (10^{-7})	1.70 × (10^{-7})	4.15 × (10^{-7})	3.85 × (10^{-7})	1.11 × (0.003)	1.11 × (0.003)	0.69 × (slower, 0.530)	0.76 × (slower, 0.590)
{NO, SymGS}	10^{-8}	3.25 × ($< 10^{-10}$)	1.25 × (10^{-7})	1.72 × ($< 10^{-10}$)	1.75 × ($< 10^{-10}$)	1.67 × (10^{-7})	1.20 × (10^{-7})	1.08 × (0.007)	0.98 × (slower, 0.701)	1.07 × (0.004)	1.07 × (0.004)	0.31 × (slower, 10^{-7})	0.28 × (slower, 10^{-7})
	h^2	9.21 × ($< 10^{-10}$)	2.34 × ($< 10^{-10}$)	3.20 × ($< 10^{-10}$)	3.42 × ($< 10^{-10}$)	9.96 × (10^{-7})	2.27 × (10^{-7})	4.00 × (10^{-7})	3.98 × (10^{-7})	—	—	—	—
{NO, SOR (1.5)}	10^{-8}	5.26 × ($< 10^{-10}$)	1.81 × ($< 10^{-10}$)	1.99 × ($< 10^{-10}$)	1.93 × ($< 10^{-10}$)	2.97 × (10^{-7})	1.22 × (10^{-7})	1.16 × (0.004)	1.05 × (0.006)	—	—	—	—
	h^2	10.3 × ($< 10^{-10}$)	3.22 × ($< 10^{-10}$)	3.28 × ($< 10^{-10}$)	3.45 × ($< 10^{-10}$)	10.9 × (10^{-7})	2.33 × (10^{-7})	4.28 × (10^{-7})	4.13 × (10^{-7})	3.10 × (10^{-7})	3.10 × (10^{-7})	1.49 × (0.052)	1.64 × (0.037)
{NO, Multigrid}	10^{-8}	2.65 × ($< 10^{-10}$)	1.25 × (10^{-7})	1.14 × (10^{-7})	1.08 × (0.011)	2.48 × (10^{-7})	1.09 × (10^{-7})	1.16 × (0.004)	1.04 × (0.060)	1.37 × (10^{-7})	1.09 × (10^{-7})	0.68 × (slower, 0.005)	0.62 × (slower, 0.005)
	h^2	1.96 × ($< 10^{-10}$)	1.88 × ($< 10^{-10}$)	2.62 × ($< 10^{-10}$)	3.42 × ($< 10^{-10}$)	2.46 × (10^{-7})	2.39 × (10^{-7})	2.35 × (10^{-7})	2.53 × (10^{-7})	2.08 × (10^{-7})	2.08 × (10^{-7})	1.96 × (10^{-7})	2.12 × (10^{-7})
	10^{-8}	1.66 × ($< 10^{-10}$)	1.27 × ($< 10^{-10}$)	1.75 × ($< 10^{-10}$)	1.60 × ($< 10^{-10}$)	1.71 × (10^{-10})	1.20 × (10^{-7})	1.72 × (10^{-7})	1.53 × (10^{-10})	1.39 × (10^{-7})	1.09 × (0.002)	1.31 × (10^{-7})	1.23 × (0.009)




Figure 4: Poisson: fraction of test runs that execute a corrector call at each iteration, for the learned router, the cost-aware oracle and HINTS ($\tau=25$), for every pairing (columns) and grid (rows; all test instances of each grid).

Table 7: Poisson: median iterations (and corrector calls) to $\varepsilon = h^2$ for HINTS ($\tau=25$), the cost-aware oracle and the learned router on every grid.

Pairing	128 × 128			256 × 256			512 × 512		
	HINTS-25	oracle	router	HINTS-25	oracle	router	HINTS-25	oracle	router
{NO, Jacobi}	25 (1)	2 (1)	2 (1)	25 (1)	6 (1)	6 (1)	65 (2)	64 (1)	64 (1)
{NO, Jacobi (0.67)}	25 (1)	2 (1)	2 (1)	25 (1)	8 (1)	8 (1)	—	—	—
{NO, GS}	26 (1)	2 (1)	2 (1)	27 (1)	4 (1)	4 (1)	38 (1)	35 (1)	35 (1)
{NO, SymGS}	26 (1)	2 (1)	2 (1)	26 (1)	3 (1)	3 (1)	—	—	—
{NO, SOR (1.5)}	50 (2)	2 (1)	2 (1)	50 (2)	4 (1)	4 (1)	50 (2)	17 (1)	17 (2)
{NO, Multigrid}	3 (0)	2 (1)	2 (1)	3 (0)	2 (1)	2 (1)	4 (0)	2 (1)	3 (1)

Table 8: Poisson: measured cost of one classical iteration c_{solver} and of one corrector call c_{NO} (residual evaluation included), and the resulting macro-action size m_{solver} (sweeps per decision; the corrector is always applied once).

Pairing	128 × 128			256 × 256			512 × 512		
	c_{solver}	c_{NO}	m_{solver}	c_{solver}	c_{NO}	m_{solver}	c_{solver}	c_{NO}	m_{solver}
{NO, Jacobi}	52 μs	635 μs	12	167 μs	1.00 ms	6	1.07 ms	3.36 ms	3
{NO, Jacobi (0.67)}	52 μs	639 μs	12	150 μs	971 μs	6	—	—	—
{NO, GS}	200 μs	643 μs	3	780 μs	1.01 ms	1	3.50 ms	3.30 ms	1
{NO, SymGS}	366 μs	635 μs	2	1.40 ms	931 μs	1	—	—	—
{NO, SOR (1.5)}	197 μs	635 μs	3	771 μs	960 μs	1	3.41 ms	3.20 ms	1
{NO, Multigrid}	1.25 ms	698 μs	1	4.67 ms	950 μs	1	21.7 ms	2.86 ms	1

Table 9: Poisson, 128×128 : iteration-based metrics of the main text—relative error after $T = 300$ iterations and AUC over the T iterations (mean and standard error over 64 instances; p : one-sided paired t -test that the method’s AUC exceeds the router’s). Bold: router better than both solver-only and HINTS.

Method	$\ e_h^{(T)}\ /\ u_h\ $	AUC	p
{NO, Jacobi}			
Solver only	7.99e-01 (4.5e-03)	2.68e+02 (8.9e-01)	$< 10^{-10}$
HINTS ($\tau=25$)	7.10e-06 (1.9e-06)	2.37e+01 (1.1e-02)	$< 10^{-10}$
Learned router (ours)	6.86e-06 (1.8e-06)	7.15e-03 (1.9e-03)	-
Cost-aware oracle	<i>6.86e-06 (1.8e-06)</i>	<i>7.15e-03 (1.9e-03)</i>	-
{NO, Jacobi (0.67)}			
Solver only	8.56e-01 (3.8e-03)	2.77e+02 (7.1e-01)	$< 10^{-10}$
HINTS ($\tau=25$)	1.15e-12 (1.0e-13)	2.38e+01 (8.3e-03)	$< 10^{-10}$
Learned router (ours)	7.57e-10 (2.9e-11)	2.49e-03 (6.3e-04)	-
Cost-aware oracle	<i>8.19e-10 (1.2e-11)</i>	<i>2.49e-03 (6.3e-04)</i>	-
{NO, GS}			
Solver only	6.53e-01 (5.2e-03)	2.42e+02 (1.2e+00)	$< 10^{-10}$
HINTS ($\tau=25$)	2.84e-11 (1.9e-12)	2.35e+01 (2.1e-02)	$< 10^{-10}$
Learned router (ours)	7.66e-10 (1.8e-11)	1.94e-03 (4.9e-04)	-
Cost-aware oracle	<i>7.08e-10 (2.5e-11)</i>	<i>1.94e-03 (4.9e-04)</i>	-
{NO, SymGS}			
Solver only	4.49e-01 (4.6e-03)	2.03e+02 (1.3e+00)	$< 10^{-10}$
HINTS ($\tau=25$)	7.14e-10 (5.4e-11)	2.30e+01 (3.5e-02)	$< 10^{-10}$
Learned router (ours)	6.50e-10 (2.4e-11)	1.46e-03 (3.7e-04)	-
Cost-aware oracle	<i>6.32e-10 (2.5e-11)</i>	<i>1.46e-03 (3.7e-04)</i>	-
{NO, SOR (1.5)}			
Solver only	3.09e-01 (3.5e-03)	1.72e+02 (1.3e+00)	$< 10^{-10}$
HINTS ($\tau=25$)	2.16e-10 (1.3e-11)	2.26e+01 (4.6e-02)	$< 10^{-10}$
Learned router (ours)	8.74e-10 (9.7e-12)	3.30e-03 (8.4e-04)	-
Cost-aware oracle	<i>8.36e-10 (1.5e-11)</i>	<i>3.29e-03 (8.4e-04)</i>	-

Table 10: Poisson, 128×128 : five independent router-training seeds per pairing—mean and standard deviation over seeds of the median work-unit time to h^2 , the share of test instances on which the seed router’s decision sequence to h^2 coincides with the main router’s, the range over seeds of the paired speedup over HINTS ($\tau=25$), and the number of seeds significantly faster ($p < 0.01$) than both HINTS ($\tau=25$) and the best fixed period.

Pairing	time to h^2 over 5 seeds	identical decisions	speedup vs. HINTS-25	seeds with $p < 0.01$
{NO, Jacobi}	0.74 ± 0.00 ms	100%	2.59–2.59×	5/5
{NO, Jacobi (0.67)}	0.75 ± 0.00 ms	100%	2.58–2.58×	5/5
{NO, GS}	0.90 ± 0.00 ms	100%	6.30–6.30×	5/5
{NO, SymGS}	1.06 ± 0.00 ms	100%	9.29–9.29×	5/5
{NO, SOR (1.5)}	0.89 ± 0.00 ms	100%	10.45–10.45×	5/5




Figure 5: Poisson, 128×128 : routing decisions (dark: corrector call) of the learned router (top row) and the cost-aware oracle (bottom row) over the first 60 iterations of every test instance, for every pairing.

912 E.3 Convection–diffusion

913 The convection–diffusion results (Tables 11 to 16 and fig. 6) mirror the Poisson ones: at 128^2 the
 914 router is the fastest deployable method in every pairing at h^2 and at 10^{-8} , reproduces the cost-aware
 915 oracle, and is $2\text{--}3\times$ faster than multigrid alone and than multigrid-preconditioned BiCGSTAB; on
 916 256^2 the same holds in every pairing; at 512^2 the router is faster than HINTS and than the best
 917 fixed period in the Jacobi, GS and SOR pairings and the {NO, multigrid} router is again the fastest
 918 iterative method, with the same mechanism as on Poisson (the Jacobi and GS pairings are limited by
 919 the modes outside the corrector’s band: the router, the oracle and HINTS all reach h^2 after 84–85
 920 Jacobi iterations). The nonsymmetric operator changes only the details: the best fixed period is $\tau=10$
 921 rather than 5 in several cells, the undamped-Jacobi pairing stalls at the near-Nyquist band below 10^{-6}
 922 for every policy, and GMRES(20) is two orders of magnitude slower than the router.

Table 11: Convection–diffusion: median wall-clock time to relative error $\varepsilon \in \{10^{-3}, h^2, 10^{-8}\}$ on the 128^2 , 256^2 and 512^2 grids for every pairing and policy, and for the classical baselines without corrector (conventions as in Table 5).

Pairing	Method	128 × 128			256 × 256			512 × 512		
		$\varepsilon=10^{-3}$	$\varepsilon=h^2$	$\varepsilon=10^{-8}$	$\varepsilon=10^{-3}$	$\varepsilon=h^2$	$\varepsilon=10^{-8}$	$\varepsilon=10^{-3}$	$\varepsilon=h^2$	$\varepsilon=10^{-8}$
{NO, Jacobi}	Solver only	1.07 s	1.51 s	2.90 s	>12.73 s ¹²⁶	>12.73 s ¹³²	>12.73 s ¹³²	>13.99 s ¹¹⁶	>13.99 s ¹¹⁶	>13.99 s ¹¹⁶
	HINTS ($\tau=25$)	3.41 ms	3.44 ms	384.4 ms	9.15 ms	9.37 ms	500.8 ms	49.2 ms	158.0 ms	1.48 s ¹⁷
	HINTS (best τ)	1.26 ms _{$\tau=5$}	1.28 ms _{$\tau=5$}	321.5 ms _{$\tau=50$}	2.80 ms _{$\tau=5$}	4.95 ms _{$\tau=10$}	475.6 ms _{$\tau=50$}	13.1 ms _{$\tau=5$}	147.4 ms _{$\tau=50$}	1.43 s ¹⁷ _{$\tau=50$}
	Greedy oracle (Alg. 1)	874 μ s	992 μ s	281.0 ms	1.57 ms	4.43 ms	436.3 ms	6.23 ms	156.1 ms	1.37 s ¹⁷
	Cost-aware oracle	979 μ s	1.15 ms	281.8 ms	1.70 ms	4.81 ms	437.5 ms	6.23 ms	154.3 ms	1.31 s ¹⁷
	Learned router (ours)	1.01 ms	1.14 ms	292.7 ms	1.77 ms	4.76 ms	435.1 ms	6.67 ms	144.3 ms	1.29 s¹⁷
{NO, Jacobi (0.67)}	Solver only	1.55 s	2.22 s	4.32 s	>12.55 s ¹³²	>12.55 s ¹³²	>12.55 s ¹³²	—	—	—
	HINTS ($\tau=25$)	3.40 ms	3.40 ms	6.53 ms	9.03 ms	9.31 ms	30.4 ms	—	—	—
	HINTS (best τ)	1.28 ms _{$\tau=5$}	1.31 ms _{$\tau=5$}	3.79 ms _{$\tau=10$}	5.68 ms _{$\tau=5$}	5.68 ms _{$\tau=10$}	30.4 ms _{$\tau=25$}	—	—	—
	Greedy oracle (Alg. 1)	885 μ s	1.10 ms	3.56 ms	1.54 ms	5.93 ms	26.8 ms	—	—	—
	Cost-aware oracle	878 μ s	1.08 ms	3.61 ms	1.58 ms	5.94 ms	27.5 ms	—	—	—
	Learned router (ours)	912 μs	1.10 ms	3.70 ms	1.61 ms	5.76 ms	26.9 ms	—	—	—
{NO, GS}	Solver only	1.15 s	1.63 s	3.14 s	18.36 s	30.05 s	>36.70 s ¹³²	>32.99 s ¹¹⁶	>32.99 s ¹¹⁶	>32.99 s ¹¹⁶
	HINTS ($\tau=25$)	7.03 ms	8.27 ms	20.5 ms	23.5 ms	31.0 ms	59.3 ms	105.1 ms	214.1 ms	504.5 ms
	HINTS (best τ)	1.90 ms _{$\tau=5$}	3.60 ms _{$\tau=5$}	6.24 ms _{$\tau=5$}	5.29 ms _{$\tau=5$}	10.5 ms _{$\tau=5$}	36.7 ms _{$\tau=10$}	22.0 ms _{$\tau=5$}	192.7 ms _{$\tau=10$}	504.5 ms _{$\tau=25$}
	Greedy oracle (Alg. 1)	887 μ s	1.23 ms	5.32 ms	1.57 ms	7.20 ms	31.1 ms	5.70 ms	184.8 ms	493.9 ms
	Cost-aware oracle	909 μ s	1.32 ms	5.46 ms	1.61 ms	7.15 ms	31.0 ms	5.97 ms	185.9 ms	489.9 ms
	Learned router (ours)	906 μs	1.35 ms	5.59 ms	1.57 ms	7.05 ms	31.4 ms	5.79 ms	179.9 ms	486.0 ms
{NO, SymGS}	Solver only	1.10 s	1.56 s	2.99 s	16.90 s	27.77 s	46.85 s	—	—	—
	HINTS ($\tau=25$)	10.7 ms	11.1 ms	21.6 ms	39.2 ms	42.3 ms	78.1 ms	—	—	—
	HINTS (best τ)	2.52 ms _{$\tau=5$}	2.96 ms _{$\tau=5$}	7.35 ms _{$\tau=5$}	7.80 ms _{$\tau=5$}	10.9 ms _{$\tau=5$}	34.8 ms _{$\tau=5$}	—	—	—
	Greedy oracle (Alg. 1)	911 μ s	1.25 ms	4.91 ms	1.58 ms	6.93 ms	28.4 ms	—	—	—
	Cost-aware oracle	905 μ s	1.27 ms	4.48 ms	1.56 ms	6.89 ms	28.4 ms	—	—	—
	Learned router (ours)	924 μs	1.28 ms	4.54 ms	1.55 ms	6.96 ms	29.6 ms	—	—	—
{NO, SOR (1.5)}	Solver only	270.0 ms	383.6 ms	737.6 ms	5.14 s	8.46 s	14.22 s	>32.91 s ¹¹⁶	>32.91 s ¹¹⁶	>32.91 s ¹¹⁶
	HINTS ($\tau=25$)	6.92 ms	13.6 ms	20.2 ms	23.2 ms	46.1 ms	69.0 ms	103.8 ms	206.6 ms	308.3 ms
	HINTS (best τ)	1.98 ms _{$\tau=5$}	3.64 ms _{$\tau=5$}	9.62 ms _{$\tau=5$}	5.22 ms _{$\tau=5$}	10.6 ms _{$\tau=5$}	31.2 ms _{$\tau=5$}	22.1 ms _{$\tau=5$}	89.3 ms _{$\tau=5$}	234.2 ms _{$\tau=10$}
	Greedy oracle (Alg. 1)	919 μ s	1.39 ms	7.66 ms	1.56 ms	5.90 ms	28.1 ms	6.82 ms	83.3 ms	212.8 ms
	Cost-aware oracle	922 μ s	1.63 ms	7.97 ms	1.57 ms	5.74 ms	27.5 ms	5.68 ms	85.0 ms	217.1 ms
	Learned router (ours)	934 μs	1.59 ms	8.18 ms	1.57 ms	5.80 ms	27.6 ms	5.64 ms	83.6 ms	216.7 ms
{NO, Multigrid}	Solver only	3.18 ms	4.71 ms	9.28 ms	10.9 ms	16.4 ms	31.8 ms	49.5 ms	90.8 ms	148.5 ms
	HINTS ($\tau=25$)	3.17 ms	4.68 ms	9.28 ms	10.9 ms	16.7 ms	31.6 ms	50.0 ms	97.6 ms	150.4 ms
	HINTS (best τ)	3.17 ms _{$\tau=25$}	4.66 ms _{$\tau=5$}	7.14 ms _{$\tau=5$}	10.9 ms _{$\tau=5$}	16.5 ms _{$\tau=5$}	23.0 ms _{$\tau=5$}	49.4 ms _{$\tau=10$}	92.5 ms _{$\tau=10$}	112.8 ms _{$\tau=5$}
	Greedy oracle (Alg. 1)	896 μ s	2.34 ms	5.57 ms	1.58 ms	7.30 ms	20.8 ms	6.18 ms	50.7 ms	103.5 ms
	Cost-aware oracle	851 μ s	2.33 ms	5.49 ms	1.55 ms	7.23 ms	19.9 ms	5.74 ms	50.9 ms	104.1 ms
	Learned router (ours)	873 μs	2.32 ms	5.50 ms	1.60 ms	7.26 ms	20.2 ms	5.63 ms	53.7 ms	102.6 ms
<i>Classical baselines without corrector (same stencil, same accounting)</i>										
{NO, Multigrid}	Multigrid V(2,2) alone	3.06 ms	4.52 ms	8.92 ms	12.3 ms	18.8 ms	36.1 ms	52.4 ms	99.7 ms	150.7 ms
	BiCGSTAB	56.0 ms	67.6 ms	93.1 ms	377.7 ms	486.6 ms	605.9 ms	3.79 s	5.56 s	6.72 s ¹¹
	BiCGSTAB (multigrid)	3.12 ms	6.17 ms	9.21 ms	12.5 ms	24.0 ms	36.1 ms	46.1 ms	93.5 ms	144.8 ms
	GMRES(20)	301.7 ms	434.1 ms	882.8 ms	1.85 s	3.00 s	5.14 s	—	—	—

Table 12: Convection–diffusion: paired median speedup of the learned router at $\varepsilon = h^2$ and 10^{-8} over HINTS ($\tau=25$), the best fixed period, multigrid alone and multigrid-preconditioned BiCGSTAB, with one-sided Wilcoxon p -values (conventions as in Table 6).

Pairing	ε	128 × 128				256 × 256				512 × 512			
		vs. HINTS-25	vs. best τ	vs. multigrid	vs. MG-Krylov	vs. HINTS-25	vs. best τ	vs. multigrid	vs. MG-Krylov	vs. HINTS-25	vs. best τ	vs. multigrid	vs. MG-Krylov
{NO, Jacobi}	h^2	2.95 × ($< 10^{-10}$)	1.34 × ($< 10^{-10}$)	3.54 × (10^{-1})	5.18 × (10^{-7})	1.92 × (10^{-7})	1.30 × (10^{-5})	4.60 × (0.000)	5.06 × (0.005)	1.17 × (0.003)	1.14 × (0.009)	0.69 × (slower, 0.174)	0.67 × (slower, 0.202)
	10^{-8}	1.33 × ($< 10^{-10}$)	1.17 × ($< 10^{-10}$)	0.03 × (slower, 10^{-10})	0.03 × (slower, 10^{-10})	1.12 × (10^{-7})	1.08 × (10^{-5})	0.08 × (slower, 10^{-7})	0.08 × (slower, 10^{-7})	0.99 × (slower, 0.990)	0.98 × (slower, 0.963)	0.22 × (slower, 0.006)	0.18 × (slower, 0.010)
{NO, Jacobi (0.67)}	h^2	3.14 × ($< 10^{-10}$)	1.32 × (10^{-10})	3.71 × ($< 10^{-10}$)	3.21 × ($< 10^{-10}$)	1.90 × (10^{-7})	1.25 × (10^{-5})	3.85 × (10^{-7})	4.42 × (10^{-7})	—	—	—	—
	10^{-8}	1.77 × ($< 10^{-10}$)	1.07 × (10^{-10})	2.37 × ($< 10^{-10}$)	2.56 × ($< 10^{-10}$)	1.06 × (0.002)	1.06 × (0.002)	1.27 × (10^{-7})	1.26 × (10^{-7})	—	—	—	—
{NO, GS}	h^2	6.76 × ($< 10^{-10}$)	2.60 × ($< 10^{-10}$)	3.29 × ($< 10^{-10}$)	4.42 × ($< 10^{-10}$)	3.52 × (10^{-7})	1.83 × (10^{-5})	3.08 × (10^{-7})	3.47 × (10^{-7})	1.05 × (0.002)	1.09 × (10^{-7})	0.56 × (slower, 0.161)	0.54 × (slower, 0.174)
	10^{-8}	3.78 × ($< 10^{-10}$)	1.17 × (10^{-10})	1.40 × ($< 10^{-10}$)	1.91 × ($< 10^{-10}$)	1.95 × (10^{-7})	1.17 × (10^{-5})	1.16 × (0.002)	1.19 × (0.003)	1.03 × (0.002)	1.03 × (0.002)	0.32 × (slower, 10^{-7})	0.29 × (slower, 10^{-7})
{NO, SymGS}	h^2	8.70 × ($< 10^{-10}$)	2.40 × ($< 10^{-10}$)	3.46 × ($< 10^{-10}$)	4.73 × ($< 10^{-10}$)	6.26 × (10^{-7})	1.69 × (10^{-5})	3.15 × (10^{-7})	3.54 × (10^{-7})	—	—	—	—
	10^{-8}	4.75 × ($< 10^{-10}$)	1.62 × ($< 10^{-10}$)	1.91 × ($< 10^{-10}$)	2.08 × ($< 10^{-10}$)	2.63 × (10^{-7})	1.18 × (10^{-5})	1.24 × (10^{-7})	1.24 × (10^{-7})	—	—	—	—
{NO, SOR (1.5)}	h^2	8.89 × ($< 10^{-10}$)	2.40 × ($< 10^{-10}$)	3.07 × ($< 10^{-10}$)	3.89 × ($< 10^{-10}$)	7.99 × (10^{-7})	1.96 × (10^{-5})	3.72 × (10^{-7})	4.15 × (10^{-7})	2.48 × (10^{-7})	1.18 × (10^{-7})	1.23 × (0.065)	1.18 × (0.088)
	10^{-8}	2.88 × ($< 10^{-10}$)	1.31 × ($< 10^{-10}$)	1.16 × (10^{-7})	1.27 × (10^{-7})	2.50 × (10^{-7})	1.19 × (10^{-5})	1.32 × (10^{-7})	1.29 × (10^{-7})	1.41 × (10^{-7})	1.08 × (10^{-7})	0.71 × (slower, 0.008)	0.64 × (slower, 0.009)
{NO, Multigrid}	h^2	2.07 × ($< 10^{-10}$)	2.01 × ($< 10^{-10}$)	2.12 × ($< 10^{-10}$)	2.84 × ($< 10^{-10}$)	2.19 × (10^{-7})	2.20 × (10^{-7})	2.56 × (10^{-7})	3.27 × (10^{-7})	1.83 × (10^{-7})	1.75 × (10^{-7})	1.44 × (10^{-7})	1.76 × (10^{-7})
	10^{-8}	1.70 × ($< 10^{-10}$)	1.28 × ($< 10^{-10}$)	1.65 × ($< 10^{-10}$)	1.74 × ($< 10^{-10}$)	1.54 × (10^{-7})	1.14 × (10^{-5})	1.70 × (10^{-7})	1.81 × (10^{-7})	1.37 × (10^{-7})	1.07 × (10^{-7})	1.42 × (10^{-7})	1.32 × (10^{-7})




Figure 6: Convection–diffusion: corrector usage per iteration of the learned router, the cost-aware oracle and HINTS ($\tau=25$) for every pairing (columns) and grid (rows).

Table 13: Convection–diffusion: median iterations (and corrector calls) to $\varepsilon = h^2$ on every grid.

Pairing	128 × 128			256 × 256			512 × 512		
	HINTS-25	oracle	router	HINTS-25	oracle	router	HINTS-25	oracle	router
{NO, Jacobi}	25 (1)	2 (1)	2 (1)	25 (1)	10 (1)	10 (1)	85 (2)	84 (1)	84 (1)
{NO, Jacobi (0.67)}	25 (1)	4 (1)	4 (1)	25 (1)	15 (1)	15 (1)	—	—	—
{NO, GS}	30 (1)	2 (1)	2 (1)	33 (1)	7 (1)	7 (1)	51 (2)	44 (1)	44 (2)
{NO, SymGS}	26 (1)	2 (1)	2 (1)	27 (1)	4 (1)	4 (1)	—	—	—
{NO, SOR (1.5)}	50 (2)	4 (1)	4 (1)	50 (2)	6 (1)	6 (1)	50 (2)	20 (1)	20 (2)
{NO, Multigrid}	3 (0)	2 (1)	2 (1)	3 (0)	2 (1)	2 (1)	4 (0)	3 (1)	3 (1)

Table 14: Convection–diffusion: measured per-iteration costs and macro-action sizes.

Pairing	128 × 128			256 × 256			512 × 512		
	c_{solver}	c_{NO}	m_{solver}	c_{solver}	c_{NO}	m_{solver}	c_{solver}	c_{NO}	m_{solver}
{NO, Jacobi}	72 μs	658 μs	9	286 μs	1.13 ms	4	1.42 ms	3.61 ms	3
{NO, Jacobi (0.67)}	71 μs	659 μs	9	288 μs	1.15 ms	4	—	—	—
{NO, GS}	234 μs	673 μs	3	855 μs	1.08 ms	1	3.94 ms	3.67 ms	1
{NO, SymGS}	378 μs	665 μs	2	1.55 ms	1.11 ms	1	—	—	—
{NO, SOR (1.5)}	223 μs	660 μs	3	813 μs	1.03 ms	1	3.90 ms	3.62 ms	1
{NO, Multigrid}	1.47 ms	737 μs	1	5.68 ms	1.06 ms	1	24.7 ms	3.60 ms	1

Table 15: Convection–diffusion, 128×128 : iteration-based metrics (conventions as in Table 9).

Method	$\ e_h^{(T)}\ /\ u_h\ $	AUC	p
{NO, Jacobi}			
Solver only	7.43e-01 (9.1e-03)	2.57e+02 (1.9e+00)	$< 10^{-10}$
HINTS ($\tau=25$)	1.55e-05 (4.2e-06)	2.36e+01 (2.7e-02)	$< 10^{-10}$
Learned router (ours)	1.50e-05 (4.1e-06)	1.55e-02 (4.1e-03)	-
Cost-aware oracle	1.50e-05 (4.1e-06)	1.55e-02 (4.1e-03)	-
{NO, Jacobi (0.67)}			
Solver only	8.11e-01 (7.9e-03)	2.69e+02 (1.5e+00)	$< 10^{-10}$
HINTS ($\tau=25$)	1.46e-12 (2.7e-13)	2.37e+01 (2.1e-02)	$< 10^{-10}$
Learned router (ours)	7.96e-10 (1.6e-11)	5.39e-03 (1.4e-03)	-
Cost-aware oracle	8.18e-10 (1.3e-11)	5.39e-03 (1.4e-03)	-
{NO, GS}			
Solver only	5.46e-01 (9.9e-03)	2.20e+02 (2.4e+00)	$< 10^{-10}$
HINTS ($\tau=25$)	8.78e-11 (4.8e-12)	2.32e+01 (5.4e-02)	$< 10^{-10}$
Learned router (ours)	6.84e-10 (2.8e-11)	4.03e-03 (1.0e-03)	-
Cost-aware oracle	6.73e-10 (2.8e-11)	4.03e-03 (1.0e-03)	-
{NO, SymGS}			
Solver only	3.77e-01 (9.1e-03)	1.84e+02 (2.7e+00)	$< 10^{-10}$
HINTS ($\tau=25$)	3.64e-12 (2.3e-13)	2.26e+01 (8.1e-02)	$< 10^{-10}$
Learned router (ours)	5.19e-10 (3.4e-11)	3.16e-03 (8.0e-04)	-
Cost-aware oracle	5.00e-10 (3.2e-11)	3.16e-03 (8.0e-04)	-
{NO, SOR (1.5)}			
Solver only	1.23e-01 (3.4e-03)	1.17e+02 (1.9e+00)	$< 10^{-10}$
HINTS ($\tau=25$)	8.06e-10 (2.3e-11)	2.12e+01 (1.2e-01)	$< 10^{-10}$
Learned router (ours)	8.08e-10 (1.3e-11)	6.49e-03 (1.7e-03)	-
Cost-aware oracle	7.77e-10 (1.8e-11)	6.49e-03 (1.7e-03)	-

Table 16: Convection–diffusion, 128×128 : five-seed router-training trials (conventions as in Table 10).

Pairing	time to h^2 over 5 seeds	identical decisions	speedup vs. HINTS-25	seeds with $p < 0.01$
{NO, Jacobi}	0.84 ± 0.00 ms	100%	$2.91\text{--}2.91\times$	5/5
{NO, Jacobi (0.67)}	0.91 ± 0.00 ms	100%	$2.67\text{--}2.67\times$	5/5
{NO, GS}	1.11 ± 0.00 ms	100%	$7.22\text{--}7.22\times$	5/5
{NO, SymGS}	1.12 ± 0.00 ms	100%	$9.10\text{--}9.10\times$	5/5
{NO, SOR (1.5)}	1.30 ± 0.00 ms	100%	$9.41\text{--}9.41\times$	5/5




Figure 7: Convection–diffusion, 128×128 : routing decisions (dark: corrector call) of the learned router (top row) and the cost-aware oracle (bottom row) over the first 60 iterations of every test instance.

923 E.4 Anisotropic diffusion

924 The anisotropic equation is where the two components of the hybrid are most complementary and
 925 where the classical solvers are weakest: the point smoothers cannot damp the x_1 -oscillatory modes

at all (Table 17: none of the five solvers alone reaches h^2 within the iteration cap on either grid), whereas the corrector, whose sensor grid keeps full resolution along x_1 , removes exactly those modes together with the smooth ones in a single call and leaves the smoother only the x_2 -oscillatory band, which it damps at the isotropic rate. Consequently the learned router reaches h^2 in under 1 ms at 128^2 for every pairing—the same absolute time as on the isotropic equations—while HINTS ($\tau=25$) needs 3–10 ms and the best fixed period 1.2–2.4 ms; at 256^2 (Jacobi, GS and SOR pairings, 16 instances) the router needs 9–10 ms against 15–30 ms for HINTS and 10–13 ms for the best period (Table 18). The router again matches the cost-aware oracle, its decisions are identical across the five training seeds (Table 22), and it spends the iterations after the single corrector call on the x_2 -oscillatory band (Figure 8 and table 19). The multigrid and Krylov baselines are not run for this equation (Section E.1).

Table 17: Anisotropic diffusion: median wall-clock time to relative error $\varepsilon \in \{10^{-3}, h^2, 10^{-8}\}$ on the 128^2 (64 instances, five pairings) and 256^2 (16 instances; Jacobi, GS and SOR pairings) grids for every policy (conventions as in Table 5; the solver-only runs are capped at 60,000 and 40,000 iterations and never reach h^2).

Pairing	Method	128×128			256×256		
		$\varepsilon=10^{-3}$	$\varepsilon=h^2$	$\varepsilon=10^{-8}$	$\varepsilon=10^{-3}$	$\varepsilon=h^2$	$\varepsilon=10^{-8}$
{NO, Jacobi}	Solver only	$>4.62 \text{ s}^{164}$	$>4.62 \text{ s}^{164}$	$>4.62 \text{ s}^{164}$	$>9.29 \text{ s}^{116}$	$>9.29 \text{ s}^{116}$	$>9.29 \text{ s}^{116}$
	HINTS ($\tau=25$)	2.90 ms	2.90 ms	67.6 ms	14.9 ms	14.9 ms	127.4 ms
	HINTS (best τ)	$1.15 \text{ ms}_{\tau=5}$	$1.15 \text{ ms}_{\tau=5}$	$58.8 \text{ ms}_{\tau=50}$	$10.2 \text{ ms}_{\tau=5}$	$10.3 \text{ ms}_{\tau=5}$	$94.6 \text{ ms}_{\tau=50}$
	Greedy oracle (Alg. 1)	801 μs	826 μs	49.6 ms	9.19 ms	9.29 ms	65.1 ms
	Cost-aware oracle	876 μs	892 μs	49.9 ms	9.26 ms	9.42 ms	65.3 ms
	Learned router (ours)	903 μs	911 μs	55.6 ms	9.35 ms	9.55 ms	65.6 ms
{NO, Jacobi (0.67)}	Solver only	$>4.52 \text{ s}^{164}$	$>4.52 \text{ s}^{164}$	$>4.52 \text{ s}^{164}$	—	—	—
	HINTS ($\tau=25$)	2.85 ms	2.85 ms	5.46 ms	—	—	—
	HINTS (best τ)	$1.16 \text{ ms}_{\tau=5}$	$1.18 \text{ ms}_{\tau=5}$	$4.21 \text{ ms}_{\tau=10}$	—	—	—
	Greedy oracle (Alg. 1)	787 μs	810 μs	3.70 ms	—	—	—
	Cost-aware oracle	802 μs	823 μs	3.81 ms	—	—	—
	Learned router (ours)	831 μs	850 μs	3.86 ms	—	—	—
{NO, GS}	Solver only	$>13.63 \text{ s}^{164}$	$>13.63 \text{ s}^{164}$	$>13.63 \text{ s}^{164}$	$>34.18 \text{ s}^{116}$	$>34.18 \text{ s}^{116}$	$>34.18 \text{ s}^{116}$
	HINTS ($\tau=25$)	6.48 ms	6.48 ms	23.7 ms	29.7 ms	29.7 ms	64.5 ms
	HINTS (best τ)	$1.73 \text{ ms}_{\tau=5}$	$1.74 \text{ ms}_{\tau=5}$	$16.5 \text{ ms}_{\tau=5}$	$12.7 \text{ ms}_{\tau=5}$	$12.8 \text{ ms}_{\tau=5}$	$37.8 \text{ ms}_{\tau=10}$
	Greedy oracle (Alg. 1)	812 μs	831 μs	5.87 ms	9.24 ms	9.40 ms	38.7 ms
	Cost-aware oracle	802 μs	833 μs	5.33 ms	9.20 ms	9.32 ms	35.1 ms
	Learned router (ours)	827 μs	852 μs	5.40 ms	9.20 ms	9.30 ms	35.9 ms
{NO, SymGS}	Solver only	$>23.15 \text{ s}^{164}$	$>23.15 \text{ s}^{164}$	$>23.15 \text{ s}^{164}$	—	—	—
	HINTS ($\tau=25$)	10.3 ms	10.3 ms	20.5 ms	—	—	—
	HINTS (best τ)	$2.40 \text{ ms}_{\tau=5}$	$2.40 \text{ ms}_{\tau=5}$	$7.45 \text{ ms}_{\tau=5}$	—	—	—
	Greedy oracle (Alg. 1)	811 μs	832 μs	5.63 ms	—	—	—
	Cost-aware oracle	826 μs	844 μs	5.59 ms	—	—	—
	Learned router (ours)	856 μs	871 μs	5.47 ms	—	—	—
{NO, SOR (1.5)}	Solver only	$>13.80 \text{ s}^{164}$	$>13.80 \text{ s}^{164}$	$>13.80 \text{ s}^{164}$	$>33.90 \text{ s}^{116}$	$>33.90 \text{ s}^{116}$	$>33.90 \text{ s}^{116}$
	HINTS ($\tau=25$)	6.62 ms	6.62 ms	17.9 ms	29.6 ms	29.7 ms	60.2 ms
	HINTS (best τ)	$1.78 \text{ ms}_{\tau=5}$	$1.80 \text{ ms}_{\tau=5}$	$12.8 \text{ ms}_{\tau=10}$	$12.7 \text{ ms}_{\tau=5}$	$12.8 \text{ ms}_{\tau=5}$	$51.7 \text{ ms}_{\tau=10}$
	Greedy oracle (Alg. 1)	813 μs	843 μs	7.60 ms	9.21 ms	9.33 ms	43.3 ms
	Cost-aware oracle	823 μs	862 μs	6.71 ms	9.12 ms	9.27 ms	34.2 ms
	Learned router (ours)	845 μs	869 μs	7.68 ms	9.18 ms	9.93 ms	44.1 ms

Table 18: Anisotropic diffusion: paired median speedup of the learned router at $\varepsilon = h^2$ and 10^{-8} over HINTS ($\tau=25$) and the best fixed period, with one-sided Wilcoxon p -values (no classical baselines without corrector are run for this equation).

Pairing	ε	128×128		256×256	
		vs. HINTS-25	vs. best τ	vs. HINTS-25	vs. best τ
{NO, Jacobi}	h^2	$3.16 \times (< 10^{-10})$	$1.25 \times (10^{-9})$	$1.58 \times (10^{-5})$	$1.11 \times (10^{-4})$
	10^{-8}	$1.44 \times (< 10^{-10})$	$1.21 \times (< 10^{-10})$	$1.93 \times (10^{-5})$	$1.80 \times (10^{-5})$
{NO, Jacobi (0.67)}	h^2	$3.31 \times (< 10^{-10})$	$1.34 \times (< 10^{-10})$	—	—
	10^{-8}	$1.43 \times (< 10^{-10})$	$1.17 \times (10^{-9})$	—	—
{NO, GS}	h^2	$7.61 \times (< 10^{-10})$	$1.99 \times (< 10^{-10})$	$3.18 \times (10^{-5})$	$1.37 \times (10^{-5})$
	10^{-8}	$4.05 \times (< 10^{-10})$	$2.74 \times (< 10^{-10})$	$2.26 \times (10^{-5})$	$1.64 \times (10^{-5})$
{NO, SymGS}	h^2	$11.8 \times (< 10^{-10})$	$2.72 \times (< 10^{-10})$	—	—
	10^{-8}	$3.71 \times (< 10^{-10})$	$1.36 \times (< 10^{-10})$	—	—
{NO, SOR (1.5)}	h^2	$7.36 \times (< 10^{-10})$	$2.04 \times (< 10^{-10})$	$2.98 \times (10^{-5})$	$1.38 \times (10^{-5})$
	10^{-8}	$2.40 \times (< 10^{-10})$	$1.82 \times (< 10^{-10})$	$1.54 \times (10^{-5})$	$1.35 \times (10^{-5})$




Figure 8: Anisotropic diffusion: corrector usage per iteration of the learned router, the cost-aware oracle and HINTS ($\tau=25$) for every pairing (columns) and grid (rows).

Table 19: Anisotropic diffusion: median iterations (and corrector calls) to $\varepsilon = h^2$.

Pairing	128 × 128			256 × 256		
	HINTS-25	oracle	router	HINTS-25	oracle	router
{NO, Jacobi}	25 (1)	1 (1)	1 (1)	25 (1)	1 (1)	1 (1)
{NO, Jacobi (0.67)}	25 (1)	1 (1)	1 (1)	—	—	—
{NO, GS}	25 (1)	1 (1)	1 (1)	25 (1)	1 (1)	1 (1)
{NO, SymGS}	25 (1)	1 (1)	1 (1)	—	—	—
{NO, SOR (1.5)}	25 (1)	1 (1)	1 (1)	25 (1)	1 (1)	1 (1)

Table 20: Anisotropic diffusion: measured per-iteration costs and macro-action sizes.

Pairing	128 × 128			256 × 256		
	c_{solver}	c_{NO}	m_{solver}	c_{solver}	c_{NO}	m_{solver}
{NO, Jacobi}	63 μs	695 μs	11	235 μs	8.89 ms	38
{NO, Jacobi (0.67)}	64 μs	692 μs	11	—	—	—
{NO, GS}	215 μs	711 μs	3	826 μs	8.91 ms	11
{NO, SymGS}	382 μs	711 μs	2	—	—	—
{NO, SOR (1.5)}	219 μs	730 μs	3	777 μs	8.80 ms	11

Table 21: Anisotropic diffusion, 128×128 : iteration-based metrics (conventions as in Table 9).

Method	$\ e_h^{(T)}\ /\ u_h\ $	AUC	p
{NO, Jacobi}			
Solver only	9.96e-01 (1.5e-04)	2.99e+02 (2.3e-02)	$< 10^{-10}$
HINTS ($\tau=25$)	7.65e-07 (2.1e-07)	2.40e+01 (1.7e-04)	$< 10^{-10}$
Learned router (ours)	7.45e-07 (2.0e-07)	6.80e-04 (1.8e-04)	-
Cost-aware oracle	7.45e-07 (2.0e-07)	6.80e-04 (1.8e-04)	-
{NO, Jacobi (0.67)}			
Solver only	9.97e-01 (1.0e-04)	3.00e+02 (1.6e-02)	$< 10^{-10}$
HINTS ($\tau=25$)	3.34e-10 (4.6e-11)	2.40e+01 (1.2e-04)	$< 10^{-10}$
Learned router (ours)	8.97e-10 (7.4e-12)	2.82e-04 (6.6e-05)	-
Cost-aware oracle	8.96e-10 (7.4e-12)	2.82e-04 (6.6e-05)	-
{NO, GS}			
Solver only	9.92e-01 (2.7e-04)	2.99e+02 (4.3e-02)	$< 10^{-10}$
HINTS ($\tau=25$)	9.79e-10 (1.4e-12)	2.40e+01 (3.4e-04)	$< 10^{-10}$
Learned router (ours)	8.15e-10 (1.3e-11)	1.95e-04 (4.6e-05)	-
Cost-aware oracle	8.02e-10 (1.5e-11)	1.95e-04 (4.6e-05)	-
{NO, SymGS}			
Solver only	9.84e-01 (4.8e-04)	2.98e+02 (7.7e-02)	$< 10^{-10}$
HINTS ($\tau=25$)	7.48e-10 (3.2e-11)	2.40e+01 (6.5e-04)	$< 10^{-10}$
Learned router (ours)	7.53e-10 (2.2e-11)	1.23e-04 (2.9e-05)	-
Cost-aware oracle	7.54e-10 (1.8e-11)	1.23e-04 (2.9e-05)	-
{NO, SOR (1.5)}			
Solver only	9.76e-01 (6.8e-04)	2.96e+02 (1.1e-01)	$< 10^{-10}$
HINTS ($\tau=25$)	9.09e-10 (1.6e-11)	2.40e+01 (9.5e-04)	$< 10^{-10}$
Learned router (ours)	8.67e-10 (9.3e-12)	2.87e-04 (6.9e-05)	-
Cost-aware oracle	8.63e-10 (9.4e-12)	2.87e-04 (6.9e-05)	-

Table 22: Anisotropic diffusion, 128×128 : five-seed router-training trials (conventions as in Table 10).

Pairing	time to h^2 over 5 seeds	identical decisions	speedup vs. HINTS-25	seeds with $p < 0.01$
{NO, Jacobi}	0.76 $\pm$ 0.00 ms	100%	2.97–2.97 $\times$	5/5
{NO, Jacobi (0.67)}	0.76 $\pm$ 0.00 ms	100%	3.01–3.01 $\times$	5/5
{NO, GS}	0.78 $\pm$ 0.00 ms	100%	7.63–7.63 $\times$	5/5
{NO, SymGS}	0.78 $\pm$ 0.00 ms	100%	12.77–12.77 $\times$	5/5
{NO, SOR (1.5)}	0.80 $\pm$ 0.00 ms	100%	7.58–7.58 $\times$	5/5




Figure 9: Anisotropic diffusion, 128×128 : routing decisions (dark: corrector call) of the learned router (top row) and the cost-aware oracle (bottom row) over the first 60 iterations of every test instance.

936 E.5 Solver ensembles

937 **When can a second classical solver help?** The cost-aware oracle is a per-decision comparison of
938 macro-actions, so a router over $\text{NO} \cup \mathcal{W}$ can only beat the best pairwise router of its members if,

along the way to the tolerance, the error passes through regimes in which *different* members give the largest reduction per unit cost. We screened this ceiling exhaustively before training any router: for every subset $\mathcal{W}$ of up to three members of {Jacobi, Jacobi (0.67), GS, SymGS, SOR (1.5), multigrid} (multigrid on the isotropic equations only), every equation and grid, we ran the cost-aware oracle over $\text{NO} \cup \mathcal{W}$ on the test instances and compared its work-unit time to that of the oracle over each member alone (Table 23). Two facts emerge. At $\varepsilon = h^2$ no ensemble beats the best single solver anywhere: one corrector call followed by one or two macro-actions of the cheapest smoother already reaches truncation-level accuracy, so there is no second regime to exploit. At $\varepsilon = 10^{-8}$, after the corrector has removed the smooth band, the smoother has to damp the whole band $|k| > 31$, whose sub-bands favour different solvers: undamped Jacobi is the cheapest per unit cost on the modes near the band edge (it contracts them by $\cos(2\pi \cdot 32/N)$ -like factors that damped Jacobi cannot match), but it does not damp the near-Nyquist modes at all, which damped Jacobi and GS remove in a few sweeps. The oracle over {Jacobi, Jacobi (0.67)}—a *nonstationary* Jacobi iteration whose damping parameter is chosen per macro-action—is therefore faster than every single solver, by 13–28% on four of the six (equation, grid) settings at 128^2 and 256^2 , and adding a third member (GS) changes nothing, i.e., the gain saturates. At 512^2 multigrid dominates every subset (its cycle removes all sub-bands at once), so the oracle over any ensemble containing it equals the oracle over {NO, multigrid}: adding members neither helps nor hurts. The “difficulty” that matters for ensembles is thus not the equation but the width of the spectrum left to the classical members after the corrector, which grows with the grid (fixed band) and with the tolerance.

Nested ensembles. Table 24 evaluates the learned router over the nested sets {Jacobi}, {Jacobi (0.67)}, {Jacobi, Jacobi (0.67)} and {Jacobi, Jacobi (0.67), GS} on every equation at 128^2 and 256^2 ; the single-member routers are the pairwise routers of Sections E.2 to E.4, evaluated in the same session on the same instances, and all members keep the macro-action sizes of their pairwise runs. The table reports live and work-unit times to h^2 , 10^{-6} and 10^{-8} for the router and its oracle, and the paired speedup at 10^{-8} over the best single-member router with its p -value (work units are the primary measure here: the nested runs share the machine with other jobs, so their live times are uniformly inflated relative to Sections E.2 to E.4, while work units are determined by the decision sequence alone). Two properties are visible. *Monotonicity*: the router over a larger set is never slower than the best of its members beyond timer noise, because the oracle ignores a member that never gives the largest reduction per unit cost and the router imitates it (adding GS to {Jacobi, Jacobi (0.67)} leaves the decision sequence unchanged); the earlier ensembles of Table 25 show the same for four- and five-member sets at h^2 . *Strict separation*: at 10^{-8} the two-member router is significantly faster than every single-member router in 2 of the 5 ensemble cells (paired Wilcoxon $p < 0.01$; speedups $1.08 \times$ – $1.14 \times$ over the best single solver), reproducing its oracle, and it is never significantly slower. The learned policy is the nonstationary damping schedule described above: undamped Jacobi macro-actions while the band-edge modes dominate the residual, damped Jacobi once the near-Nyquist modes do, which the router detects from the decay rate of the residual over the last macro-action.

Table 23: Oracle-level screening of ensembles: for every equation and grid, the fastest single solver and the fastest ensemble of up to three members under the cost-aware oracle (median work-unit time over 16 test instances), at $\varepsilon = h^2$ and 10^{-8} ; the ratio is the speedup of the ensemble over the best single solver, and the last column is the ratio for {Jacobi, Jacobi (0.67)}.

Equation	N	$\varepsilon = h^2$			$\varepsilon = 10^{-8}$			{Jacobi, Jacobi (0.67)}
		best single	best ensemble	ratio	best single	best ensemble	ratio	
Poisson	128^2	Jacobi (772 μ s)	{Jacobi, GS} (772 μ s)	1.00 $\times$	Jacobi (0.67) (2.51 ms)	{Jacobi, Jacobi(0.67)} (2.45 ms)	1.02 $\times$	1.02 $\times$
	256^2	Jacobi (2.86 ms)	{Jacobi, Jacobi(0.67)} (2.86 ms)	1.00 $\times$	Jacobi (0.67) (14.2 ms)	{Jacobi, Jacobi(0.67)} (11.1 ms)	1.28$\times$	1.28 $\times$
	512^2	Multigrid (36.5 ms)	{Jacobi, Multigrid} (36.5 ms)	1.00 $\times$	Multigrid (90.8 ms)	{Jacobi, Multigrid} (90.8 ms)	1.00 $\times$	–
Convection–diffusion	128^2	Jacobi (923 μ s)	{Jacobi, GS} (923 μ s)	1.00 $\times$	Jacobi (0.67) (3.06 ms)	{Jacobi, Jacobi(0.67)} (2.71 ms)	1.13$\times$	1.13 $\times$
	256^2	Jacobi (5.49 ms)	{Jacobi, Jacobi(0.67)} (5.49 ms)	1.00 $\times$	Multigrid (24.1 ms)	{Jacobi, Jacobi(0.67)} (19.3 ms)	1.25$\times$	1.25 $\times$
	512^2	Multigrid (54.3 ms)	{Jacobi, Multigrid} (54.3 ms)	1.00 $\times$	Multigrid (103.7 ms)	{Jacobi, Multigrid} (103.7 ms)	1.00 $\times$	–
Anisotropic diffusion	128^2	Jacobi (765 μ s)	{Jacobi, Jacobi(0.67)} (765 μ s)	1.00 $\times$	Jacobi (0.67) (3.58 ms)	{Jacobi, Jacobi(0.67)} (3.06 ms)	1.17$\times$	1.17 $\times$

Table 24: Nested ensembles: median time to $\varepsilon \in \{h^2, 10^{-6}, 10^{-8}\}$ of the learned router (live and work units) and of the cost-aware oracle (work units) over $\text{NO} \cup \mathcal{W}$, for the single members and the nested two- and three-member sets, all evaluated in one session on the same instances (64 at 128^2 , 32 at 256^2 , 16 for anisotropic diffusion at 256^2). The last columns give the paired speedup at 10^{-8} over the best single-member router (resp. oracle) with the one-sided Wilcoxon p -value; bold: $\geq 1.1 \times$ with $p < 0.01$.

Equation	N	$\mathcal{W}$	$\varepsilon = h^2$			$\varepsilon = 10^{-6}$			$\varepsilon = 10^{-8}$			$\varepsilon = 10^{-8}$: vs. best single solver	
			router	router (WU)	oracle (WU)	router	router (WU)	oracle (WU)	router	router (WU)	oracle (WU)	router	oracle
Poisson	128^2	{Jacobi}	1.48 ms	743 μ s	732 μ s	5.49 ms	2.21 ms	2.19 ms	342.7 ms ¹¹⁸	116.1 ms ¹¹⁸	115.2 ms	—	—
		{Jacobi(0.67)}	1.45 ms	747 μ s	737 μ s	2.13 ms	1.16 ms	1.15 ms	4.46 ms	2.33 ms	2.31 ms	(best single)	—
		{GS}	1.50 ms	901 μ s	891 μ s	3.60 ms	2.56 ms	2.53 ms	6.62 ms	4.41 ms	4.37 ms	—	—
		{Jacobi, Jacobi(0.67)}	1.37 ms	747 μ s	735 μ s	1.94 ms	1.16 ms	1.15 ms	4.93 ms	2.59 ms	2.56 ms	$1.00 \times (0.911)$	$1.00 \times (10^{-5})$
		{Jacobi, Jacobi(0.67), GS}	1.46 ms	749 μ s	738 μ s	2.42 ms	1.21 ms	1.15 ms	5.46 ms	2.59 ms	2.57 ms	$0.92 \times (\text{slower}, 10^{-8})$	$1.00 \times (10^{-4})$
Convection-diffusion	128^2	{Jacobi}	1.31 ms	842 μ s	832 μ s	11.3 ms	6.00 ms	5.95 ms	368.9 ms ¹²⁵	213.9 ms ¹²⁵	212.2 ms	—	—
		{Jacobi(0.67)}	1.44 ms	913 μ s	903 μ s	2.39 ms	1.52 ms	1.51 ms	4.45 ms	2.94 ms	2.92 ms	(best single)	—
		{GS}	1.58 ms	1.11 ms	1.10 ms	3.81 ms	2.85 ms	2.82 ms	6.13 ms	4.94 ms	4.90 ms	—	—
		{Jacobi, Jacobi(0.67)}	1.75 ms	879 μ s	903 μ s	3.12 ms	1.49 ms	1.47 ms	5.23 ms	2.70 ms	2.67 ms	$1.08 \times (< 10^{-10})$	$1.11 \times (< 10^{-10})$
		{Jacobi, Jacobi(0.67), GS}	1.37 ms	887 μ s	911 μ s	2.55 ms	1.61 ms	1.48 ms	4.39 ms	2.85 ms	2.68 ms	$0.99 \times (\text{slower}, 0.155)$	$1.10 \times (10^{-10})$
Anisotropic diffusion	128^2	{Jacobi}	1.28 ms	762 μ s	752 μ s	2.04 ms	1.36 ms	1.19 ms	68.7 ms ¹¹⁰	41.2 ms ¹¹⁰	40.9 ms	—	—
		{Jacobi(0.67)}	1.45 ms	760 μ s	750 μ s	2.57 ms	1.46 ms	1.42 ms	5.31 ms	3.26 ms	3.23 ms	(best single)	—
		{Jacobi, Jacobi(0.67)}	1.28 ms	762 μ s	751 μ s	2.08 ms	1.40 ms	1.19 ms	4.61 ms	2.87 ms	2.84 ms	$1.14 \times (< 10^{-10})$	$1.18 \times (< 10^{-10})$

Larger ensembles at truncation level. Table 25 repeats the ensembles of Section 7 ({Jacobi, GS} up to all five solvers) at 128^2 and $\varepsilon = h^2$, where the screening predicts no gain: the ensemble router matches the best pairwise router of its members to within about $\pm 20\%$ in live time (ratio of medians $0.82 \times$ – $1.14 \times$ over the 12 cells) and within 1.5% in work units, without knowing in advance which member is fastest, and is $380 \times$ – $> 10^3 \times$ faster than the fastest member alone; Table 26 lists how its iterations are distributed over the members. Its oracle is up to 25% faster in a few cells because, at h^2 , several cost-equalised smoothers are near-ties and the oracle’s choice among them depends on the spectral content of the current error, which the router’s residual-norm features reflect only indirectly (agreement with the oracle 55–83% on its own rollouts against 88–99% for pairwise routers). More capacity or imitation data does not change this: routers with twice the hidden width, three times the DAgger rounds and twice the oracle rollouts (Table 28) take the same number of iterations and corrector calls to h^2 as the default routers on 53%–91% of the test instances and are significantly faster in 2 and significantly slower in 3 of the 8 cells ($p < 0.05$); the residual disagreements are near-ties whose resolution does not affect the time to h^2 . Deciding *when* to call the corrector—the decision that dominates the wall-clock gains—is learned equally well in the ensembles.

Table 25: Ensembles of Section 7 at 128×128 and $\varepsilon = h^2$: median live time of the best member solver alone, of the best pairwise router among the members (paired speedup of the ensemble router over it in parentheses), and of the router and the cost-aware oracle over $\text{NO} \cup \mathcal{W}$ (64 instances; pairwise baselines from the pairwise runs on the same instances).

Equation	$\mathcal{W}$	Best solver only	Best pairwise router	Router($\text{NO} \cup \mathcal{W}$)	Oracle($\text{NO} \cup \mathcal{W}$)
Poisson	{Jacobi, GS}	1.15 s (Jacobi)	987 μ s (Jacobi)	964 μ s ($1.09 \times$ vs pairwise)	834 μ s
	{Jacobi, GS, SymGS}	1.15 s (Jacobi)	987 μ s (Jacobi)	951 μ s ($1.07 \times$ vs pairwise)	854 μ s
	{Jacobi, GS, SymGS, Jacobi(0.67)}	1.15 s (Jacobi)	912 μ s (Jacobi(0.67))	868 μ s ($0.98 \times$ vs pairwise)	842 μ s
	{Jacobi, GS, SymGS, Jacobi(0.67), SOR(1.5)}	600.9 ms (SOR(1.5))	912 μ s (Jacobi(0.67))	1.12 ms ($0.86 \times$ vs pairwise)	900 μ s
ConvDiff	{Jacobi, GS}	1.51 s (Jacobi)	1.14 ms (Jacobi)	1.25 ms ($1.09 \times$ vs pairwise)	974 μ s
	{Jacobi, GS, SymGS}	1.51 s (Jacobi)	1.14 ms (Jacobi)	1.01 ms ($1.16 \times$ vs pairwise)	996 μ s
	{Jacobi, GS, SymGS, Jacobi(0.67)}	1.51 s (Jacobi)	1.10 ms (Jacobi(0.67))	996 μ s ($1.00 \times$ vs pairwise)	1,07 ms
	{Jacobi, GS, SymGS, Jacobi(0.67), SOR(1.5)}	383.6 ms (SOR(1.5))	1.10 ms (Jacobi(0.67))	1.01 ms ($1.06 \times$ vs pairwise)	1,08 ms
AnisoDiff	{Jacobi, GS}	>4.62 s [†] (none)	852 μ s (GS)	898 μ s ($1.01 \times$ vs pairwise)	870 μ s
	{Jacobi, GS, SymGS}	>4.62 s [†] (none)	852 μ s (GS)	880 μ s ($1.03 \times$ vs pairwise)	857 μ s
	{Jacobi, GS, SymGS, Jacobi(0.67)}	>4.52 s [†] (none)	850 μ s (Jacobi(0.67))	880 μ s ($1.00 \times$ vs pairwise)	858 μ s
	{Jacobi, GS, SymGS, Jacobi(0.67), SOR(1.5)}	>4.52 s [†] (none)	850 μ s (Jacobi(0.67))	890 μ s ($1.00 \times$ vs pairwise)	863 μ s

Table 26: Ensembles at 128×128 : mean (standard deviation over instances) fraction of the iterations up to the h^2 crossing that the ensemble router spends in each member and in the corrector.

Equation	$\mathcal{W}$	Jacobi	GS	SymGS	Jacobi(0.67)	SOR(1.5)	DeepONet
Poisson	{Jacobi, GS}	0.228 (0.365)	0.144 (0.216)	—	—	—	0.628 (0.386)
	{Jacobi, GS, SymGS}	0.237 (0.337)	0.043 (0.134)	0.086 (0.165)	—	—	0.634 (0.378)
	{Jacobi, GS, SymGS, Jacobi(0.67)}	0.235 (0.297)	0.008 (0.062)	0.101 (0.207)	0.000 (0.000)	—	0.656 (0.352)
	{Jacobi, GS, SymGS, Jacobi(0.67), SOR(1.5)}	0.094 (0.195)	0.020 (0.051)	0.182 (0.265)	0.046 (0.125)	0.007 (0.023)	0.650 (0.366)
ConvDiff	{Jacobi, GS}	0.089 (0.255)	0.310 (0.337)	—	—	—	0.600 (0.373)
	{Jacobi, GS, SymGS}	0.279 (0.354)	0.067 (0.156)	0.055 (0.103)	—	—	0.599 (0.374)
	{Jacobi, GS, SymGS, Jacobi(0.67)}	0.305 (0.340)	0.000 (0.000)	0.022 (0.038)	0.110 (0.214)	—	0.563 (0.410)
	{Jacobi, GS, SymGS, Jacobi(0.67), SOR(1.5)}	0.166 (0.302)	0.028 (0.127)	0.000 (0.000)	0.233 (0.375)	0.000 (0.000)	0.573 (0.399)
AnisoDiff	{Jacobi, GS}	0.173 (0.320)	0.000 (0.000)	—	—	—	0.827 (0.320)
	{Jacobi, GS, SymGS}	0.052 (0.153)	0.000 (0.000)	0.097 (0.229)	—	—	0.851 (0.273)
	{Jacobi, GS, SymGS, Jacobi(0.67)}	0.000 (0.000)	0.000 (0.000)	0.000 (0.000)	0.185 (0.339)	—	0.815 (0.339)
	{Jacobi, GS, SymGS, Jacobi(0.67), SOR(1.5)}	0.000 (0.000)	0.000 (0.000)	0.000 (0.000)	0.185 (0.339)	0.000 (0.000)	0.815 (0.339)

Table 27: Ensembles at 128×128 : one-sided Wilcoxon p -values at $\varepsilon = h^2$ for the ensemble router against the fastest member solver alone (censored runs at their lower bound), against the best pairwise router in both directions, and against the cost-aware oracle over the same ensemble (alternative: the ensemble router is slower).

Equation	$\mathcal{W}$	vs. best solver only	vs. best pairwise router (faster)	vs. best pairwise router (slower)	vs. oracle($\text{NO} \cup \mathcal{W}$) (slower)
Poisson	$\{\text{Jacobi}, \text{GS}\}$	$< 10^{-10}$	10^{-7}	1.000	0.040
	$\{\text{Jacobi}, \text{GS}, \text{SymGS}\}$	$< 10^{-10}$	0.001	0.999	0.001
	$\{\text{Jacobi}, \text{GS}, \text{SymGS}, \text{Jacobi}(0.67)\}$	$< 10^{-10}$	0.856	0.144	10^{-4}
	$\{\text{Jacobi}, \text{GS}, \text{SymGS}, \text{Jacobi}(0.67), \text{SOR}(1.5)\}$	$< 10^{-10}$	1.000	10^{-8}	10^{-8}
ConvDiff	$\{\text{Jacobi}, \text{GS}\}$	$< 10^{-10}$	10^{-4}	1.000	10^{-5}
	$\{\text{Jacobi}, \text{GS}, \text{SymGS}\}$	$< 10^{-10}$	10^{-10}	1.000	0.492
	$\{\text{Jacobi}, \text{GS}, \text{SymGS}, \text{Jacobi}(0.67)\}$	$< 10^{-10}$	0.783	0.217	0.001
	$\{\text{Jacobi}, \text{GS}, \text{SymGS}, \text{Jacobi}(0.67), \text{SOR}(1.5)\}$	$< 10^{-10}$	0.047	0.953	0.566
AnisoDiff	$\{\text{Jacobi}, \text{GS}\}$	$< 10^{-10}$	0.315	0.685	0.014
	$\{\text{Jacobi}, \text{GS}, \text{SymGS}\}$	$< 10^{-10}$	0.055	0.945	10^{-4}
	$\{\text{Jacobi}, \text{GS}, \text{SymGS}, \text{Jacobi}(0.67)\}$	$< 10^{-10}$	0.551	0.449	0.064
	$\{\text{Jacobi}, \text{GS}, \text{SymGS}, \text{Jacobi}(0.67), \text{SOR}(1.5)\}$	$< 10^{-10}$	0.633	0.367	0.012

Table 28: Control on router capacity for the ensembles (128×128 , 64 instances): the default ensemble router (hidden width 64, 2 DAgger rounds, 128 oracle rollouts) against one with hidden width 128, 6 DAgger rounds and 256 rollouts. Times to h^2 are median work units net of the per-decision cost; “same decisions” is the share of instances on which both routers take the same number of iterations and corrector calls to h^2 ; p is the one-sided Wilcoxon p -value on the paired net times; “agreement” is each router’s agreement with the oracle on its own rollouts during training.

Equation	$\mathcal{W}$	default router	larger router	oracle	larger / default (medians)	same decisions	p (larger faster / slower)	agreement (default / larger)
Poisson	$\{\text{Jacobi}, \text{GS}\}$	858 μs	758 μs	758 μs	0.884 $\times$	84%	0.933 / 0.067	80% / 89%
	$\{\text{Jacobi}, \text{GS}, \text{SymGS}\}$	742 μs	742 μs	742 μs	1.000 $\times$	89%	0.525 / 0.494	73% / 73%
	$\{\text{Jacobi}, \text{GS}, \text{SymGS}, \text{Jacobi}(0.67)\}$	740 μs	740 μs	740 μs	1.000 $\times$	61%	0.958 / 0.042	55% / 76%
	$\{\text{Jacobi}, \text{GS}, \text{SymGS}, \text{Jacobi}(0.67), \text{SOR}(1.5)\}$	732 μs	732 μs	732 μs	1.000 $\times$	62%	10^{-4} / 0.999	65% / 64%
ConvDiff	$\{\text{Jacobi}, \text{GS}\}$	1.07 ms	887 μs	854 μs	0.829 $\times$	69%	0.005 / 0.995	83% / 86%
	$\{\text{Jacobi}, \text{GS}, \text{SymGS}\}$	854 μs	890 μs	856 μs	1.042 $\times$	91%	0.995 / 0.005	58% / 65%
	$\{\text{Jacobi}, \text{GS}, \text{SymGS}, \text{Jacobi}(0.67)\}$	888 μs	888 μs	932 μs	1.000 $\times$	56%	0.459 / 0.541	57% / 70%
	$\{\text{Jacobi}, \text{GS}, \text{SymGS}, \text{Jacobi}(0.67), \text{SOR}(1.5)\}$	892 μs	892 μs	936 μs	1.000 $\times$	53%	$1.000 / 10^{-8}$	65% / 74%

992 E.6 Overheads and amortisation

993 Table 29 lists the measured cost of every operation as a function of N , including one decision of
994 the LSTM router of Section 7 (full-state input of size $4N^2$, hidden size 256, 4 layers): at 128^2 one
995 LSTM decision costs 2.4 ms, about 46 Jacobi sweeps, which is why the wall-clock study uses the
996 scalar-feature router ($\approx 5 \mu\text{s}$, independent of N). The learned router’s decisions therefore account
997 for well under 1% of its solve time, and its corrector calls for 50–70% of the time to h^2 (one or two
998 calls). Table 30 accounts for training: the corrector (shared with HINTS) is fit in minutes, the router
999 in seconds to a few minutes per pairing, and the break-even number of solves after which the router’s
1000 own training cost is recovered relative to HINTS ($\tau=25$) with the same corrector is in the thousands
1001 at 128^2 and in the hundreds on the larger grids.

Table 29: Measured single-thread cost of one operation as a function of the grid size (Poisson; median of repeated calls). Multigrid uses a direct solve below 64^2 and is not listed there.

Operation	$N = 32$	$N = 64$	$N = 128$	$N = 256$	$N = 512$
Jacobi iteration	14 μs	25 μs	53 μs	191 μs	1.07 ms
GS iteration	25 μs	64 μs	209 μs	820 μs	3.74 ms
SymGS iteration	38 μs	109 μs	378 μs	1.48 ms	6.85 ms
Multigrid V(2,2) cycle	—	390 μs	1.35 ms	5.27 ms	25.35 ms
DeepONet corrector call	—	—	663 μs	1.13 ms	3.62 ms
Router decision (ours)	5 μs	5 μs	5 μs	5 μs	5 μs
LSTM router decision (Section 7)	236 μs	696 μs	2.43 ms	9.55 ms	37.46 ms

Table 30: Training-cost amortisation for the GS pairing: corrector training (data generation and least-squares fit, CPU), router training (oracle rollouts and DAGger rounds), median wall-clock saved per solve at $\varepsilon = h^2$ relative to HINTS ($\tau=25$) and to the solver alone, and the number of solves after which the router’s training cost is recovered relative to HINTS.

Equation	N	Corrector data + fit	Router training (GS pairing)	Saving per solve vs HINTS ($\tau=25$)	Break-even solves (router)	Saving per solve vs solver only
Poisson	128^2	82 s	8 s	5.64 ms	1419	1.83 s
	256^2	98 s	34 s	18.1 ms	1883	28.82 s
	512^2	504 s	229 s	11.4 ms	20139	—
Convection–diffusion	128^2	79 s	8 s	6.92 ms	1157	1.63 s
	256^2	106 s	33 s	24.0 ms	1376	30.05 s
	512^2	493 s	254 s	34.2 ms	7423	—

E.7 Verification of the theoretical assumptions

Theorem 4.1 and Propositions 4.2–4.3 assume that every error propagation map $G_j = I - C_j \circ \mathcal{L}_h^a$ is Lipschitz with constant $\rho_j < 1$ (respectively ≤ 1 with a shared eigenbasis), zero-preserving, and—for postfix monotonicity—invertible; Theorem 5.1 assumes that the per-step errors $\tilde{c}_j$ are bounded above and that at least one solver cannot annihilate the error in one step. Table 31 verifies these conditions numerically for every operation of the study on the mean-free subspace (the constant mode is the null space of the periodic operator and is excluded from the data and from every iteration): the Euclidean Lipschitz constant $\|G_j\|_2$ (power iteration on $G_j^\top G_j$; exact from the Fourier symbol for Jacobi), the Lipschitz constant $\|G_j\|_A = \|A_s^{1/2} G_j A_s^{-1/2}\|_2$ in the energy norm of the symmetric part A_s of the discrete operator, the same for the macro-action $G_j^{m_j}$, the spectral radius, the smallest observed contraction $\sigma_{\min} = \min \|G_j x\|/\|x\|$, $\|C_j(0)\|$, and the commutator $\|G_j G_{\text{NO}} - G_{\text{NO}} G_j\|$ on random vectors. Table 32 evaluates the quantities of Theorem 4.1 and Theorem 5.1 along the cost-aware oracle’s own path to $\varepsilon = 10^{-8}$ (a horizon of tens of macro-actions; to h^2 the path has one or two steps).

What holds, and in which norm. (i) Every classical solver is zero-preserving exactly and contracts in the sense that matters for convergence: the spectral radius of G_j is below 1 in every setting (at most 0.999997), and so is the energy-norm Lipschitz constant of the solver and of its macro-action (at most 0.999912). The *Euclidean* Lipschitz constant, however, exceeds 1 for Gauss–Seidel, SOR and symmetric GS (up to 1.89 for SOR), whose iteration matrices are non-normal; the proofs of Propositions 4.2–4.3 use only the triangle inequality and the composition rule for Lipschitz constants, so they hold verbatim with $\|\cdot\|_2$ replaced by any norm, and in the energy norm the Lipschitz condition is satisfied by all classical solvers in all settings. For (damped) Jacobi both constants agree (the maps are Fourier-diagonal); the values are close to 1 because the slowest modes (the smoothest ones, and the near-Nyquist mode for undamped Jacobi) are barely damped by a point smoother—which is exactly why the corrector is needed. (ii) The corrector’s map is zero-preserving, contracts every band mode to 1.7×10^{-3} or better (on average to 10^{-4} – 10^{-6}) and leaves the modes outside the band untouched, so its Lipschitz constant is exactly 1 (not < 1): Theorem 4.2 in its strict form does not apply to the corrector, but Theorem 4.3 does, and its second condition—a shared eigenbasis—holds exactly for the ensembles $\{\text{NO}, \text{Jacobi}\}$, $\{\text{NO}, \text{Jacobi} (0.67)\}$ and $\{\text{NO}, \text{Jacobi}, \text{Jacobi} (0.67)\}$ (all maps are diagonal in the Fourier basis; measured commutators $\approx 10^{-6}$, at the corrector’s float32 precision), so h is supermodular for them and Theorem 4.1 holds with $\alpha = 1$; for GS, SOR, symmetric GS and multigrid the commutators are small but nonzero (10^{-3} – 10^{-1}). (iii) Invertibility fails for most maps in exact arithmetic: a Gauss–Seidel, SOR or symmetric-GS sweep annihilates an error concentrated at the first grid point ($\sigma_{\min} = 0$), undamped Jacobi annihilates the mode $(N/4, N/4)$, and the corrector annihilates the band; postfix monotonicity is therefore not guaranteed, which Section 4 anticipates (“our experiments show that greedy routers remain effective even when invertibility is not met”). (iv) The worst-case ratio $\alpha(O)$ of Theorem 4.2, which only uses the Lipschitz constants, is loose here: along the oracle’s path to 10^{-8} (Table 32) $\sum_i \rho_{O_i}^2$ is close to T because each ρ is dominated by the slowest mode, so the bound ranges from 1 (the multigrid pairing, whose cycle contracts everything) to several hundred, and is infinite on the few instances where $\sum_i \rho_{O_i}^2 \geq T$. The empirically measured supermodularity ratio of Equation (7) along the same path—with S the oracle’s prefix and S' its suffix—is at most 1.000 (median 0.44) in every setting, and at most 1.000 on the short paths to h^2 : the objective behaves supermodularly on the states the greedy rule actually visits, which is the regime in which Theorem 4.1 gives the $(1 - e^{-1})$ guarantee. (v) For Theorem 5.1, the per-step errors are bounded by the current error ($\max_j \tilde{c}_j / \|e^{(t)}\|^2 \leq 1$ at every visited state) and at least one solver

1047 leaves a positive fraction of it ($\min_t \max_j \tilde{c}_j / \|e^{(t)}\|^2 > 0$), so the surrogate’s Bayes consistency
1048 applies.

Table 31: Numerical verification of the assumptions of Propositions 4.2–4.3 for every operation of the study (mean-free subspace). $\|\cdot\|_2$: Euclidean Lipschitz constant of the error propagation map; $\|\cdot\|_A$: Lipschitz constant in the energy norm of the symmetric part of $\mathcal{L}_h$, for one iteration and for the macro-action of m_j iterations; ρ : spectral radius; $\sigma_{\min}$: smallest observed contraction (a unit error at the first grid point for the triangular sweeps, the minimum over Fourier modes otherwise; for the corrector, the largest remaining fraction of a band mode); $\|C_j(0)\|$: zero preservation; last column: commutator with the corrector’s map on random unit vectors (Proposition 4.3). For (damped) Jacobi the constants are exact from the Fourier symbol (the value with the Nyquist modes included; without them undamped Jacobi has $\|G\|_2 = \rho = 0.9994$ at 128^2).

Equation	Operation	m_j	$\ I - C_j \mathcal{L}_h\ _2$	$\ I - C_j \mathcal{L}_h\ _A$	$\ (I - C_j \mathcal{L}_h)^{m_j}\ _A$	$\rho(I - C_j \mathcal{L}_h)$	$\sigma_{\min}$	$\ C_j(0)\ $	$\ [G_j, G_{\text{NO}}]\ $
Poisson, 128^2	DeepONet corrector	1	1.000000	1.000000	1.000000	1.000000	4.4e-05 (band)	0	–
	Jacobi	12	0.999398	0.999398	0.992797	0.999398	0	0	5.7e-06
	Jacobi (0.67)	12	0.999596	0.999596	0.995168	0.999596	1.4e-04	0	4.1e-06
	GS	3	1.021515	0.998595	0.996368	0.998797	0	0	3.1e-02
	SymGS	2	1.006607	0.997396	0.995187	0.997623	0	0	2.2e-02
	SOR (1.5)	3	1.449417	0.996467	0.989336	0.996411	7.2e-01	0	8.5e-02
	Multigrid	1	0.098203	0.097860	0.097860	0.096877	1.9e-03	0	8.8e-03
Poisson, 256^2	DeepONet corrector	1	1.000000	1.000000	1.000000	1.000000	4.4e-05 (band)	0	–
	Jacobi	6	0.999849	0.999849	0.999097	0.999849	0	0	1.3e-06
	Jacobi (0.67)	6	0.999899	0.999899	0.999395	0.999899	5.6e-05	0	1.2e-06
	GS	1	1.022053	0.998925	0.998925	0.999699	0	0	1.1e-02
	SymGS	1	1.005470	0.998839	0.998839	0.999401	0	0	9.4e-03
	SOR (1.5)	1	1.451367	0.998352	0.998352	0.999098	7.2e-01	0	3.8e-02
	Multigrid	1	0.102361	0.102263	0.102263	0.101344	1.4e-03	0	4.2e-03
Poisson, 512^2	DeepONet corrector	1	1.000000	1.000000	1.000000	1.000000	2.1e-05 (band)	0	–
	Jacobi	3	0.999962	0.999962	0.999887	0.999962	0	0	6.1e-07
	GS	1	1.022202	0.998983	0.998983	0.999925	0	0	4.5e-03
	SOR (1.5)	1	1.451920	0.999043	0.999043	0.999774	7.2e-01	0	1.7e-02
	Multigrid	1	0.103470	0.103434	0.103434	0.102511	1.0e-03	0	1.3e-03
Convection–diffusion, 128^2	DeepONet corrector	1	1.000000	1.000000	1.000000	1.000000	1.5e-05 (band)	0	–
	Jacobi	9	0.999398	0.999398	0.994593	0.999398	0	0	3.1e-06
	Jacobi (0.67)	9	0.999596	0.999596	0.996374	0.999596	1.4e-04	0	2.2e-06
	GS	3	1.025726	0.998460	0.995816	0.998589	0	0	3.1e-02
	SymGS	2	1.010112	0.997353	0.995168	0.997571	0	0	2.3e-02
	SOR (1.5)	3	1.554544	0.996981	0.983568	0.993973	7.5e-01	0	9.1e-02
	Multigrid	1	0.076967	0.074905	0.074905	0.074609	1.5e-03	0	6.9e-03
Convection–diffusion, 256^2	DeepONet corrector	1	1.000000	1.000000	1.000000	1.000000	1.1e-05 (band)	0	–
	Jacobi	4	0.999849	0.999849	0.999398	0.999849	0	0	7.6e-07
	Jacobi (0.67)	4	0.999899	0.999899	0.999596	0.999899	5.6e-05	0	6.8e-07
	GS	1	1.024025	0.998934	0.998934	0.999674	0	0	1.2e-02
	SymGS	1	1.008004	0.998846	0.998846	0.999398	0	0	9.8e-03
	SOR (1.5)	1	1.497088	0.998192	0.998192	0.998846	7.3e-01	0	4.1e-02
	Multigrid	1	0.084356	0.083809	0.083809	0.083459	1.2e-03	0	3.7e-03
Convection–diffusion, 512^2	DeepONet corrector	1	1.000000	1.000000	1.000000	1.000000	9.3e-06 (band)	0	–
	Jacobi	3	0.999962	0.999962	0.999887	0.999962	0	0	3.4e-07
	GS	1	1.023154	0.998983	0.998983	0.999922	0	0	4.6e-03
	SOR (1.5)	1	1.473282	0.999075	0.999075	0.999745	7.3e-01	0	1.8e-02
	Multigrid	1	0.091267	0.091201	0.091201	0.090840	9.4e-04	0	1.2e-03
Anisotropic diffusion, 128^2	DeepONet corrector	1	1.000000	1.000000	1.000000	1.000000	4.4e-04 (band)	0	–
	Jacobi	11	0.999988	0.999988	0.999989	0.999988	1.1e-16	0	2.4e-05
	Jacobi (0.67)	11	0.999992	0.999992	0.999912	0.999992	4.4e-03	0	1.6e-05
	GS	3	1.131668	0.999528	0.999537	0.999976	0	0	5.7e-02
	SymGS	2	1.102789	0.999489	0.999255	0.999953	0	0	5.5e-02
	SOR (1.5)	3	1.889774	0.999393	0.999471	0.999928	8.3e-01	0	1.5e-01
Anisotropic diffusion, 256^2	DeepONet corrector	1	1.000000	1.000000	1.000000	1.000000	1.7e-03 (band)	0	–
	Jacobi	38	0.999997	0.999997	0.999887	0.999997	1.1e-16	0	9.7e-05
	GS	11	1.131845	0.999245	0.999485	0.999994	0	0	4.0e-02
	SOR (1.5)	11	1.890225	0.999352	0.999473	0.999982	8.3e-01	0	1.1e-01

Table 32: Quantities of Theorem 4.1, Proposition 4.2 and Theorem 5.1 along the cost-aware oracle’s path to $\varepsilon = 10^{-8}$ (8 test instances; medians unless stated): number of macro-actions T , $\sum_i \rho_{O_i}^2$ with the energy-norm constants of the chosen macro-actions, the resulting $\alpha(O)$ of Proposition 4.2 (median over the instances on which $T - \sum_i \rho_{O_i}^2 > 0$; ∞ otherwise), the empirical supermodularity ratio $\hat{\alpha}$ of Equation (7) evaluated with S the oracle’s prefix and S' its suffix (maximum and median over prefixes and instances), and the normalised bounds of Theorem 5.1: the largest ratio $\max_j \tilde{c}_j / \|e^{(t)}\|^2$ over visited states ($\bar{E}$) and the smallest value of $\max_j \tilde{c}_j / \|e^{(t)}\|^2$ ($E_{\min}$).

Equation	Ensemble	T	$\sum_i \rho_{O_i}^2$	$\alpha(O)$ (Prop. 4.2)	$\hat{\alpha}$ max	$\hat{\alpha}$ median	$\max_t \max_j \tilde{c}_j / \ e^{(t)}\ ^2$	$\min_t \max_j \tilde{c}_j / \ e^{(t)}\ ^2$
Poisson, 128 ²	{Jacobi}	52	51.775	6	1.000	0.420	1.000	6.3e-02
	{Jacobi(0.67)}	4	3.981	207	1.000	0.598	1.000	9.0e-02
	{GS}	6	5.971	138	1.000	0.498	1.000	1.4e-01
	{SymGS}	5	4.971	139	1.000	0.501	1.000	1.1e-01
	{SOR(1.5)}	9	8.852	27	1.000	0.376	1.000	1.2e-01
	{Multigrid}	4	1.029	1	1.000	0.417	1.000	9.6e-06
	{Jacobi, Jacobi(0.67)}	4	3.976	167	1.000	0.612	1.000	9.0e-02
	{Jacobi, Jacobi(0.67), GS, SymGS, SOR(1.5)}	4	3.976	167	1.000	0.612	1.000	1.4e-01
Poisson, 256 ²	{Jacobi}	10	9.486	297	1.000	0.459	1.000	3.6e-01
	{Jacobi(0.67)}	12	11.988	331	1.000	0.302	1.000	4.7e-01
	{GS}	26	26.447	76	1.000	0.201	1.000	6.3e-01
	{SymGS}	14	14.471	138	1.000	0.255	1.000	5.1e-01
	{SOR(1.5)}	24	23.928	55	1.000	0.207	1.000	5.4e-01
	{Multigrid}	4	1.037	1	1.000	0.418	1.000	1.1e-05
	{Jacobi, Jacobi(0.67)}	9	8.988	322	1.000	0.389	1.000	4.3e-01
Poisson, 512 ²	{Jacobi}	55	54.988	322 (∞ on 1/8)	1.000	0.236	1.000	7.8e-01
	{GS}	86	86.328	22 (∞ on 1/8)	1.000	0.154	1.000	8.4e-01
	{SOR(1.5)}	41	40.925	52 (∞ on 1/8)	1.000	0.166	1.000	7.4e-01
	{Multigrid}	4	1.043	1 (∞ on 1/8)	1.000	0.419	1.000	1.0e-05
Convection-diffusion, 128 ²	{Jacobi}	156	154.834	2	1.000	0.391	1.000	7.9e-02
	{Jacobi(0.67)}	5	4.978	184	1.000	0.503	1.000	1.2e-01
	{GS}	6	5.967	120	1.000	0.494	1.000	1.2e-01
	{SymGS}	5	4.971	138	1.000	0.501	1.000	1.0e-01
	{SOR(1.5)}	9	8.772	18	1.000	0.359	1.000	7.5e-02
	{Multigrid}	4	1.017	1	1.000	0.417	1.000	8.2e-06
	{Jacobi, Jacobi(0.67)}	4	3.982	222	1.000	0.659	1.000	1.2e-01
	{Jacobi, Jacobi(0.67), GS, SymGS, SOR(1.5)}	4	3.982	222	1.000	0.615	1.000	1.2e-01
Convection-diffusion, 256 ²	{Jacobi}	26	25.971	185	1.000	0.464	1.000	5.0e-01
	{Jacobi(0.67)}	20	19.486	284	1.000	0.226	1.000	5.2e-01
	{GS}	28	28.444	71	1.000	0.186	1.000	5.9e-01
	{SymGS}	16	16.467	120	1.000	0.221	1.000	5.1e-01
	{SOR(1.5)}	24	23.921	50	1.000	0.207	1.000	5.2e-01
	{Multigrid}	4	1.021	1	1.000	0.418	1.000	2.8e-06
	{Jacobi, Jacobi(0.67)}	14	13.986	285	1.000	0.276	1.000	5.2e-01
Convection-diffusion, 512 ²	{Jacobi}	62	62.486	293	1.000	0.196	1.000	7.6e-01
	{GS}	94	93.813	21	1.000	0.138	1.000	8.4e-01
	{SOR(1.5)}	42	41.926	54	1.000	0.158	1.000	6.8e-01
	{Multigrid}	4	1.033	1	1.000	0.377	1.000	2.0e-06
Anisotropic diffusion, 128 ²	{Jacobi}	10	9.998	1386 (∞ on 1/8)	1.000	0.481	1.000	3.5e-01
	{Jacobi(0.67)}	5	4.999	7585 (∞ on 1/8)	1.000	0.504	1.000	3.5e-01
	{GS}	7	6.995	865 (∞ on 1/8)	1.000	0.500	1.000	3.5e-01
	{SymGS}	6	6.494	672 (∞ on 1/8)	1.000	0.435	1.000	3.5e-01
	{SOR(1.5)}	9	8.993	540 (∞ on 1/8)	1.000	0.397	1.000	3.5e-01
	{Jacobi, Jacobi(0.67)}	4	4.000	9130 (∞ on 1/8)	1.000	0.652	1.000	3.5e-01
	{Jacobi, Jacobi(0.67), GS, SymGS, SOR(1.5)}	4	4.000	9130 (∞ on 1/8)	1.000	0.639	1.000	3.5e-01
Anisotropic diffusion, 256 ²	{Jacobi}	3	3.000	17652 (∞ on 2/8)	1.000	0.608	1.000	2.9e-01
	{GS}	3	2.999	3888 (∞ on 2/8)	1.000	0.658	1.000	3.9e-01
	{SOR(1.5)}	4	3.498	1897 (∞ on 2/8)	1.000	0.578	1.000	2.2e-01

F Limitations

A primary limitation of our approach is the higher on-time computational cost requires to learn the greedy routing strategy. In contrast, the simpler baselines such as HINTS incur no training overhead. While this cost is amortized during inference time and leads to improved solution quality, it may be prohibitive in settings with limited computational resources or when rapid deployment is required. On the deployment side, Section E shows that the gains survive cost accounting: with a corrector that covers the smooth half of the spectrum and a classical smoother that covers the rest, one or two corrector calls suffice and the router’s advantage over fixed schedules is uniform across solvers, equations and tolerances. The corrector we use is specific to the constant-coefficient periodic setting (a fixed Fourier trunk basis on a coarse sensor grid), its cost grows with the fourth power of the sensor-grid size, and undamped Jacobi remains limited far below truncation level by the near-Nyquist modes that neither component damps. Our wall-clock study is restricted to single-thread CPU execution; with a corrector band that does not grow with the grid, the gains over fixed schedules shrink at 512² for the point smoothers, whereas pairing the corrector with geometric multigrid restores them (Section E). Ensembles of classical solvers pay off only where the error passes through regimes that favour different members; at truncation level a single well-chosen smoother suffices.

Additionally, the effectiveness of the learned router depends on the neural operator exhibiting heterogeneous performance across samples. When the neural operators produces predictions of relatively uniform quality, the benefit of adaptive routing diminishes and it is more practical to set a

1068 fixed schedule based on the quality of the neural operator. In practice, achieving such heterogeneity
1069 can be sensitive to training choices, and neural operators themselves can be finicky to train.

1070 **G LLM Usage**

1071 LLMs, specifically ChatGPT and Gemini, supported the writing process in an iterative manner. We
1072 drafted paragraphs and asked the models for feedback on grammar and clarity. We then incorporated
1073 selected suggestions into the writing and repeated this process until we were satisfied with the writing.

1074 The code developed for the experiments was written by the authors with the help of occasional code
1075 completions. The central components (e.g., the hybrid solver implementation and the greedy-router
1076 training pipelines) were implemented exclusively by the authors.

1077 All substantive intellectual contributions, which include ideas, theorems, and analyses, are our own.
1078 LLMs were occasionally used to verify the correctness of proofs, but all proof strategies originated
1079 from the authors and relevant literature.

1080 NeurIPS Paper Checklist

1081 The checklist is designed to encourage best practices for responsible machine learning research,
1082 addressing issues of reproducibility, transparency, research ethics, and societal impact. Do not remove
1083 the checklist: **The papers not including the checklist will be desk rejected.** The checklist should
1084 follow the references and follow the (optional) supplemental material. The checklist does NOT count
1085 towards the page limit.

1086 Please read the checklist guidelines carefully for information on how to answer these questions. For
1087 each question in the checklist:

- 1088 • You should answer [Yes], [No], or [N/A].
- 1089 • [N/A] means either that the question is Not Applicable for that particular paper or the
1090 relevant information is Not Available.
- 1091 • Please provide a short (1–2 sentence) justification right after your answer (even for [N/A]).

1092 **The checklist answers are an integral part of your paper submission.** They are visible to the
1093 reviewers, area chairs, senior area chairs, and ethics reviewers. You will also be asked to include it
1094 (after eventual revisions) with the final version of your paper, and its final version will be published
1095 with the paper.

1096 The reviewers of your paper will be asked to use the checklist as one of the factors in their evaluation.
1097 While [Yes] is generally preferable to [No], it is perfectly acceptable to answer [No] provided a
1098 proper justification is given (e.g., error bars are not reported because it would be too computationally
1099 expensive” or “we were unable to find the license for the dataset we used”). In general, answering
1100 [No] or [N/A] is not grounds for rejection. While the questions are phrased in a binary way, we
1101 acknowledge that the true answer is often more nuanced, so please just use your best judgment and
1102 write a justification to elaborate. All supporting evidence can appear either in the main paper or the
1103 supplemental material, provided in appendix. If you answer [Yes] to a question, in the justification
1104 please point to the section(s) where related material for the question can be found.

1105 IMPORTANT, please:

- 1106 • **Delete this instruction block, but keep the section heading “NeurIPS Paper Checklist”,**
- 1107 • **Keep the checklist subsection headings, questions/answers and guidelines below.**
- 1108 • **Do not modify the questions and only use the provided macros for your answers.**

1109 1. Claims

1110 Question: Do the main claims made in the abstract and introduction accurately reflect the
1111 paper’s contributions and scope?

1112 Answer: [Yes]

1113 Justification: The abstract clearly delineates the paper’s main contributions and scope, and
1114 the introduction reiterates these claims in a structured, bullet-point form. All stated claims
1115 are directly supported by the theoretical and experimental results presented in the paper.

1116 Guidelines:

- 1117 • The answer [N/A] means that the abstract and introduction do not include the claims
1118 made in the paper.
- 1119 • The abstract and/or introduction should clearly state the claims made, including the
1120 contributions made in the paper and important assumptions and limitations. A [No] or
1121 [N/A] answer to this question will not be perceived well by the reviewers.
- 1122 • The claims made should match theoretical and experimental results, and reflect how
1123 much the results can be expected to generalize to other settings.
- 1124 • It is fine to include aspirational goals as motivation as long as it is clear that these goals
1125 are not attained by the paper.

1126 2. Limitations

1127 Question: Does the paper discuss the limitations of the work performed by the authors?

1128 Answer: [Yes]

1129 Justification: The paper discusses key limitations in Appendix E.

1130 Guidelines:

- 1131 • The answer [N/A] means that the paper has no limitation while the answer [No] means
1132 that the paper has limitations, but those are not discussed in the paper.
- 1133 • The authors are encouraged to create a separate “Limitations” section in their paper.
- 1134 • The paper should point out any strong assumptions and how robust the results are to
1135 violations of these assumptions (e.g., independence assumptions, noiseless settings,
1136 model well-specification, asymptotic approximations only holding locally). The authors
1137 should reflect on how these assumptions might be violated in practice and what the
1138 implications would be.
- 1139 • The authors should reflect on the scope of the claims made, e.g., if the approach was
1140 only tested on a few datasets or with a few runs. In general, empirical results often
1141 depend on implicit assumptions, which should be articulated.
- 1142 • The authors should reflect on the factors that influence the performance of the approach.
1143 For example, a facial recognition algorithm may perform poorly when image resolution
1144 is low or images are taken in low lighting. Or a speech-to-text system might not be
1145 used reliably to provide closed captions for online lectures because it fails to handle
1146 technical jargon.
- 1147 • The authors should discuss the computational efficiency of the proposed algorithms
1148 and how they scale with dataset size.
- 1149 • If applicable, the authors should discuss possible limitations of their approach to
1150 address problems of privacy and fairness.
- 1151 • While the authors might fear that complete honesty about limitations might be used by
1152 reviewers as grounds for rejection, a worse outcome might be that reviewers discover
1153 limitations that aren’t acknowledged in the paper. The authors should use their best
1154 judgment and recognize that individual actions in favor of transparency play an impor-
1155 tant role in developing norms that preserve the integrity of the community. Reviewers
1156 will be specifically instructed to not penalize honesty concerning limitations.

1157 3. Theory assumptions and proofs

1158 Question: For each theoretical result, does the paper provide the full set of assumptions and
1159 a complete (and correct) proof?

1160 Answer: [Yes]

1161 Justification: All theoretical results include clearly stated assumptions in the statement
1162 of theorems. Complete proofs are provided in Appendices A and B with appropriate
1163 cross-referencing from the main text.

1164 Guidelines:

- 1165 • The answer [N/A] means that the paper does not include theoretical results.
- 1166 • All the theorems, formulas, and proofs in the paper should be numbered and cross-
1167 referenced.
- 1168 • All assumptions should be clearly stated or referenced in the statement of any theorems.
- 1169 • The proofs can either appear in the main paper or the supplemental material, but if
1170 they appear in the supplemental material, the authors are encouraged to provide a short
1171 proof sketch to provide intuition.
- 1172 • Inversely, any informal proof provided in the core of the paper should be complemented
1173 by formal proofs provided in appendix or supplemental material.
- 1174 • Theorems and Lemmas that the proof relies upon should be properly referenced.

1175 4. Experimental result reproducibility

1176 Question: Does the paper fully disclose all the information needed to reproduce the main ex-
1177 perimental results of the paper to the extent that it affects the main claims and/or conclusions
1178 of the paper (regardless of whether the code and data are provided or not)?

1179 Answer: [Yes]

Justification: The paper provides sufficient detail to reproduce all experiments, including training procedures and implementation details described in the main text and Appendix C. Additionally, we release code with a comprehensive README to support the reproducibility of our results.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: We provide open access to code with detailed instructions in a comprehensive README, including environment setup and commands to train routers and reproduce the main experimental results. These files will automatically generate the data needed for training. Data generation has also been documented in Appendix C.

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.

- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [\[Yes\]](#)

Justification: The paper specifies all key training and evaluation details in the Section 6 and Appendix C.

Guidelines:

- The answer [\[N/A\]](#) means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [\[Yes\]](#)

Justification: The paper reports standard deviations in the main results tables, and further provides statistical significance via t-test p values in Appendix D.

Guidelines:

- The answer [\[N/A\]](#) means that the paper does not include experiments.
- The authors should answer [\[Yes\]](#) if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

1288 Answer: [Yes]

1289 Justification: The paper provides details on compute resources in Appendix C.

1290 Guidelines:

- 1291 • The answer [N/A] means that the paper does not include experiments.
- 1292 • The paper should indicate the type of compute workers CPU or GPU, internal cluster,
- 1293 or cloud provider, including relevant memory and storage.
- 1294 • The paper should provide the amount of compute required for each of the individual
- 1295 experimental runs as well as estimate the total compute.
- 1296 • The paper should disclose whether the full research project required more compute
- 1297 than the experiments reported in the paper (e.g., preliminary or failed experiments that
- 1298 didn't make it into the paper).

1299 **9. Code of ethics**

1300 Question: Does the research conducted in the paper conform, in every respect, with the

1301 NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

1302 Answer: [Yes]

1303 Justification: This paper does not involve any human subjects and all data has been simu-

1304 lated. This research focuses on numerical methods for physical systems and does not raise

1305 concerns related to safety, security, discrimination, surveillance, deception & harassment,

1306 environment, human rights, and bias & fairness.

1307 Guidelines:

- 1308 • The answer [N/A] means that the authors have not reviewed the NeurIPS Code of
- 1309 Ethics.
- 1310 • If the authors answer [No], they should explain the special circumstances that require a
- 1311 deviation from the Code of Ethics.
- 1312 • The authors should make sure to preserve anonymity (e.g., if there is a special consid-
- 1313 eration due to laws or regulations in their jurisdiction).

1314 **10. Broader impacts**

1315 Question: Does the paper discuss both potential positive societal impacts and negative

1316 societal impacts of the work performed?

1317 Answer: [Yes]

1318 Justification: The paper discusses broader impacts in the introduction section. No direct

1319 negative societal impacts are anticipated.

1320 Guidelines:

- 1321 • The answer [N/A] means that there is no societal impact of the work performed.
- 1322 • If the authors answer [N/A] or [No], they should explain why their work has no societal
- 1323 impact or why the paper does not address societal impact.
- 1324 • Examples of negative societal impacts include potential malicious or unintended uses
- 1325 (e.g., disinformation, generating fake profiles, surveillance), fairness considerations
- 1326 (e.g., deployment of technologies that could make decisions that unfairly impact specific
- 1327 groups), privacy considerations, and security considerations.
- 1328 • The conference expects that many papers will be foundational research and not tied
- 1329 to particular applications, let alone deployments. However, if there is a direct path to
- 1330 any negative applications, the authors should point it out. For example, it is legitimate
- 1331 to point out that an improvement in the quality of generative models could be used to
- 1332 generate Deepfakes for disinformation. On the other hand, it is not needed to point out
- 1333 that a generic algorithm for optimizing neural networks could enable people to train
- 1334 models that generate Deepfakes faster.
- 1335 • The authors should consider possible harms that could arise when the technology is
- 1336 being used as intended and functioning correctly, harms that could arise when the
- 1337 technology is being used as intended but gives incorrect results, and harms following
- 1338 from (intentional or unintentional) misuse of the technology.

- 1339 • If there are negative societal impacts, the authors could also discuss possible mitigation
1340 strategies (e.g., gated release of models, providing defenses in addition to attacks,
1341 mechanisms for monitoring misuse, mechanisms to monitor how a system learns from
1342 feedback over time, improving the efficiency and accessibility of ML).

1343 11. Safeguards

1344 Question: Does the paper describe safeguards that have been put in place for responsible
1345 release of data or models that have a high risk for misuse (e.g., pre-trained language models,
1346 image generators, or scraped datasets)?

1347 Answer: [N/A]

1348 Justification: The paper does not release any models or datasets at all. It focuses on method-
1349 ological contributions for numerical solvers using simulated data. Therefore, safeguards are
1350 not applicable.

1351 Guidelines:

- 1352 • The answer [N/A] means that the paper poses no such risks.
1353 • Released models that have a high risk for misuse or dual-use should be released with
1354 necessary safeguards to allow for controlled use of the model, for example by requiring
1355 that users adhere to usage guidelines or restrictions to access the model or implementing
1356 safety filters.
1357 • Datasets that have been scraped from the Internet could pose safety risks. The authors
1358 should describe how they avoided releasing unsafe images.
1359 • We recognize that providing effective safeguards is challenging, and many papers do
1360 not require this, but we encourage authors to take this into account and make a best
1361 faith effort.

1362 12. Licenses for existing assets

1363 Question: Are the creators or original owners of assets (e.g., code, data, models), used in
1364 the paper, properly credited and are the license and terms of use explicitly mentioned and
1365 properly respected?

1366 Answer: [N/A]

1367 Justification: The paper does not use any external datasets or pretrained models. All
1368 experiments are conducted using simulated data and models trained in-house.

1369 Guidelines:

- 1370 • The answer [N/A] means that the paper does not use existing assets.
1371 • The authors should cite the original paper that produced the code package or dataset.
1372 • The authors should state which version of the asset is used and, if possible, include a
1373 URL.
1374 • The name of the license (e.g., CC-BY 4.0) should be included for each asset.
1375 • For scraped data from a particular source (e.g., website), the copyright and terms of
1376 service of that source should be provided.
1377 • If assets are released, the license, copyright information, and terms of use in the
1378 package should be provided. For popular datasets, `paperswithcode.com/datasets`
1379 has curated licenses for some datasets. Their licensing guide can help determine the
1380 license of a dataset.
1381 • For existing datasets that are re-packaged, both the original license and the license of
1382 the derived asset (if it has changed) should be provided.
1383 • If this information is not available online, the authors are encouraged to reach out to
1384 the asset’s creators.

1385 13. New assets

1386 Question: Are new assets introduced in the paper well documented and is the documentation
1387 provided alongside the assets?

1388 Answer: [Yes]

1389 Justification: We release code for our method which is documented with a detailed
1390 README.

1391 Guidelines:

1392 • The answer [N/A] means that the paper does not release new assets.

1393 • Researchers should communicate the details of the dataset/code/model as part of their

1394 submissions via structured templates. This includes details about training, license,

1395 limitations, etc.

1396 • The paper should discuss whether and how consent was obtained from people whose

1397 asset is used.

1398 • At submission time, remember to anonymize your assets (if applicable). You can either

1399 create an anonymized URL or include an anonymized zip file.

1400 **14. Crowdsourcing and research with human subjects**

1401 Question: For crowdsourcing experiments and research with human subjects, does the paper

1402 include the full text of instructions given to participants and screenshots, if applicable, as

1403 well as details about compensation (if any)?

1404 Answer: [N/A]

1405 Justification: The paper does not involve crowdsourcing or any research with human subjects.

1406 All experiments are conducted using simulated data.

1407 Guidelines:

1408 • The answer [N/A] means that the paper does not involve crowdsourcing nor research

1409 with human subjects.

1410 • Including this information in the supplemental material is fine, but if the main contribu-

1411 tion of the paper involves human subjects, then as much detail as possible should be

1412 included in the main paper.

1413 • According to the NeurIPS Code of Ethics, workers involved in data collection, curation,

1414 or other labor should be paid at least the minimum wage in the country of the data

1415 collector.

1416 **15. Institutional review board (IRB) approvals or equivalent for research with human**

1417 **subjects**

1418 Question: Does the paper describe potential risks incurred by study participants, whether

1419 such risks were disclosed to the subjects, and whether Institutional Review Board (IRB)

1420 approvals (or an equivalent approval/review based on the requirements of your country or

1421 institution) were obtained?

1422 Answer: [N/A]

1423 Justification: This paper does not involve human subjects. All experiments are conducted

1424 using simulated data. Therefore, IRB approval is not applicable.

1425 Guidelines:

1426 • The answer [N/A] means that the paper does not involve crowdsourcing nor research

1427 with human subjects.

1428 • Depending on the country in which research is conducted, IRB approval (or equivalent)

1429 may be required for any human subjects research. If you obtained IRB approval, you

1430 should clearly state this in the paper.

1431 • We recognize that the procedures for this may vary significantly between institutions

1432 and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the

1433 guidelines for their institution.

1434 • For initial submissions, do not include any information that would break anonymity (if

1435 applicable), such as the institution conducting the review.

1436 **16. Declaration of LLM usage**

1437 Question: Does the paper describe the usage of LLMs if it is an important, original, or

1438 non-standard component of the core methods in this research? Note that if the LLM is used

1439 only for writing, editing, or formatting purposes and does *not* impact the core methodology,

1440 scientific rigor, or originality of the research, declaration is not required.

1441 Answer: [N/A]

1442 Justification: LLMs are primarily used for writing and editing. Further details are provided
1443 in Appendix F.
1444 Guidelines:
1445 • The answer [N/A] means that the core method development in this research does not
1446 involve LLMs as any important, original, or non-standard components.
1447 • Please refer to our LLM policy in the NeurIPS handbook for what should or should not
1448 be described.